



Cloud Native Architecture at Netflix

Yow December 2013 (Brisbane)

Adrian Cockcroft

@adrianco @NetflixOSS

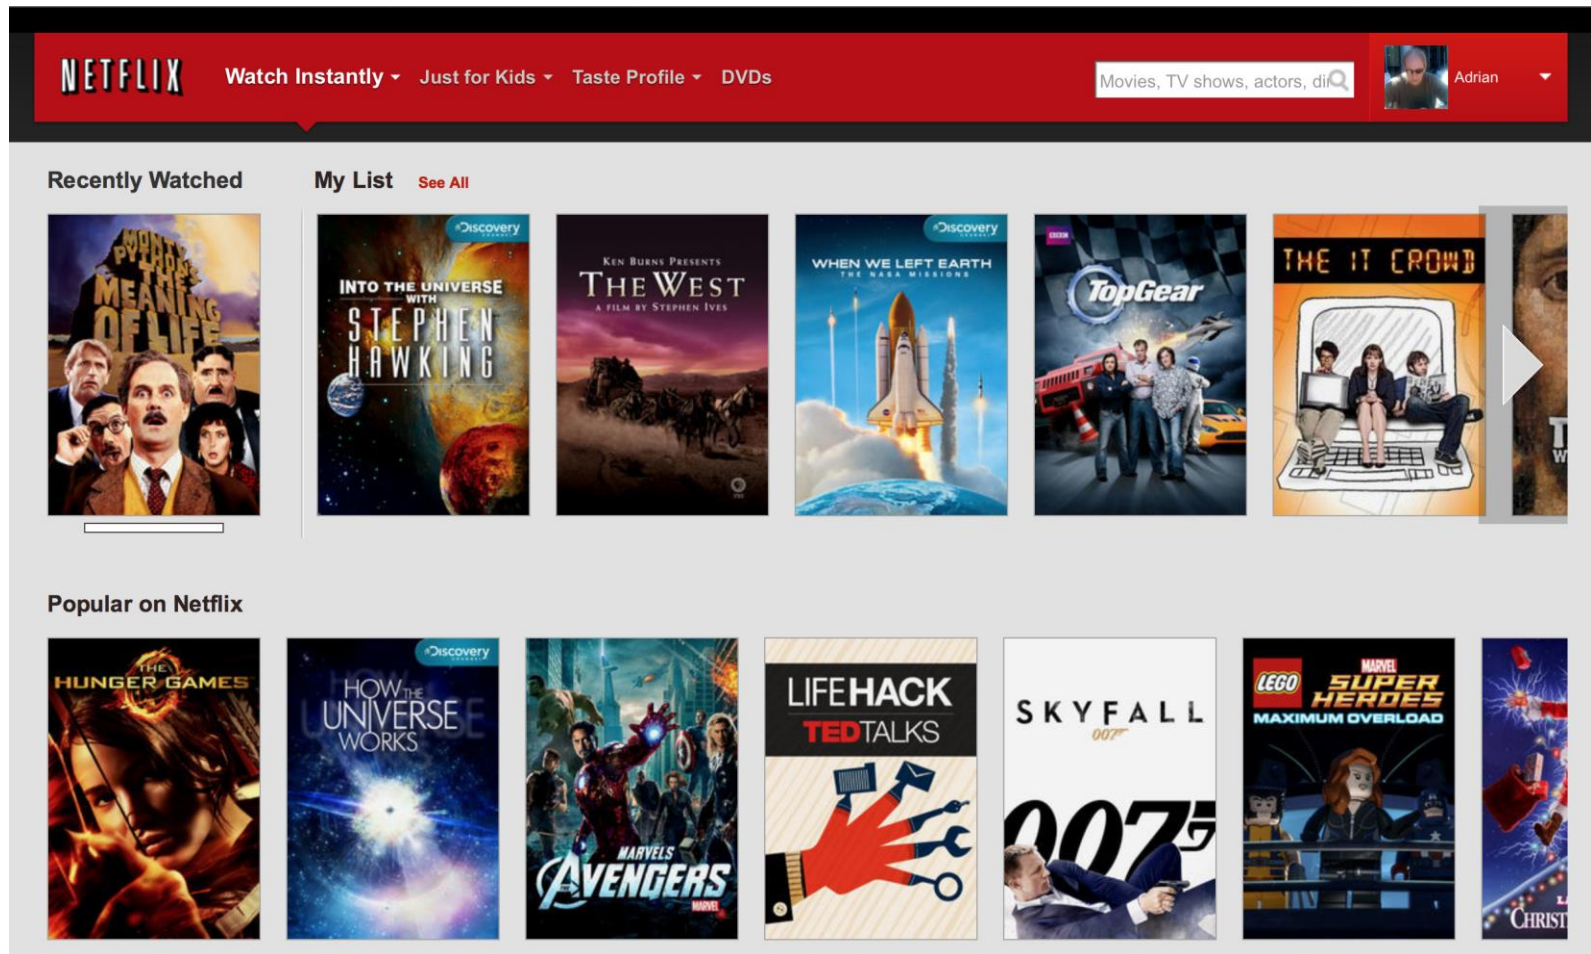
<http://www.linkedin.com/in/adriancockcroft>

Netflix History (current size)

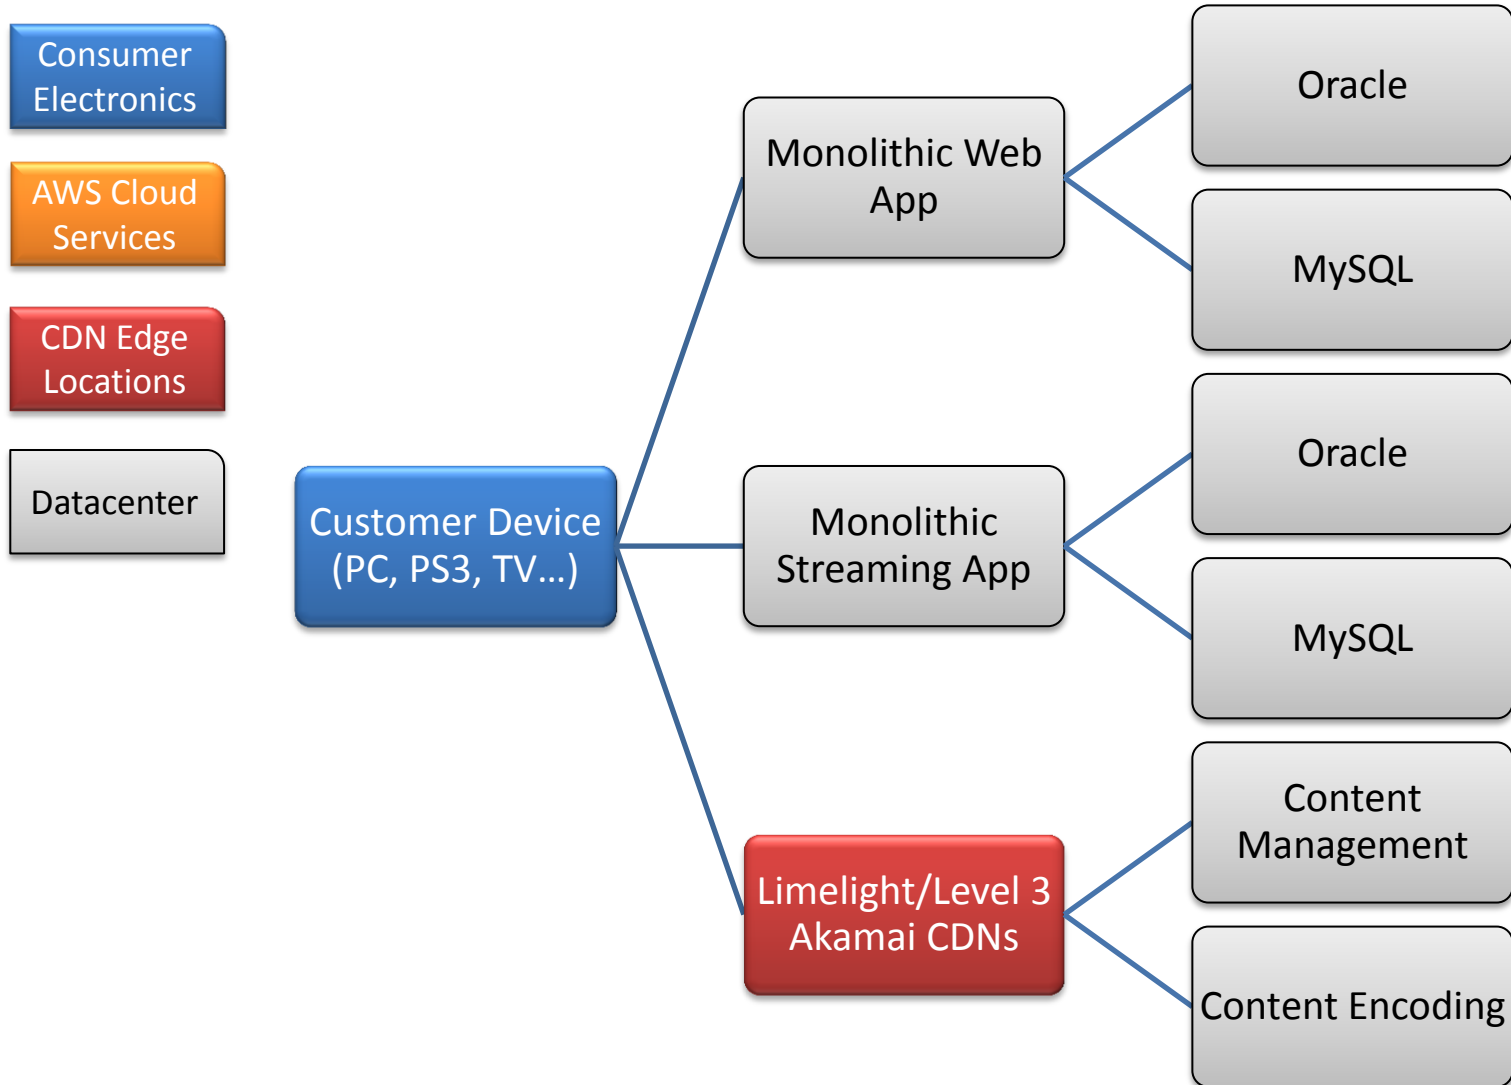
- 1998 DVD Shipping Service in USA (~7M users)
- 2007 Streaming video in USA (~31M users)
- International streaming video (~9M users)
 - 2010 Canada
 - 2011 Latin America
 - 2012 UK and Ireland
 - 2012 Nordics
 - 2013 Netherlands

Netflix Member Web Site Home Page

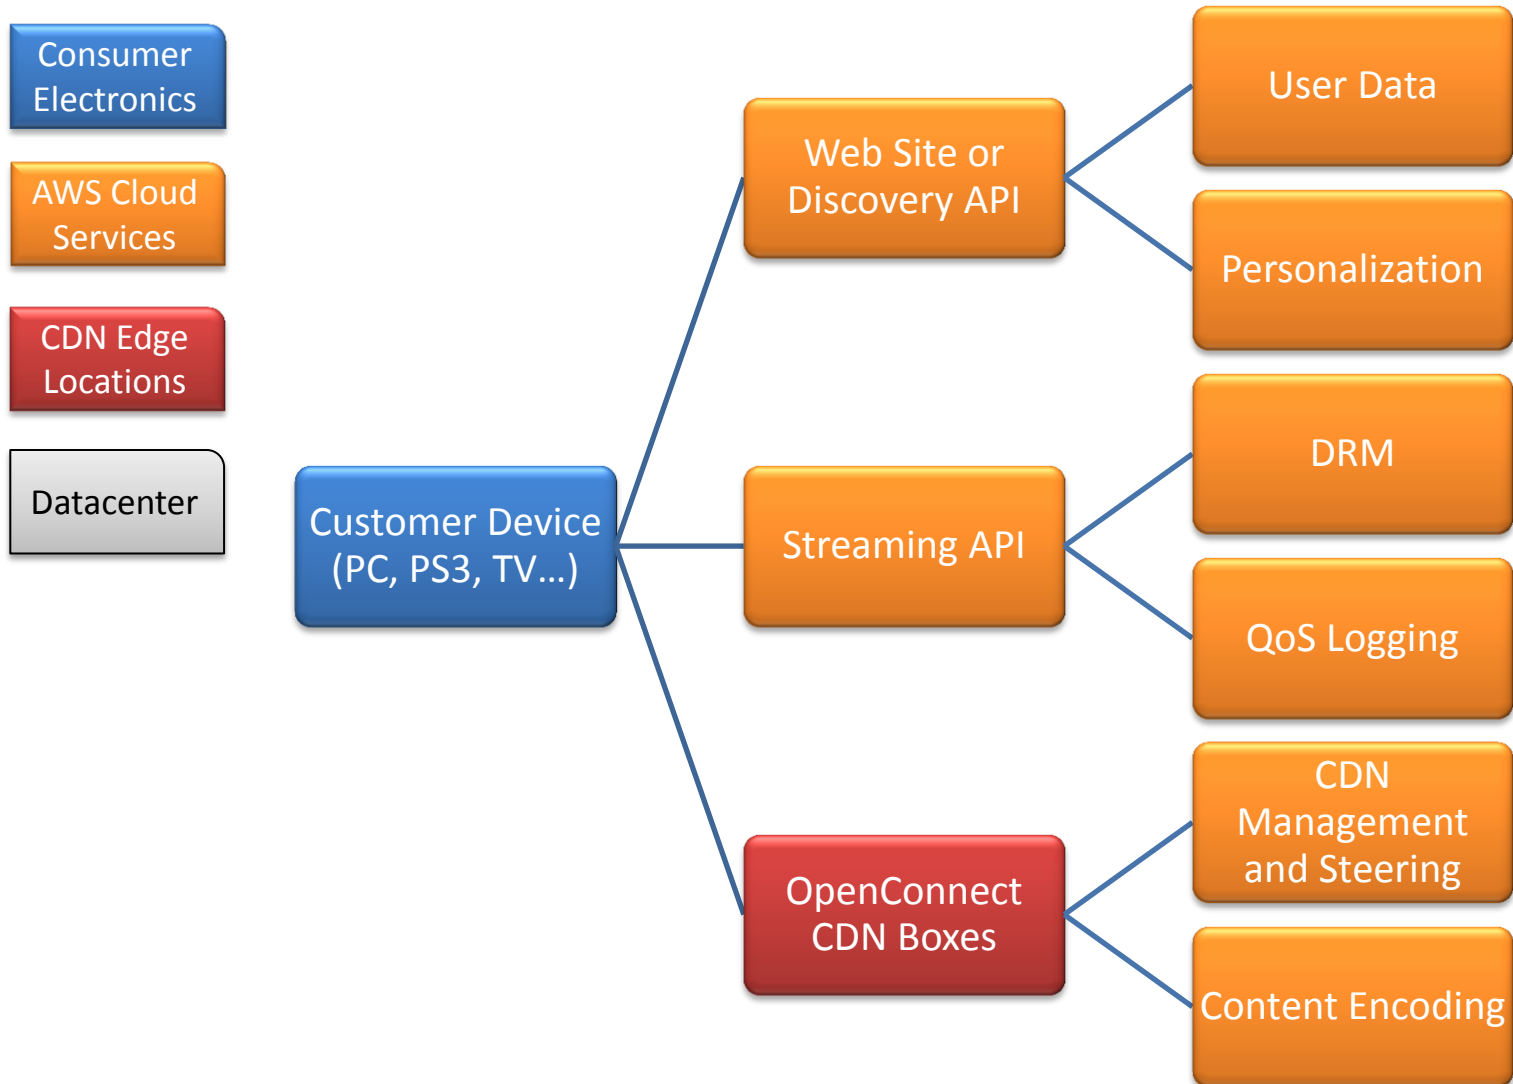
Personalization Driven – How Does It Work?



How Netflix Used to Work



How Netflix Streaming Works Today



2H 2013

1H 2013

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	36.35%	Netflix	31.62%	Netflix	28.18%
2	HTTP	6.03%	YouTube	18.69%	YouTube	16.78%
3	SSL	5.87%	HTTP	9.74%	HTTP	9.26%
4	Netflix	4.44%	BitTorrent	4.05%	BitTorrent	7.39%
5	YouTube	3.63%	iTunes	3.27%	iTunes	2.91%
6	Skype	2.76%	MPEG - Other	2.60%	SSL	2.54%
7	QVoD	2.55%	SSL	2.05%	MPEG - Other	2.32%
8	Facebook	1.54%	Amazon Video	1.61%	Amazon Video	1.48%
9	FaceTime	1.44%	Facebook	1.31%	Facebook	1.34%
10	Dropbox	1.39%	Hulu	1.29%	Hulu	1.15%
		66.00%		76.23%		73.35%

 sandvine

Table 2 - Top 10 Peak Period Applications - North America, Fixed Access

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	34.81%	Netflix	32.25%	Netflix	28.88%
2	HTTP	7.53%	YouTube	17.11%	YouTube	15.43%
3	SSL	5.81%	HTTP	11.11%	HTTP	10.66%
4	Netflix	5.38%	BitTorrent	5.57%	BitTorrent	9.23%
5	Skype	4.88%	MPEG	2.58%	SSL	2.39%
6	YouTube	3.71%	Hulu	2.41%	MPEG	2.30%
7	Facebook	1.71%	iTunes	1.90%	Hulu	2.16%
8	Apple Photostream	1.34%	SSL	1.89%	iTunes	1.71%
9	Dropbox	1.21%	Flash Video	1.72%	Flash Video	1.53%
10	Carbonite	0.99%	Facebook	1.48%	Facebook	1.52%
Top 10		67.38%		78.03%		75.82%

 sandvine

Table 2 - Top 10 Peak Period Applications - North America, Fixed Access

2H 2012

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	36.8%	Netflix	33.0%	Netflix	28.8%
2	HTTP	9.83%	YouTube	14.8%	YouTube	13.1%
3	Skype	4.76%	HTTP	12.0%	HTTP	11.7%
4	Netflix	4.51%	BitTorrent	5.89%	BitTorrent	10.3%
5	SSL	3.73%	iTunes	3.92%	iTunes	3.43%
6	YouTube	2.70%	MPEG	2.22%	SSL	2.23%
7	PPStream	1.65%	Flash Video	2.21%	MPEG	2.05%
8	Facebook	1.62%	SSL	1.97%	Flash Video	2.01%
9	Apple PhotoStream	1.46%	Amazon Video	1.75%	Facebook	1.50%
10	Dropbox	1.17%	Facebook	1.48%	RTMP	1.41%
Top 10		68.24%	Top 10	79.01%	Top 10	76.54%


 sandvine

Table 3 - Top 10 Peak Period Applications (North America, Fixed Access)

Netflix Scale

- Tens of thousands of instances on AWS
 - Typically 4 core, 30GByte, Java business logic
 - Thousands created/removed every day
- Thousands of Cassandra NoSQL nodes on AWS
 - Many hi1.4xl - 8 core, 60Gbyte, 2TByte of SSD
 - 65 different clusters, over 300TB data, triple zone
 - Over 40 are multi-region clusters (6, 9 or 12 zone)
 - Biggest 288 m2.4xl – over 300K rps, 1.3M wps



Reactions over time

2009 “You guys are crazy! Can’t believe it”

2010 “What Netflix is doing won’t work”

2011 “It only works for ‘Unicorns’ like Netflix”

2012 “We’d like to do that but can’t”

2013 “We’re on our way using Netflix OSS code”



YOW! Workshop

175 slides of Netflix Architecture

See bit.ly/netflix-workshop

A whole day...



This Talk

Abstract the principles from the
architecture



Objectives:

Scalability

Availability

Agility

Efficiency



Principles:

Immutability

Separation of Concerns

Anti-fragility

High trust organization

Sharing



Outcomes:

- Public cloud – scalability, agility, sharing
- Micro-services – separation of concerns
- De-normalized data – separation of concerns
- Chaos Engines – anti-fragile operations
- Open source by default – agility, sharing
- Continuous deployment – agility, immutability
- DevOps – high trust organization, sharing
- Run-what-you-wrote – anti-fragile development

When to use public cloud?





adrian cockcroft @adrianco

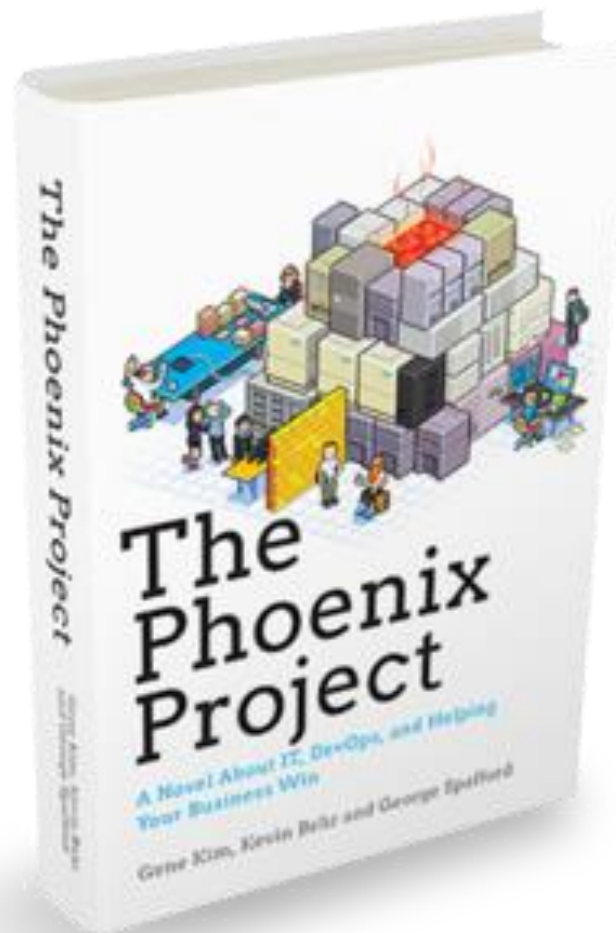
10 Apr

Baffling-late-adopters as a Service

Retweeted by Andrew Clay Shafer

Expand

"This is the IT swamp draining manual for anyone who is neck deep in alligators." -
Adrian Cockcroft, Cloud Architect at Netflix



Goal of Traditional IT:
Reliable hardware
running stable software

SCALE

Breaks hardware

...*SPEED*

Breaks software

SPEED at

SCALE

Breaks everything

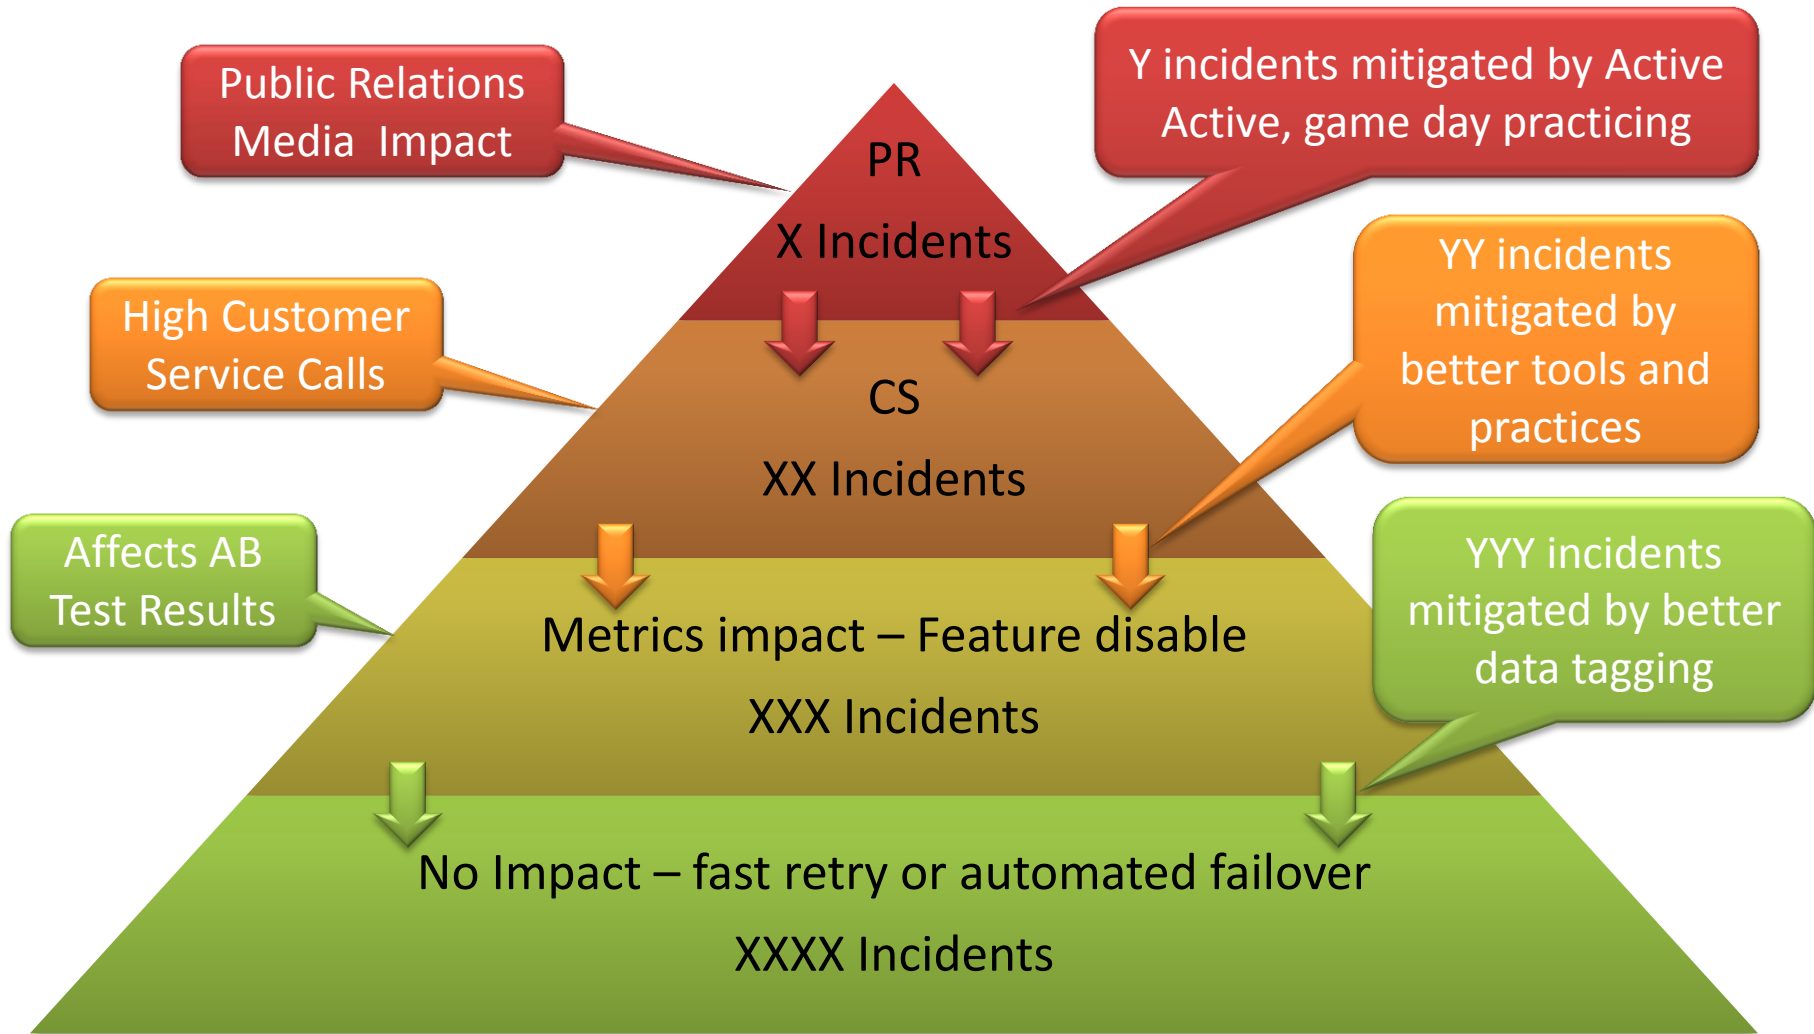
The STRANGE WORLD of the FUTURE

*STRANDED without video!
No way to fill their empty hours!
They were victims of...*

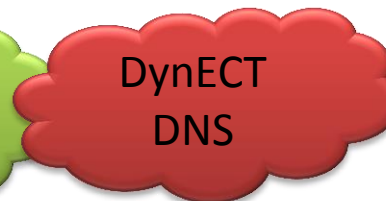
THE CLOUD OF BROKEN STREAMS

CREATED WITH PULP-O-MIZER COVER MAKER

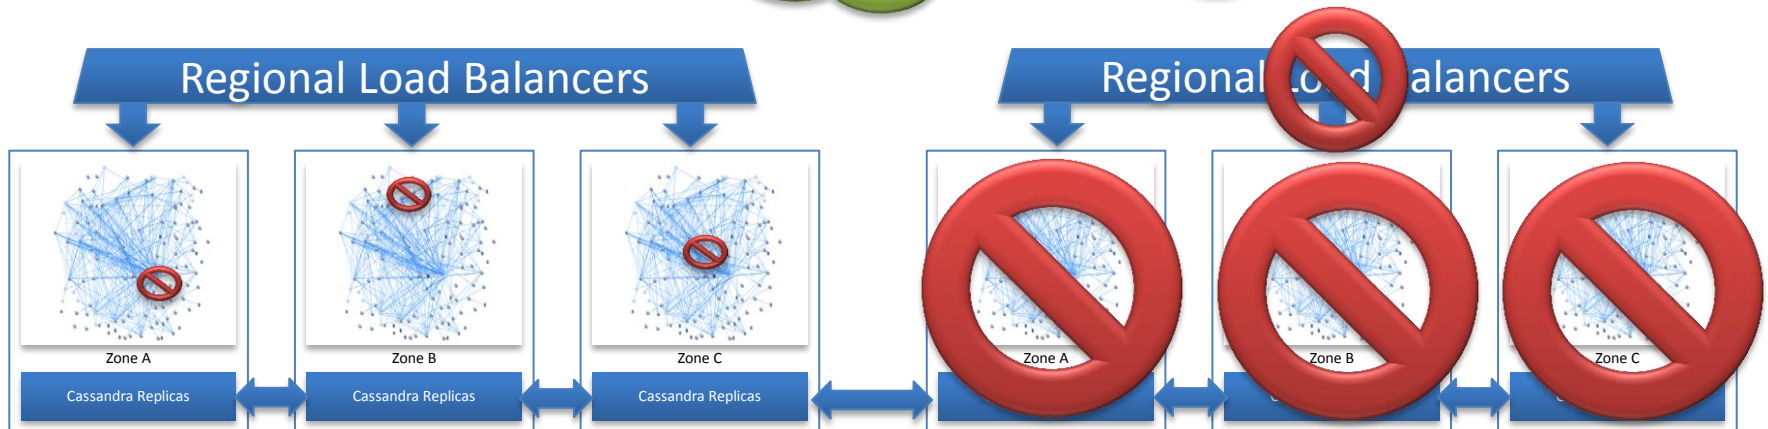
Incidents – Impact and Mitigation



Web Scale Architecture



DNS
Automation

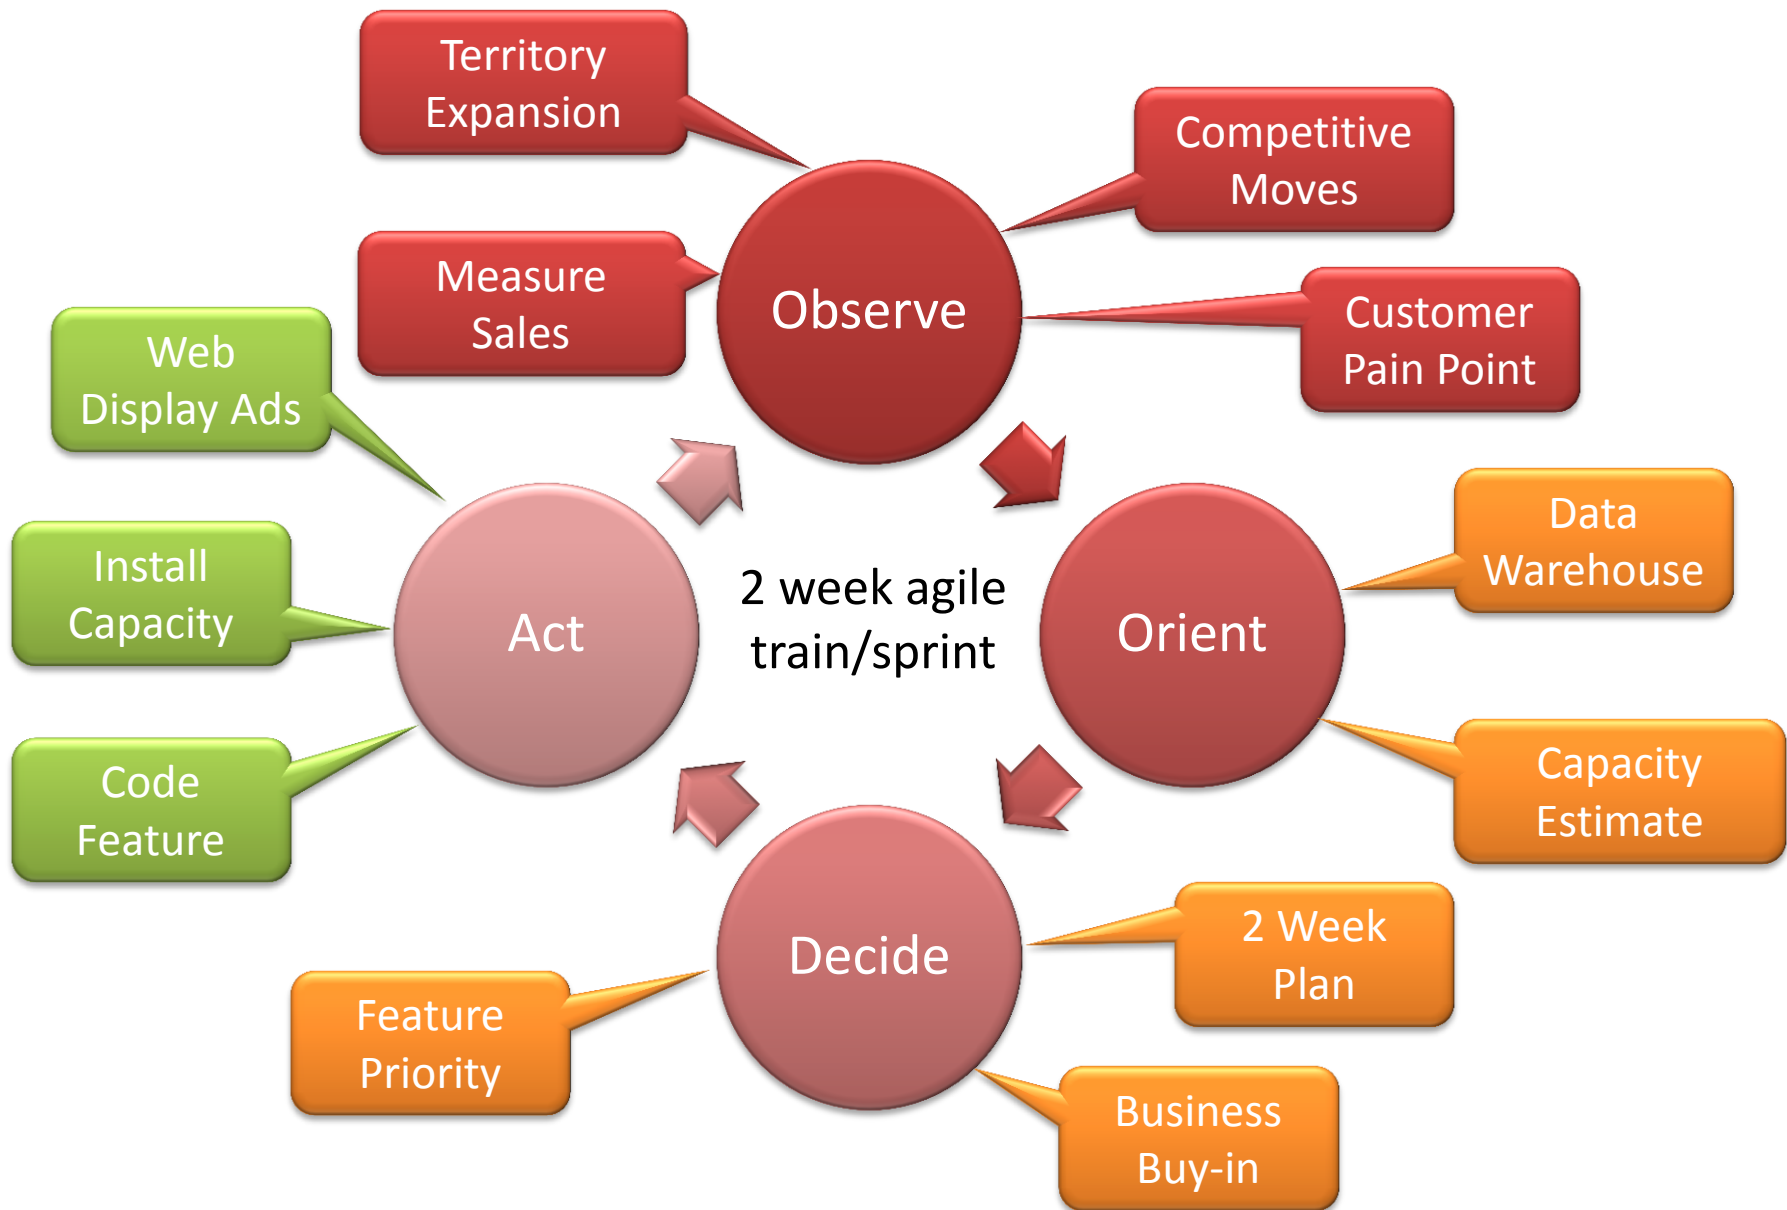


“Get inside your adversaries'
OODA loop to disorient them”

Colonel Boyd, USAF

“Agile” vs. “Continuous Delivery”

Speed Wins



2 Week Train Model Hand-Off Steps

Product Manager – 2 days



```
graph TD; A[Product Manager – 2 days] --> B[Developer – 2 days coding, 2 days meetings]; B --> C[QA Integration Team – 3 days]; C --> D[Operations Deploy Team – 4 days]; D --> E[BI Analytics Team – 1 day];
```

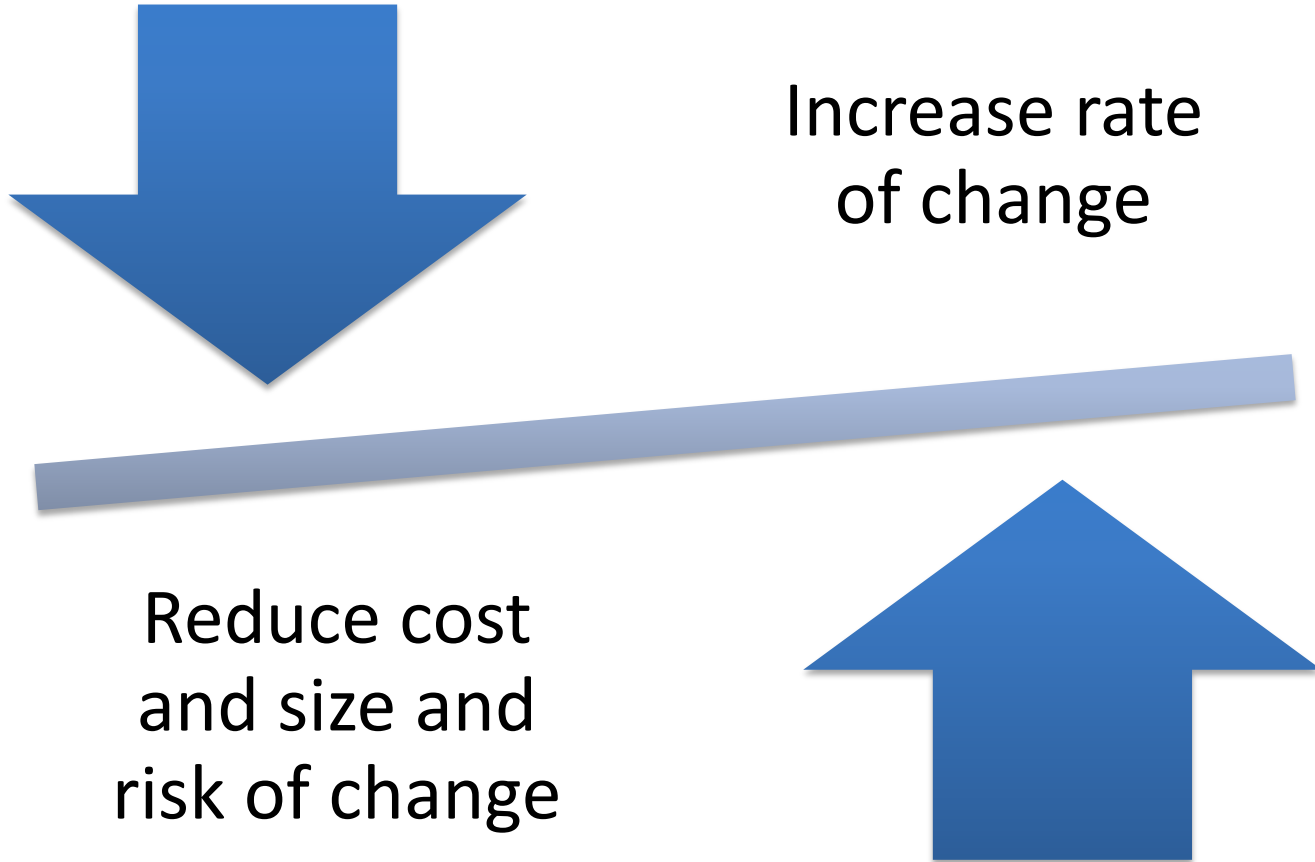
Developer – 2 days coding, 2 days meetings

QA Integration Team – 3 days

Operations Deploy Team – 4 days

BI Analytics Team – 1 day

What's Next?





Cloud Native

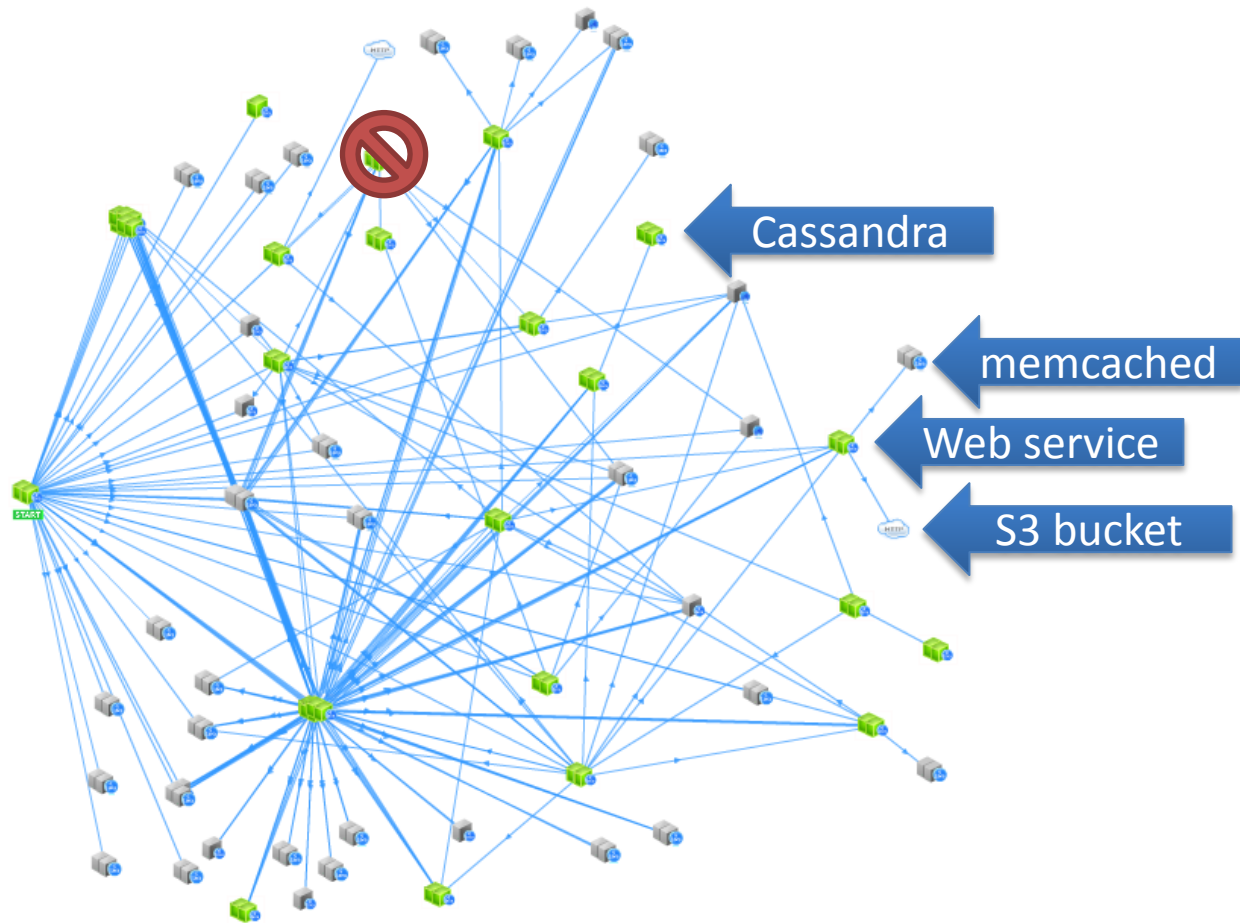
Construct a highly agile and highly available service from ephemeral and assumed broken components

Cloud Native Microservices

Each icon is
three to a few
hundred
instances
across three
AWS zones

Start Here →

Each microservice
is updated
independently
and continuously
at whatever rate
seems appropriate





Continuous Deployment

No time for handoff to IT



Developer Self Service

Freedom and Responsibility



Developers run what they wrote

Root access *and* pagerduty



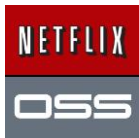
IT is a Cloud API

DEVops automation



Github all the things!

Leverage social coding



Netflix.github.com

NETFLIX

OSS

Netflix Open Source Software Center


Repositories

Commit Timeline

Community

Around the Web

AMIs



Getting Started

Welcome to the Netflix Open Source Software Center. To begin, we recommend working with our **RSS Reader application**. See [this walkthrough](#) on AWS Answers to get up and running quickly.

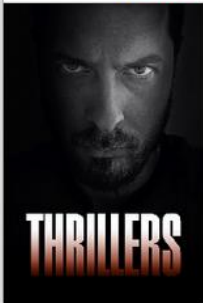
After you've tackled that, check out the **IBM ACME Air** and **Flux Capacitor** apps.

Also, be sure to join our mailing lists and follow us [@NetflixOSS](#) to stay up to date.

Our Repositories

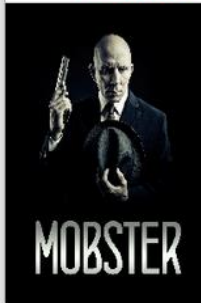
Availability

HYSTRIX




THRILLERS

SIMIANARMY



MOBSTER

TURBINE



FOREIGN

35 public repos

28 members

Mailing Lists

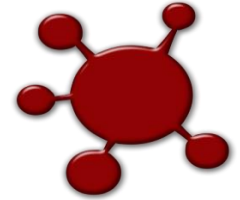
We maintain mailing lists for some of our open source projects. Click on the following links to subscribe and/or maintain your subscriptions.

- Asgard Mailing List
- Astyanax Mailing List
- Curator Mailing List
- Exhibitor Mailing List
- SimianArmy Mailing List

Stay in Touch

35 repos
36 repos
37 today

Karyon Microservice Template



Bootstrapping

Dependency & Lifecycle management via Governor.

Service registry via Eureka.

Property management via Archaius

Hooks for Latency Monkey injection testing

Preconfigured status page and heathcheck servlets

```
localhost:8989/hello-netflix-oss/rest/v1/hello/to/newbee
{
  Message: "Hello newbie from Netflix OSS"
}
```

Karyon Microservice Status Page

Eureka discovery service metadata

Environment variables

JMX

Versions



PLATFORMSERVICE v1.0 us-east-1_prod

Asyanax Counters Debug Data Diagnostics Discovery Environment Info Jars JMX NWS Properties SLA Tracers

Refresh Machine Readable filter

Filter: [x] [+] [-]

com.netflix.servo.name=LBStats_AvailableZones,*

Attributes

Refresh

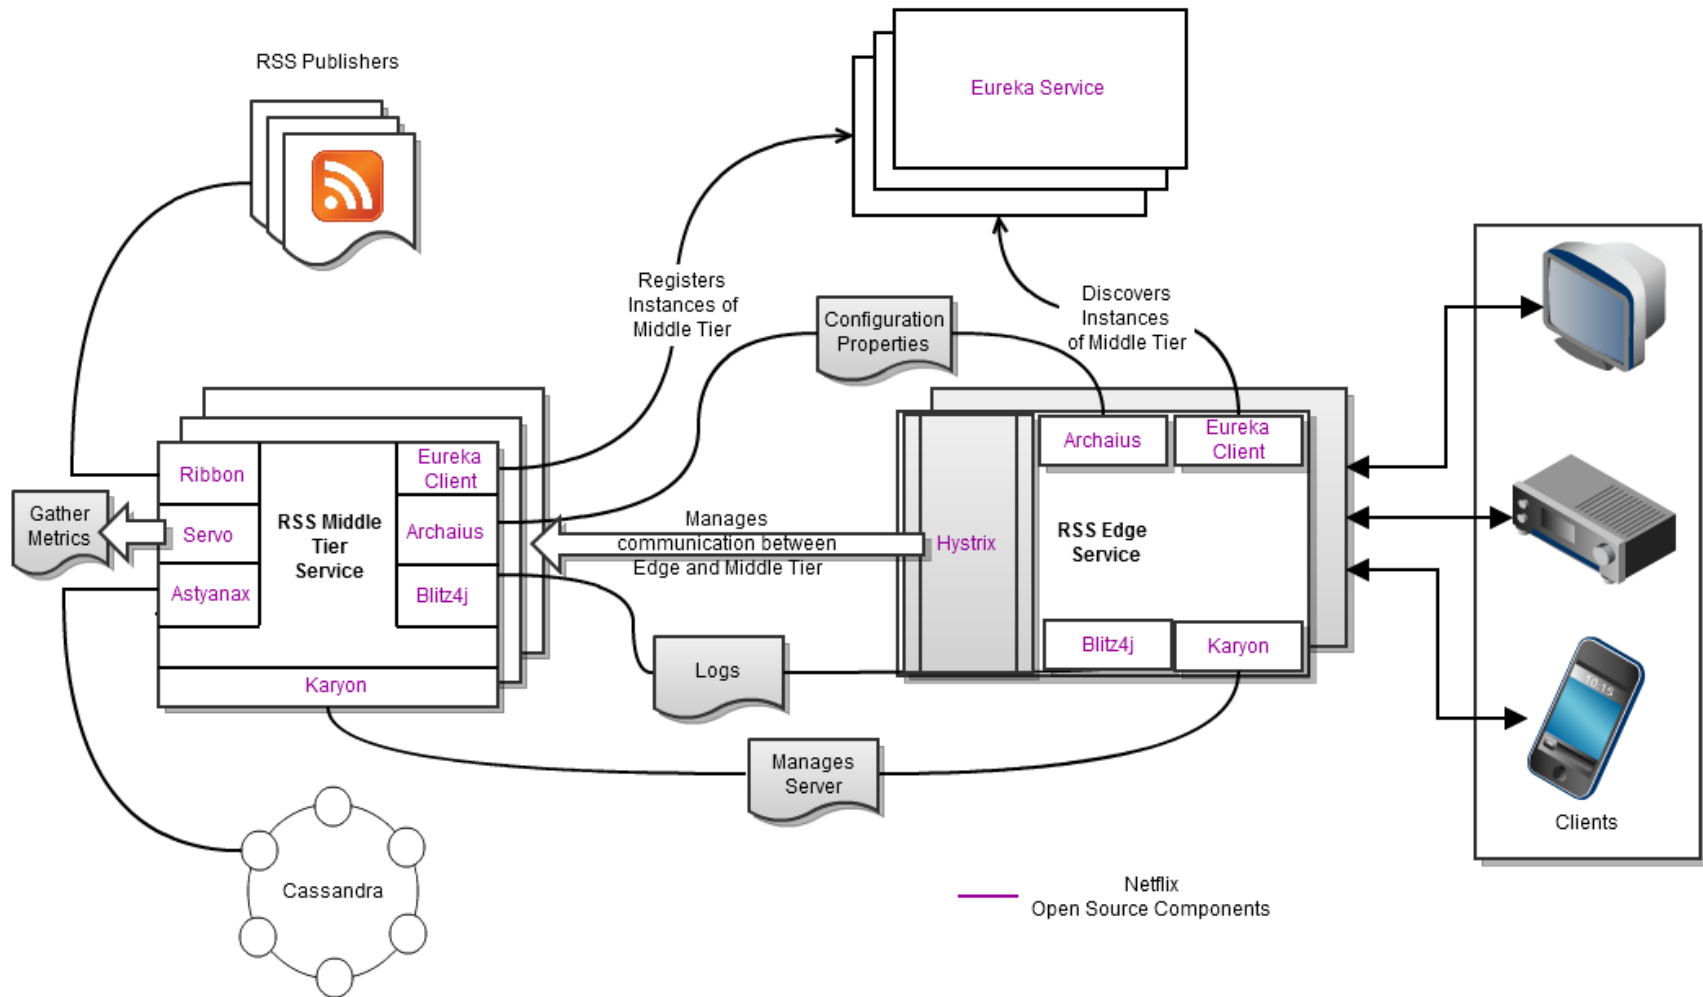
Key	Name	Value
{\"class\":\"LoadBalancerStats\", \"id\":\"akmsclient\", \"level\":\"INFO\", \"type\":\"INFORMATIONAL\"}	value	[us-east-1e, us-east-1d, us-east-1c]
{\"class\":\"LoadBalancerStats\", \"id\":\"atlas_publish\", \"level\":\"INFO\", \"type\":\"INFORMATIONAL\"}	value	[us-east-1e, us-east-1d, us-east-1c]
{\"class\":\"LoadBalancerStats\", \"id\":\"chukwatrackerclient\", \"level\":\"INFO\", \"type\":\"INFORMATIONAL\"}	value	[us-east-1e, us-east-1d, us-east-1c]
{\"class\":\"LoadBalancerStats\", \"id\":\"defaultRestClient\", \"level\":\"INFO\", \"type\":\"INFORMATIONAL\"}	value	[]
{\"class\":\"LoadBalancerStats\", \"id\":\"epic_publish\", \"level\":\"INFO\", \"type\":\"INFORMATIONAL\"}	value	[us-east-1e, us-east-1d, us-east-1c]
{\"class\":\"LoadBalancerStats\", \"id\":\"epicplugin\", \"level\":\"INFO\", \"type\":\"INFORMATIONAL\"}	value	[us-east-1e, us-east-1d, us-east-1c]

Showing 1 to 6 of 6 entries

Error: Visible: 898 Total: 898 Last updated: Mon Mar 04 2013 23:13:52

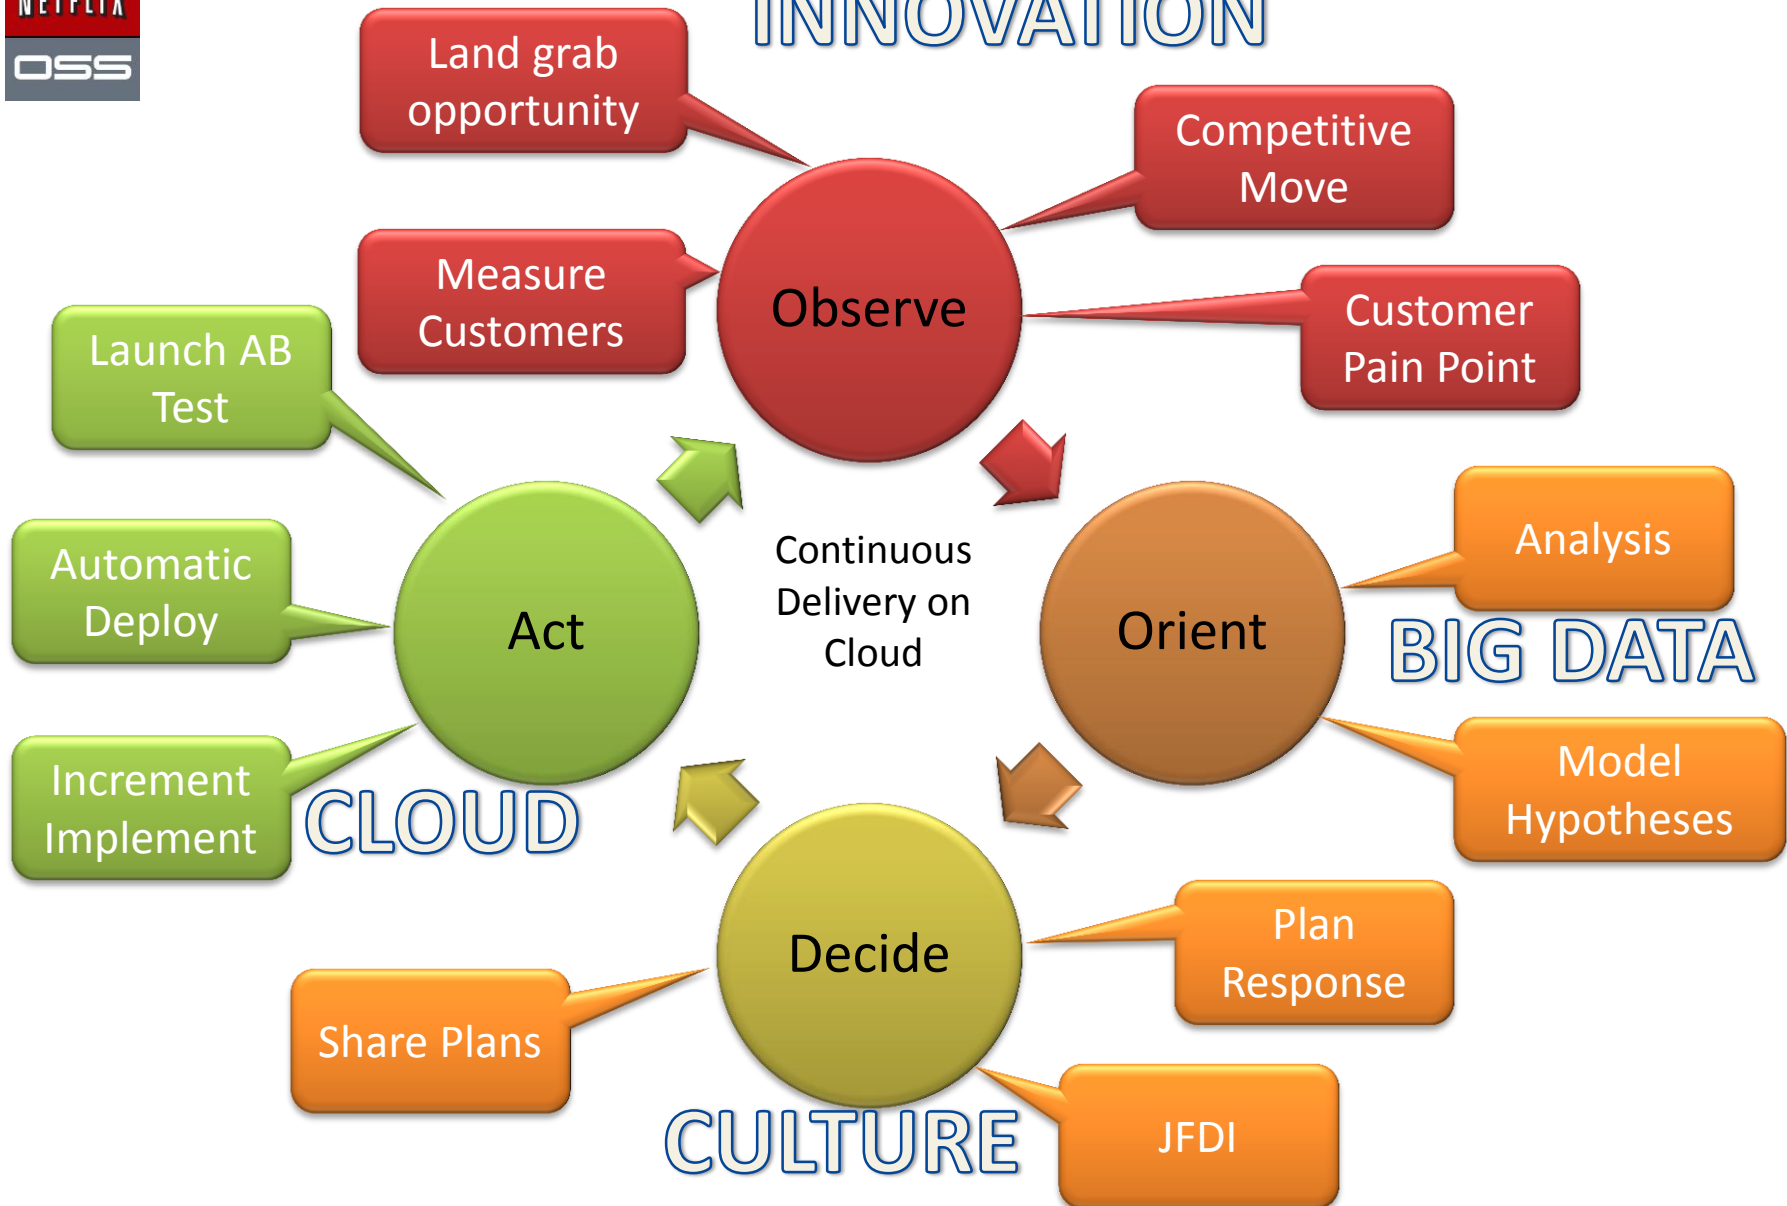
Conformity Monkey Support

Sample Application – RSS Reader



Putting it all together...

INNOVATION





Continuous Deploy Hand-Off

Product Manager - 2 days

A/B test setup and enable

Self service hypothesis test results

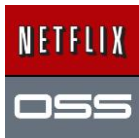


Developer – 2 days

Automated test

Automated deploy, on call

Self service analytics



Continuous Deploy Automation

Check in code, Jenkins build

Bake AMI, launch in test env

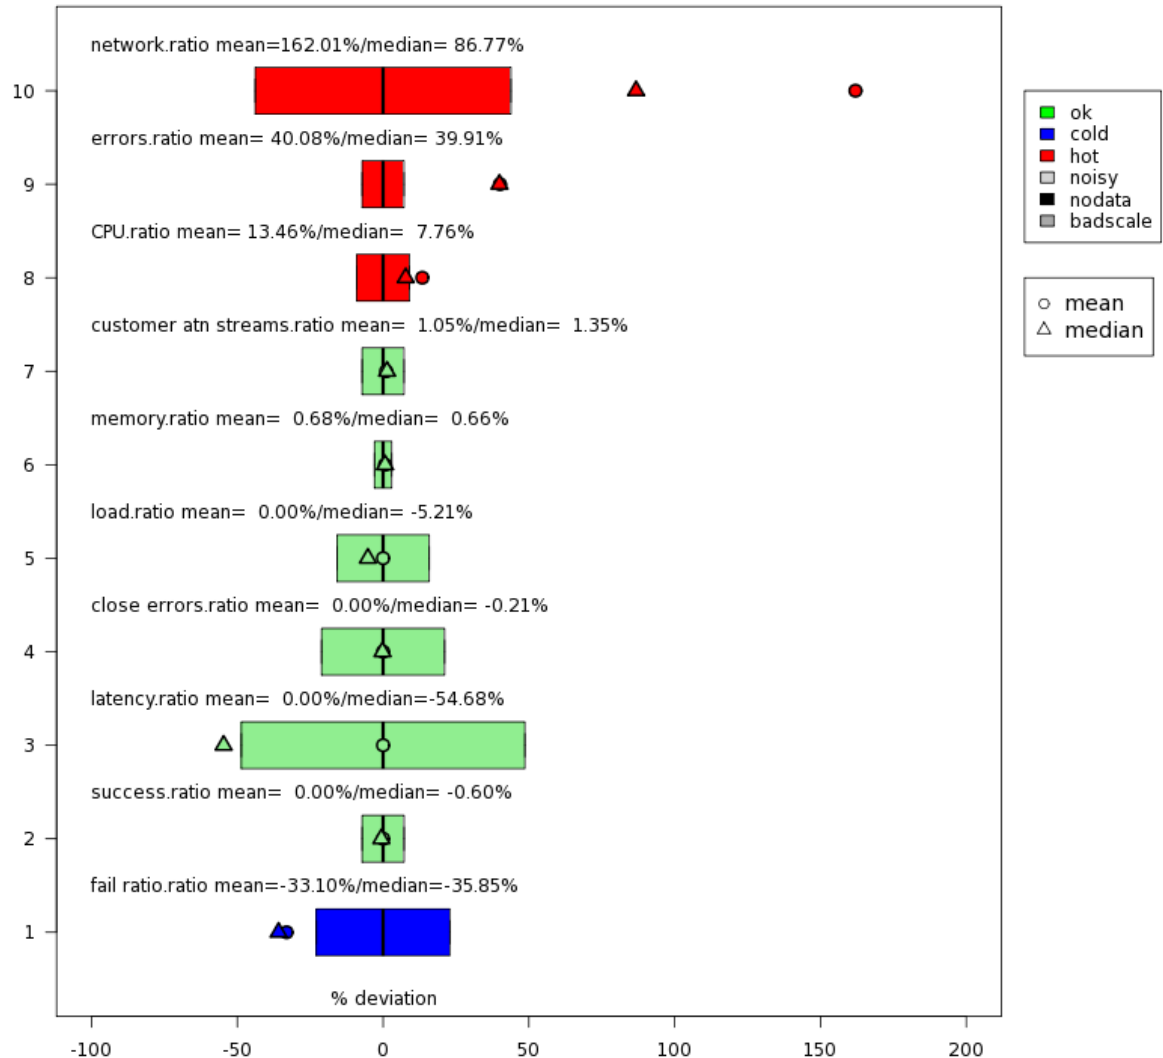
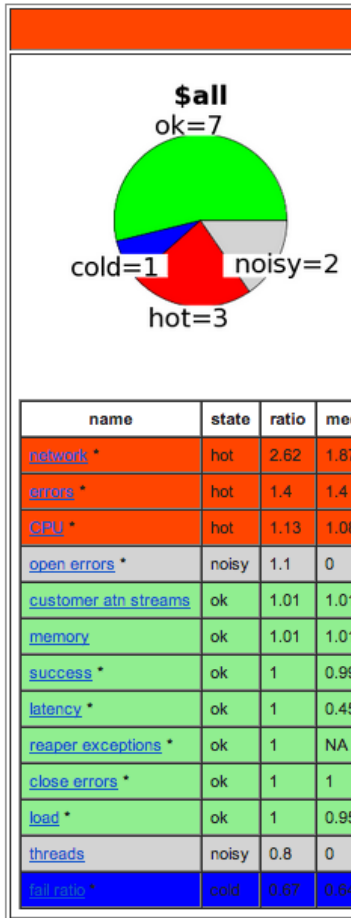
Functional and performance test

Production canary test

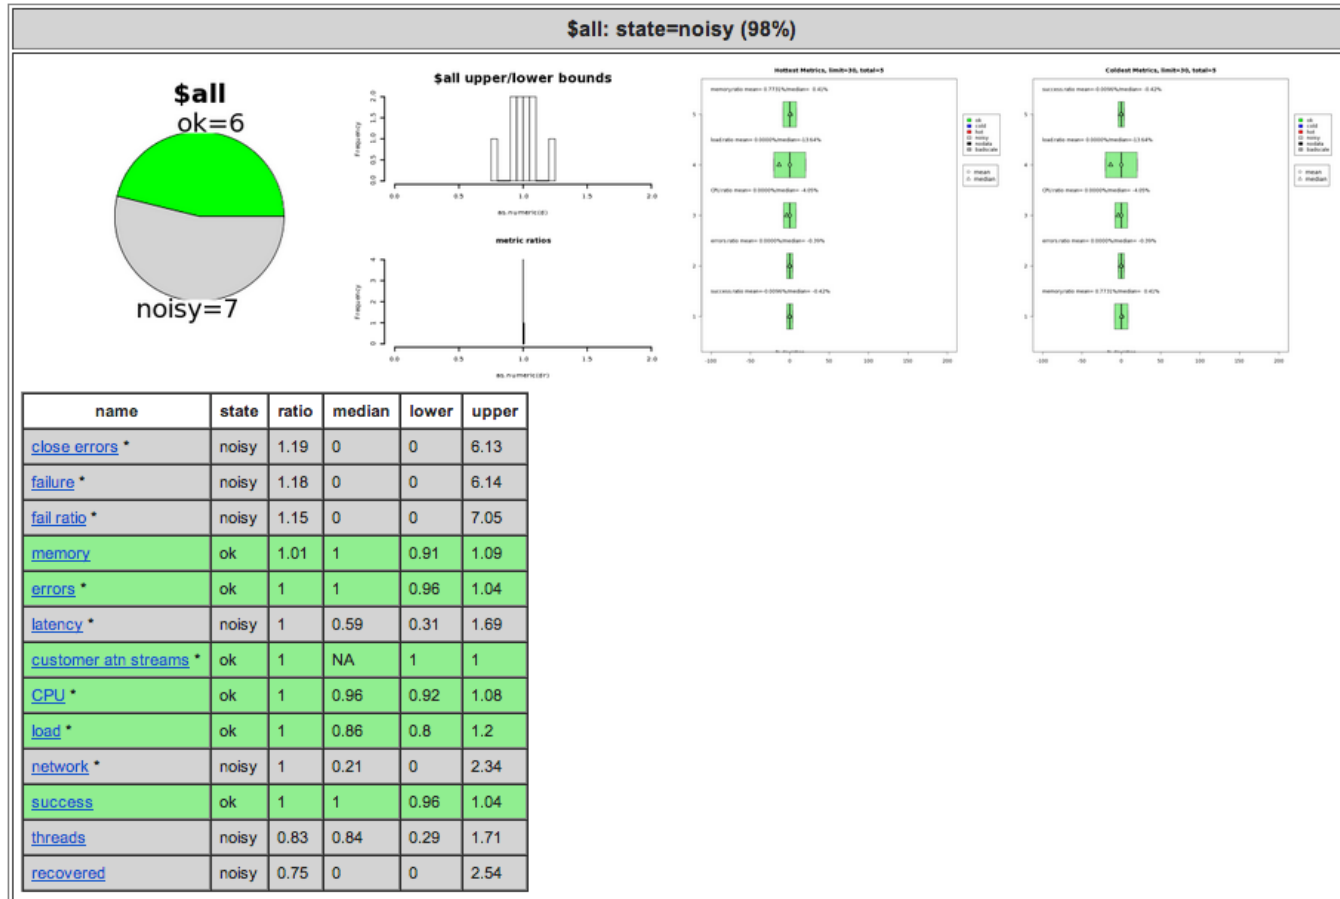
Production red/black push

Bad Canary Signature

Hottest Metrics, limit=30, total=10

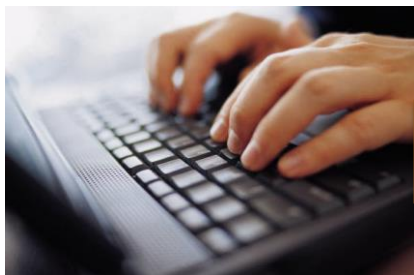


Happy Canary Signature



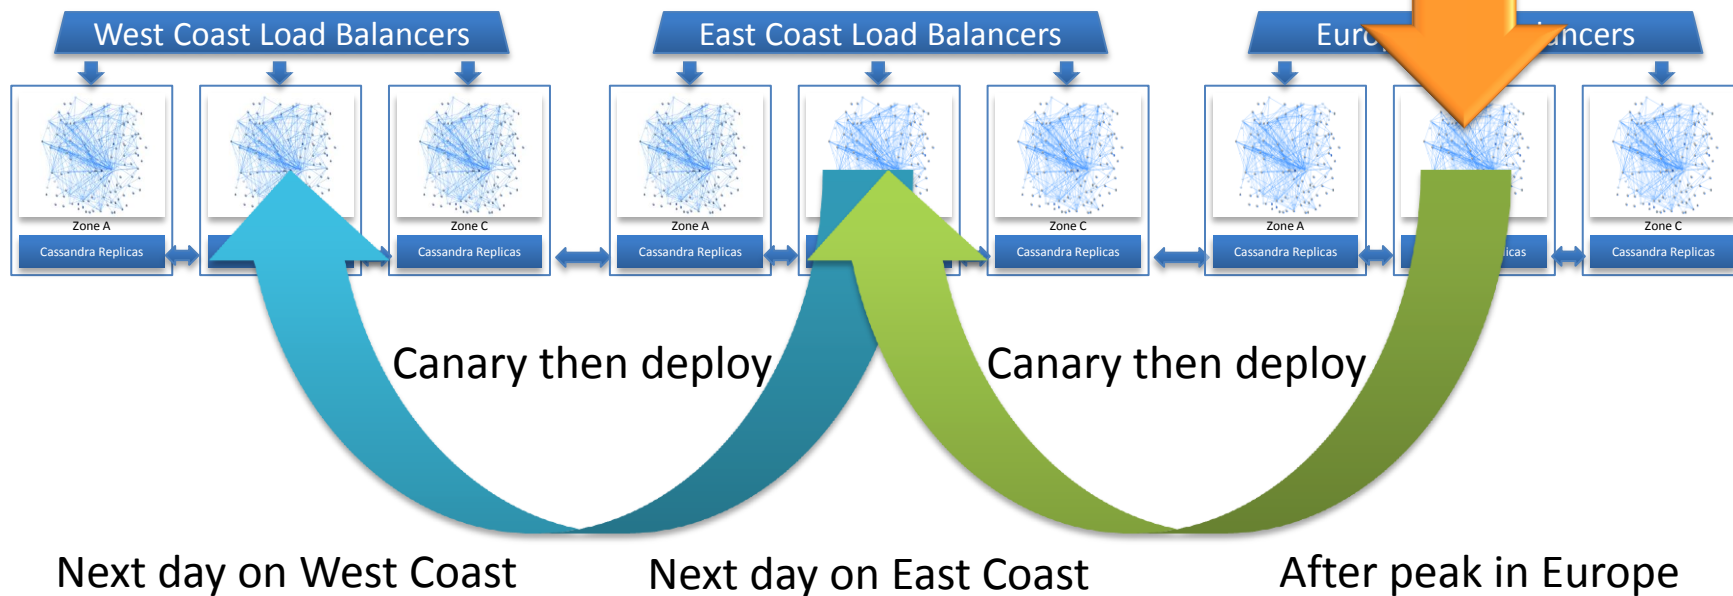
Global Deploy Automation

Afternoon in California



Night-time in Europe

If passes test suite, canary then deploy



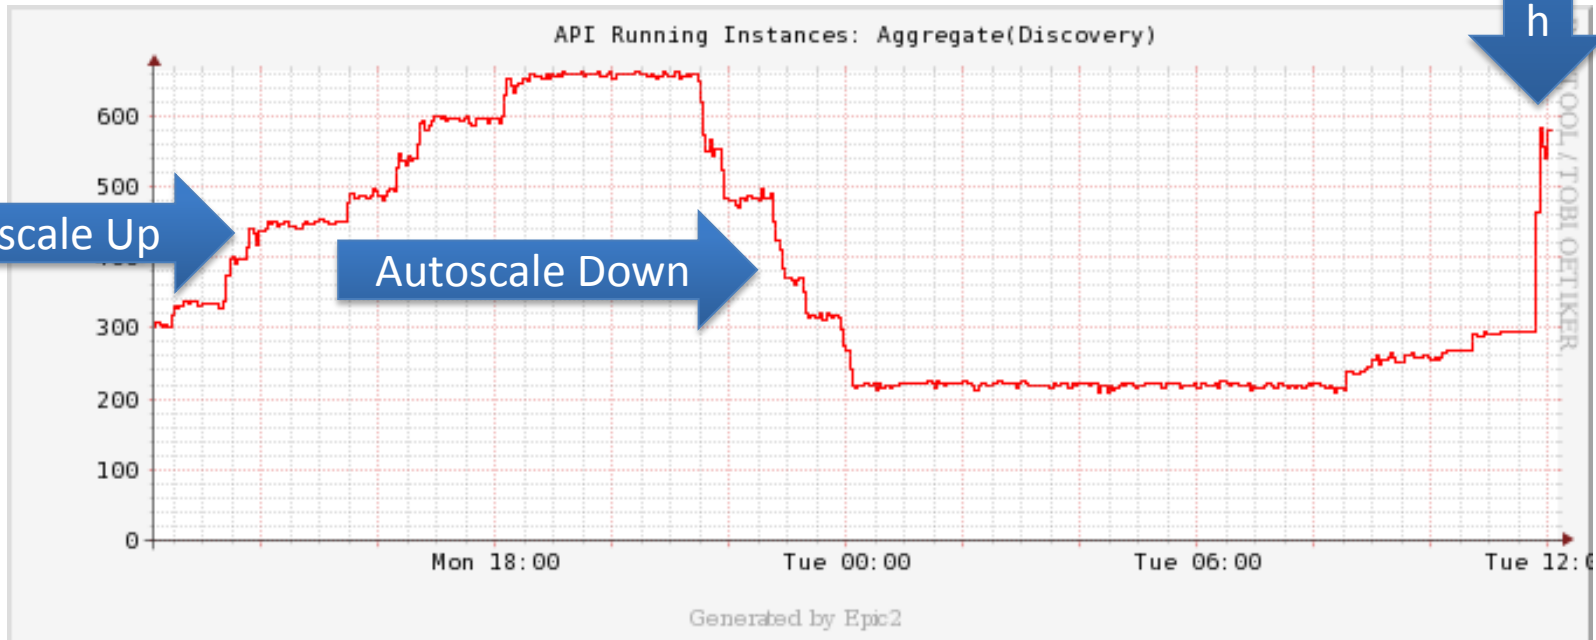
Ephemeral Instances

- Largest services are autoscaled
- Average lifetime of an instance is 36 hours

P
u
s
h

Autoscale Up

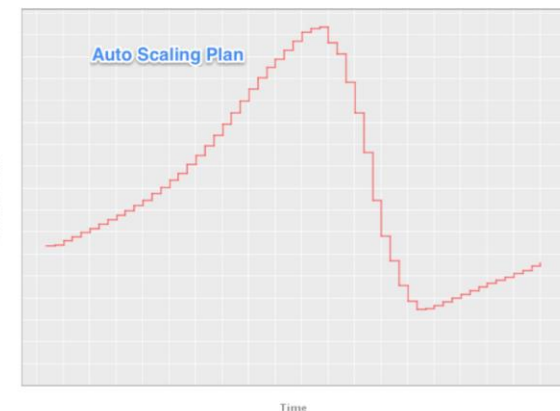
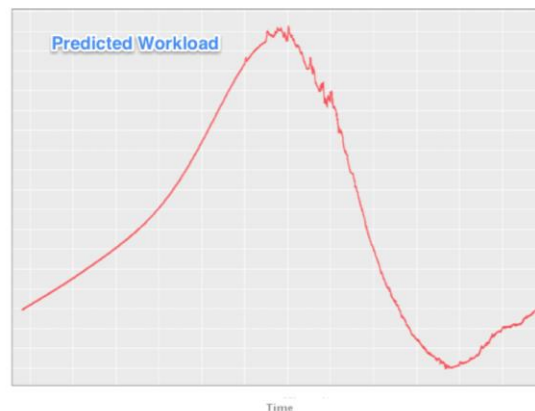
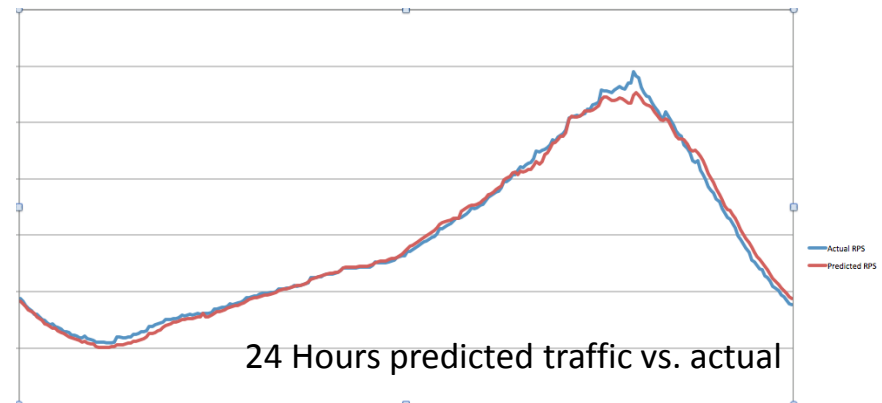
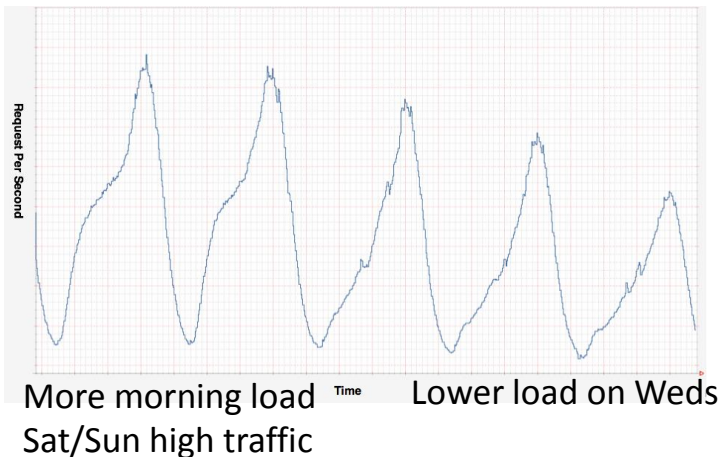
Autoscale Down





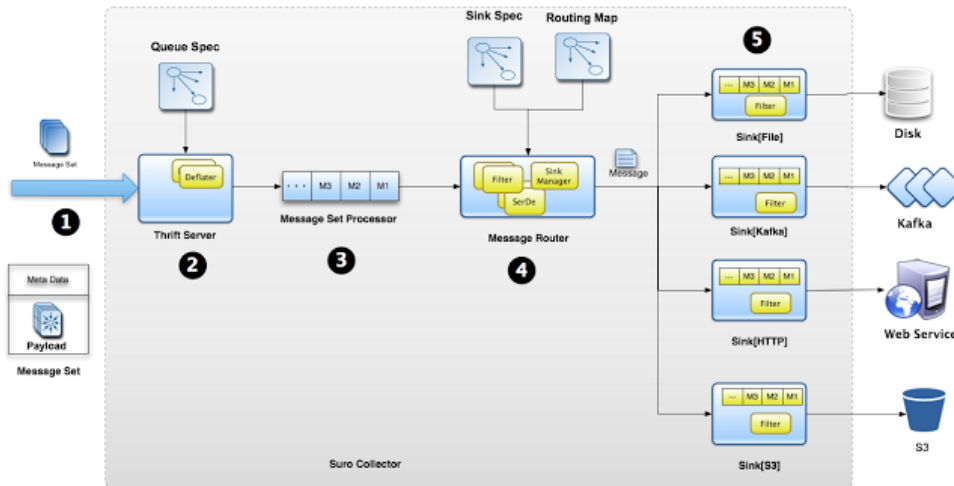
Scryer - Predictive Auto-scaling

See techblog.netflix.com



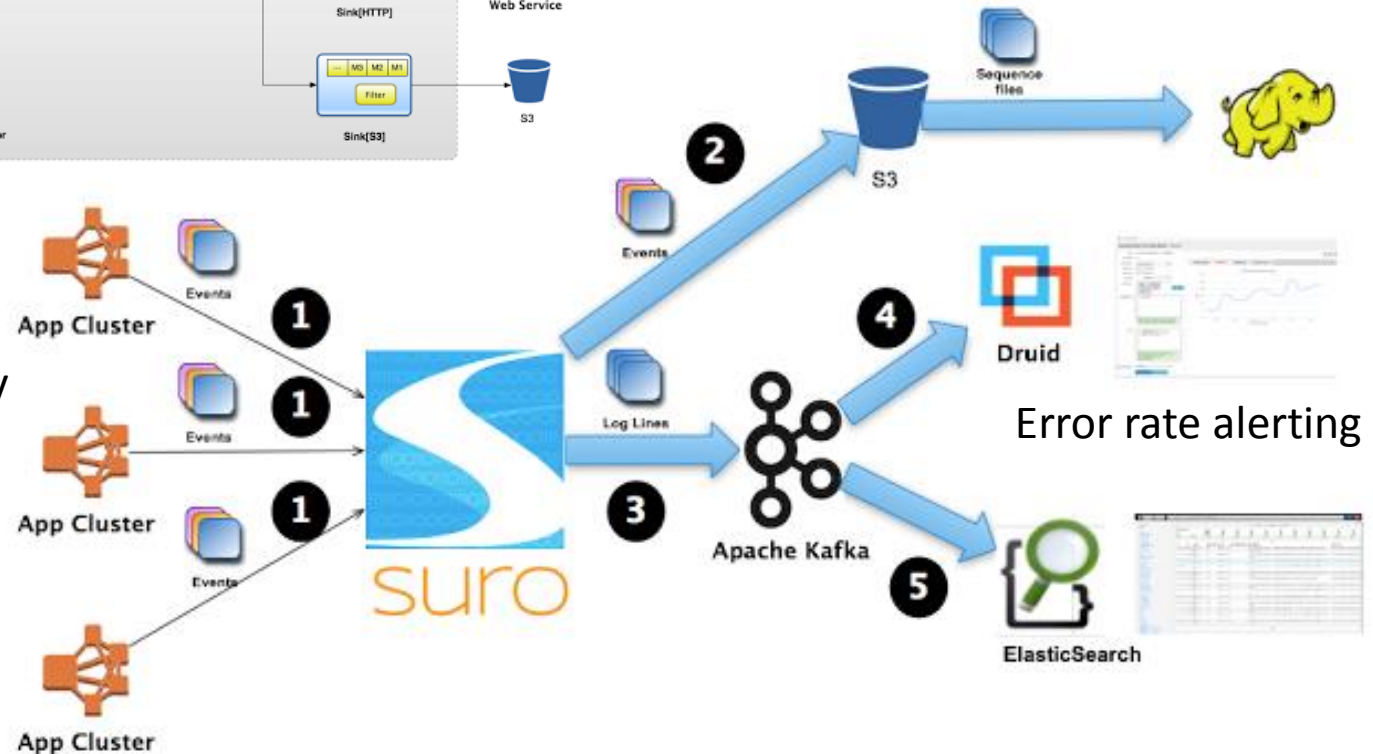
FFT based prediction driving AWS Autoscaler to plan minimum capacity

Suro Event Pipeline

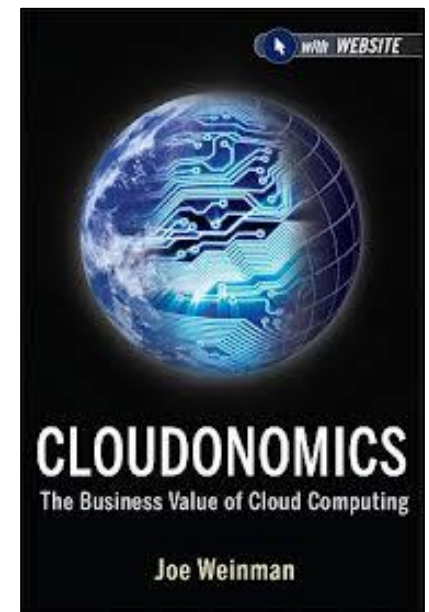
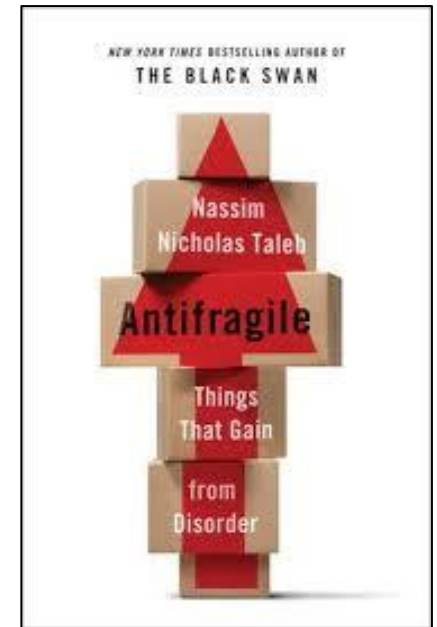
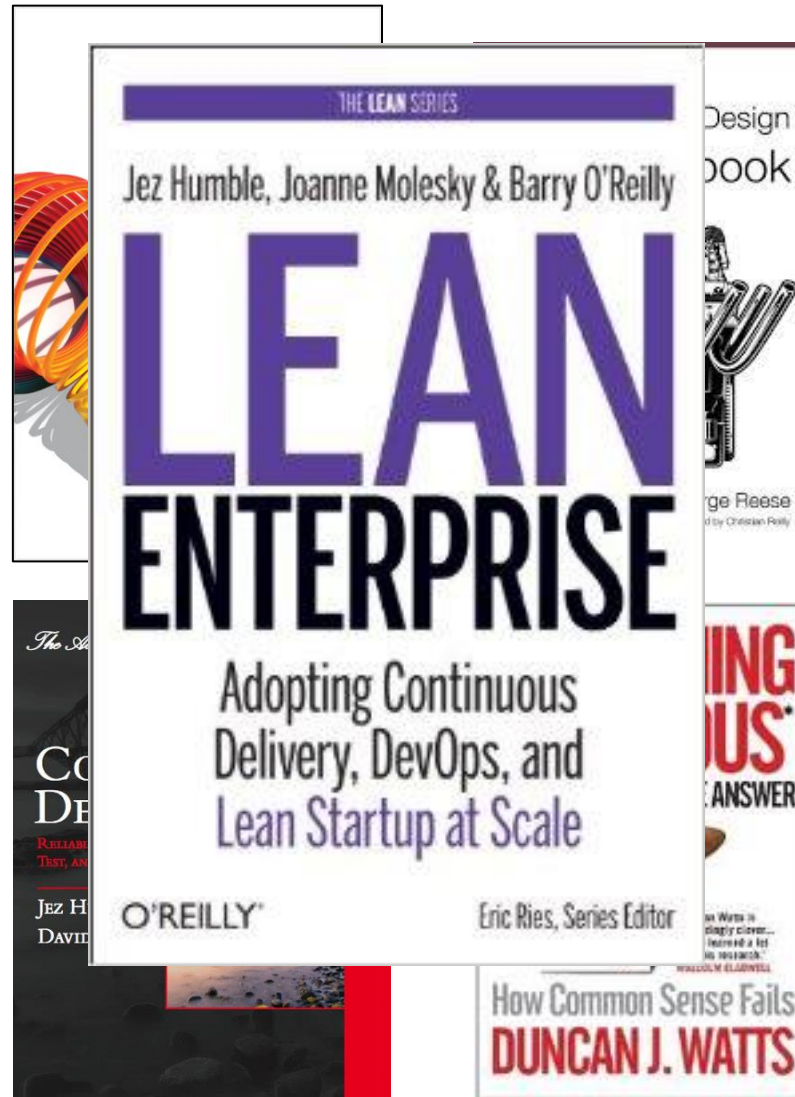
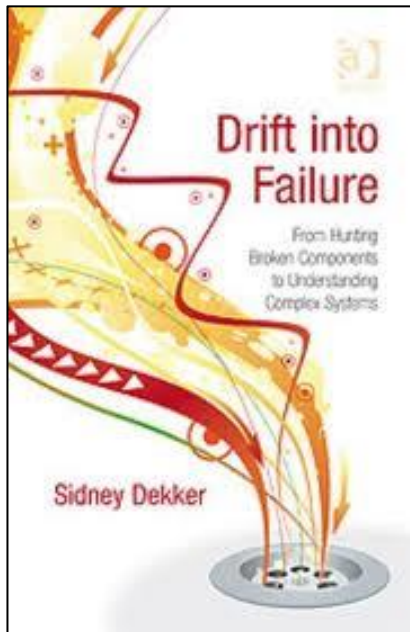
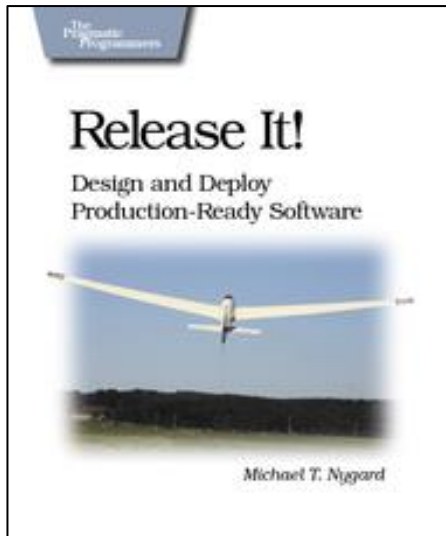


Cloud native, dynamic,
configurable offline and
realtime data sinks

1.5 Million events/s
80 Billion events/day



Inspiration





Principles Revisited:

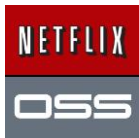
Immutability

Separation of Concerns

Anti-fragility

High trust organization

Sharing



Takeaway

Speed Wins

Assume Broken

Cloud Native Automation

Github is the “app store” and resumé

@adriano @NetflixOSS

<http://netflix.github.com>