



BDT310

Big Data Architectural Patterns and Best Practices on AWS

Siva Raghupathy, Principal Solutions Architect, Amazon Web Services

October 2015

What to Expect from the Session

Big data challenges

How to simplify big data processing

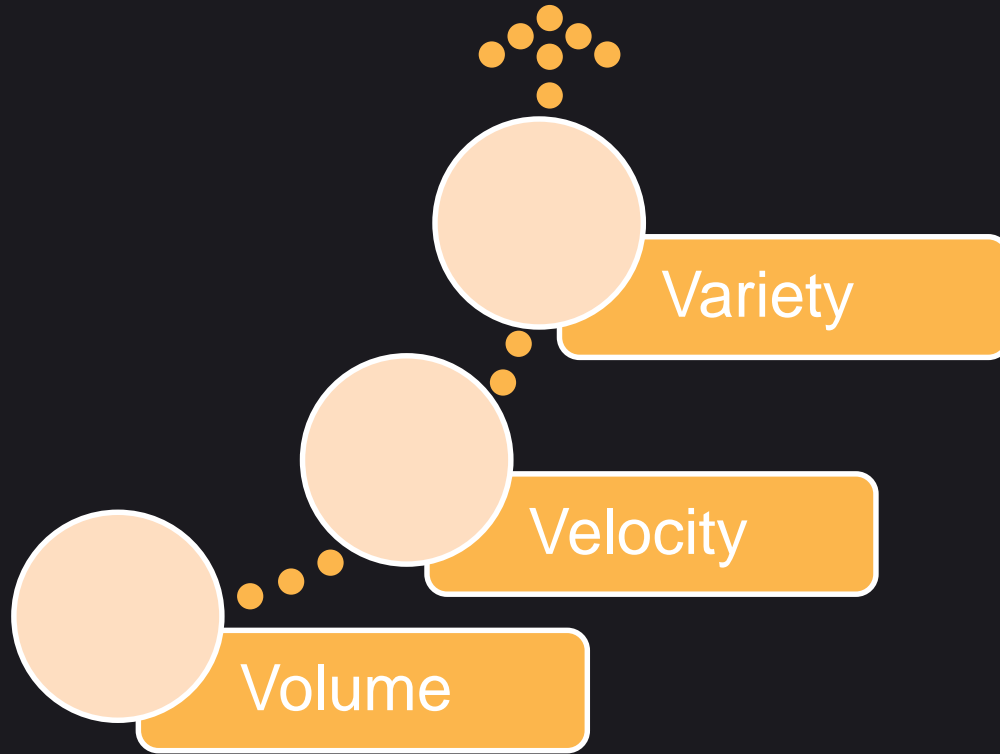
What technologies should you use?

- Why?
- How?

Reference architecture

Design patterns

Ever Increasing Big Data



Big Data Evolution

Batch



Real-time



Prediction

Report



Alerts



Forecast



Plethora of Tools



EMR



S3



DynamoDB



SQS



Amazon
Redshift




Amazon
Glacier



RDS



ElastiCache



Cassandra



Amazon Kinesis



Kinesis-
enabled
app



Data Pipeline



CloudSearch



Flink




Lambda



ML



DynamoDB
Streams

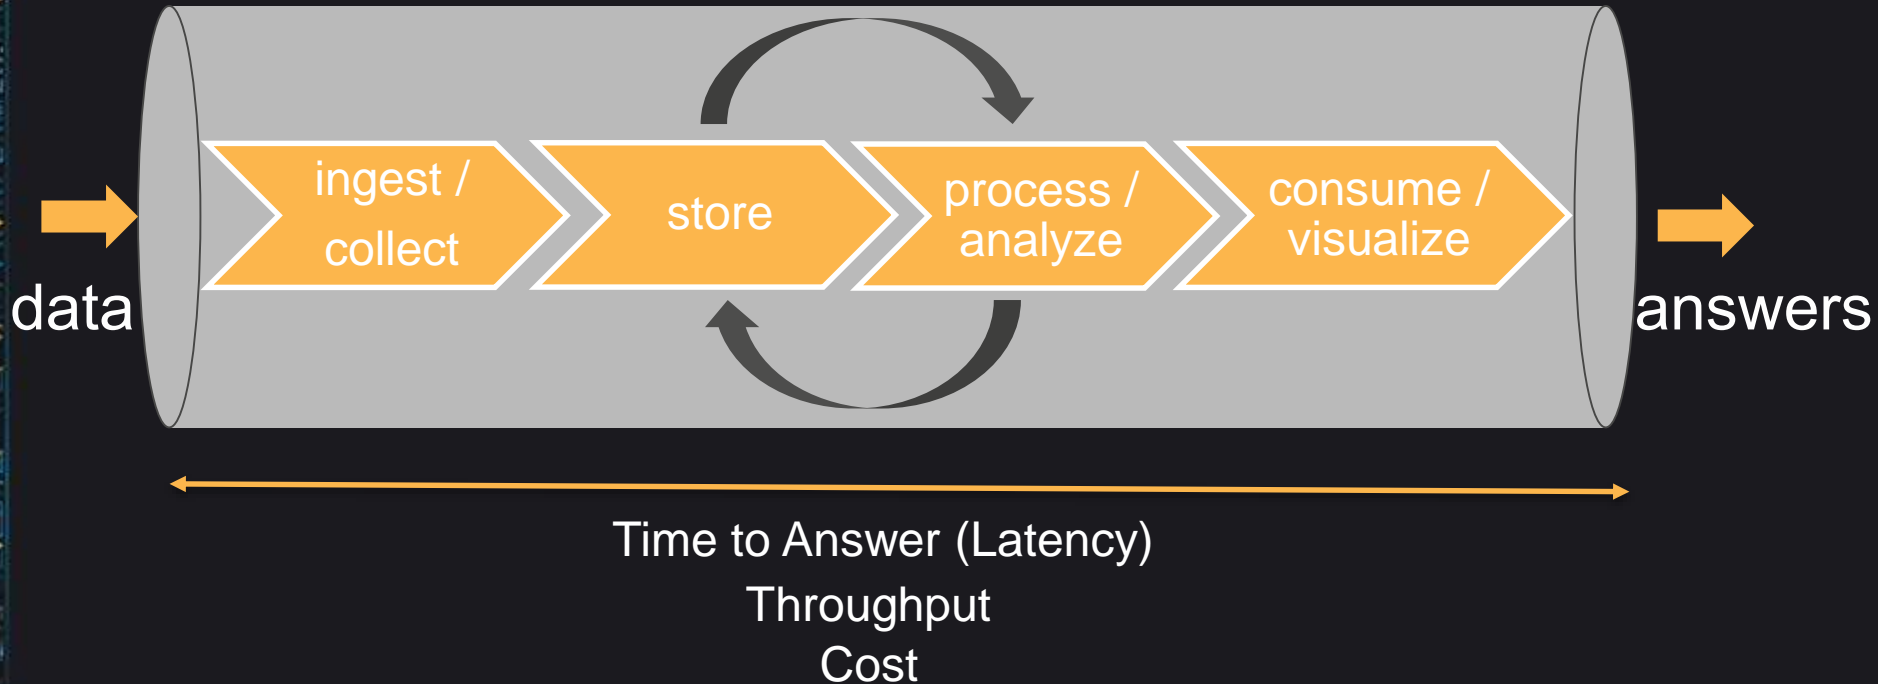


Is there a reference architecture ?
What tools should I use ?
How ?
Why ?

Architectural Principles

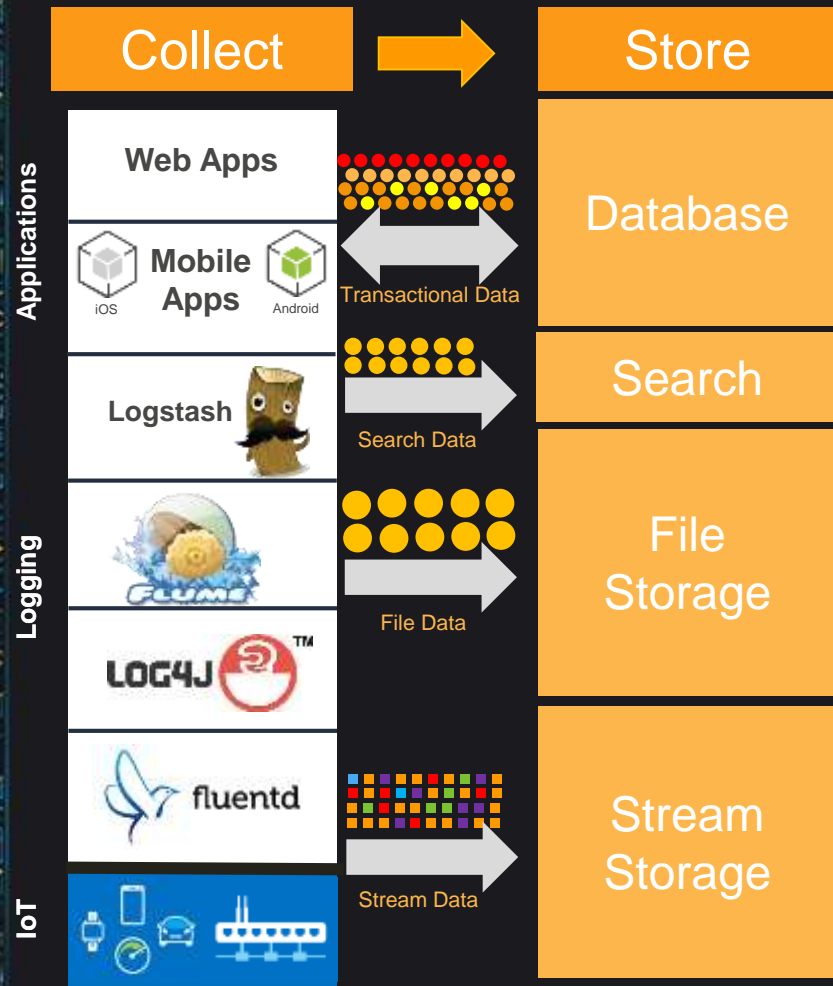
- Decoupled “data bus”
 - Data → Store → Process → Answers
- Use the right tool for the job
 - Data structure, latency, throughput, access patterns
- Use Lambda architecture ideas
 - Immutable (append-only) log, batch/speed/serving layer
- Leverage AWS managed services
 - No/low admin
- Big data ≠ big cost

Simplify Big Data Processing





Collect /
Ingest



Types of Data

- Transactional
 - Database reads & writes (OLTP)
 - Cache
- Search
 - Logs
 - Streams
- File
 - Log files (/var/log)
 - Log collectors & frameworks
- Stream
 - Log records
 - Sensors & IoT data



Store




Collect



Store

Applications

Web Apps



Database

Transactional Data

Logstash



Search

Search Data

Logging



File Storage

File Data



Stream Storage

Apache Kafka



Amazon Kinesis



Amazon DynamoDB



Stream Data

IoT



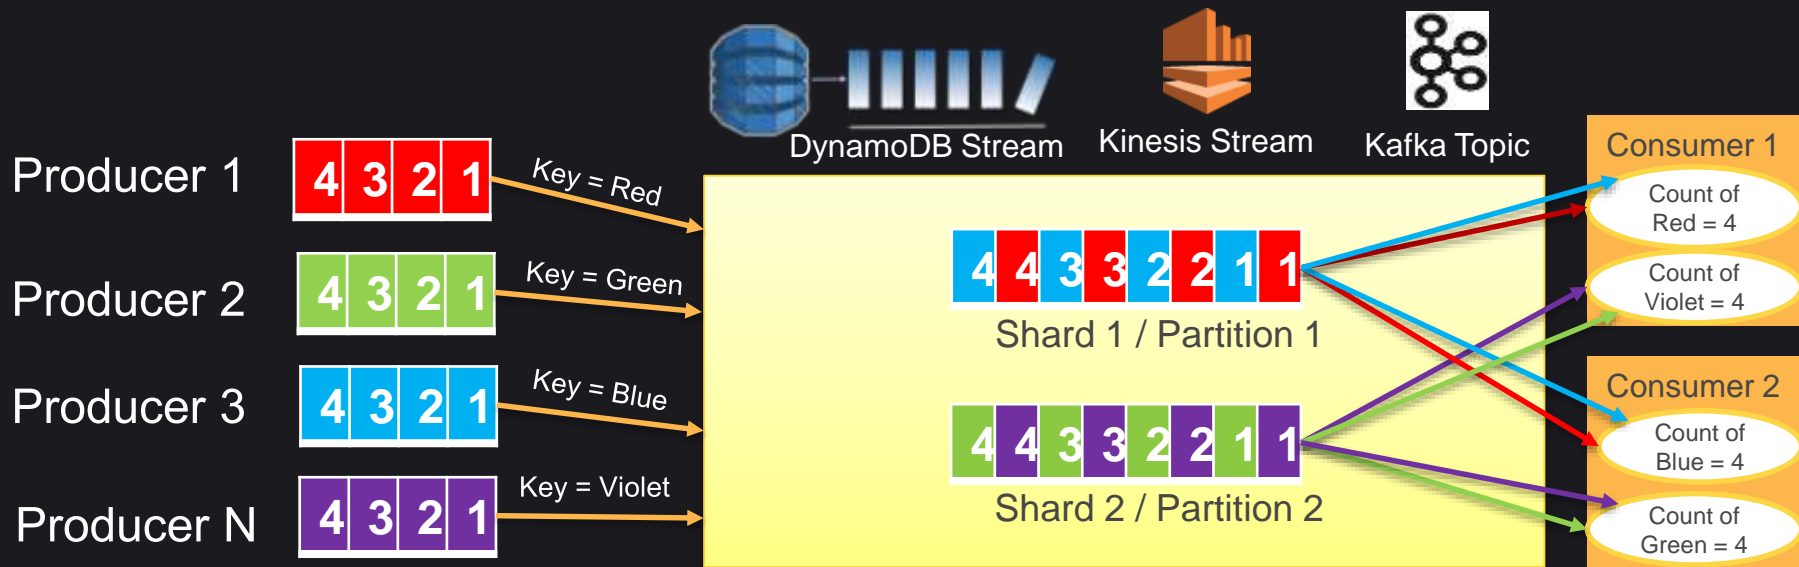
Stream Storage

Stream Storage Options

- AWS managed services
 - Amazon Kinesis → streams
 - DynamoDB Streams → table + streams
 - Amazon SQS → queue
 - Amazon SNS → pub/sub
- Unmanaged
 - Apache Kafka → stream

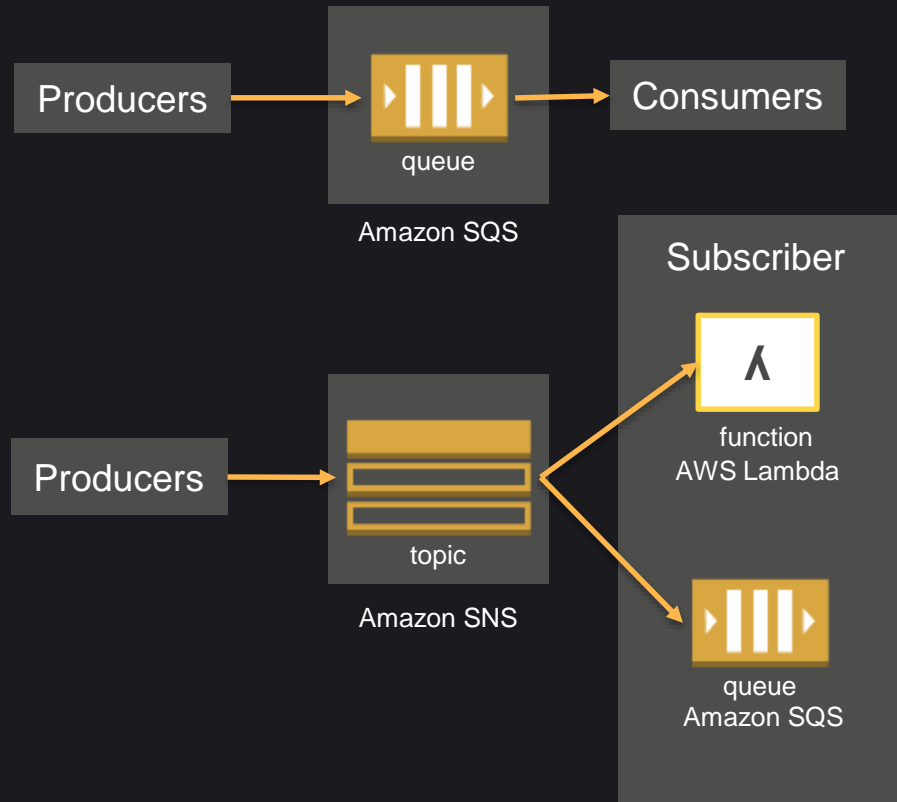
Why Stream Storage?

- Decouple producers & consumers
- Persistent buffer
- Collect multiple streams
- Preserve client ordering
- Streaming MapReduce
- Parallel consumption



What About Queues & Pub/Sub ?

- Decouple producers & consumers/subscribers
- Persistent buffer
- Collect multiple streams
- **No** client ordering
- **No** parallel consumption for Amazon SQS
 - Amazon SNS can route to multiple queues or Λ functions
- **No** streaming MapReduce



Which stream storage should I use?

	Amazon Kinesis	DynamoDB Streams	Amazon SQS Amazon SNS	Kafka
Managed	Yes	Yes	Yes	No
Ordering	Yes	Yes	No	Yes
Delivery	at-least-once	exactly-once	at-least-once	at-least-once
Lifetime	7 days	24 hours	14 days	Configurable
Replication	3 AZ	3 AZ	3 AZ	Configurable
Throughput	No Limit	No Limit	No Limit	~ Nodes
Parallel Clients	Yes	Yes	No (SQS)	Yes
MapReduce	Yes	Yes	No	Yes
Record size	1MB	400KB	256KB	Configurable
Cost	Low	Higher(table cost)	Low-Medium	Low (+admin)



Collect




Store

Applications

Web Apps

Mobile Apps

iOS Android



Database


Logging

Logstash

ELUMET

LOG4J

fluentd



Search

File Storage

HDFS

Amazon S3

Amazon Glacier

IoT

IoT devices



Stream Storage

Apache Kafka

Amazon Kinesis

Amazon DynamoDB



Why Is Amazon S3 Good for Big Data?



- Natively supported by big data frameworks (Spark, Hive, Presto, etc.)
- No need to run compute clusters for storage (unlike HDFS)
- Can run transient Hadoop clusters & Amazon EC2 Spot instances
- Multiple distinct (Spark, Hive, Presto) clusters can use the same data
- Unlimited number of objects
- Very high bandwidth – no aggregate throughput limit
- Highly available – can tolerate AZ failure
- Designed for 99.999999999% durability
- Tired-storage (Standard, IA, Amazon Glacier) via life-cycle policy
- Secure – SSL, client/server-side encryption at rest
- Low cost

What about HDFS & Amazon Glacier?

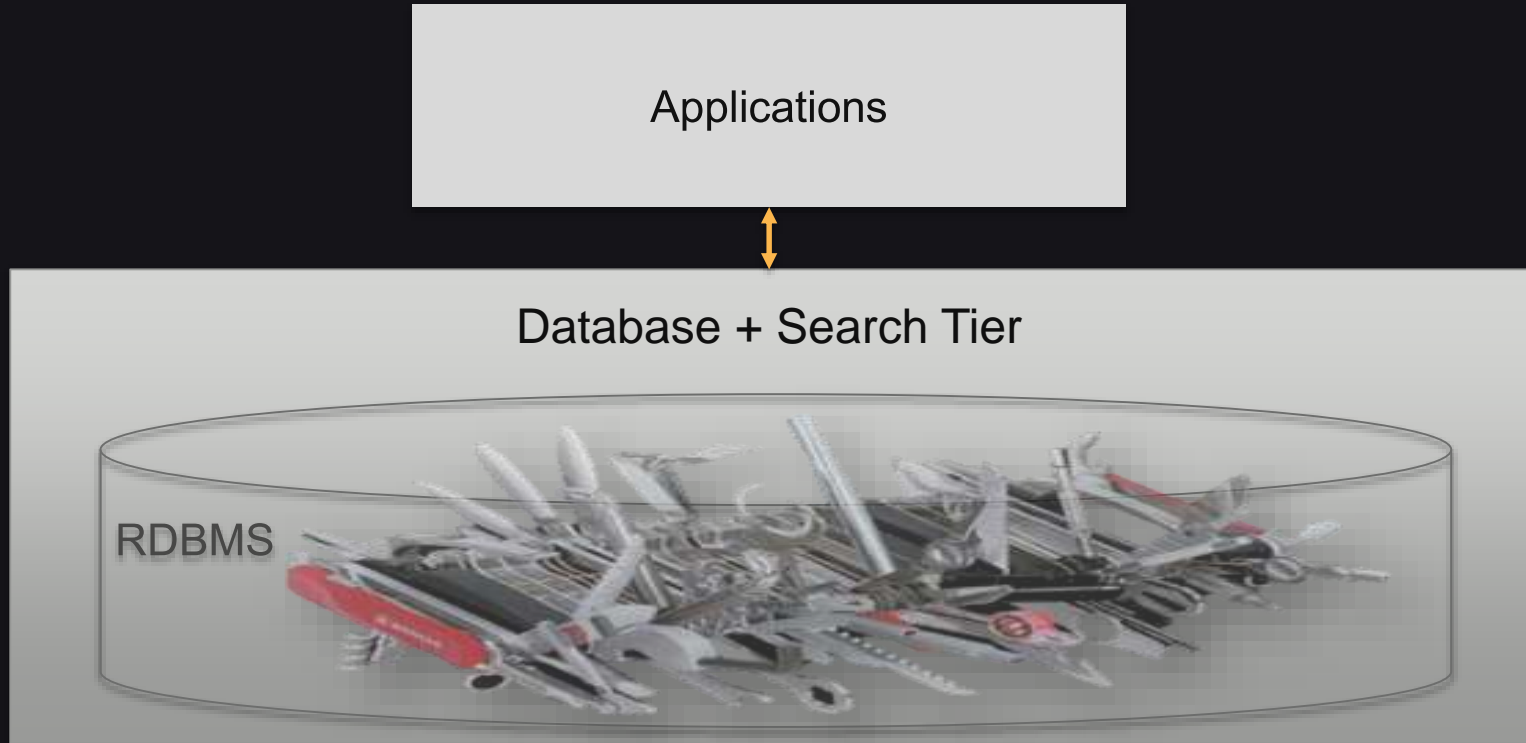
- Use HDFS for very frequently accessed (hot) data
- Use Amazon S3 Standard for frequently accessed data
- Use Amazon S3 Standard – IA for infrequently accessed data
- Use Amazon Glacier for archiving cold data



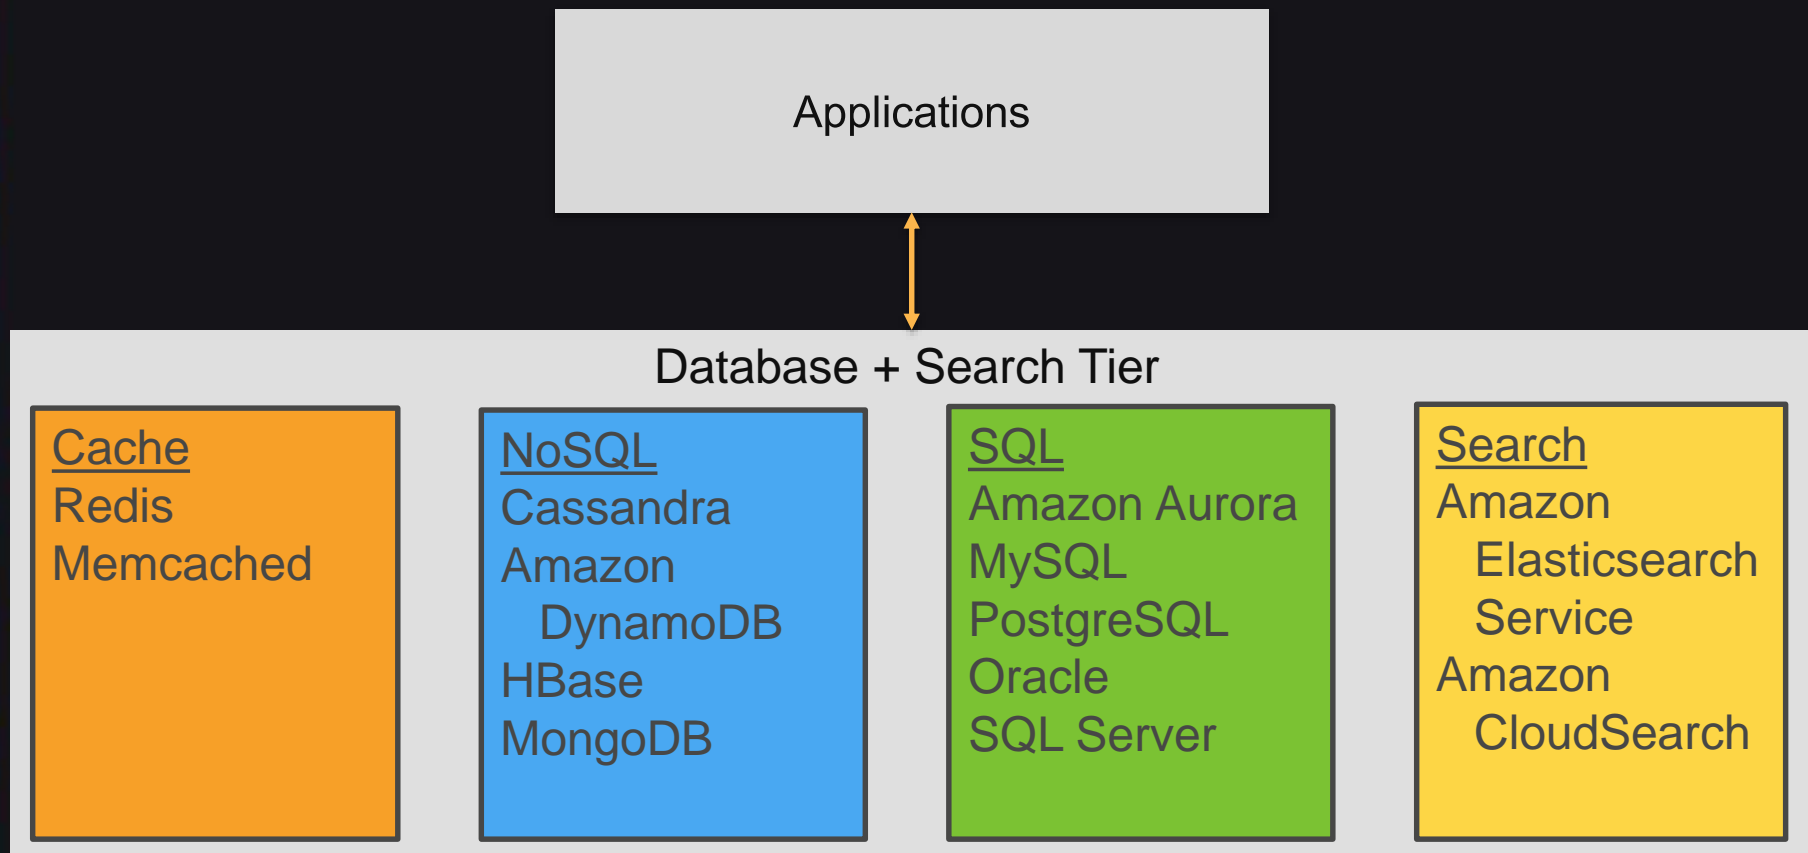


Database + Search Tier

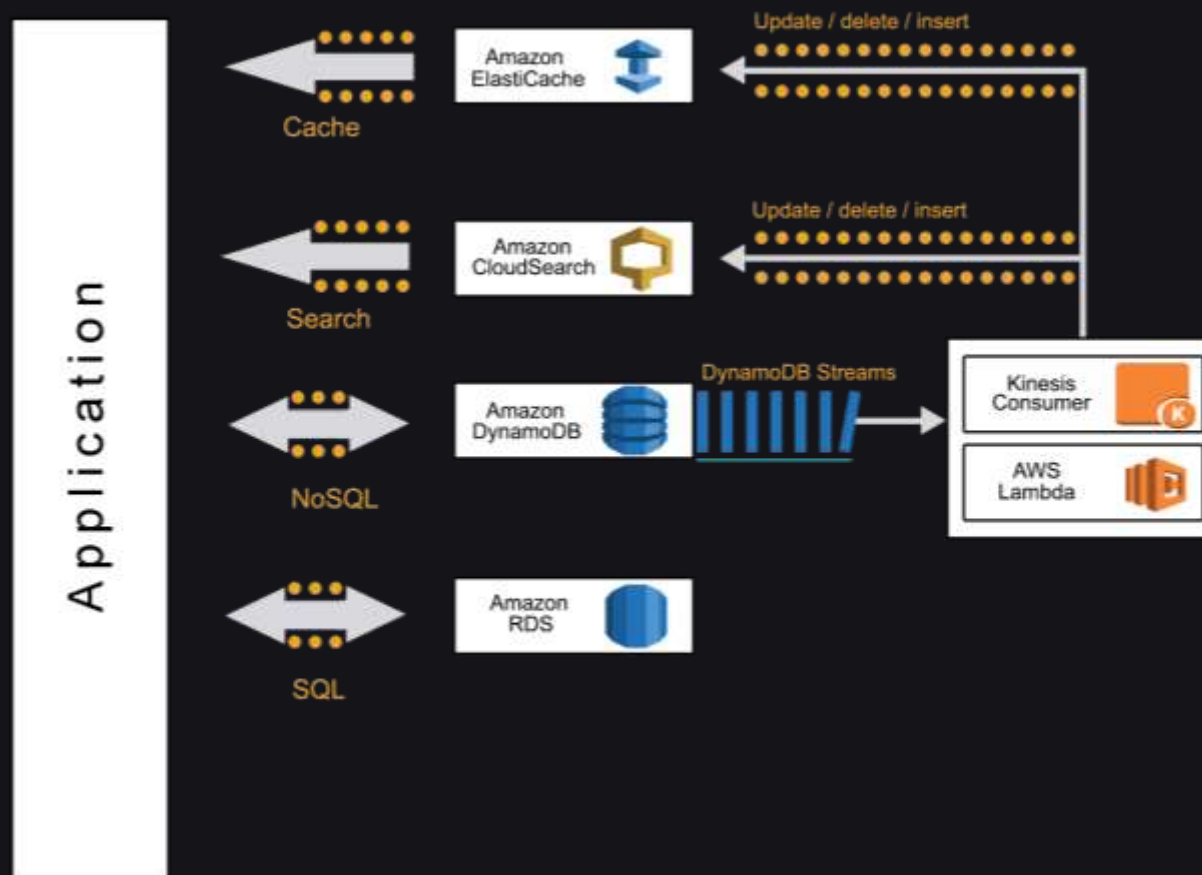
Database + Search Tier Anti-pattern



Best Practice — Use the Right Tool for the Job



Materialized Views



What Data Store Should I Use?


- Data structure → Fixed schema, JSON, key-value
- Access patterns → Store data in the format you will access it
- Data / access characteristics → Hot, warm, cold
- Cost → Right cost

Data Structure and Access Patterns

Access Patterns	What to use?
Put/Get (Key, Value)	Cache, NoSQL
Simple relationships → 1:N, M:N	NoSQL
Cross table joins, transaction, SQL	SQL
Faceting, Search	Search

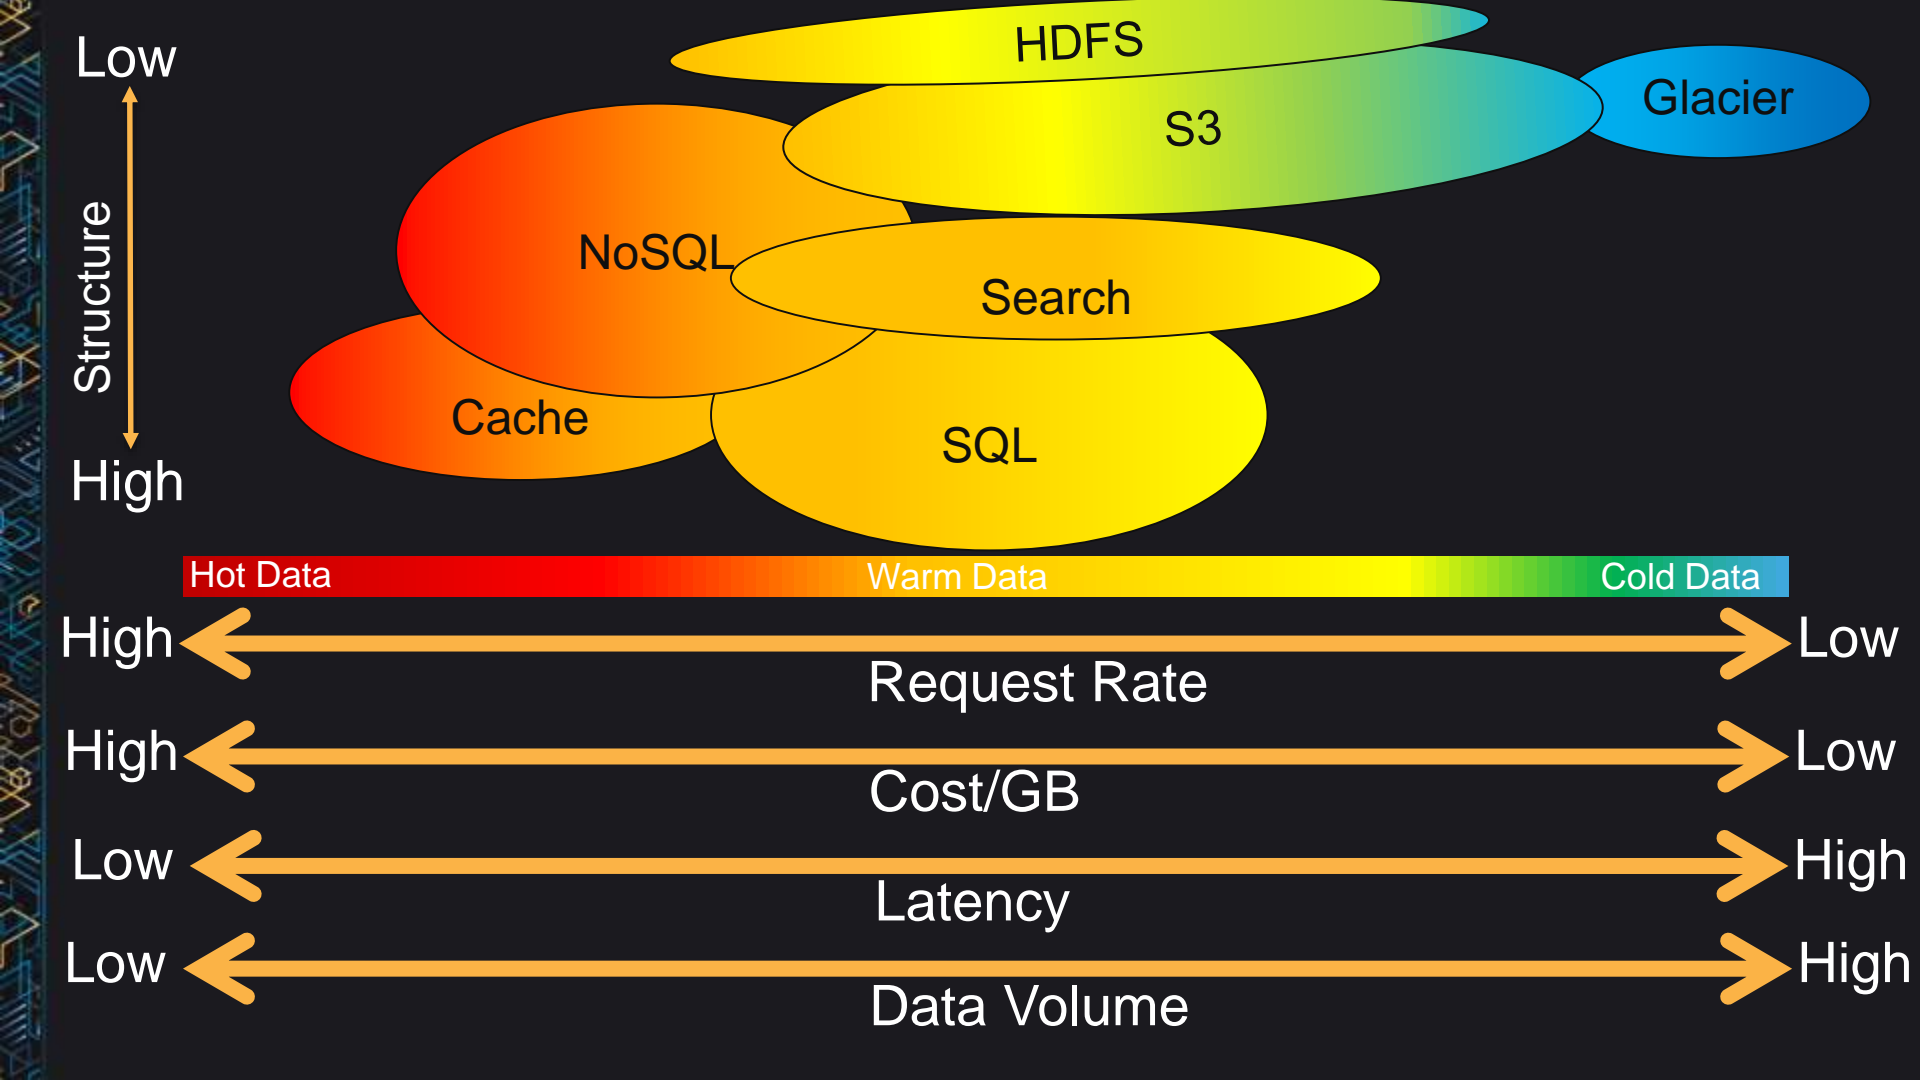
Data Structure	What to use?
Fixed schema	SQL, NoSQL
Schema-free (JSON)	NoSQL, Search
(Key, Value)	Cache, NoSQL

What Is the Temperature of Your Data / Access ?



Data / Access Characteristics: Hot, Warm, Cold

	Hot Data	Warm Data	Cold Data
	Hot	Warm	Cold
Volume	MB–GB	GB–TB	PB
Item size	B–KB	KB–MB	KB–TB
Latency	ms	ms, sec	min, hrs
Durability	Low–High	High	Very High
Request rate	Very High	High	Low
Cost/GB	\$\$-\$	\$-¢¢	¢



What Data Store Should I Use?

	Hot Data			Warm Data			Cold Data
	Amazon ElastiCache	Amazon DynamoDB	Amazon Aurora	Amazon Elasticsearch	Amazon EMR (HDFS)	Amazon S3	Amazon Glacier
Average latency	ms	ms	ms, sec	ms,sec	sec,min,hrs	ms,sec,min (~ size)	hrs
Data volume	GB	GB–TBs (no limit)	GB–TB (64 TB Max)	GB–TB	GB–PB (~nodes)	MB–PB (no limit)	GB–PB (no limit)
Item size	B-KB	KB (400 KB max)	KB (64 KB)	KB (1 MB max)	MB-GB	KB-GB (5 TB max)	GB (40 TB max)
Request rate	High - Very High	Very High (no limit)	High	High	Low – Very High	Low – Very High (no limit)	Very Low
Storage cost GB/month	\$\$	¢¢	¢¢	¢¢	¢	¢	¢/10
Durability	Low - Moderate	Very High	Very High	High	High	Very High	Very High
	Hot Data			Warm Data			Cold Data

Cost Conscious Design

Example: Should I use Amazon S3 or Amazon DynamoDB?

“I’m currently scoping out a project that will greatly increase my team’s use of Amazon S3. Hoping you could answer some questions. The current iteration of the design calls for **many small files**, perhaps up to a **billion during peak**. The **total size** would be on the order of **1.5 TB per month**...”

Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
300	2048	1483	777,600,000

Cost Conscious Design

Example: Should I use Amazon S3 or Amazon DynamoDB?



<https://calculator.s3.amazonaws.com/index.html>

Amazon S3 or Amazon DynamoDB?

Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
300	2,048	1,483	777,600,000

Amazon DynamoDB is a high performance non-relational database service that is easy to set up, operate, and scale. It is designed to address the core problems of database management, performance, scalability, and reliability. It also provides predictable high performance and low latency at scale.

Indexed Data Storage:

Dataset Size:

1483 GB

Provisioned Throughput Capacity *:

Item Size (All attributes):

2 KB

Number of items read per second:

0 Reads/Second

Read Consistency:

☒ Strongly Consistent

☐ Eventually Consistent (cheaper)

Number of items written per second:

300 Writes/Second

Amazon S3 is storage for the Internet. It is designed to make web-scale computing easier for developers.

Storage:

Storage:

1483 GB

Reduced Redundancy Storage:

0 GB

Requests:

PUT/COPY/POST/LIST Requests:

77760000 Requests

GET and Other Requests:

0 Requests

Amazon S3 Service (US-East)

\$ 3932.27

Storage:

\$ 44.27

Put/List Requests:

\$ 3888.00

Amazon DynamoDB Service (US-East)

\$ 644.30

Provisioned Throughput Capacity:

\$ 261.69


Indexed Data Storage:

\$ 382.61



amazon
webservices

SIMPLE MONTHLY CALCULATOR




Amazon DynamoDB

use

<u>Amazon S3 Service (US-East)</u>		\$	3932.27
Storage:	\$	44.27	
Put/List Requests:	\$	3888.00	
<u>Amazon DynamoDB Service (US-East)</u>		\$	644.30
Provisioned Throughput Capacity:	\$	261.69	
Indexed Data Storage:	\$	382.61	
DynamoDB Streams:	\$	0.00	

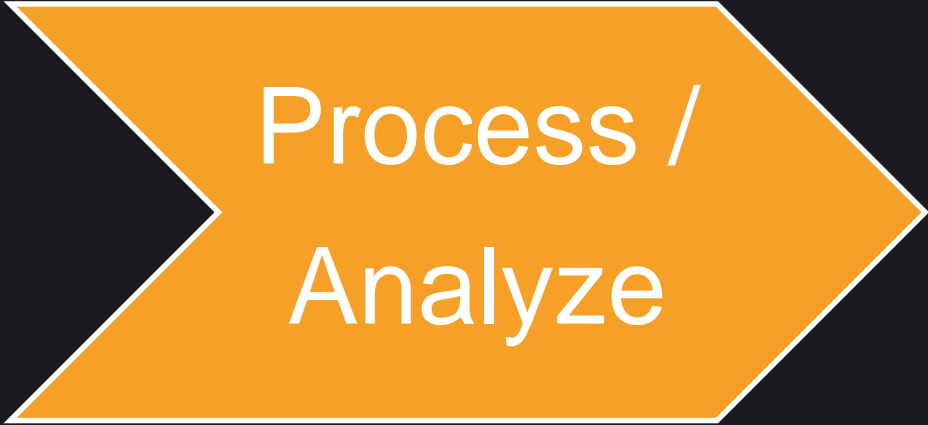

	Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
<u>Scenario 1</u>	300	2,048	1,483	777,600,000
<u>Scenario 2</u>	300	32,768	23,730	777,600,000

use



Amazon S3

<u>Amazon S3 Service (US-East)</u>		\$	4588.55
Storage:	\$	700.55	
Put/List Requests:	\$	3888.00	
<u>Amazon DynamoDB Service (US-East)</u>		\$	10131.40
Provisioned Throughput Capacity:	\$	4187.04	
Indexed Data Storage:	\$	5944.36	
DynamoDB Streams:	\$	0.00	



Process /
Analyze

Applications



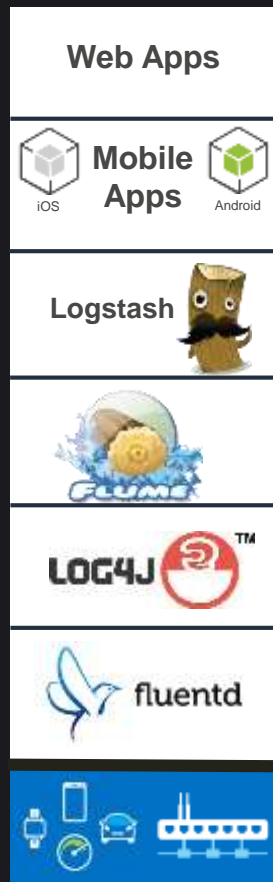
Logging

IoT

Collect

Store

Analyze



Interactive

Batch

Stream Processing



Amazon Elastic MapReduce

Analyze

Process / Analyze

Analysis of data is a process of inspecting, cleaning, transforming, and modeling data with the goal of discovering useful information, suggesting conclusions, and supporting decision-making.

Examples

- Interactive dashboards → **Interactive analytics**
- Daily/weekly/monthly reports → **Batch analytics**
- Billing/fraud alerts, 1 minute metrics → **Real-time analytics**
- Sentiment analysis, prediction models → **Machine learning**



Interactive Analytics

Takes large amount of (warm/cold) data

Takes **seconds** to get answers back

Example: Self-service dashboards



Batch Analytics

Takes large amount of (warm/cold) data

Takes **minutes or hours** to get answers back

Example: Generating daily, weekly, or monthly reports

Real-Time Analytics

Take small amount of hot data and ask questions

Takes short amount of time (milliseconds or seconds) to get your answer back

- Real-time (event)
 - Real-time response to events in data streams
 - Example: Billing/Fraud Alerts
- Near real-time (micro-batch)
 - Near real-time operations on small batches of events in data streams
 - Example: 1 Minute Metrics

Predictions via Machine Learning

ML gives computers the ability to learn without being explicitly programmed

Machine Learning Algorithms:

- Supervised Learning ← “teach” program
 - Classification ← Is this transaction fraud? (Yes/No)
 - Regression ← Customer Life-time value?
- Unsupervised Learning ← let it learn by itself
 - Clustering ← Market Segmentation

Analysis Tools and Frameworks

Machine Learning

- Mahout, Spark ML, Amazon ML

Interactive Analytics

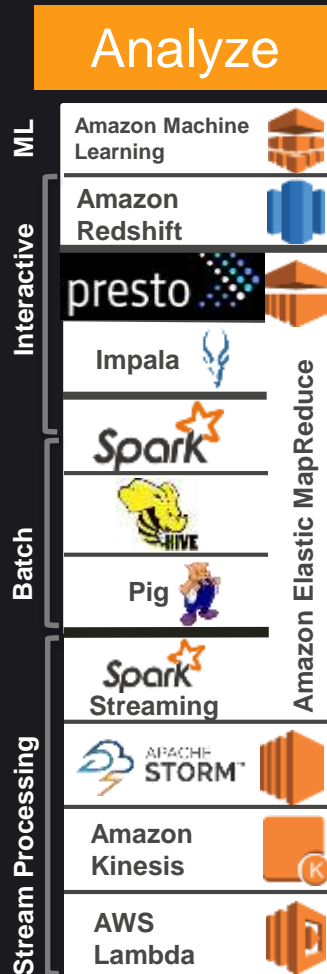
- Amazon Redshift, Presto, Impala, Spark

Batch Processing

- MapReduce, Hive, Pig, Spark

Stream Processing

- Micro-batch: Spark Streaming, KCL, Hive, Pig
- Real-time: Storm, AWS Lambda, KCL



What Stream Processing Technology Should I Use?

	Spark Streaming	Apache Storm	Amazon Kinesis Client Library	AWS Lambda	Amazon EMR (Hive, Pig)
Scale / Throughput	~ Nodes	~ Nodes	~ Nodes	Automatic	~ Nodes
Batch or Real-time	Real-time	Real-time	Real-time	Real-time	Batch
Manageability	Yes (Amazon EMR)	Do it yourself	Amazon EC2 + Auto Scaling	AWS managed	Yes (Amazon EMR)
Fault Tolerance	Single AZ	Configurable	Multi-AZ	Multi-AZ	Single AZ
Programming languages	Java, Python, Scala	Any language via Thrift	Java, via MultiLangDaemon (.Net, Python, Ruby, Node.js)	Node.js, Java	Hive, Pig, Streaming languages

Low

Low

Low

High

Query Latency (Low is better)

What Data Processing Technology Should I Use?

	Amazon Redshift	Impala	Presto	Spark	Hive
Query Latency	Low	Low	Low	Low	Medium (Tez) – High (MapReduce)
Durability	High	High	High	High	High
Data Volume	1.6 PB Max	~Nodes	~Nodes	~Nodes	~Nodes
Managed	Yes	Yes (EMR)	Yes (EMR)	Yes (EMR)	Yes (EMR)
Storage	Native	HDFS / S3A*	HDFS / S3	HDFS / S3	HDFS / S3
SQL Compatibility	High	Medium	High	Low (SparkSQL)	Medium (HQL)

Low

Low

Low

Medium

High

Query Latency (Low is better)

What about ETL?



Data Integration

Reduce the effort to move, cleanse, synchronize, manage, and automatize data related processes.



Attunity CloudBeam



Informatica Cloud



Matillion ETL for Redshift



snapLogic

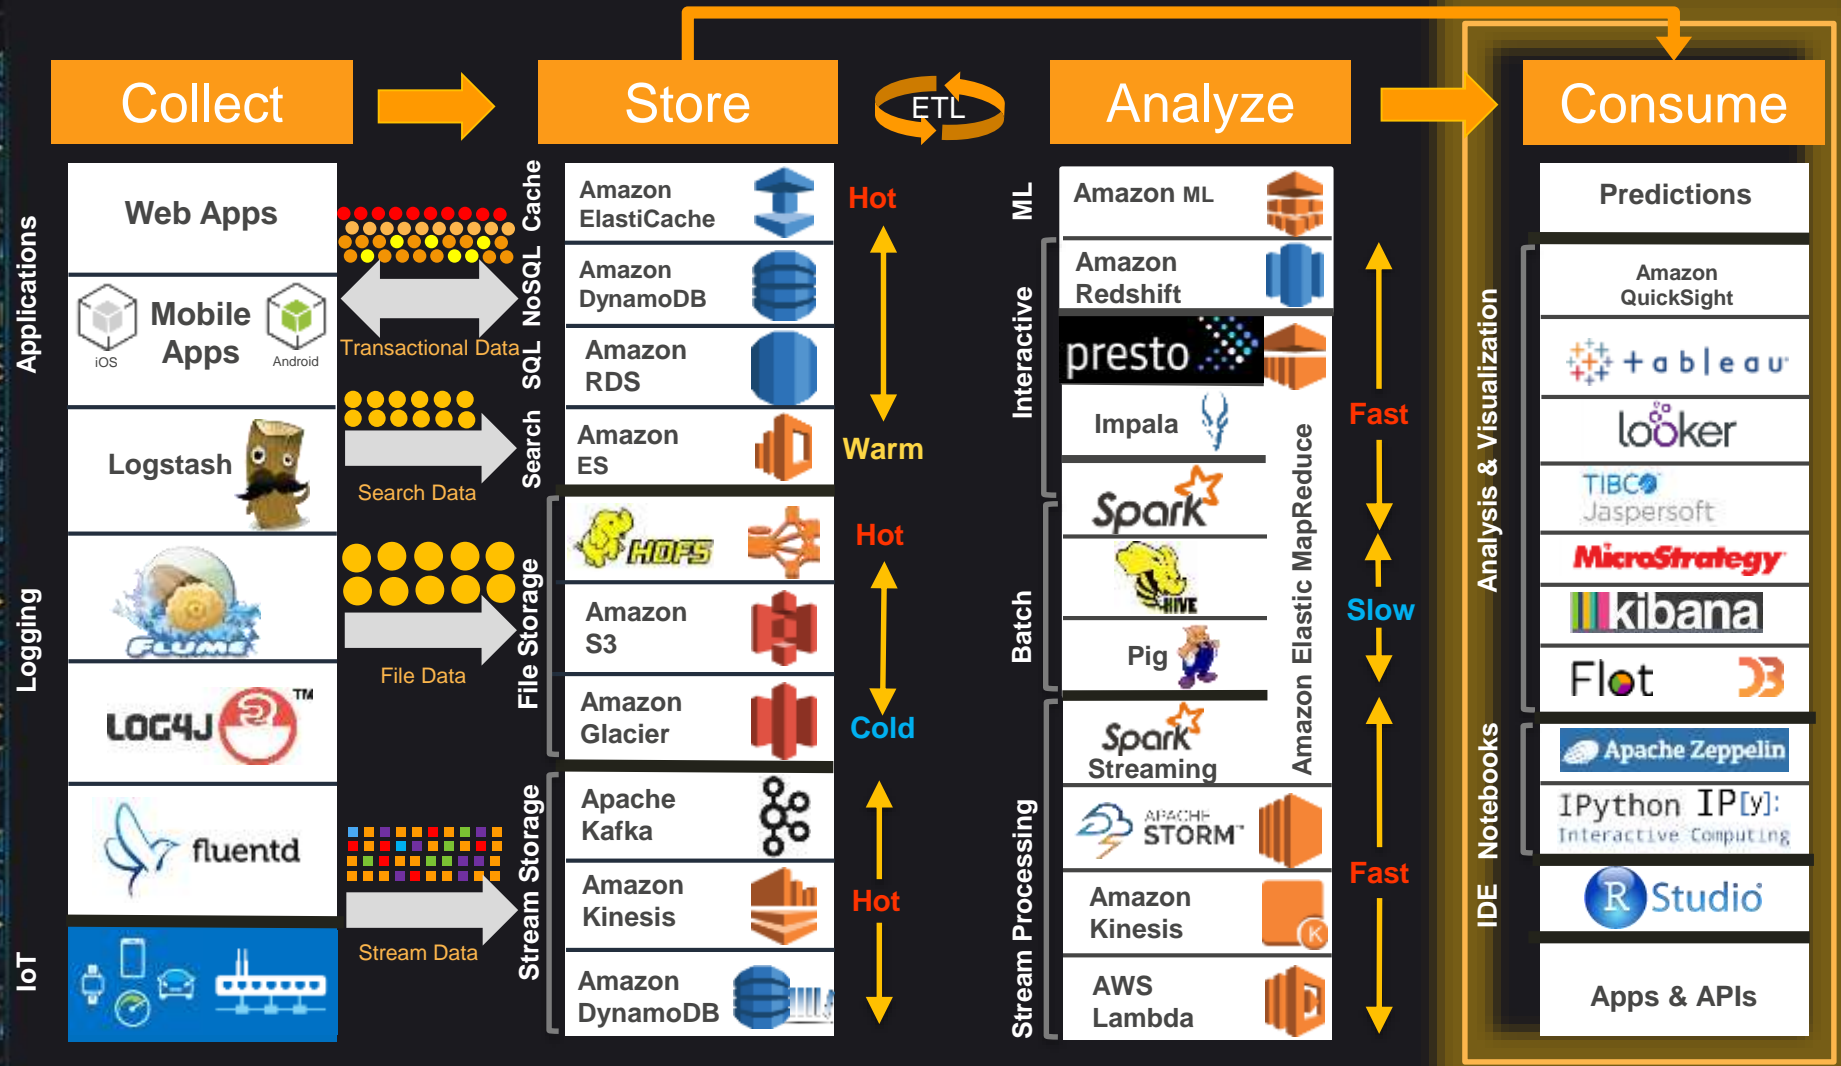


alteryx

<https://aws.amazon.com/big-data/partner-solutions/>

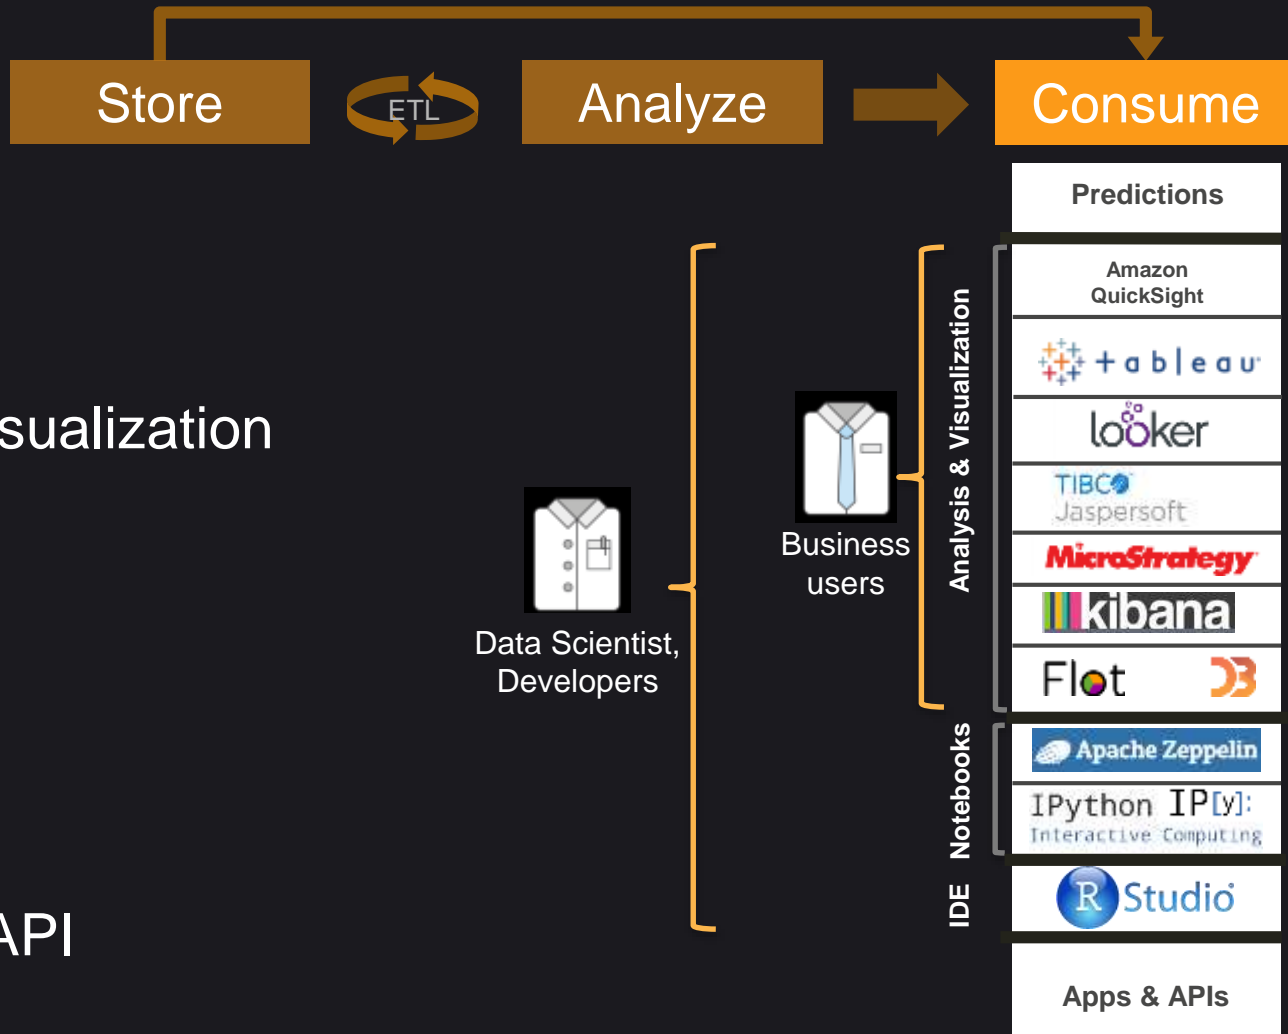


Consume /
Visualize



Consume

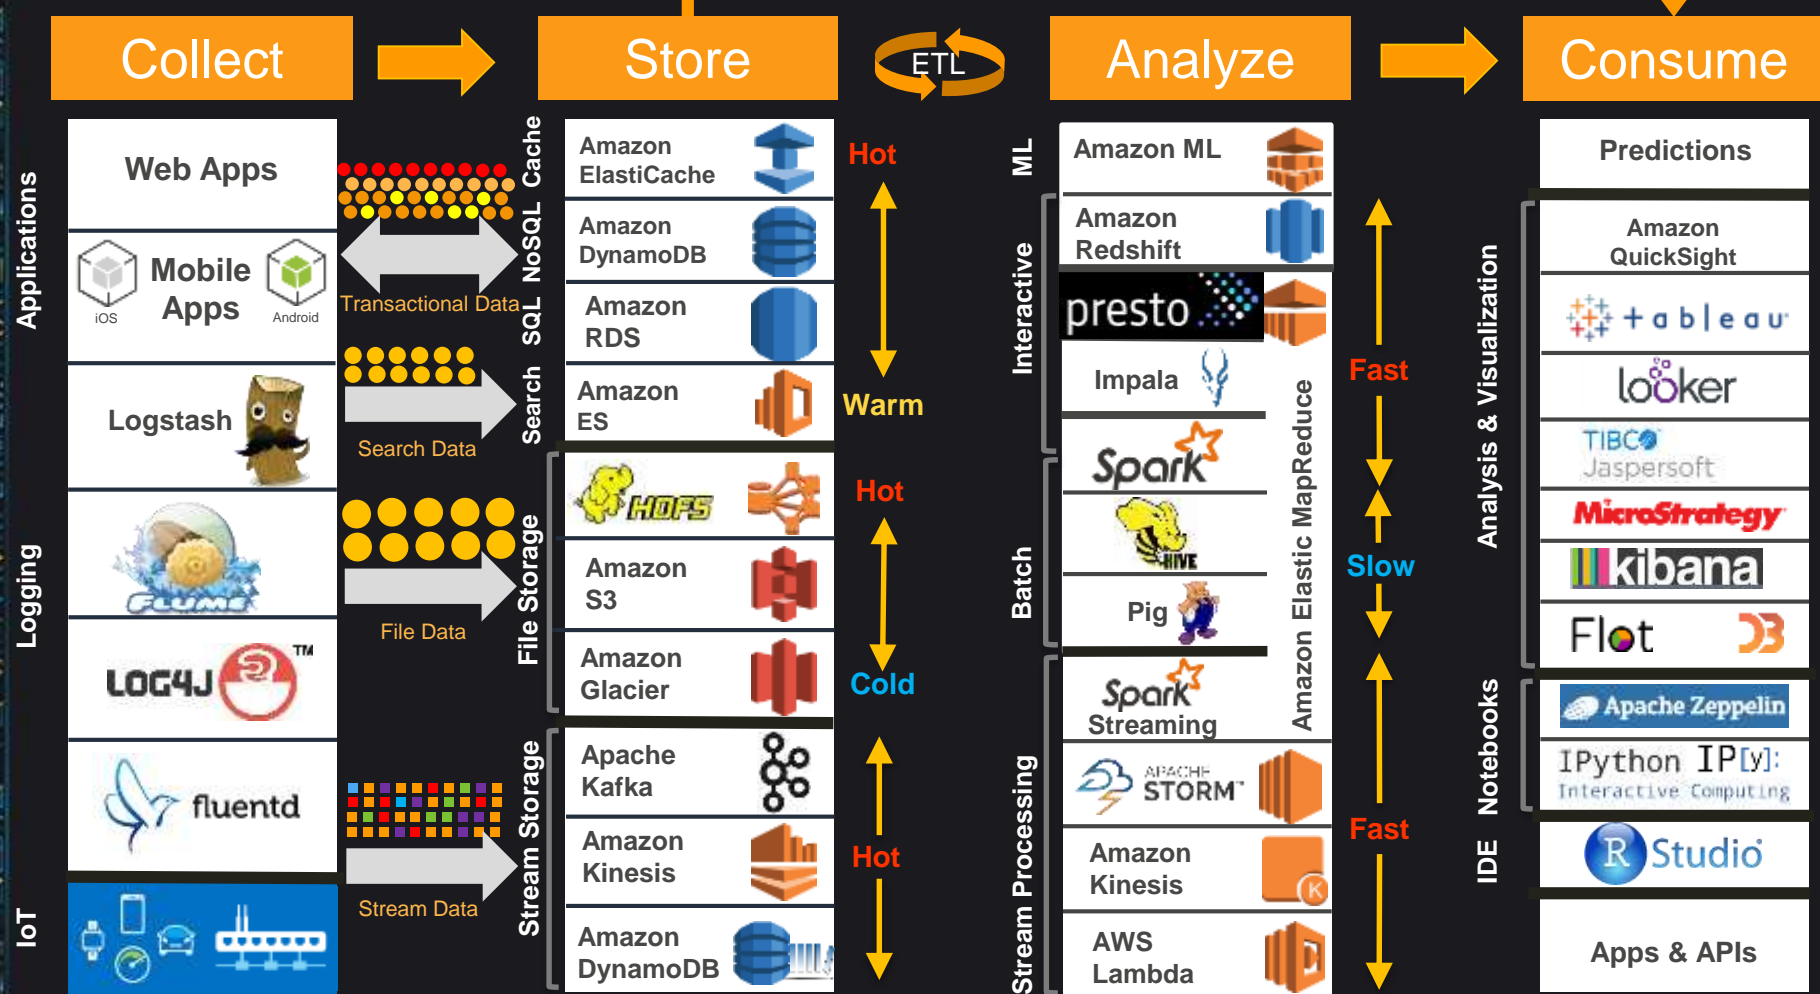
- Predictions
- Analysis and Visualization
- Notebooks
- IDE
- Applications & API





Putting It All Together

Reference Architecture





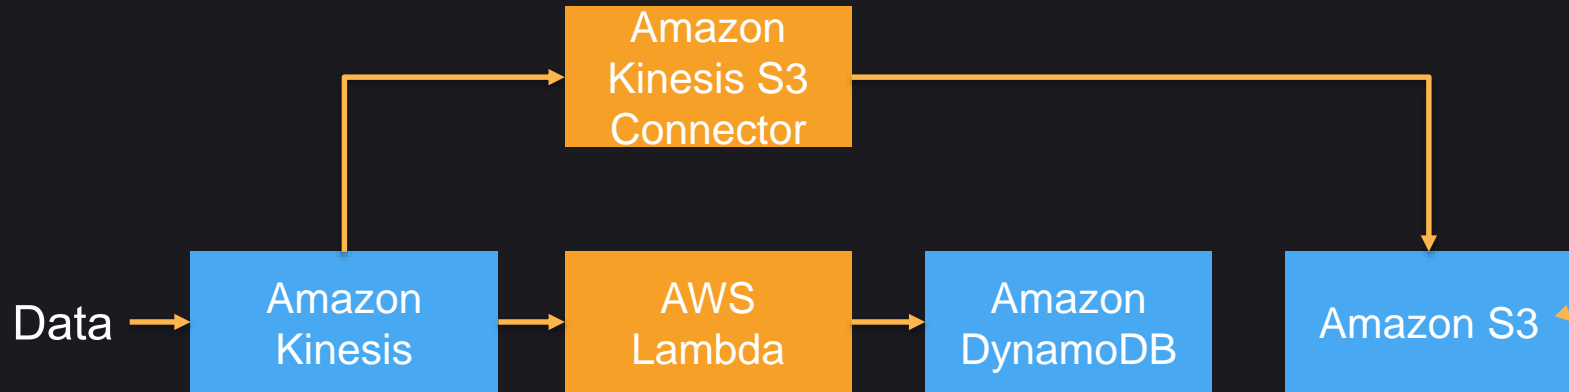
Design Patterns

Multi-Stage Decoupled “Data Bus”

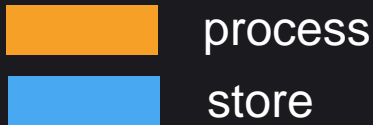
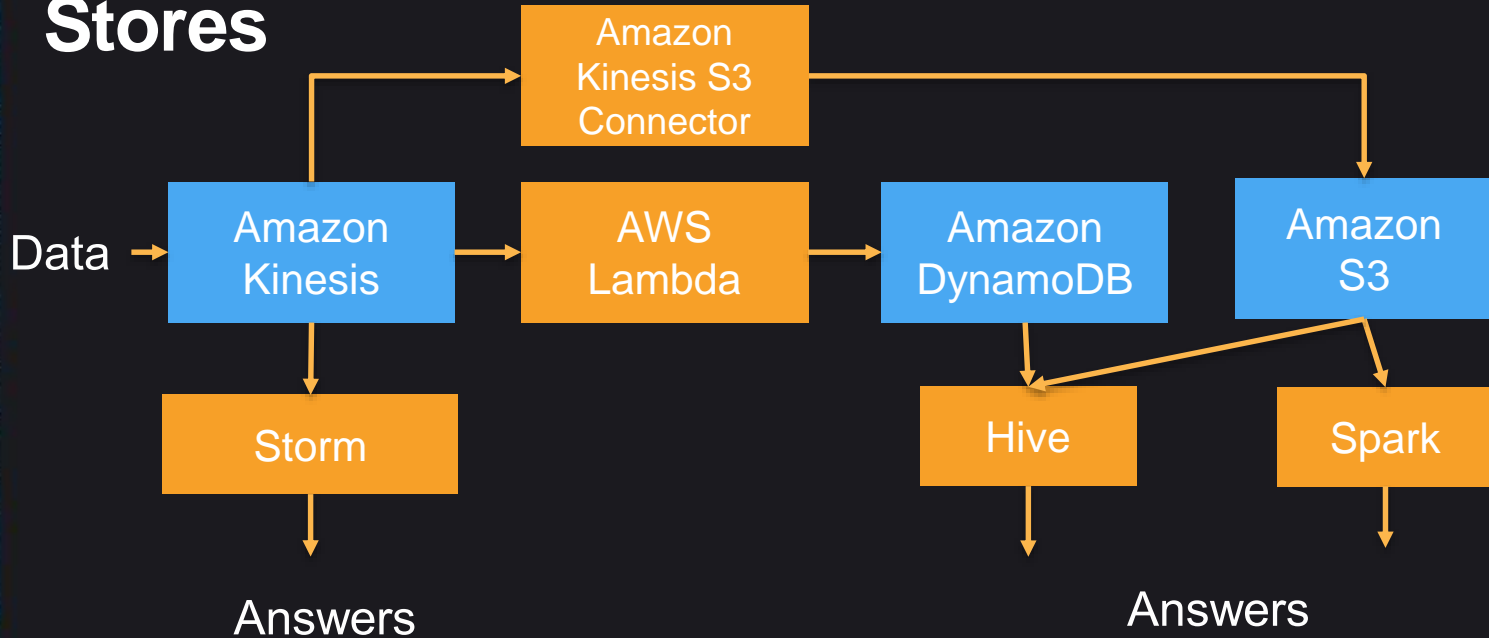
- Multiple stages
- Storage decoupled from processing



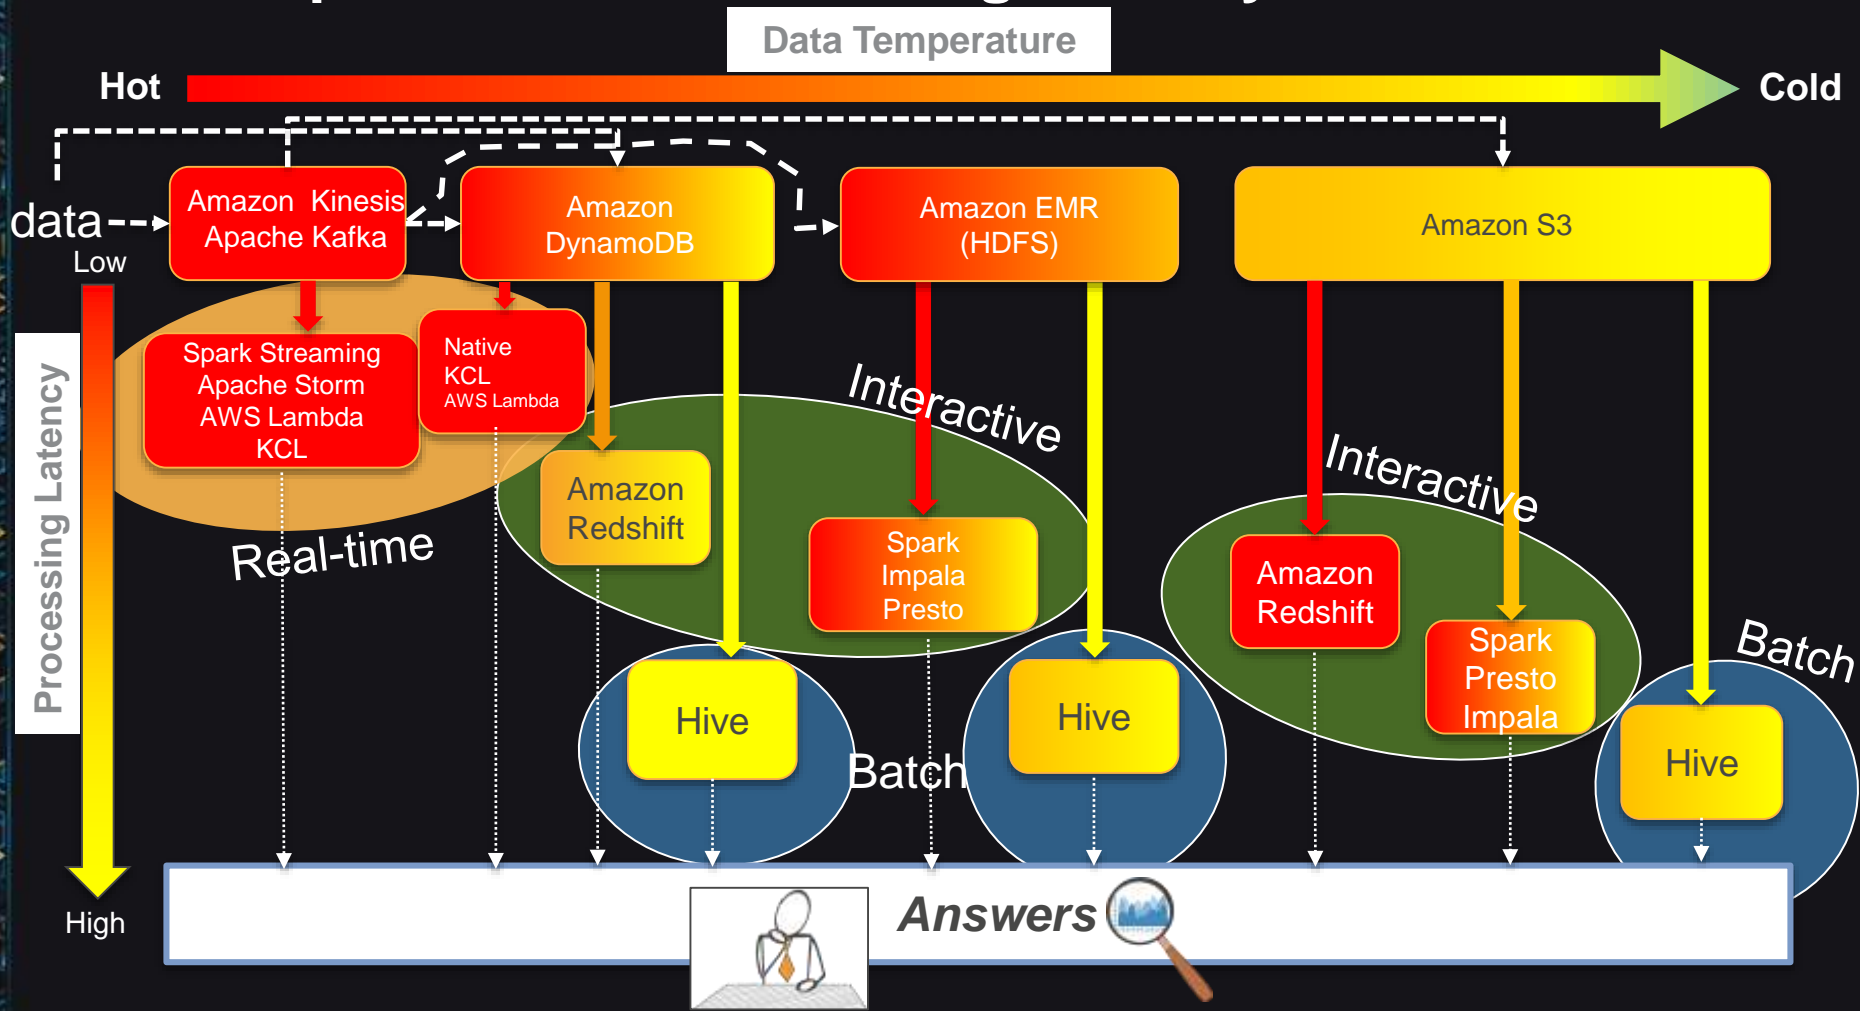
Multiple Processing Applications (or Connectors) Can Read from or Write to Multiple Data Stores



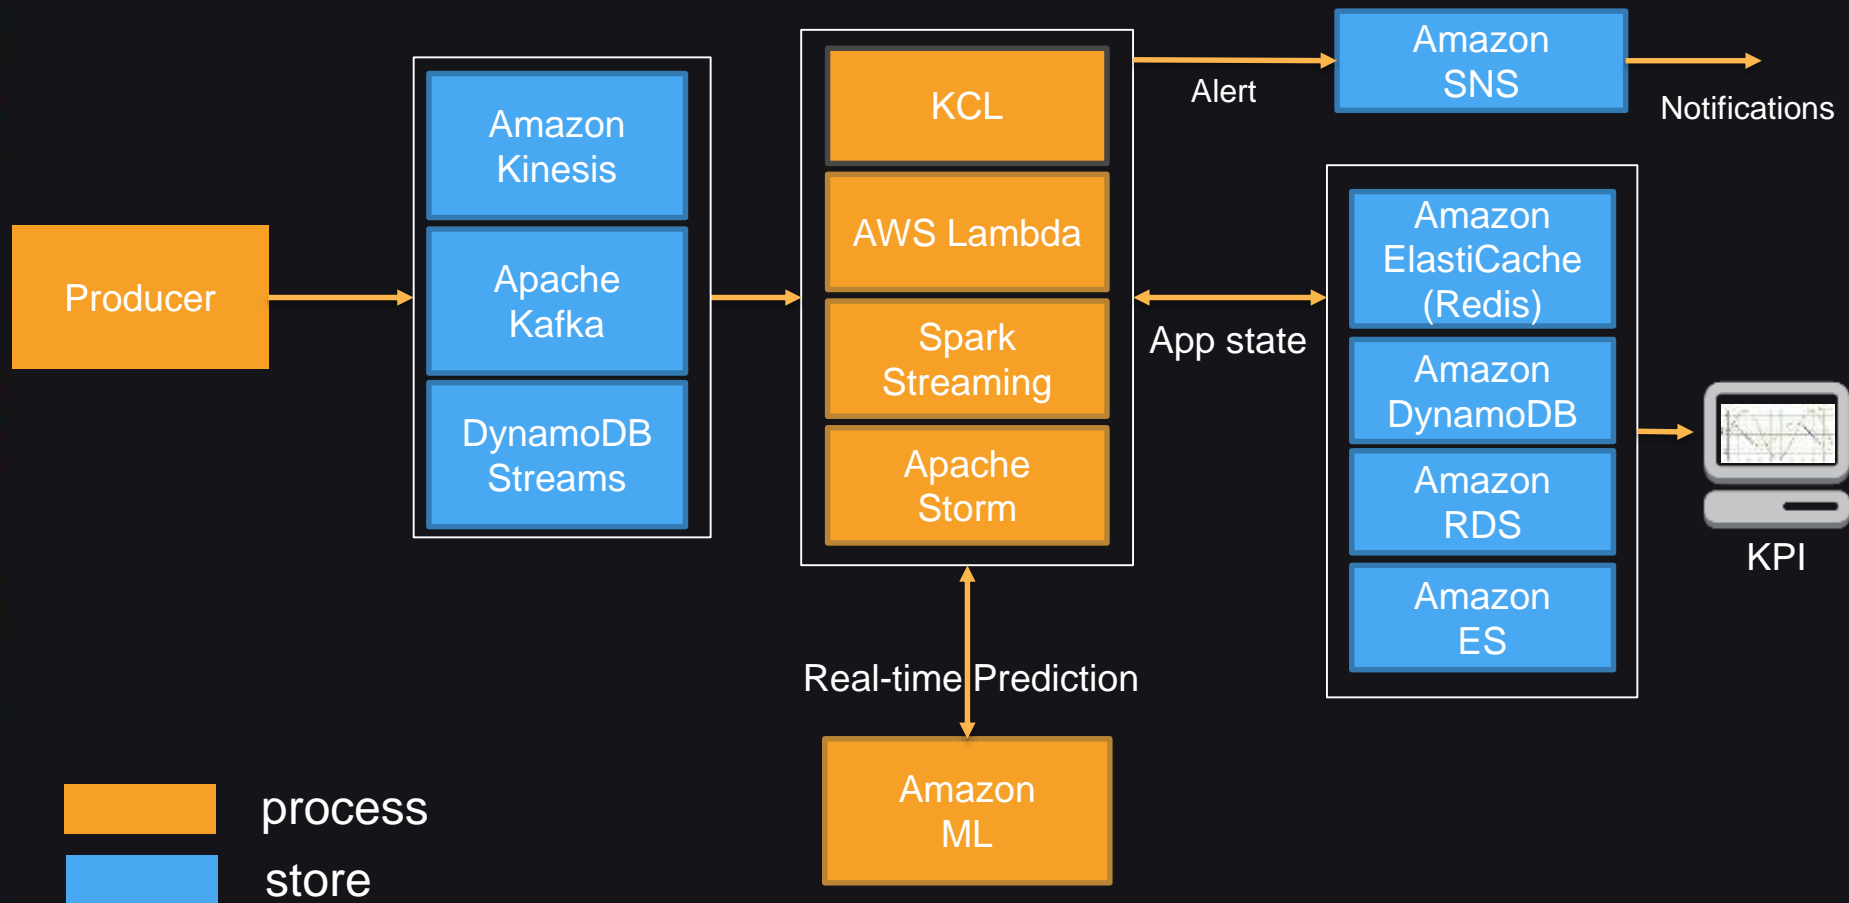
Processing Frameworks (KCL, Storm, Hive, Spark, etc.) Could Read from Multiple Data Stores



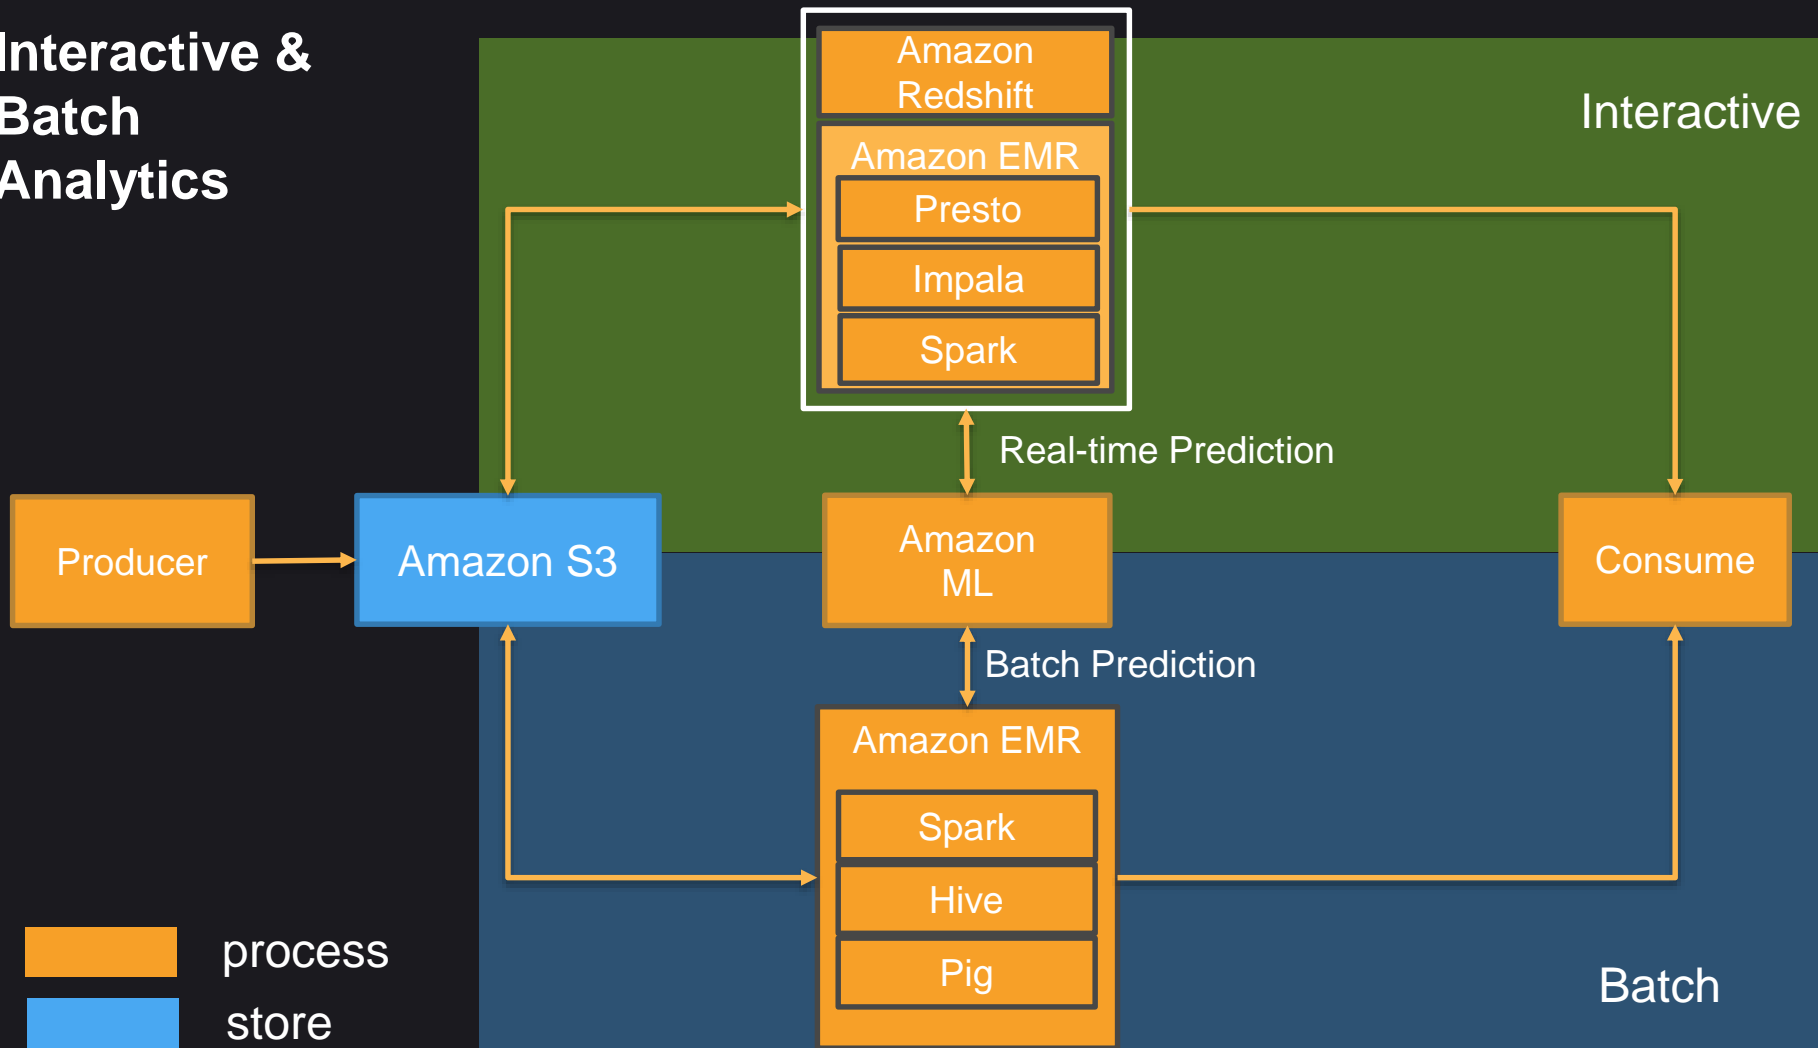
Data Temperature vs Processing Latency



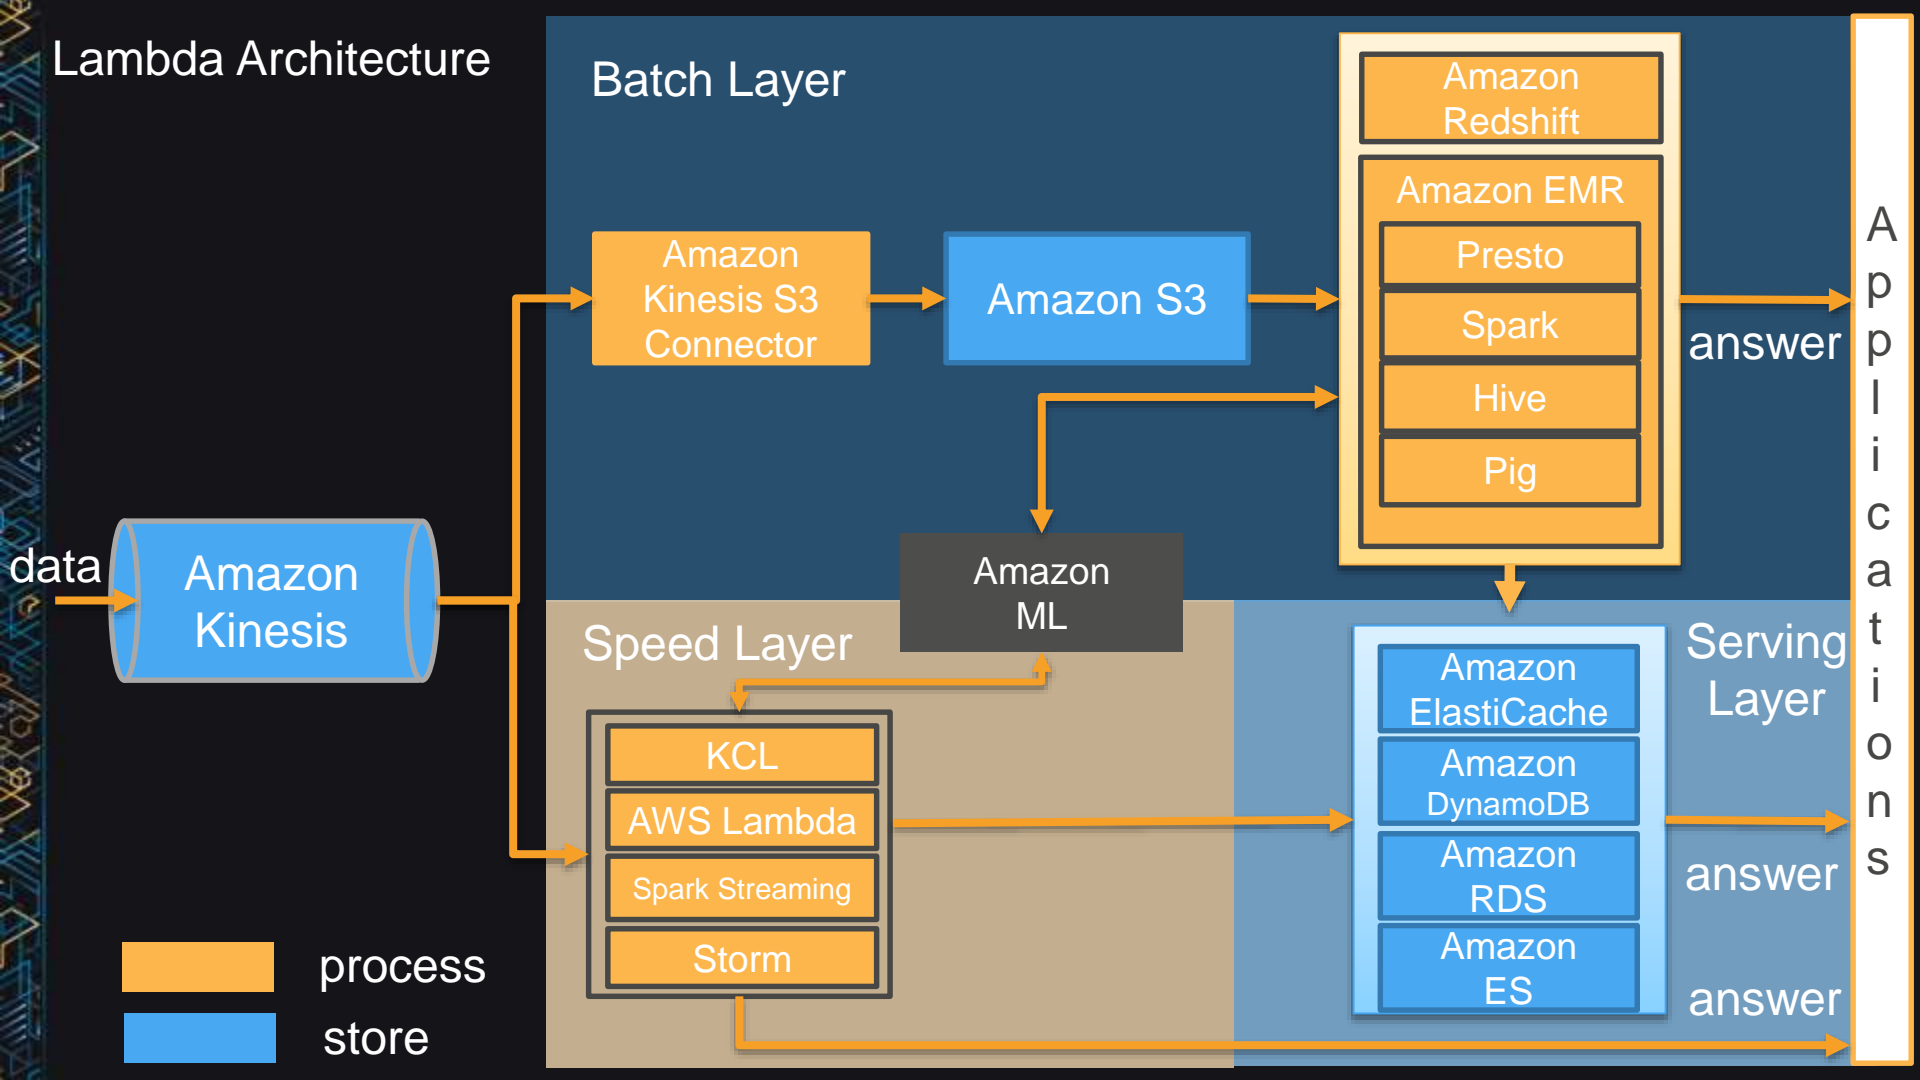
Real-time Analytics



Interactive & Batch Analytics



Lambda Architecture



Summary

- Build decoupled “data bus”
 - Data → Store ↔ Process → Answers
- Use the right tool for the job
 - Latency, throughput, access patterns
- Use Lambda architecture ideas
 - Immutable (append-only) log, batch/speed/serving layer
- Leverage AWS managed services
 - No/low admin
- Be cost conscious
 - Big data ≠ big cost



**Remember to complete
your evaluations!**



Thank you!