

Cloud Application Architecture Patterns

Developing Microservice Applications for PaaS

Matt Stine (@mstine)

Cloud Foundry Platform Engineer

matt.stine@gmail.com

<http://www.mattstine.com>



The Cloud

Software
as a Service

This Talk

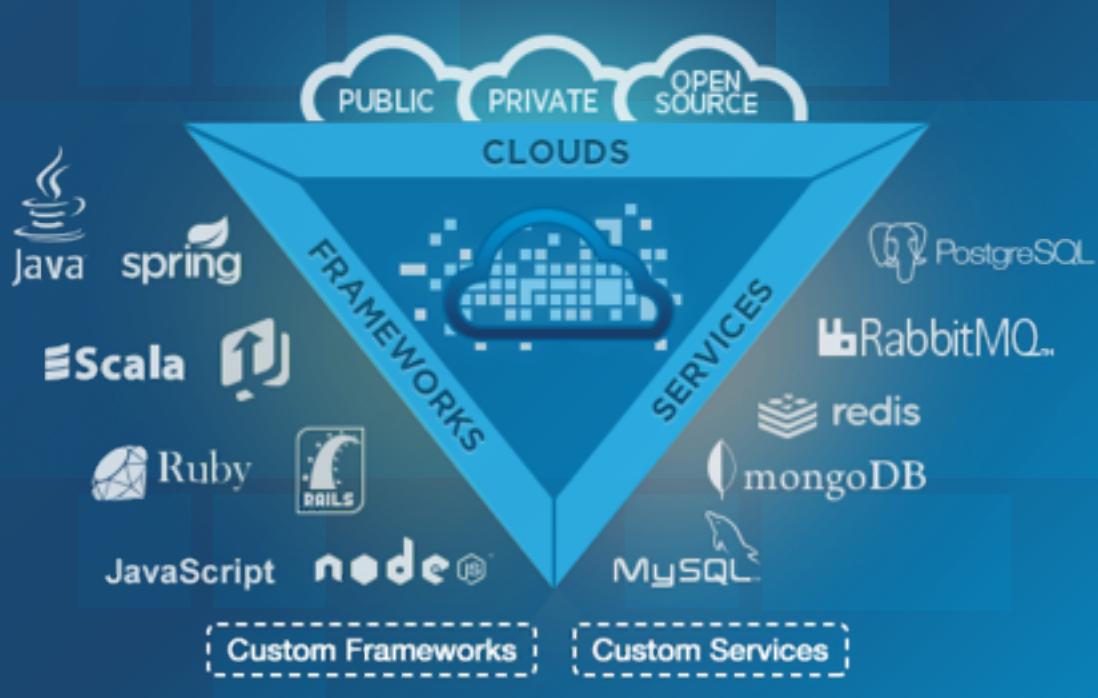
Platform as a Service

Infrastructure as a Service

PaaS Platforms

- Heroku (<http://www.heroku.com>)
- OpenShift (<https://www.openshift.com>)
- Google App Engine (<https://appengine.google.com>)
- AWS Elastic Beanstalk (<https://aws.amazon.com/elasticbeanstalk>)
- Microsoft Azure Cloud Services (<http://azure.microsoft.com/en-us/services/cloud-services/>)
- Cloud Foundry (<http://www.cloudfoundry.org>)
 - Pivotal Web Services (<https://run.pivotal.io>)
 - IBM BlueMix (<https://ace.ng.bluemix.net>)
 - Anynines (<http://www.anynines.com>)

Cloud Foundry PaaS



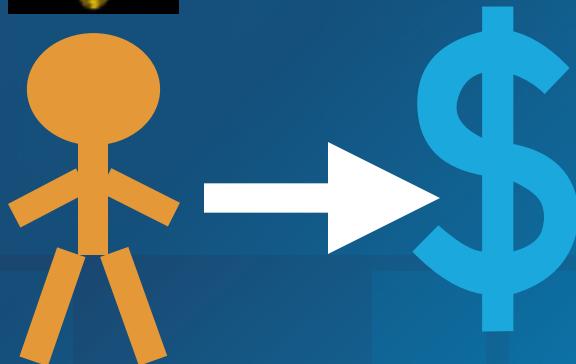
Spring (<https://spring.io>)



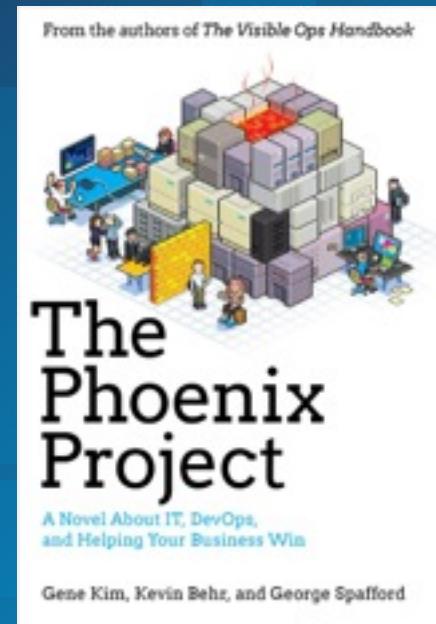
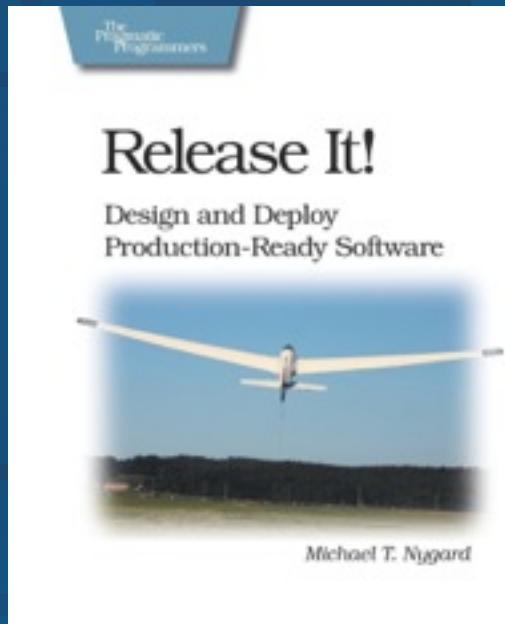
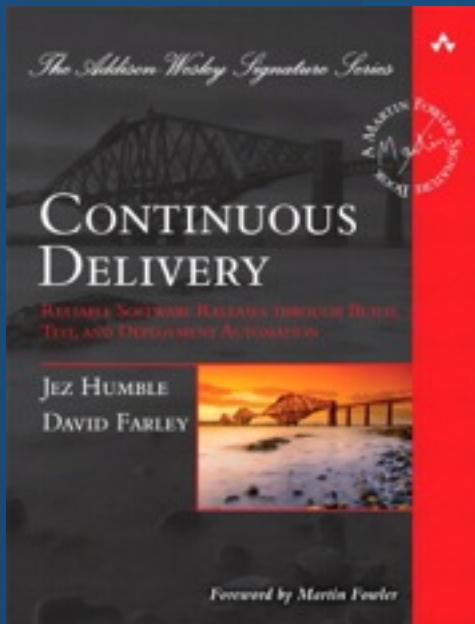
Architecting for Continuous Delivery



Architecting for Continuous Delivery



Continuous Delivery - How?



Warner Music: Software Factories



Warner Software Factory Platform

- New applications and major updates
 - **Before:** 6 months, team of 10 developers
 - **After:** 6 weeks, same team
 - **Speed/Agility:** 400% faster on new platform
 - **HR Hard Savings:** \$1.1M per application update delivered

Iterative Development



Horizontal Scale



Slow/Expensive

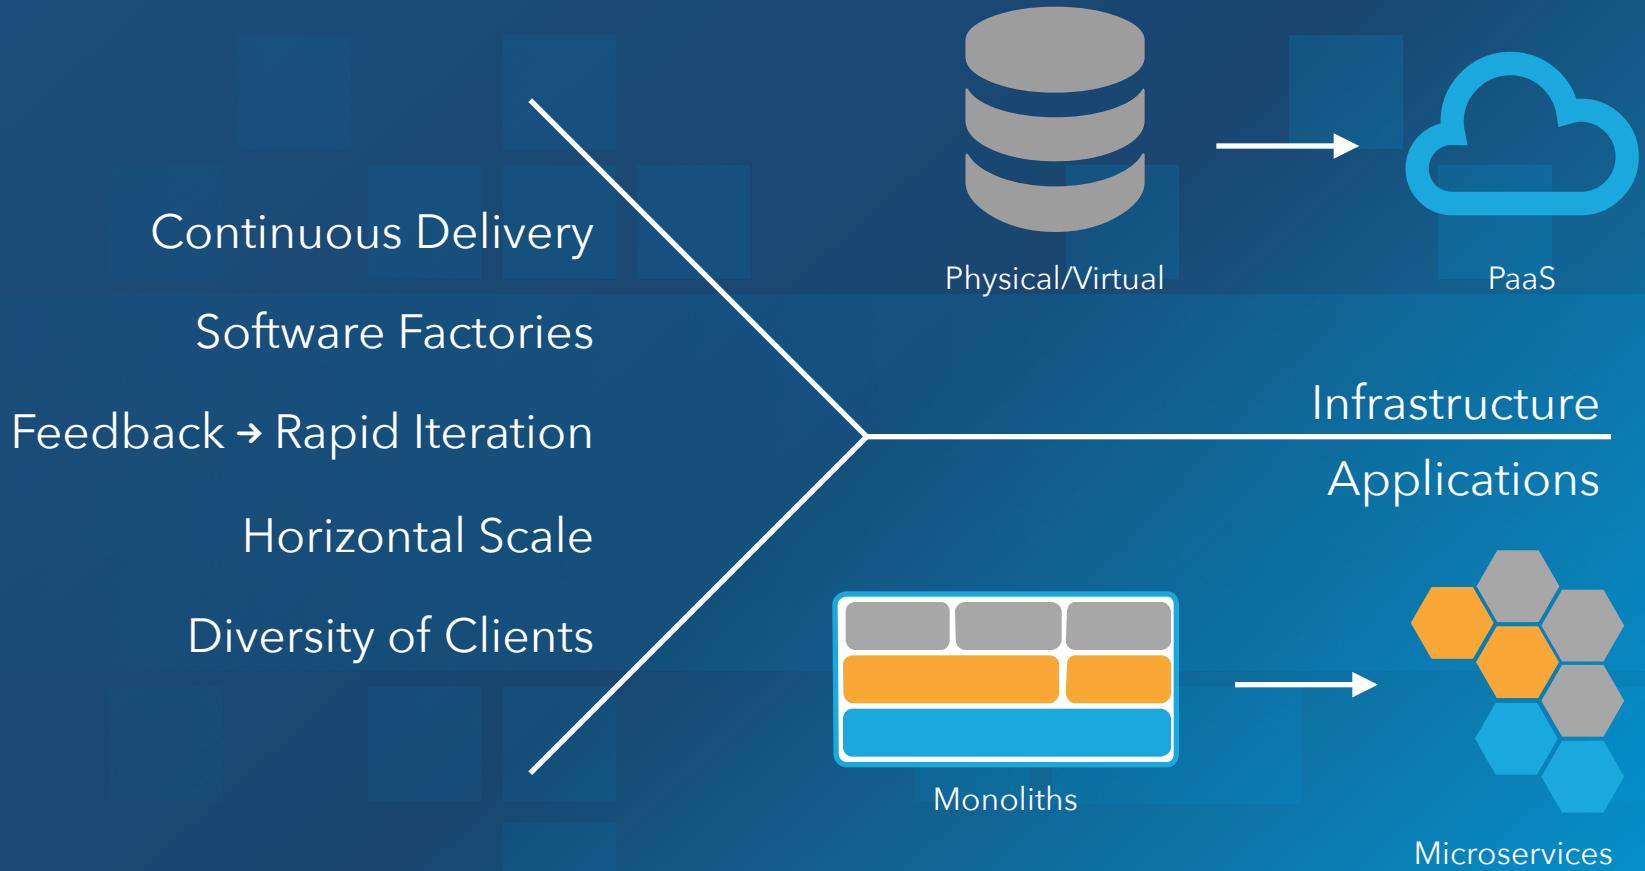
Fast/Cheap

Diversity of Clients

In January 2014, mobile devices accounted for 55% of Internet usage in the United States. Apps made up 47% of Internet traffic and 8% of traffic came from mobile browsers.



<http://money.cnn.com/2014/02/28/technology/mobile/mobile-apps-internet/>

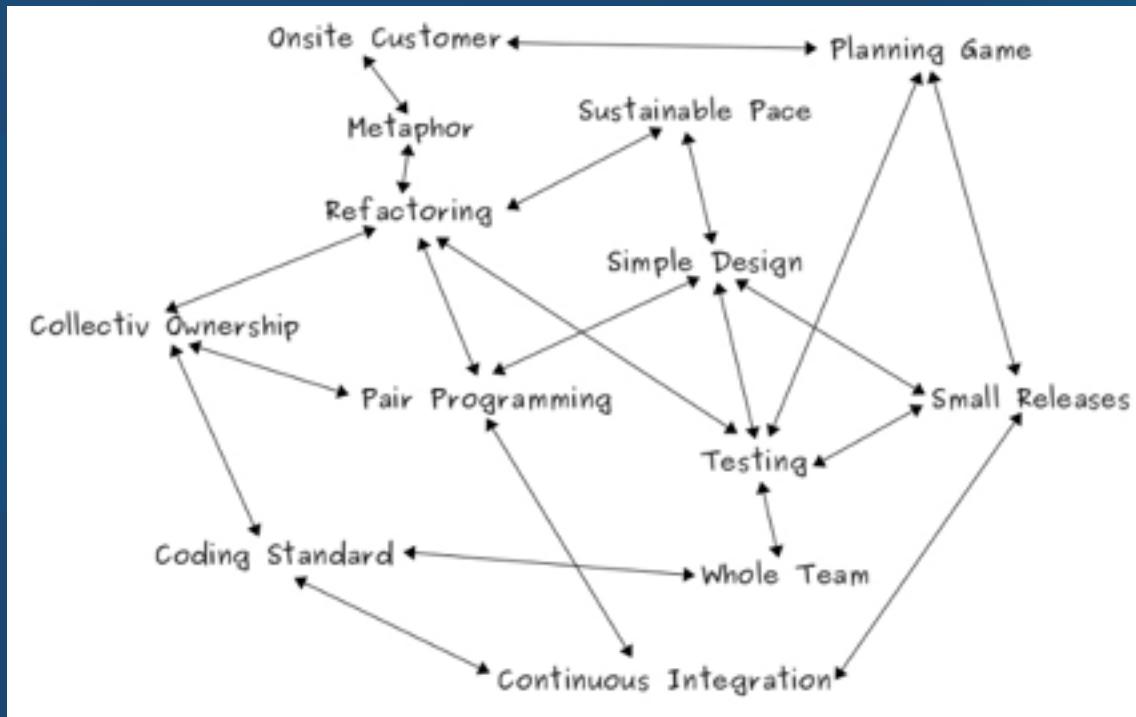


New Architectural Constraints

- Platforms like CF and Heroku optimizes for 12 Factor Linux applications
- Microservices: a radical departure from traditional monolithic applications
- In both cases, the enterprise is forced to “think different.”



How XP Practices Support Each Other



A Symbiotic Relationship...



Patterns

- Microservice
- API Gateway
- Stateless/Shared-Nothing
- Configuration/Service Consumption
- Fault Tolerance

Pattern:

Microservice



Simple vs. Easy

- Simple

- *sim-plex*

- one fold/braid

- vs **complex**

- Easy

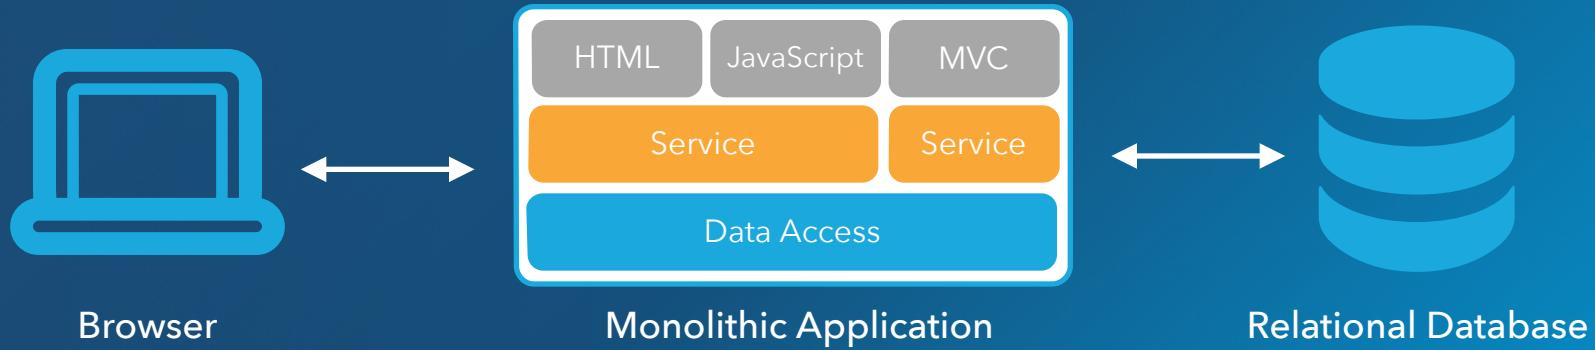
- *ease < aise < adjacens*

- lie near

- vs hard



Monolithic Architecture

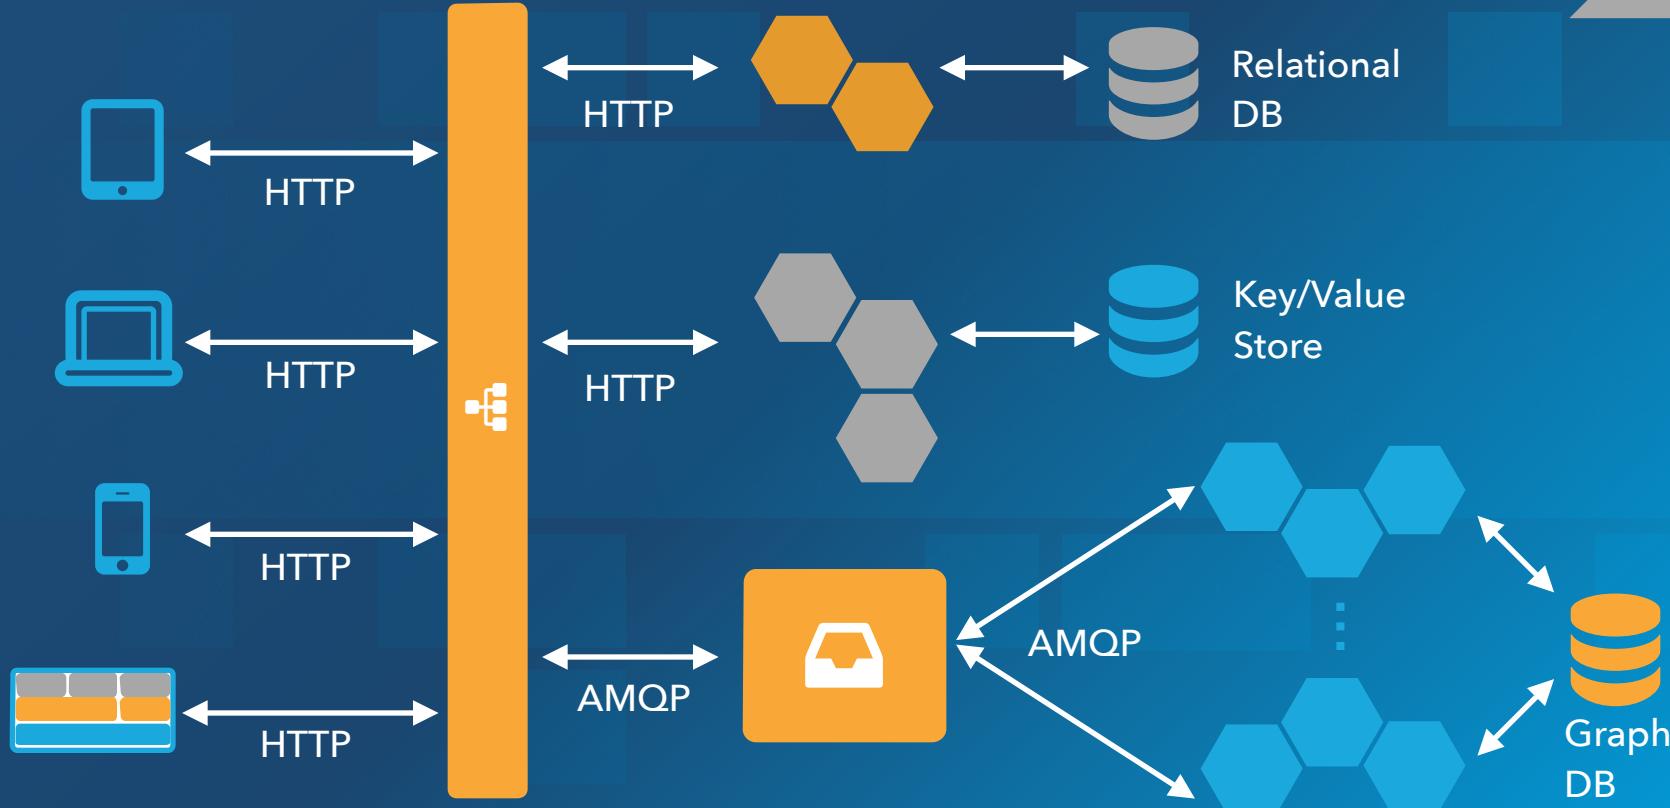




Monolithic Architectures

- Complex / Easy
- Modularity Dependent Upon Language / Frameworks
- Change Cycles Tightly Coupled / Obstacle to Frequent Deploys
- Inefficient Scaling
- Can Be Intimidating to New Developers
- Obstacle to Scaling Development
- Requires Long-Term Commitment to Technical Stack

Microservice Architecture





Microservice Architectures

- Simple / Challenging
- Modularity Based on Component Services
- Change Cycles Decoupled / Enable Frequent Deploys
- Efficient Scaling
- Individual Components Less Intimidating to New Developers
- Enables Scaling of Development
- Eliminates Long-Term Commitment to Technical Stack

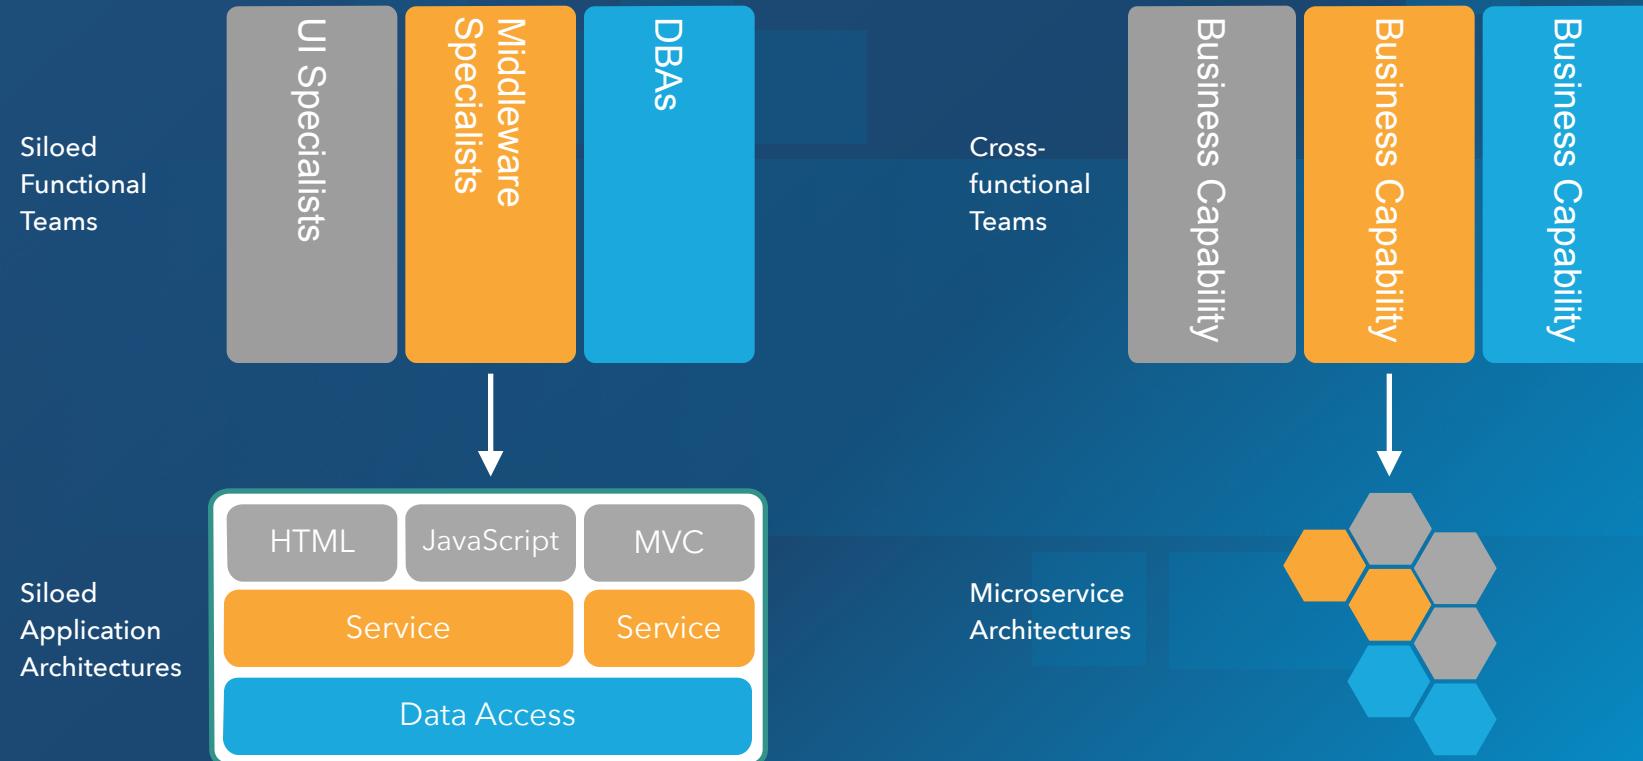


Conway's Law

Any organization that designs a system (defined broadly) will produce a design whose structure is a copy of the organization's communication structure.

Melvyn Conway, 1967

Organize Around Business Capabilities

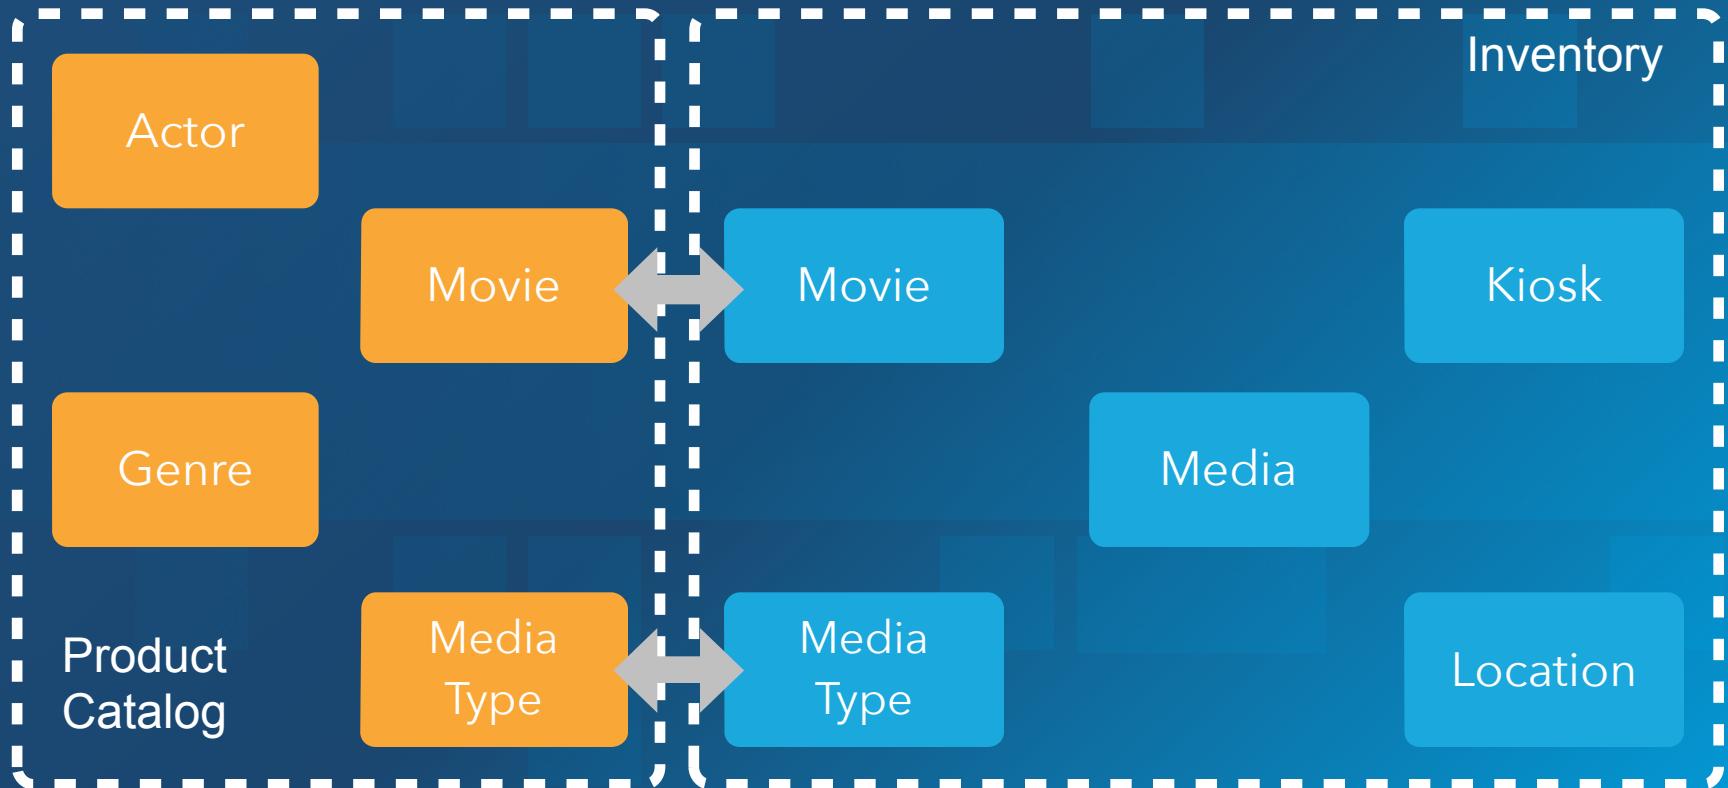




Partitioning Strategies

- By Noun (e.g. product info service)
- By Verb (e.g. shipping service)
- Single Responsibility Principle
(http://programmer.97things.oreilly.com/wiki/index.php/The_Single_Responsibility_Principle)
- Bounded Context (<http://martinfowler.com/bliki/BoundedContext.html>)

Bounded Contexts





UNIX Pipes and Filters



```
cut -d" " -f1 < access.log | sort | uniq -c | sort -rn | less
```



Choreography over Orchestration





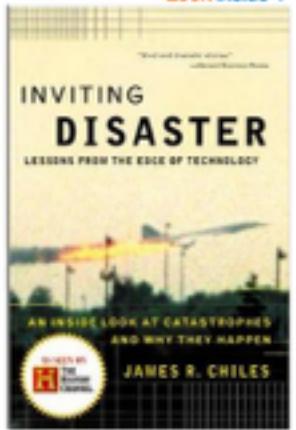
Challenges of Microservices

- Distributed System
- Remote Calls More Expensive Than In-process Calls
- Eventual Consistency
- Features Spanning Multiple Services
- Dependency Management / API Versioning
- Refactoring Module Boundaries

Pattern:

API Gateway





Look Inside

Inviting Disaster: Lessons From the Edge of Technology Paperback

by James R. Chiles (Author)

4.5 out of 5 stars - 58 customer reviews

See all 7 formats and editions

Kindle
\$10.23

Hardcover
From \$0.01

Paperback
\$12.08 ✓Prime

37 Used from \$0.01
14 New from \$23.99

50 Used from \$0.71
16 New from \$6.16
2 Collectible from \$9.00

Combining captivating storytelling with eye-opening findings, *Inviting Disaster* delves inside some of history's worst catastrophes in order to show how increasingly "smart" systems leave us wide open to human tragedy.

Weaving a dramatic narrative that explains how breakdowns in these systems result in such disasters as the chain reaction crash of the Air France Concorde to the meltdown at the Chernobyl Nuclear Power Station, Chiles vividly demonstrates how the battle between man and machine may be escalating beyond manageable limits -- and why we all have a stake in its outcome.

Included in this edition is a special introduction providing a behind-the-scenes look at the World Trade Center catastrophe. Combining firsthand accounts of employees' escapes with an in-depth look at the

+ Read more

Share



Buy Now

✓Prime

\$12.08

Qty: 1

List Price: \$16.00

Save: \$3.91 (24%)

Only 16 left in stock (more on the way).

Ship from and sold by Amazon.com.

Gift wrap available.



Add to Cart

Sign in to turn on 1-click ordering

Want it tomorrow, May 27? Order within 2 hrs 22 mins and choose One-Day Shipping at checkout. Details

Buy Used

✓Prime

\$9.22

Add to Wish List

Have one to sell?

Sell on Amazon

Frequently Bought Together



Price for all three: \$53.26

Add all three to Cart

Add all three to Wish List

Show availability and shipping details

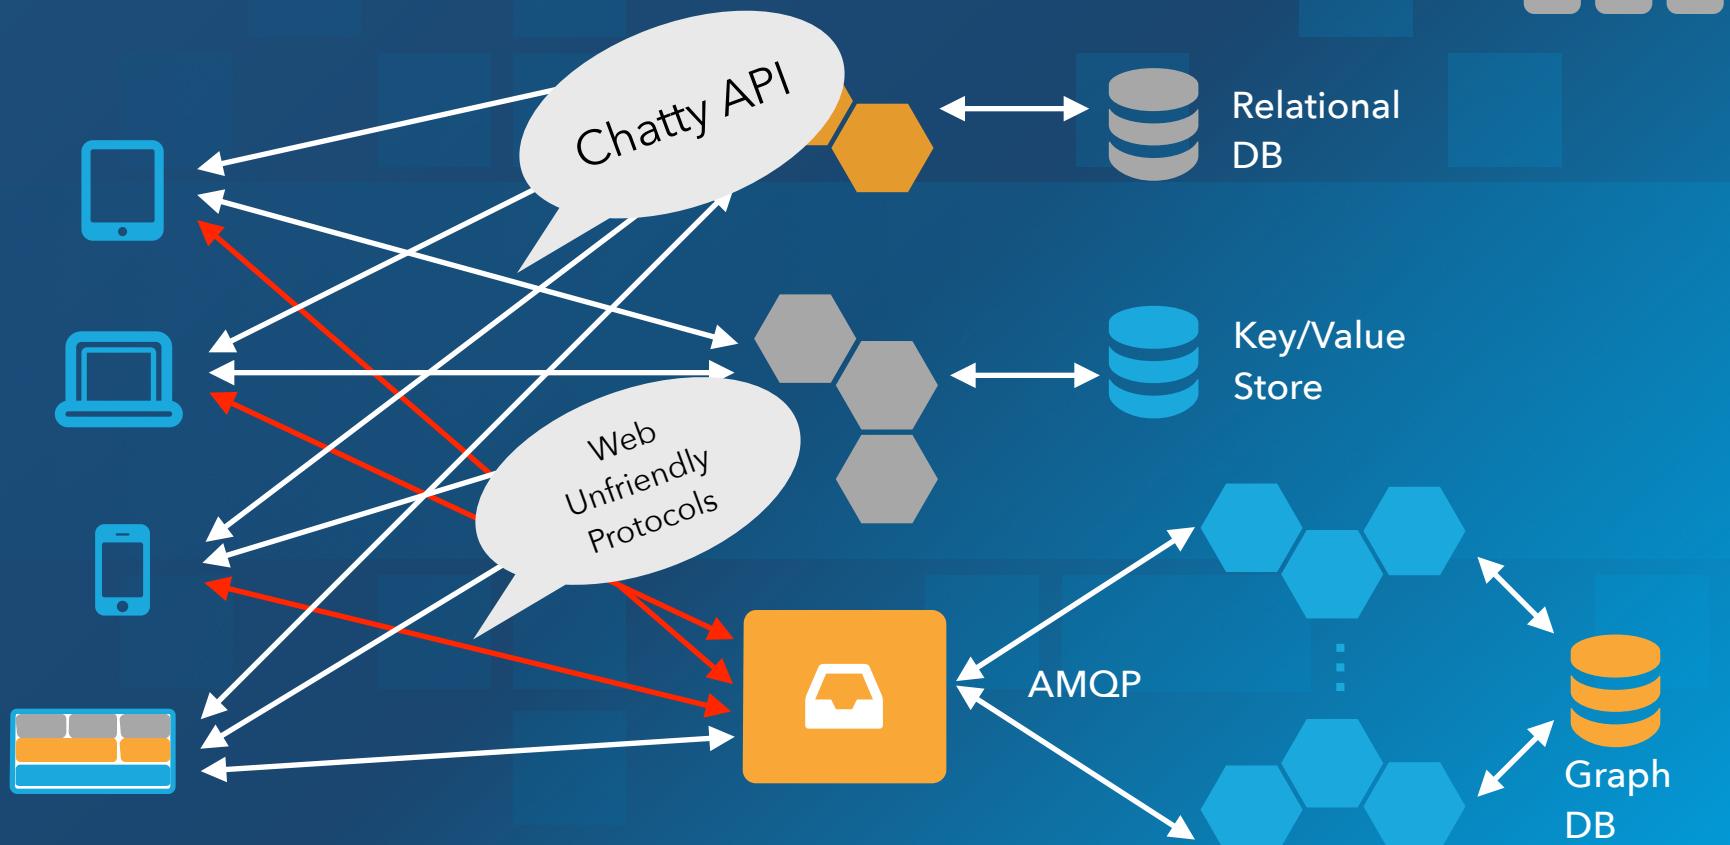
- This item: Inviting Disaster: Lessons From the Edge of Technology by James R. Chiles Paperback \$12.08
- The Logic Of Failure: Recognizing And Avoiding Error In Complex Situations by Detrich Damer Paperback \$12.26
- Normal Accidents: Living with High-Risk Technologies by Charles Perrow Paperback \$28.82

How many microservices?

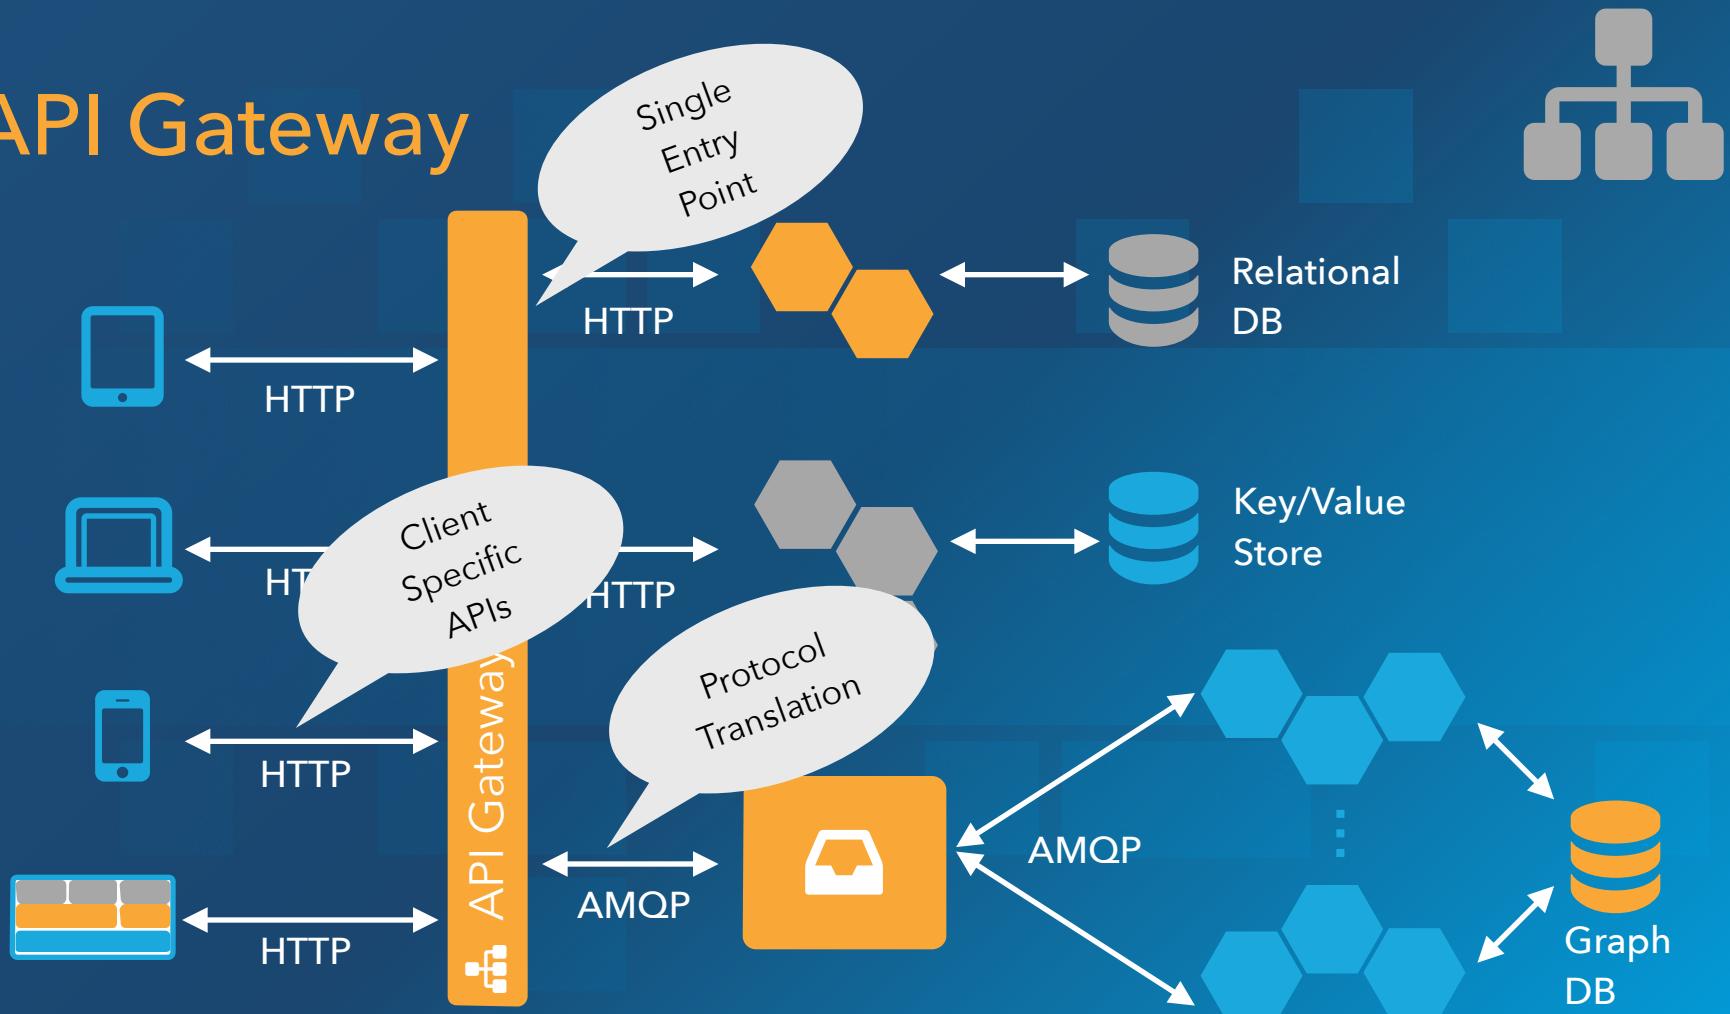
ELEVEN



Direct Connect



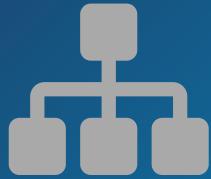
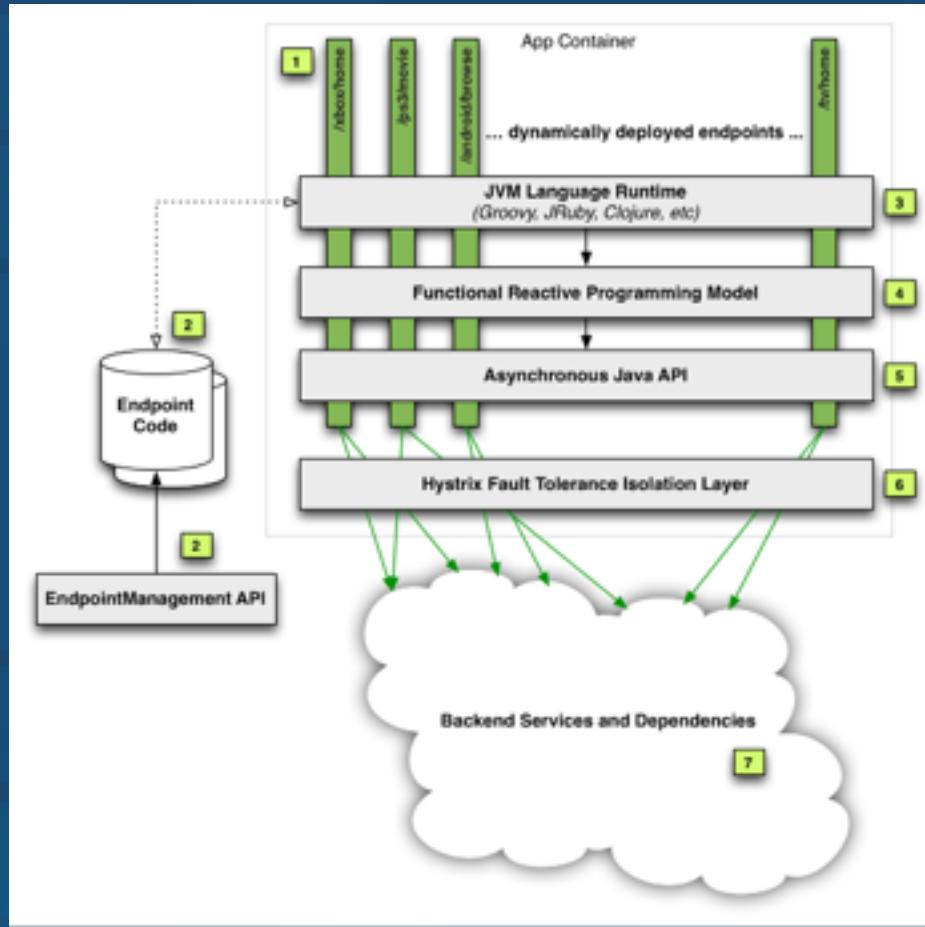
API Gateway



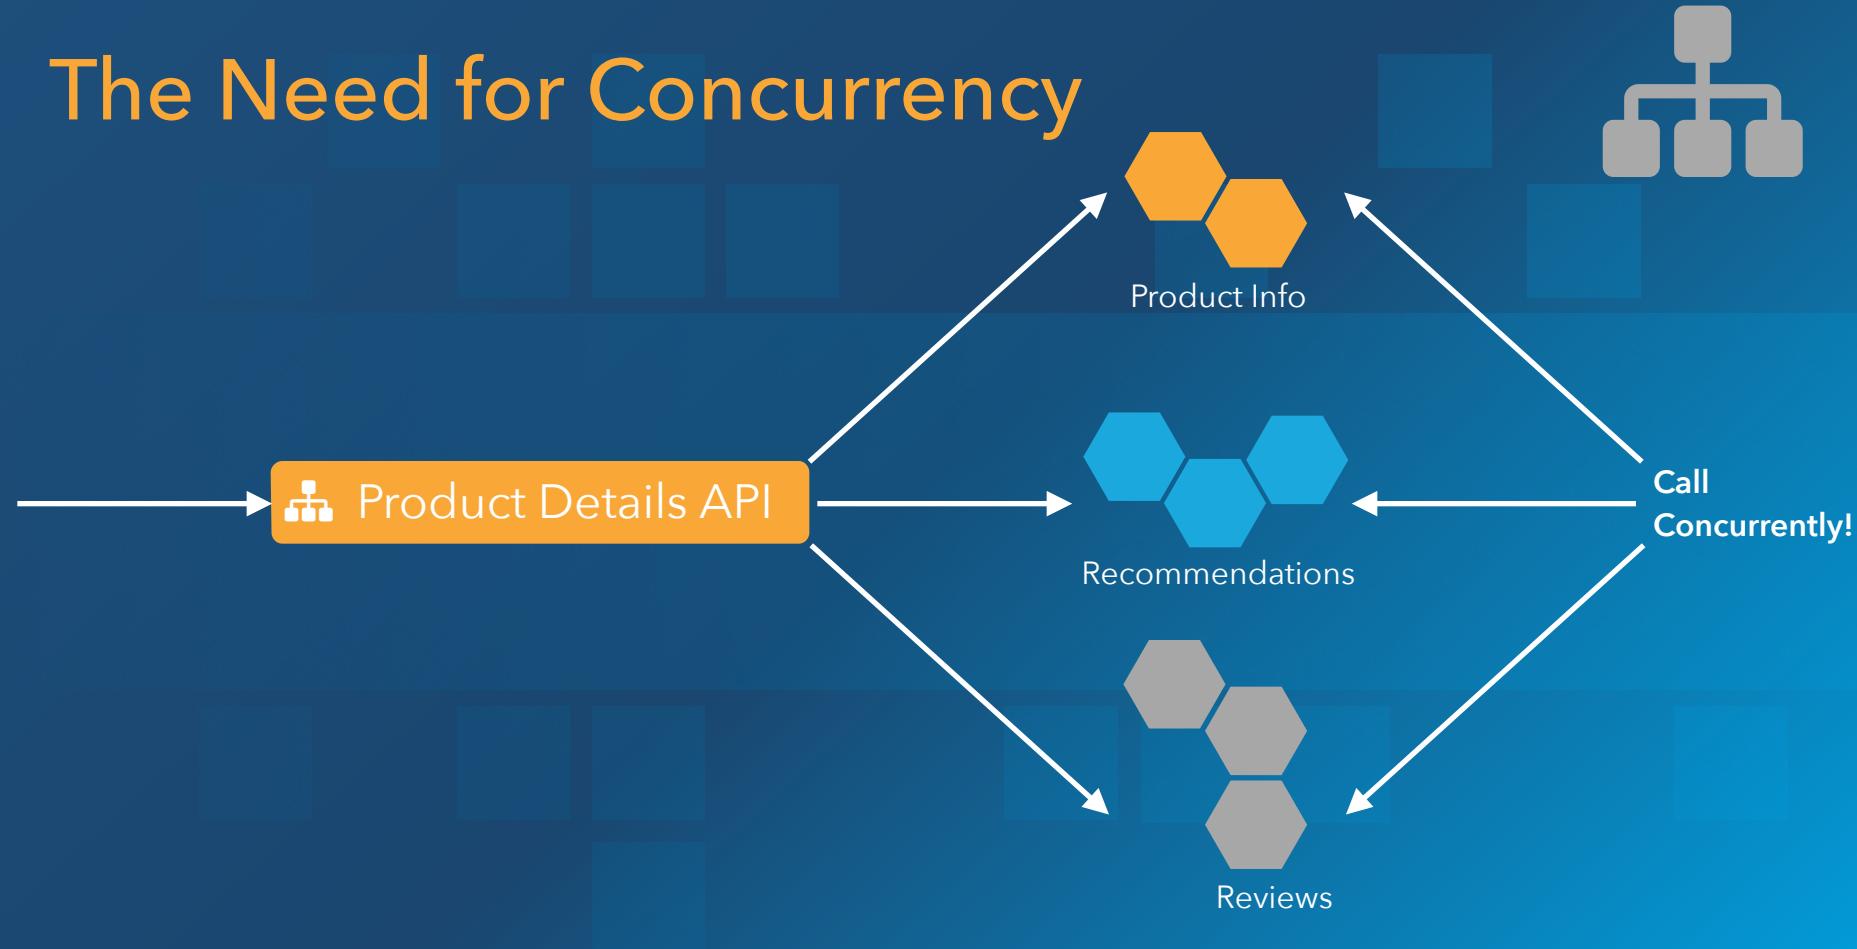
Importance for Mobile

- Latency
 - You can't wait for 2 second or 1 MB responses
- Round Trips
 - Too much network == KILL THE BATTERY
- Aggregation/Transformation
 - Device-specific responses (manufacturer/device/type/size/etc.)
 - Different Form Factors == Different Use Cases

Netflix API Gateway



The Need for Concurrency



Building API Gateways



- **Reactor** (<https://spring.io/blog/2013/11/12/it-can-t-just-be-big-data-it-has-to-be-fast-data-reactor-1-0-goes-ga>)
- **RxJava** (<https://github.com/Netflix/RxJava>)
- **Vert.x** (<http://vertx.io/>)
- **Go** (<http://golang.org/>)
- **Node.js** (<http://nodejs.org/>)
- **Pivotal CF Mobile Services** (<https://network.pivotal.io/products/p-api-gateway>)

Pattern:

Stateless/Shared-Nothing



Elasticity



http://www.flickr.com/photos/karen_d/2944127077



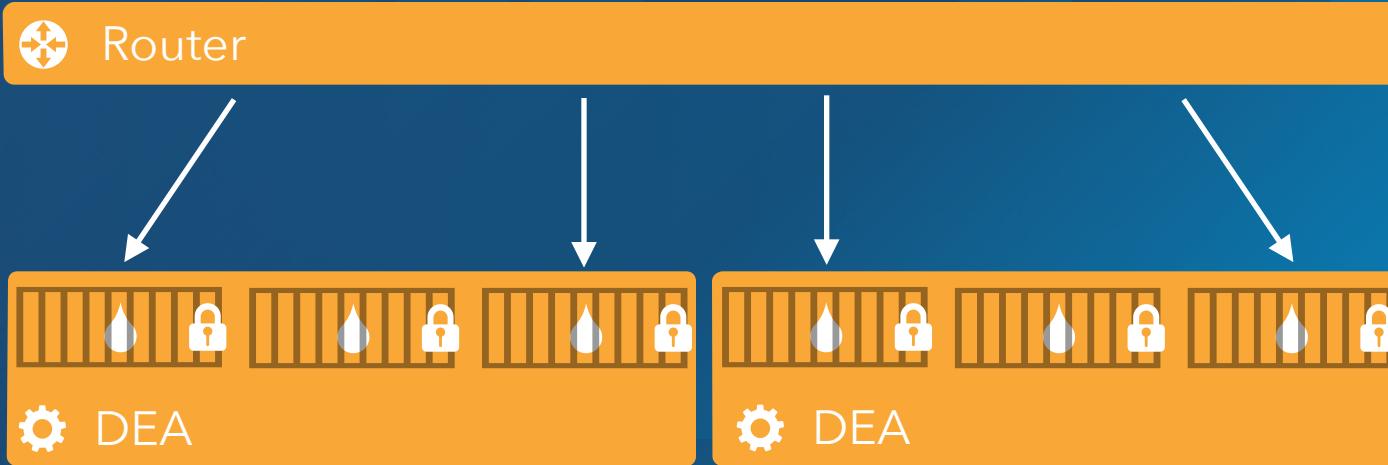
Ephemerality



<http://www.flickr.com/photos/smathur/852322080>

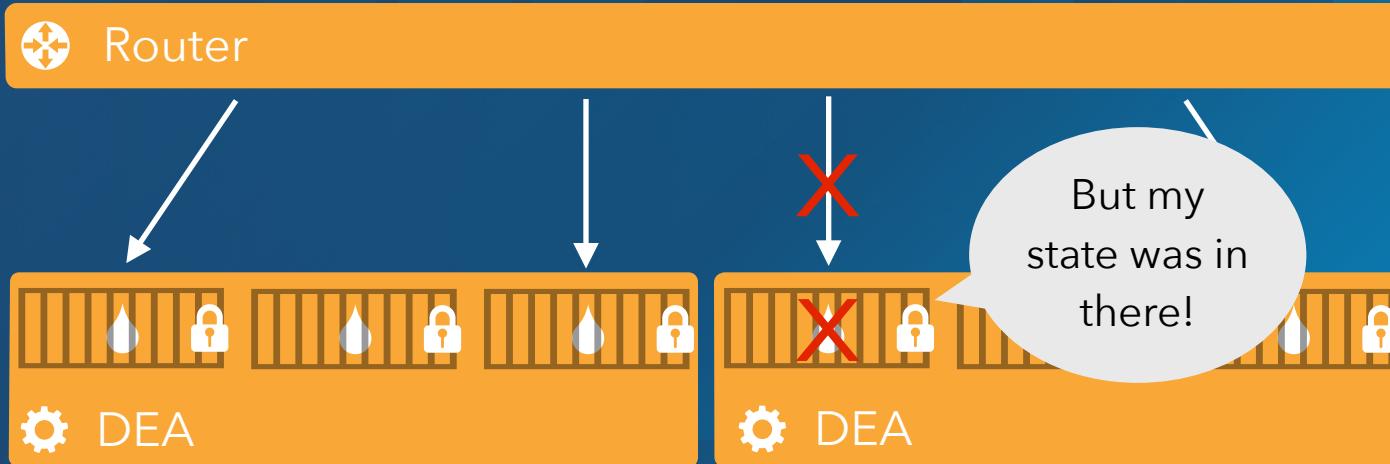


Why is state a problem?





Why is state a problem?



The Usual Suspects



6'6"

6'0"

5'6"

5'0"

4'6"

4'0"

A large, dense word cloud centered around the words "User" and "Sessions". The words are repeated in various sizes and colors (blue, orange, yellow) across the page. The background features faint horizontal lines.

The image is a dense collage of text elements, primarily in purple, blue, yellow, and orange, set against a white background. The words are arranged in a non-linear, overlapping fashion. Key words include 'Cache' (large, purple), 'Custom Application' (large, blue), 'App' (large, yellow), 'Frame' (large, blue and orange), and 'works' (large, orange). Smaller words interspersed include 'Application', 'Custom', 'Cache', 'App', 'Frame', 'works', 'Frameworks', and 'Frame'. The overall effect is a dynamic, abstract representation of software development concepts.



Option #1: Push State to the Client

Cookies, HTML 5, Single Page Applications (SPA)

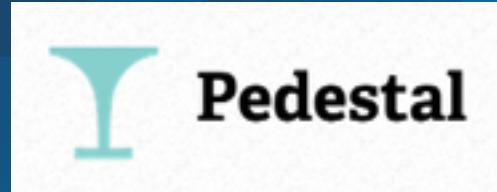


BACKBONE.JS



cujoJS

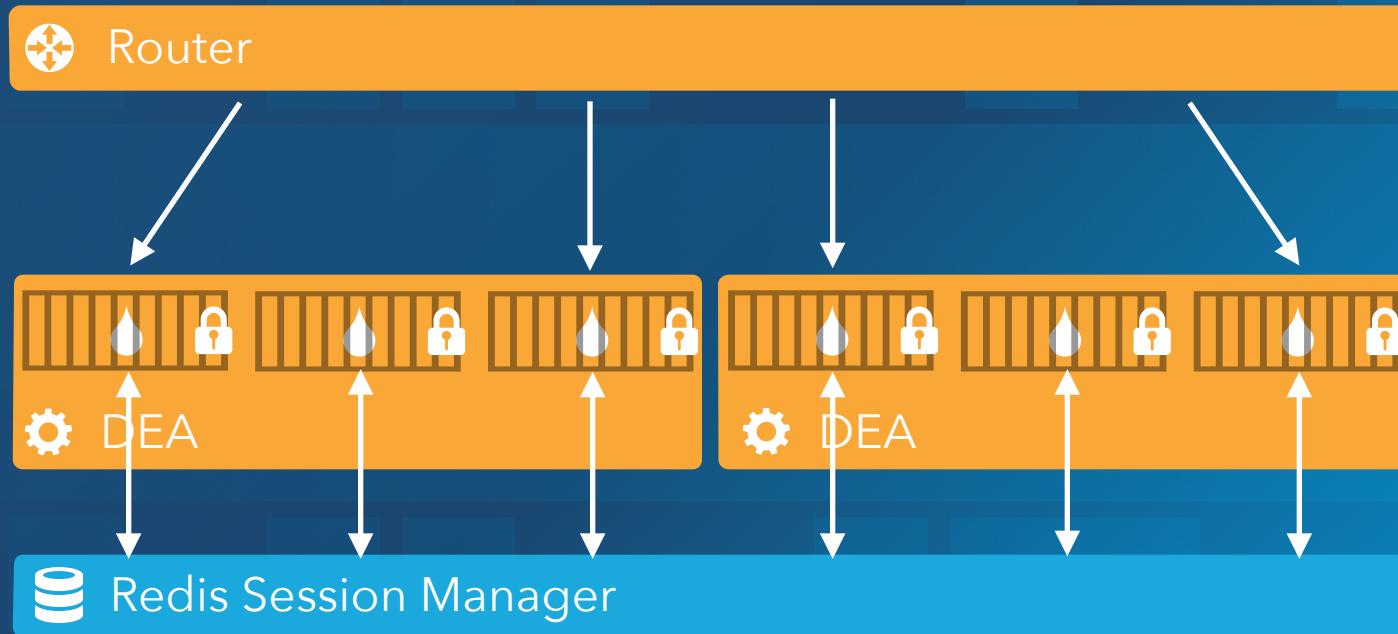
Knockout.



ANGULARJS
by Google



Option #2: External Services



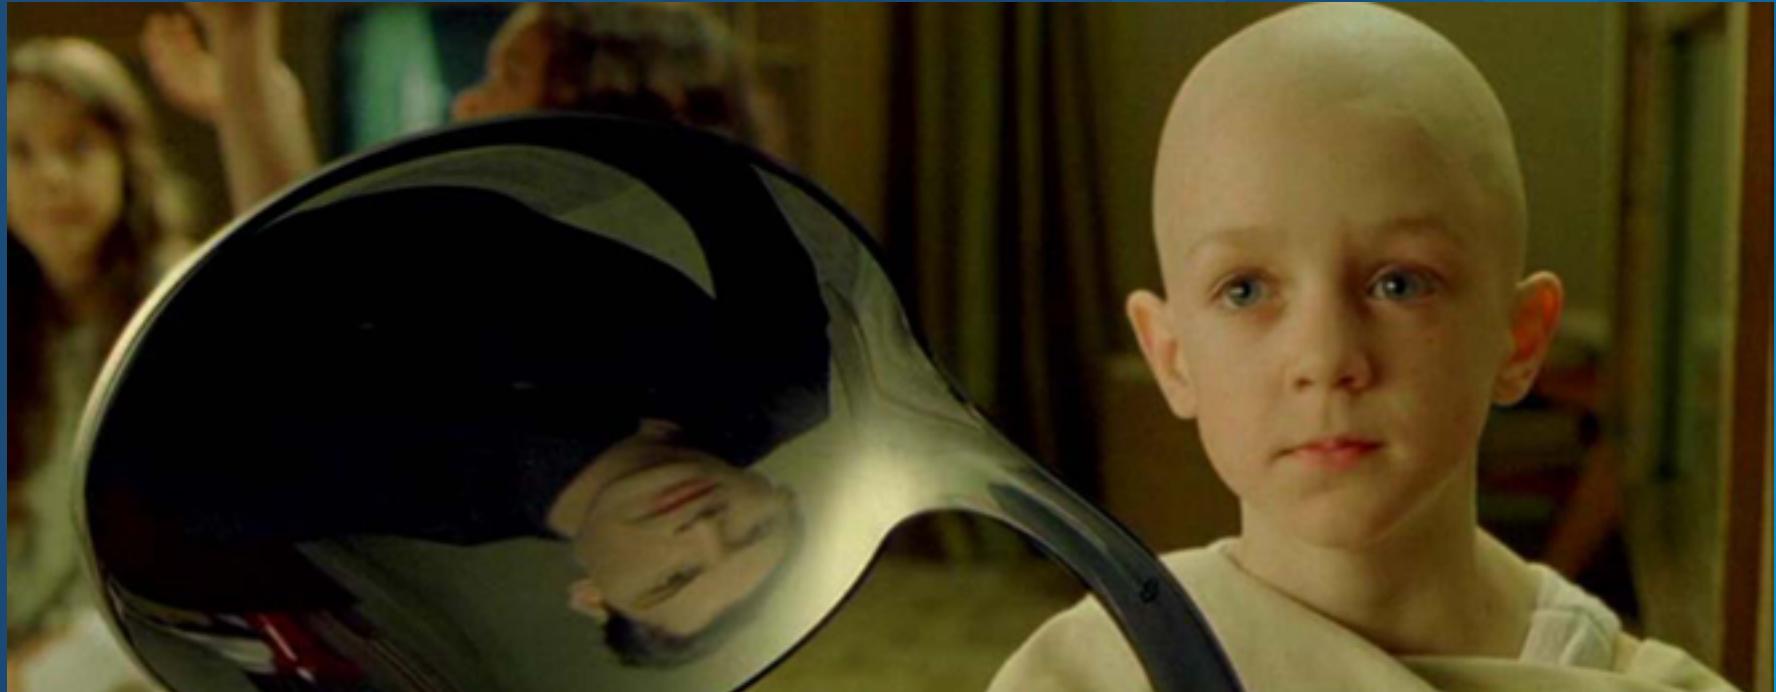


Option #2: External Services

- Spring Session
 - Container-independent session abstraction
 - <https://github.com/spring-projects/spring-session>
- Redis + CF Java Buildpack
 - Create a CF service containing a name, label, or tag with "session-replication" as a substring.
 - <https://github.com/cloudfoundry/java-buildpack/blob/master/docs/container-tomcat.md#session-replication>



There is no file system...





It's not shared...

 Router

HTTP POST - 200 OK

 DEA

 DEA

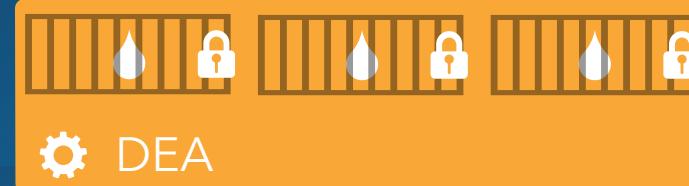




It's not shared...

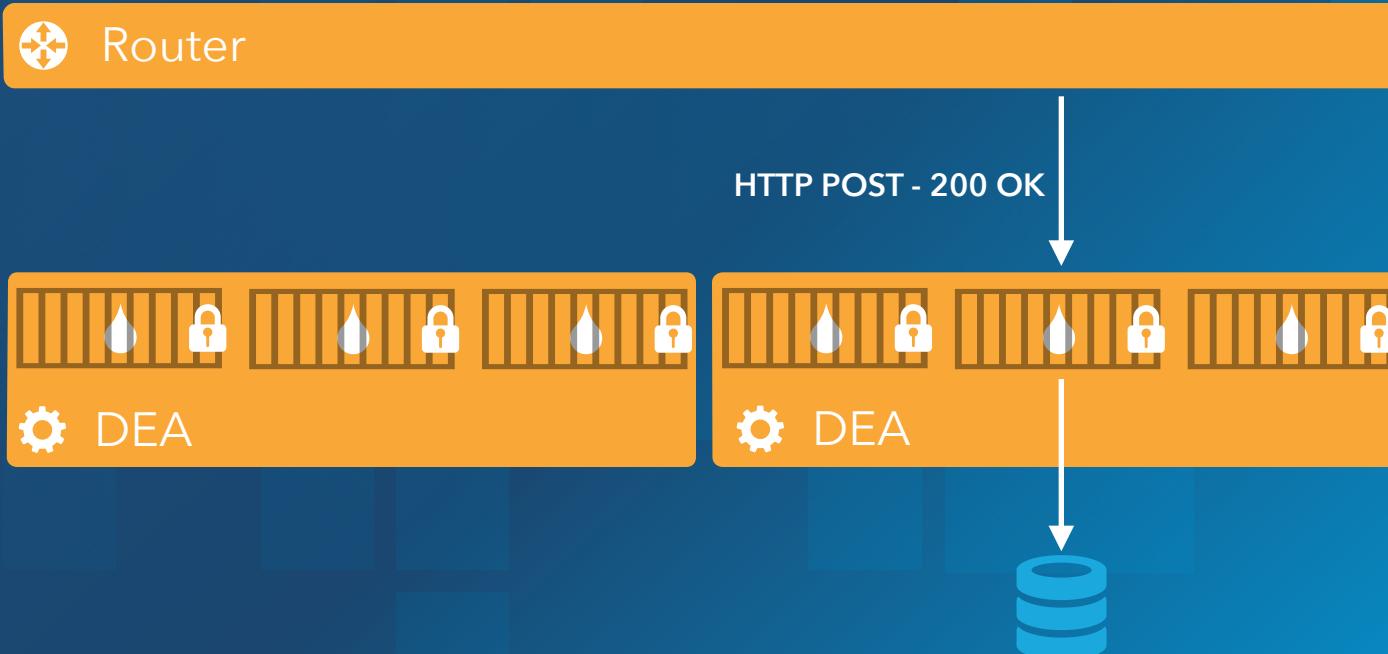
Router

HTTP GET - 404 NOT FOUND





It's not persistent...





It's not persistent...



Router



DEA

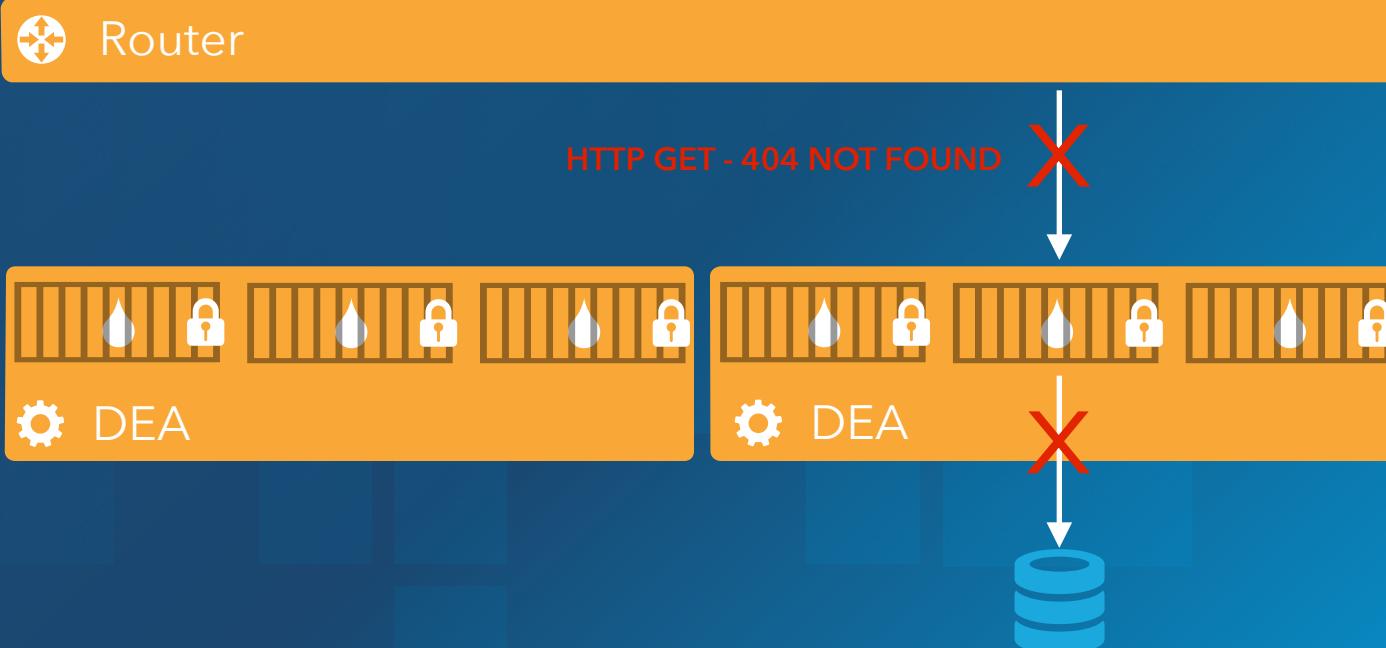


DEA





It's not persistent...





Use a persistent store!

- Amazon S3 (<http://aws.amazon.com/s3>)
- Google Cloud Storage (<https://cloud.google.com/products/cloud-storage>)
- OpenStack Swift (<http://swift.openstack.org>)
- Hadoop HDFS (http://hadoop.apache.org/docs/r1.2.1/hdfs_design.html)
- Riak CS for Pivotal CF (<https://network.pivotal.io/products/p-riakcs>)



PUBLIC



cloudfoundry-samples / cf-s3-demo



Unwatch

▼

3



Star

0



Fork

1

Basic demo using S3 as a user-provided service — Edit

14 commits

1 branch

0 releases

1 contributor



branch: master ▾

cf-s3-demo /

update for multiple manifests

mstline authored 2 months ago

latest commit c8e699d86b

gradle add Gradle wrapper 5 months ago

src Port from MongoDB to MySQL. Upgrade to Spring Boot 1.0.0.BUILD-SNAP... 2 months ago

.gitignore update for multiple manifests 2 months ago

README.md Port from MongoDB to MySQL. Upgrade to Spring Boot 1.0.0.BUILD-SNAP... 2 months ago

build.gradle remove commented dependency 2 months ago

gradlew add Gradle wrapper 5 months ago

gradlew.bat add Gradle wrapper 5 months ago

Code

Issues

1

Pull Requests

1

Wiki

Pulse

Graphs

Network

Settings

HTTPS clone URL

<https://github.com/cloudfoundry-samples/cf-s3-demo><https://github.com/cloudfoundry-samples/cf-s3-demo>



```
@Configuration
@Profile("cloud")
public class CloudConfig extends AbstractCloudConfig {

    @Bean
    public DataSource dataSource() {
        // Default pool size to 4 connections to support ClearDB Spark (free)
        PooledServiceConnectorConfig.PoolConfig poolConfig =
            new PooledServiceConnectorConfig.PoolConfig(4, 200);

        DataSourceConfig config = new DataSourceConfig(poolConfig, new DataSourceConfig.
            ConnectionConfig(""));

        return connectionFactory().dataSource(config);
    }

    @Bean
    public S3 s3() {
        return connectionFactory().service(S3.class);
    }

}
```



```
public class S3 {  
  
    private AmazonS3 amazonS3;  
    private String bucket;  
  
    public S3(AmazonS3 amazonS3, String bucket) {  
        this.amazonS3 = amazonS3;  
        this.bucket = bucket;  
    }  
  
    public S3File createS3FileObject(String id, String name, File file) {  
        return new S3File(id, bucket, name, file);  
    }  
  
    public void put(S3File file) {  
        PutObjectRequest putObjectRequest = new PutObjectRequest(bucket, file.getActualFileName(),  
            file.getFile());  
        putObjectRequest.withCannedAcl(CannedAccessControlList.PublicRead);  
        amazonS3.putObject(putObjectRequest);  
    }  
  
    public void delete(S3File file) {  
        amazonS3.deleteObject(bucket, file.getActualFileName());  
    }  
}
```



```
@RequestMapping(value = "/upload", method = RequestMethod.POST)
public String handleFileUpload(@RequestParam("file") MultipartFile file) {

    String id = UUID.randomUUID().toString();
    File uploadedFile = new File(file.getOriginalFilename());

    try {
        byte[] bytes = file.getBytes();
        BufferedOutputStream stream = new BufferedOutputStream(new FileOutputStream(uploadedFile));
        stream.write(bytes);
        stream.close();
    } catch (IOException e) {
        throw new RuntimeException("Failed to upload file!", e);
    }

    S3File s3File = s3.createS3FileObject(id, file.getOriginalFilename(), uploadedFile);

    s3.put(s3File);
    log.info(s3File.getName() + " put to S3.");
    repository.save(s3File);
    log.info(s3File.getName() + " record saved to MySQL.");
    uploadedFile.delete();
    log.info(s3File.getFile().getAbsolutePath() + " is deleted.");

    return "redirect:/";
}
```



```
@RequestMapping(value = "/delete/{id}", method = RequestMethod.GET)
public String deleteFile(@PathVariable String id) {

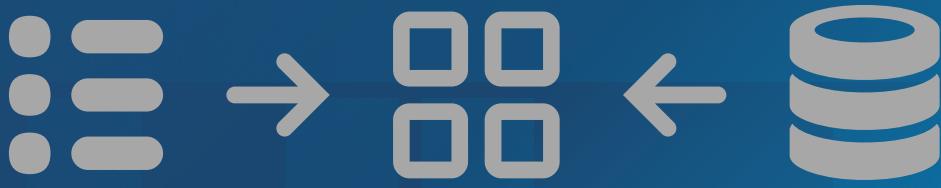
    S3File s3File = repository.findOne(id);

    repository.delete(s3File);
    log.info(s3File.getId() + " deleted from MySQL.");
    s3.delete(s3File);
    log.info(s3File.getActualFileName() + " deleted from S3 bucket.");

    return "redirect:/";
}
```

Pattern:

Configuration/Service Consumption





THE TWELVE-FACTOR APP

<http://12factor.net/>

What is configuration?



- Resource handles to databases and other backing services
- Credentials to external sources (e.g. S3, Twitter, ...)
- Per-deploy values (e.g. canonical hostname for deploy)
- ANYTHING that's likely to vary between deploys (dev, test, stage, prod)

Where NOT to store it:



- In the **CODE** (Obvious)
- In **PROPERTIES FILES** (That's code...)
- In the **BUILD** (ONE build, MANY deploys)
- In the **APP SERVER** (e.g. JNDI datasources)



Store it in the
ENVIRONMENT!

```
TMPDIR=/home/vcap/tmp
VCAP_APP_PORT=61863
USER=vcap
VCAP_APPLICATION={"instance_id": "b3e92a6fc443436888a525d100c91a12",
                  "instance_index": 0, "host": "0.0.0.0", "port": 61863, "started_at": "2013-12-04 01:52:01 +0000", "started_at_timestamp": 1386121921,
                  "start": "2013-12-04 01:52:01+0000", "state_timestamp": 1386121921,
                  "limits": {"mem": 512, "disk": 1024, "fds": 16384},
                  "application_version": "09c9bfe9-c14e-4fcb-8ad8-9fcfd4b854893",
                  "application_name": "cf-s3-demo", "application_uris": ["cf-s3-demo.cfapps.io"], "version": "09c9bfe9-c14e-4fcb-8ad8-9fcfd4b854893",
                  "name": "cf-s3-demo", "uris": ["cf-s3-demo.cfapps.io"], "users": null}
PATH=/bin:/usr/bin
PWD=/home/vcap
VCAP_SERVICES={"mongolab-n/a": [{"name": "mongo-cf-s3", "label": "mongolab-n/a",
                                    "tags": ["mongodb", "document"], "plan": "sandbox", "credentials": {"uri": "mongodb://user:****@ds053708.mongolab.com:53708/CloudFoundry_shageik2_iijfo9ve"}}, {"name": "s3-bucket-service", "label": "user-provided", "tags": [], "credentials": {"awsAccessKey": "*****", "awsSecretKey": "*****", "bucket": "cf-s3-demo"}, "syslog_drain_url": ""}]}
SHLVL=1
HOME=/home/vcap/app
PORT=61863
VCAP_APP_HOST=0.0.0.0
DATABASE_URL=
MEMORY_LIMIT=512m
_= /usr/bin/env
```



Cloud Foundry Environment

When am I done?



When "...the codebase could be made open source at any moment, without compromising any credentials."

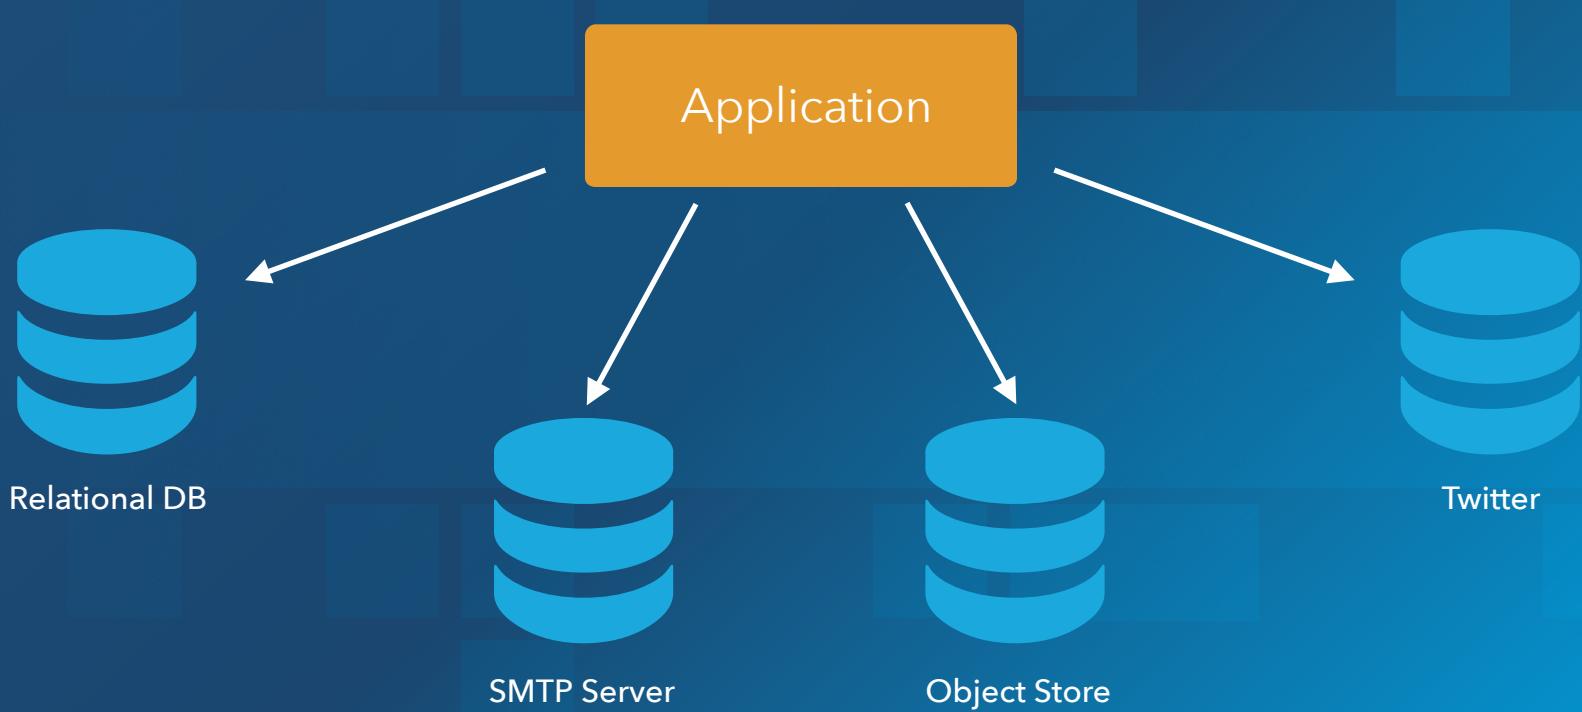
<http://12factor.net/config>

Why environment variables?



- Easy to change
- Little chance of being “checked in” to VCS
- Language/OS-agnostic standard

Backing Services



```
TMPDIR=/home/vcap/tmp
VCAP_APP_PORT=61863
USER=vcap
VCAP_APPLICATION={"instance_id": "b3e92a6fc443436888a525d100c91a12",
                  "instance_index": 0, "host": "0.0.0.0", "port": 61863, "started_at": "2013-12-04 01:52:01 +0000", "started_at_timestamp": 1386121921,
                  "start": "2013-12-04 01:52:01+0000", "state_timestamp": 1386121921,
                  "limits": {"mem": 512, "disk": 1024, "fds": 16384},
                  "application_version": "09c9bfe9-c14e-4fcb-8ad8-9fcfd4b854893",
                  "application_name": "cf-s3-demo", "application_uris": ["cf-s3-demo.cfapps.io"], "version": "09c9bfe9-c14e-4fcb-8ad8-9fcfd4b854893",
                  "name": "cf-s3-demo", "uris": ["cf-s3-demo.cfapps.io"], "users": null}
PATH=/bin:/usr/bin
PWD=/home/vcap
VCAP_SERVICES={"mongolab-n/a": [{"name": "mongo-cf-s3", "label": "mongolab-n/a",
                                    "tags": ["mongodb", "document"], "plan": "sandbox", "credentials": {"uri": "mongodb://user:****@ds053708.mongolab.com:53708/CloudFoundry_shageik2_iijfo9ve"}}, "user-provided": [{"name": "s3-bucket-service", "label": "user-provided", "tags": [], "credentials": {"awsAccessKey": "*****", "awsSecretKey": "*****", "bucket": "cf-s3-demo"}, "syslog_drain_url": ""}]}}
SHLVL=1
HOME=/home/vcap/app
PORT=61863
VCAP_APP_HOST=0.0.0.0
DATABASE_URL=
MEMORY_LIMIT=512m
_= /usr/bin/env
```



Cloud Foundry Environment



spring

DOCS

GUIDES

PROJECTS

BLOG

FORUM



PROJECTS

Spring Cloud Connectors

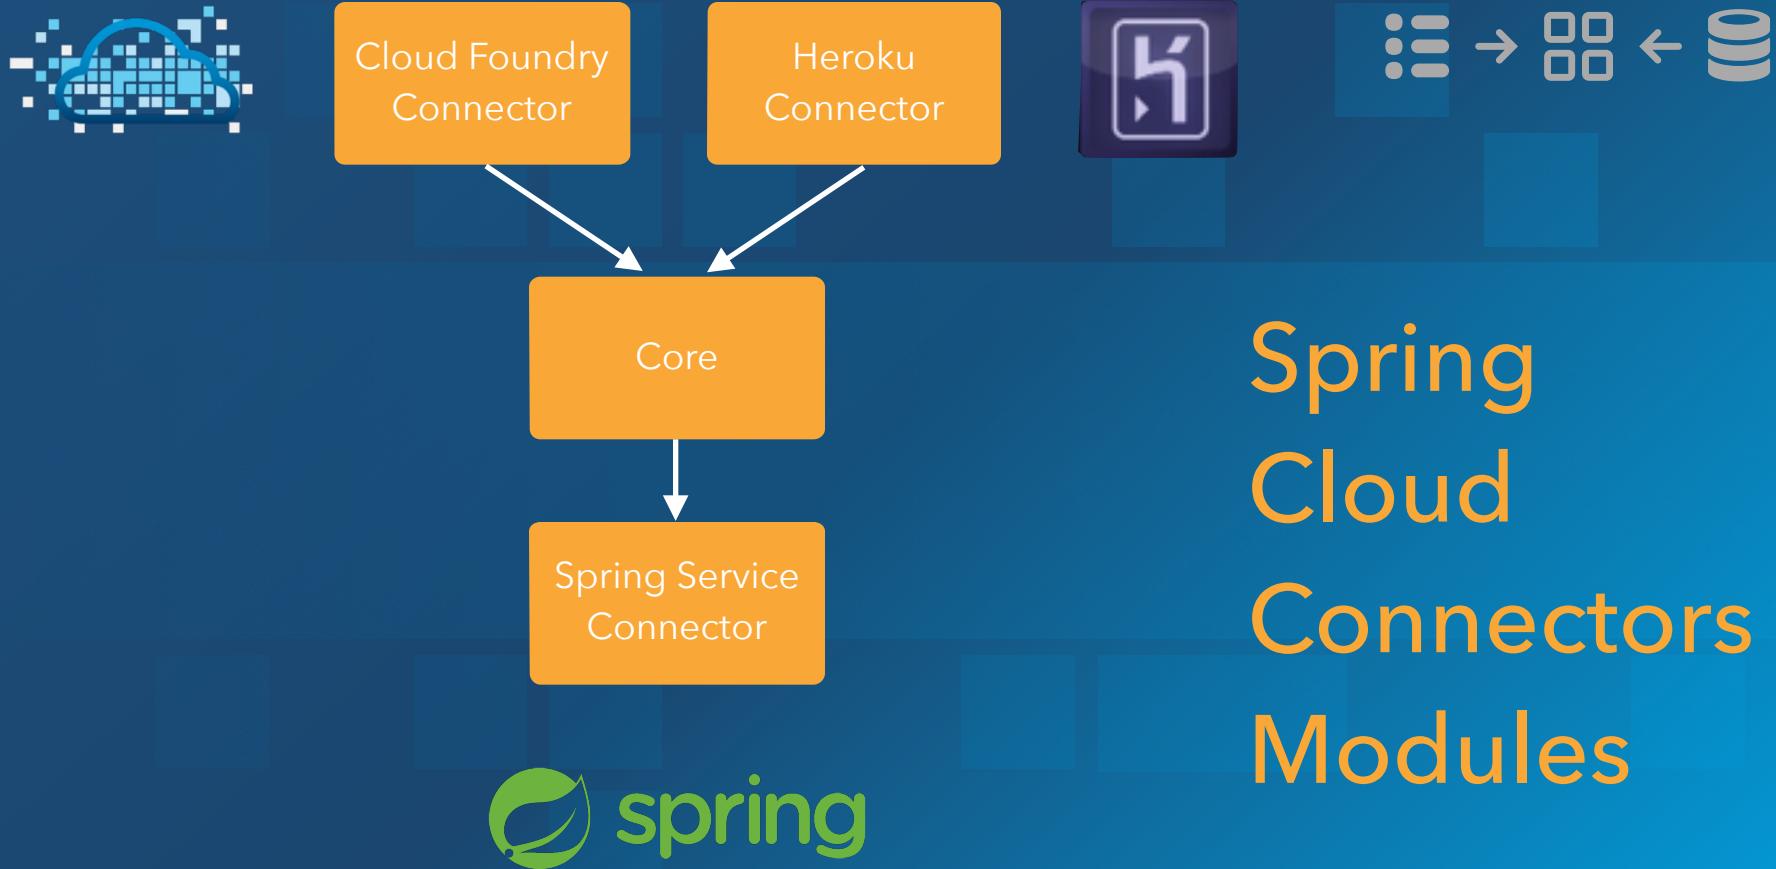


Spring Cloud Connectors simplifies connecting to services and gaining operating environment awareness in cloud platforms like Cloud Foundry and Heroku. Special support for Spring application through Java and XML config makes it trivial for apps to connect to cloud services. Designed for extensibility, you can use one of the existing cloud connectors (Cloud Foundry and Heroku) or write one for your cloud platform. While supporting commonly used services (relational databases, MongoDB, Redis, Rabbit) out of the box, it allows extending it to your own services. Neither of these require modifying Spring Cloud itself; all you need to do is add jars for your extensions to your classpath.

QUICK START

<http://cloud.spring.io/spring-cloud-connectors/>

Spring Cloud Connectors Modules



Spring Service Beans

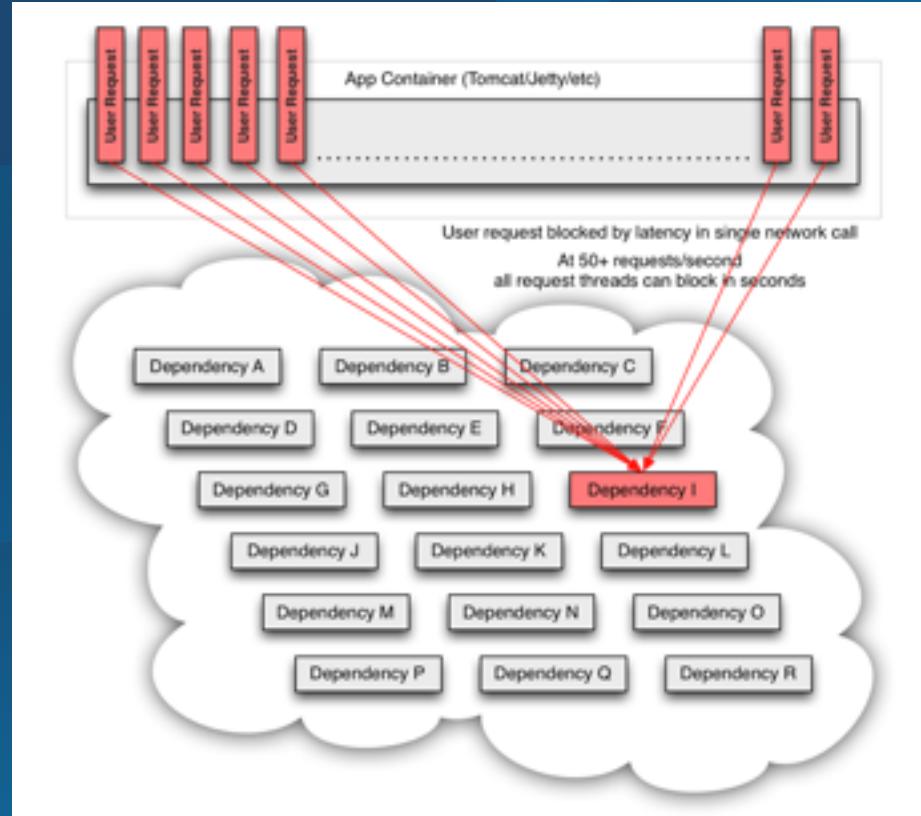
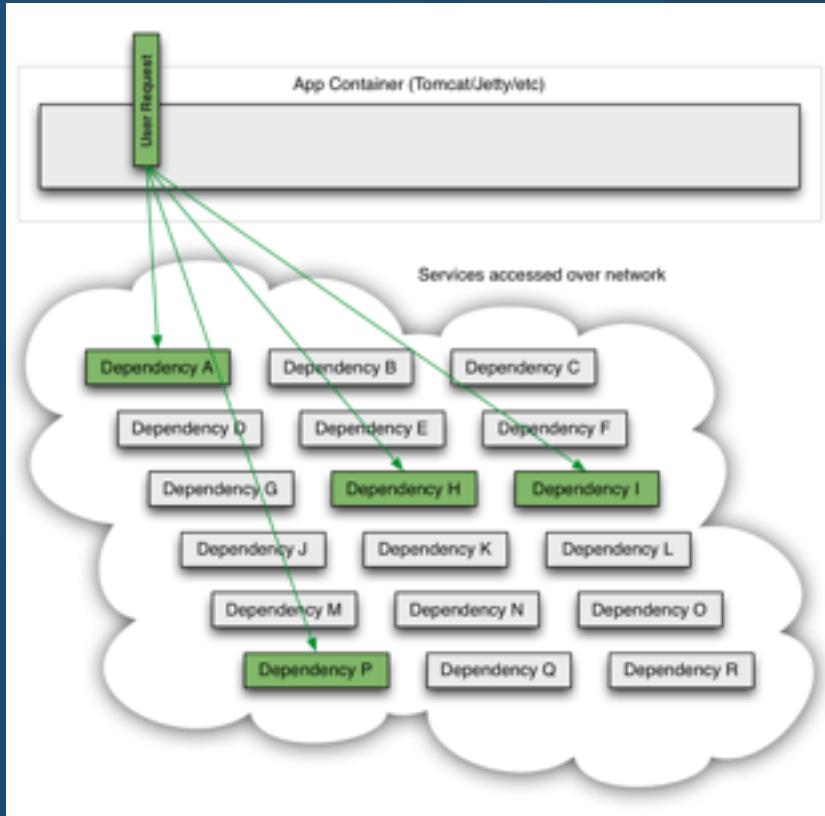


```
public class CloudConfig extends AbstractCloudConfig {  
    @Bean  
    public ConnectionFactory rabbitConnectionFactory() {  
        return connectionFactory.rabbitConnectionFactory();  
    }  
  
    @Bean  
    public DataSource dataSource() {  
        return connectionFactory().dataSource();  
    }  
  
    @Bean  
    public S3 s3() {  
        return connectionFactory().service(S3.class);  
    }  
}
```

Pattern: Fault Tolerance



Fault Tolerance at Netflix

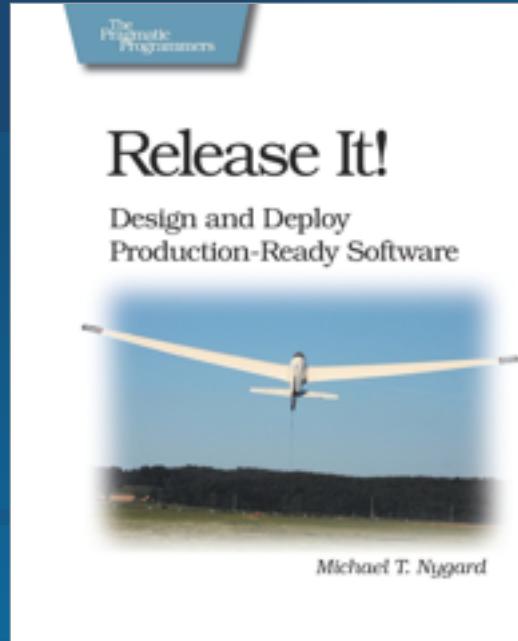




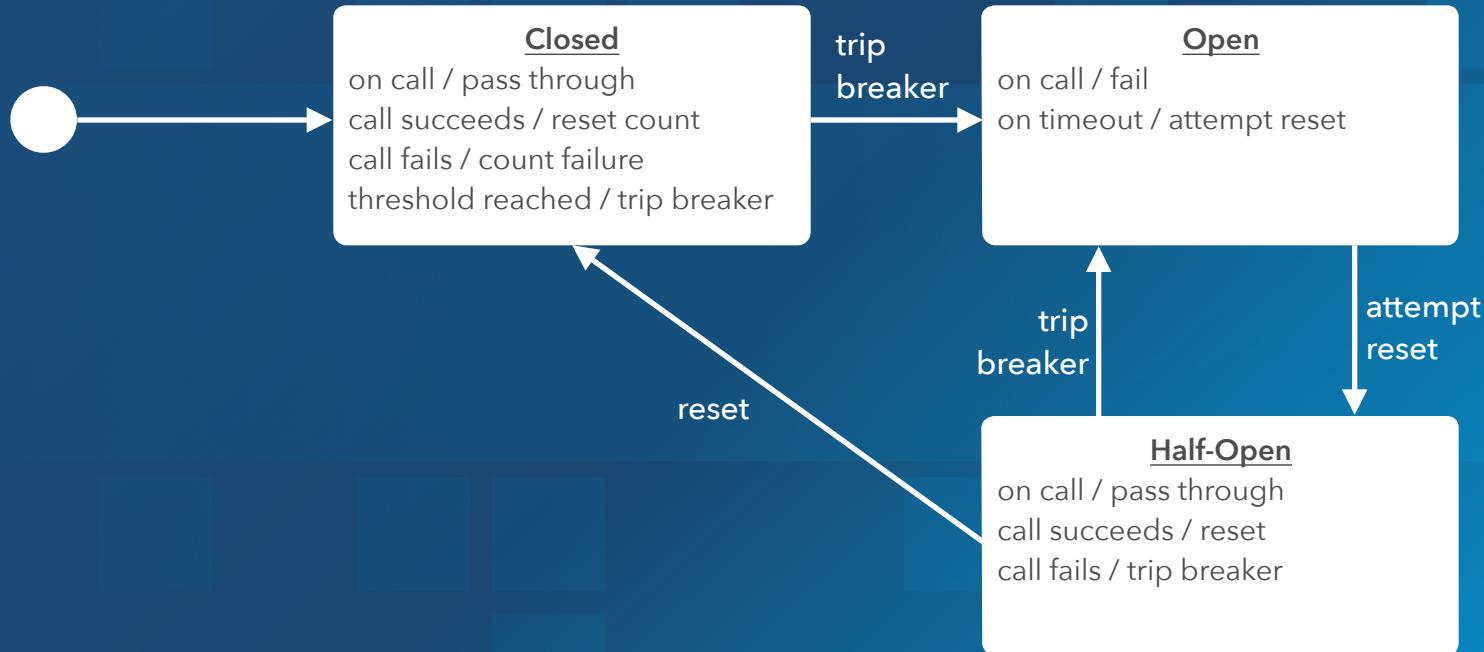
Without taking steps to ensure fault tolerance, 30 dependencies each with 99.99% uptime would result in 2+ hours downtime/month ($99.99\%^{30} = 99.7\%$ uptime = 2+ hours downtime in a month).

<http://techblog.netflix.com/2012/02/fault-tolerance-in-high-volume.html>

Circuit Breaker



Circuit Breaker State Machine





HYSTRIX

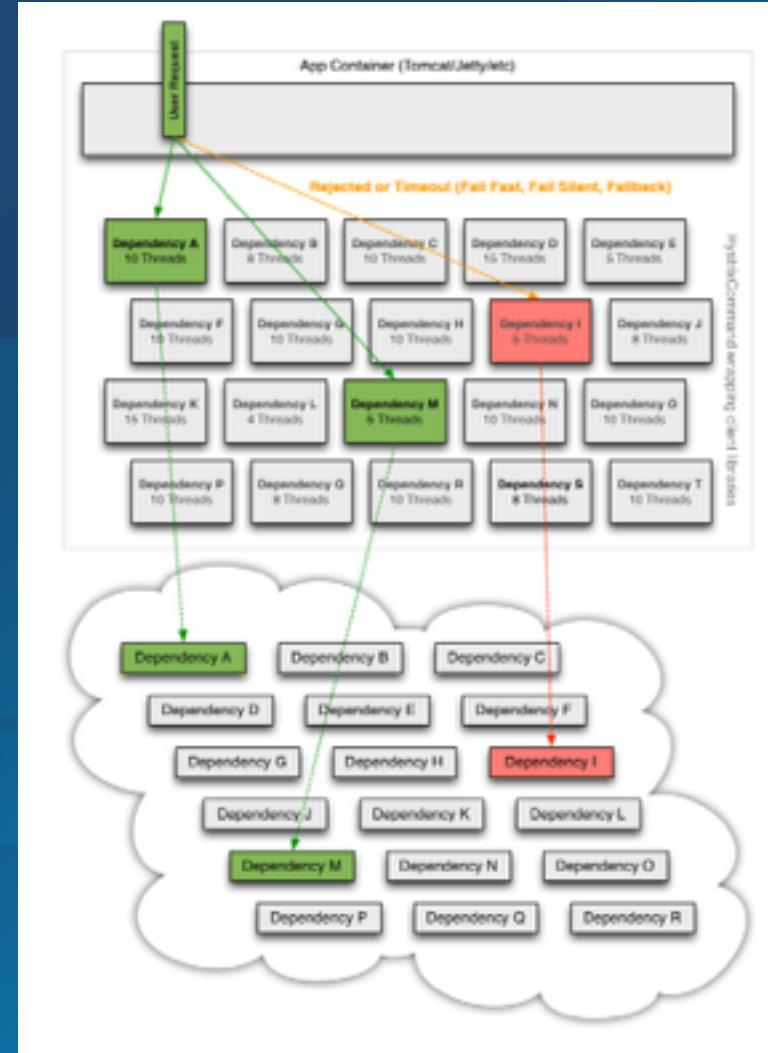
DEFEND YOUR APP

<https://github.com/Netflix/Hystrix>

Hystrix



- Latency and Fault Tolerance
- Realtime Operations
- Concurrency



Hello Hystrix!



```
public class CommandHelloWorld extends HystrixCommand<String> {  
  
    private final String name;  
  
    public CommandHelloWorld(String name) {  
        super(HystrixCommandGroupKey.Factory.asKey("ExampleGroup"));  
        this.name = name;  
    }  
  
    @Override  
    protected String run() {  
        return "Hello " + name + "!";  
    }  
}
```

Consuming Hystrix Commands



```
String s = new CommandHelloWorld("Bob").execute();  
  
Future<String> s = new CommandHelloWorld("Bob").queue();  
  
Observable<String> s = new CommandHelloWorld("Bob").observe();  
  
s.subscribe(new Action1<String>() {  
  
    @Override  
    public void call(String s) {  
        // value emitted here  
    }  
});
```

Fail Fast



```
public class CommandThatFailsFast extends HystrixCommand<String> {  
  
    private final boolean throwException;  
  
    public CommandThatFailsFast(boolean throwException) {  
        super(HystrixCommandGroupKey.Factory.asKey("ExampleGroup"));  
        this.throwException = throwException;  
    }  
  
    @Override  
    protected String run() {  
        if (throwException) {  
            throw new RuntimeException("failure from CommandThatFailsFast");  
        } else {  
            return "success";  
        }  
    }  
}
```

Fail Silent



```
public class CommandThatFailsSilently extends HystrixCommand<String> {  
  
    private final boolean throwException;  
  
    public CommandThatFailsSilently(boolean throwException) {  
        super(HystrixCommandGroupKey.Factory.asKey("ExampleGroup"));  
        this.throwException = throwException;  
    }  
  
    @Override  
    protected String run() {  
        if (throwException) {  
            throw new RuntimeException("failure from CommandThatFailsFast");  
        } else {  
            return "success";  
        }  
    }  
  
    @Override  
    protected String getFallback() {  
        return null;  
    }  
}
```

Static Fallback



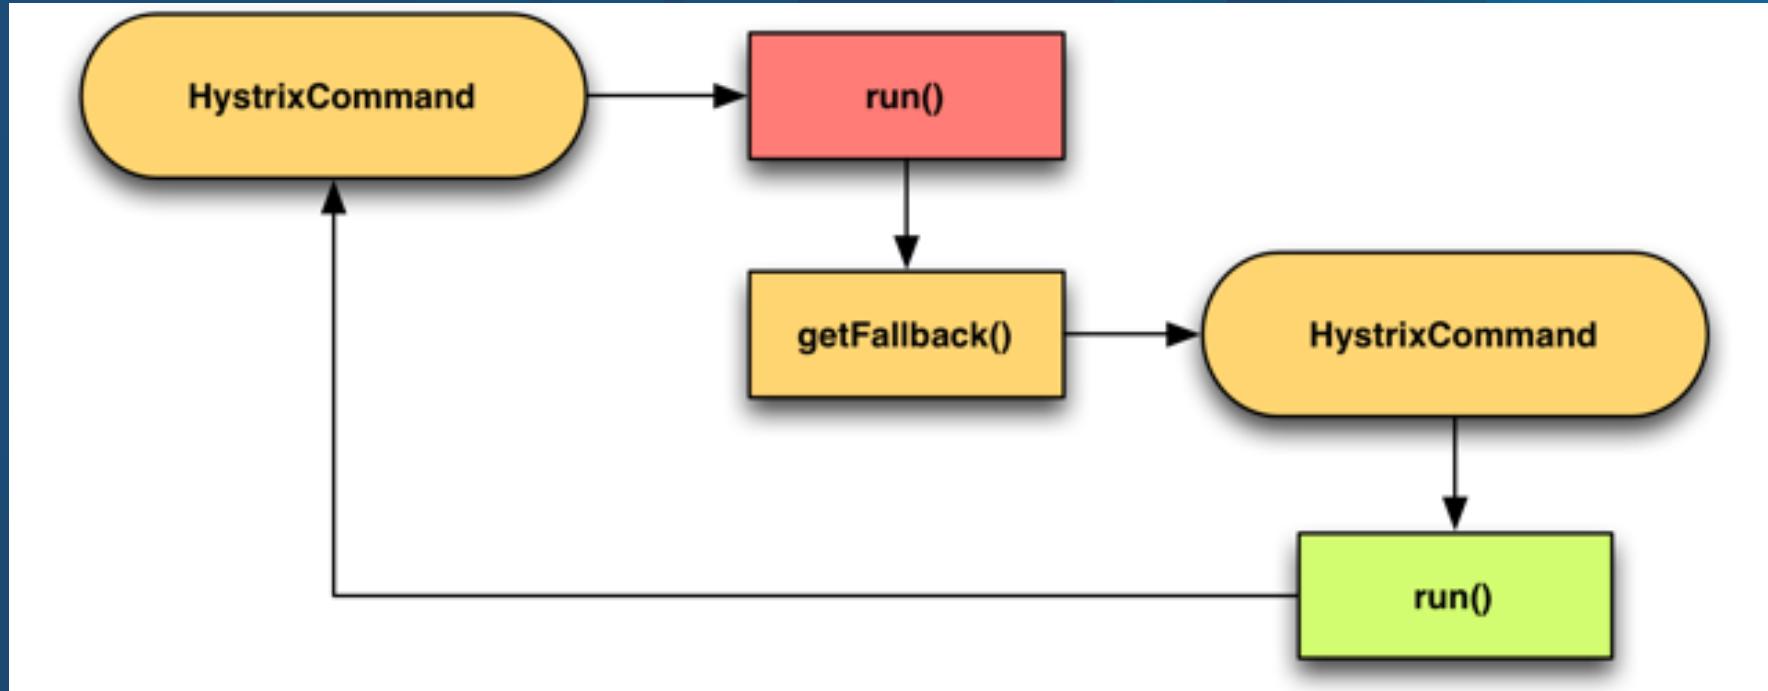
```
@Override  
protected Boolean getFallback() {  
    return true;  
}
```



Stubbed Fallback

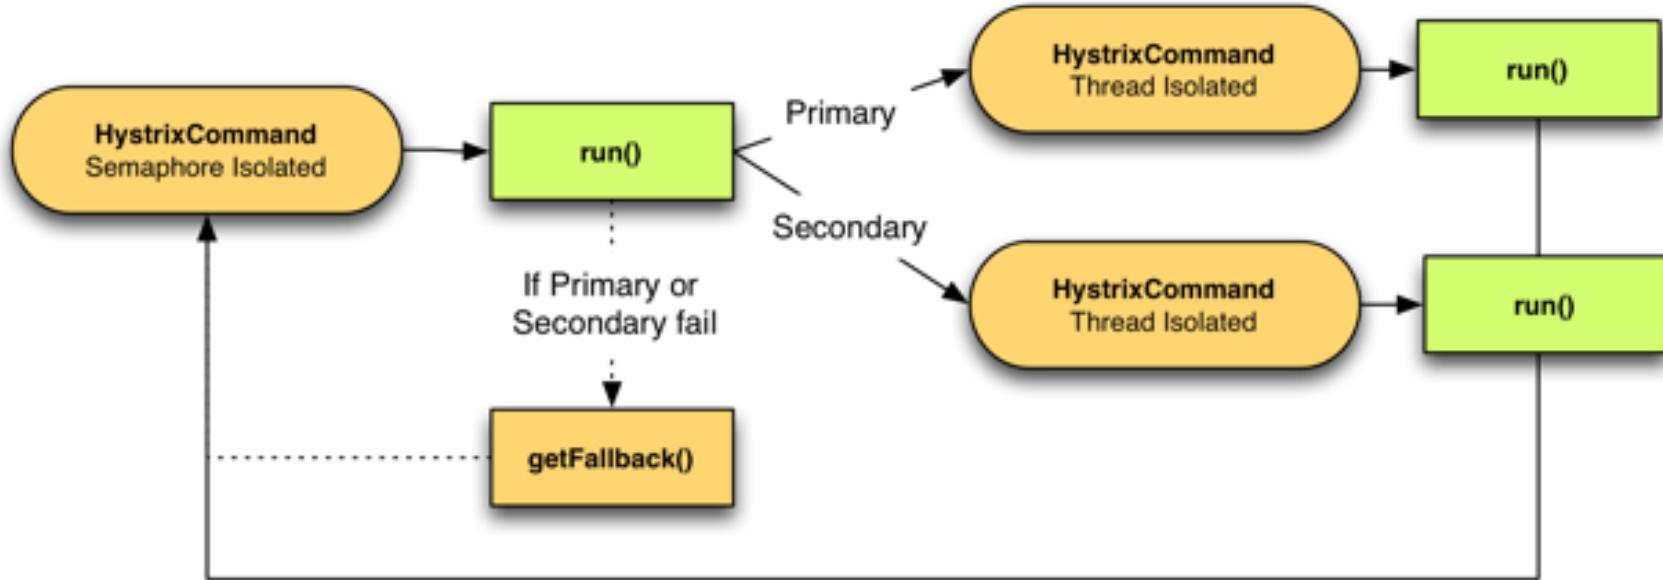
```
@Override  
protected UserAccount run() {  
    // fetch UserAccount from remote service  
    //           return UserAccountClient.getAccount(customerId);  
    throw new RuntimeException("forcing failure for example");  
}  
  
@Override  
protected UserAccount getFallback() {  
    /**  
     * Return stubbed fallback with some static defaults, placeholders,  
     * and an injected value 'countryCodeFromGeoLookup' that we'll use  
     * instead of what we would have retrieved from the remote service.  
     */  
    return new UserAccount(customerId, "Unknown Name",  
                           countryCodeFromGeoLookup, true, true, false);  
}
```

Fallback: Cache via Network





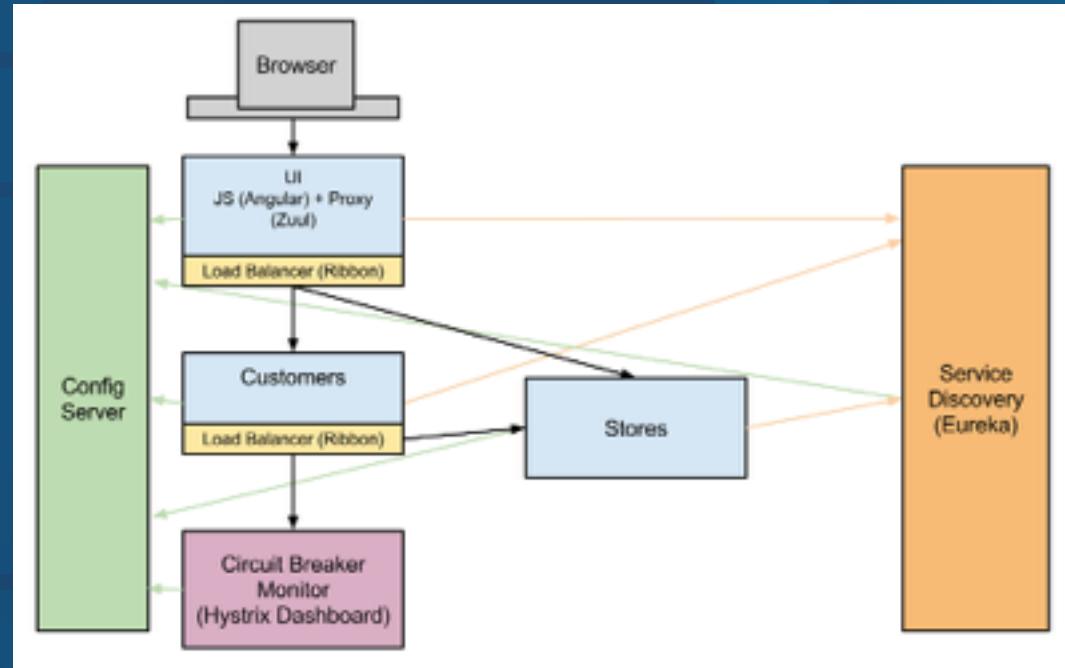
Primary + Secondary with Fallback



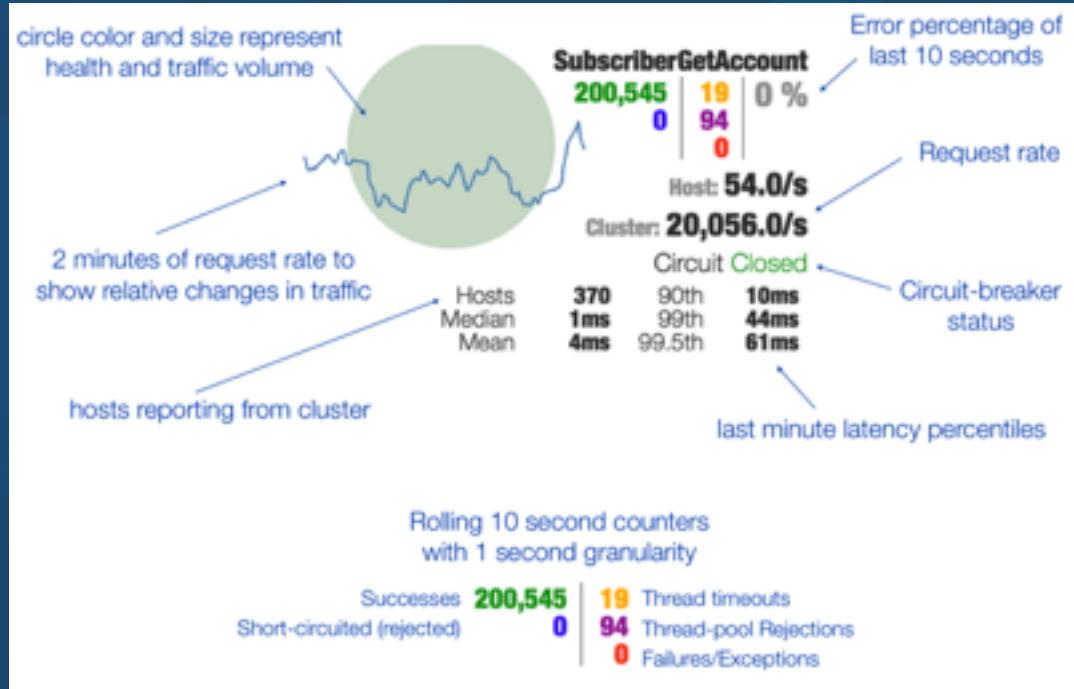
QUICK DEMO!



NETFLIX | OSS

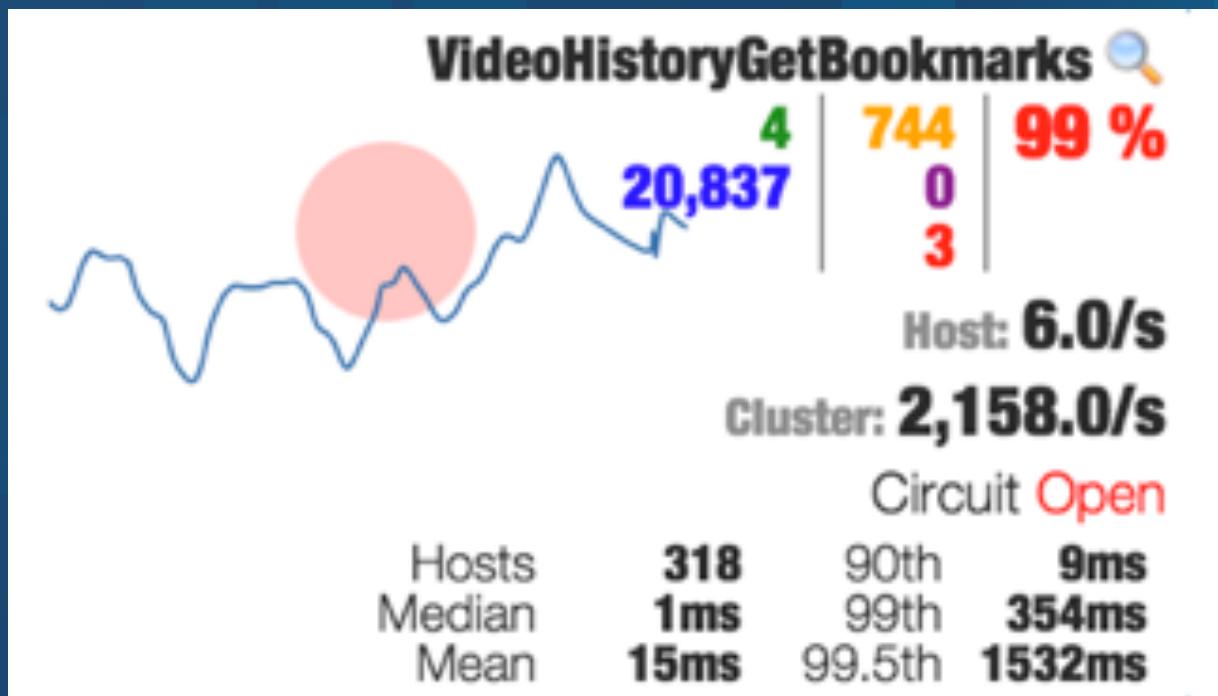


Hystrix Dashboard



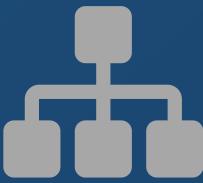


A Failing Circuit





Microservice



API Gateway



Stateless/
Shared-Nothing

Cloud Application Architecture Patterns

Thank You!!!

Matt Stine (@mstine)

Cloud Foundry Platform Engineer

matt.stine@gmail.com

<http://www.mattstine.com>



Configuration/
Service Consumption



Fault
Tolerance