



# Parallelizing CI using Docker Swarm-Mode


Akihiro Suda <suda.akihiro@lab.ntt.co.jp>  
NTT Software Innovation Center

# Who am I



<https://github.com/AkihiroSuda>

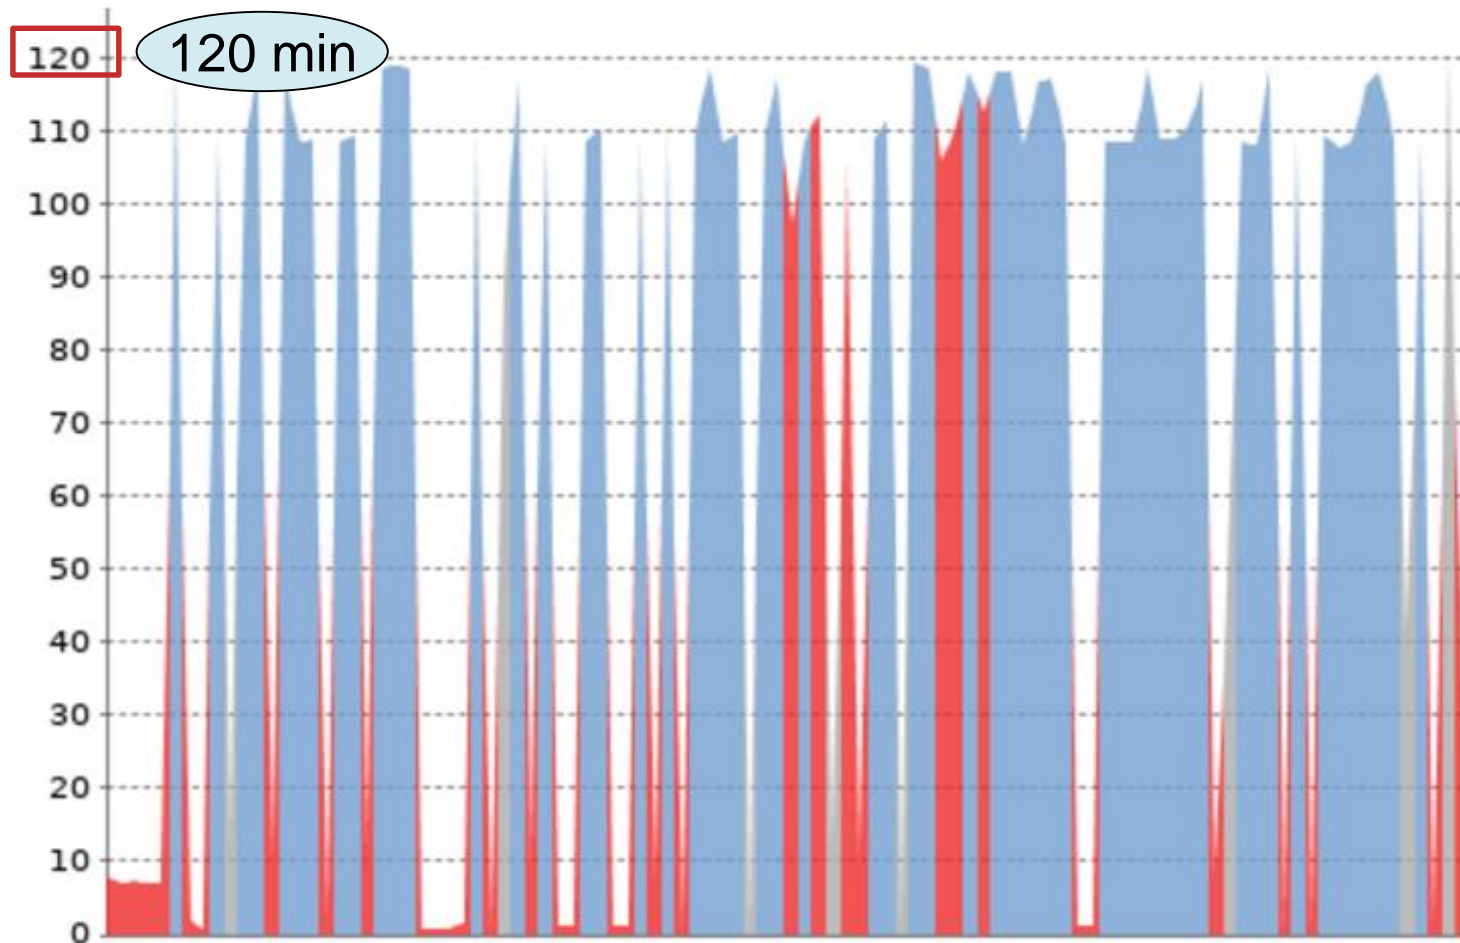
- **Software Engineer at NTT Corporation**
- **Docker core maintainer**
- **Previous talks at FLOSS community**
  - FOSDEM 2016
  - ApacheCon Core North America 2016
  - ...



# A problem in Docker project: CI is slow

capture: March 3, 2017





<https://jenkins.dockerproject.org/job/Docker-PRs/buildTimeTrend>



# How about other FLOSS projects?




capture: March 3, 2017



# Why slow CI matters?

- **Blocker for reviewing/merging patches**
- **Discourages developers from writing tests**
- **Discourages developers from enabling additional testing features (e.g. race detector)**




Poor implementation quality &  
Slow development cycle

# Solution: Parallelize CI ?

e.g.

- ``go test -parallel N``
- ``mvn --threads N``
- ``parallel``
- ...



But parallelization is not enough ☹

## • No isolation


- Concurrent test tasks may race for certain shared resources (e.g. files under ``/tmp``, TCP port, ...)

## • Poor scalability

- CPU/RAM resource limitation
- I/O parallelism limitation

# Solution: Parallelize CI across Docker Swarm-mode

- Docker provides isolation
- Docker Swarm-mode provides scalability, plus allows you to manage multiple Docker nodes as if they were a single node



# Chunking

- For ideal isolation, each of the test functions should be encapsulated into independent containers
- But this is not optimal in actual due to setup/teardown code

```
func TestMain(m *testing.M) {
```

```
    setUp()
```

```
    m.Run()
```

```
    tearDown()
```

```
}
```

```
func TestFoo(t *testing.T)
```

```
func TestBar(t *testing.T)
```

redundantly executed

container

```
setUp()
```

```
testFoo.Run()
```

```
tearDown()
```

container

```
setUp()
```


```
testBar.Run()
```

```
tearDown()
```



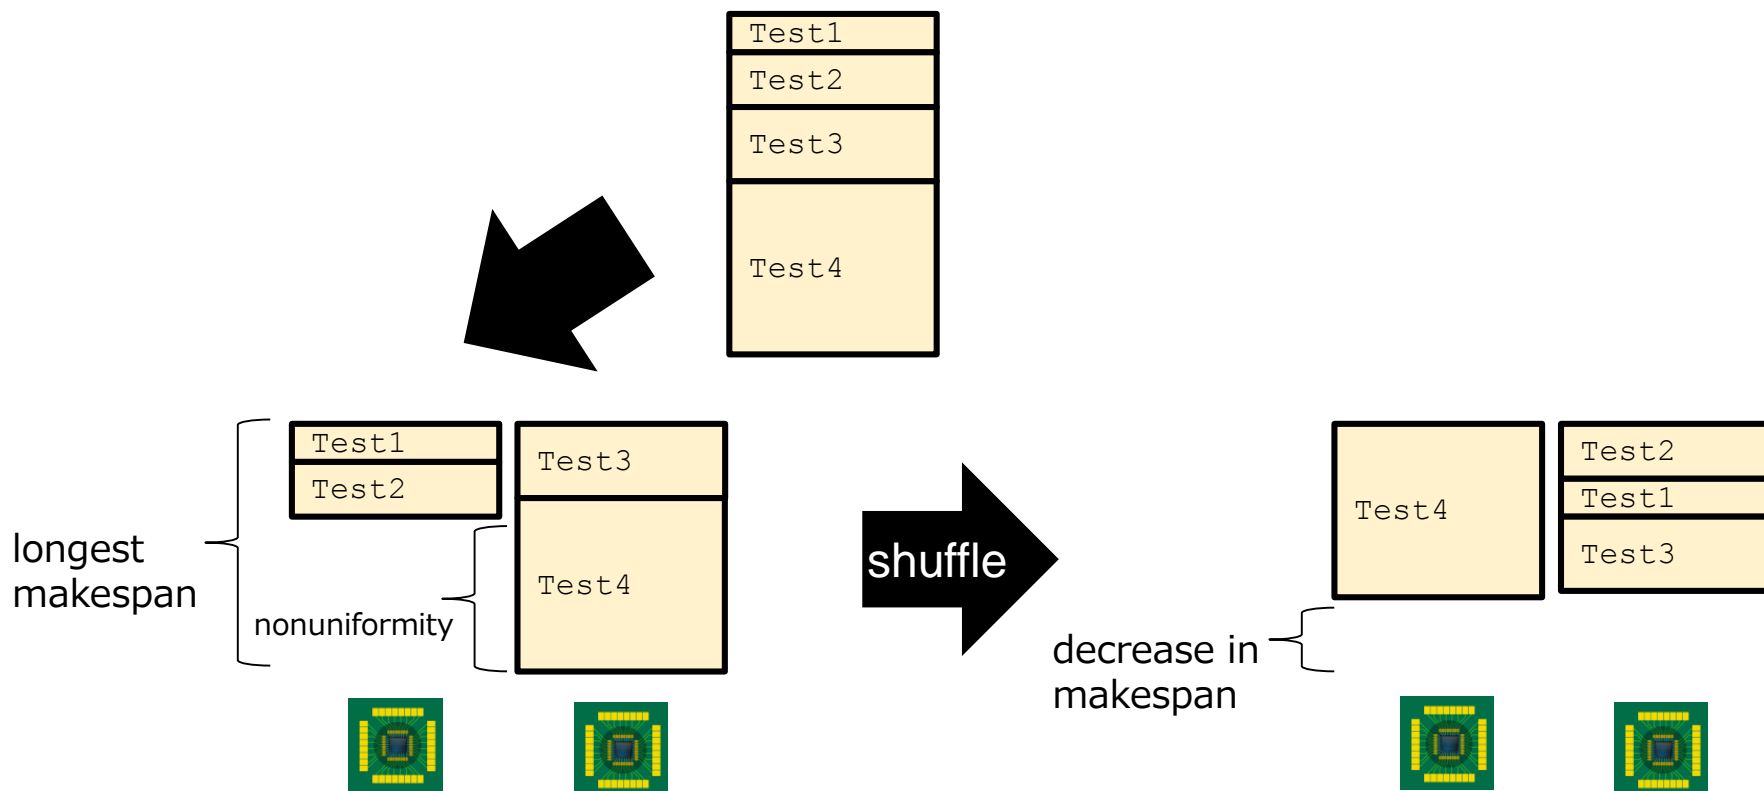
# Chunking

- So we execute a chunk of multiple test functions sequentially in a single container



# Shuffling

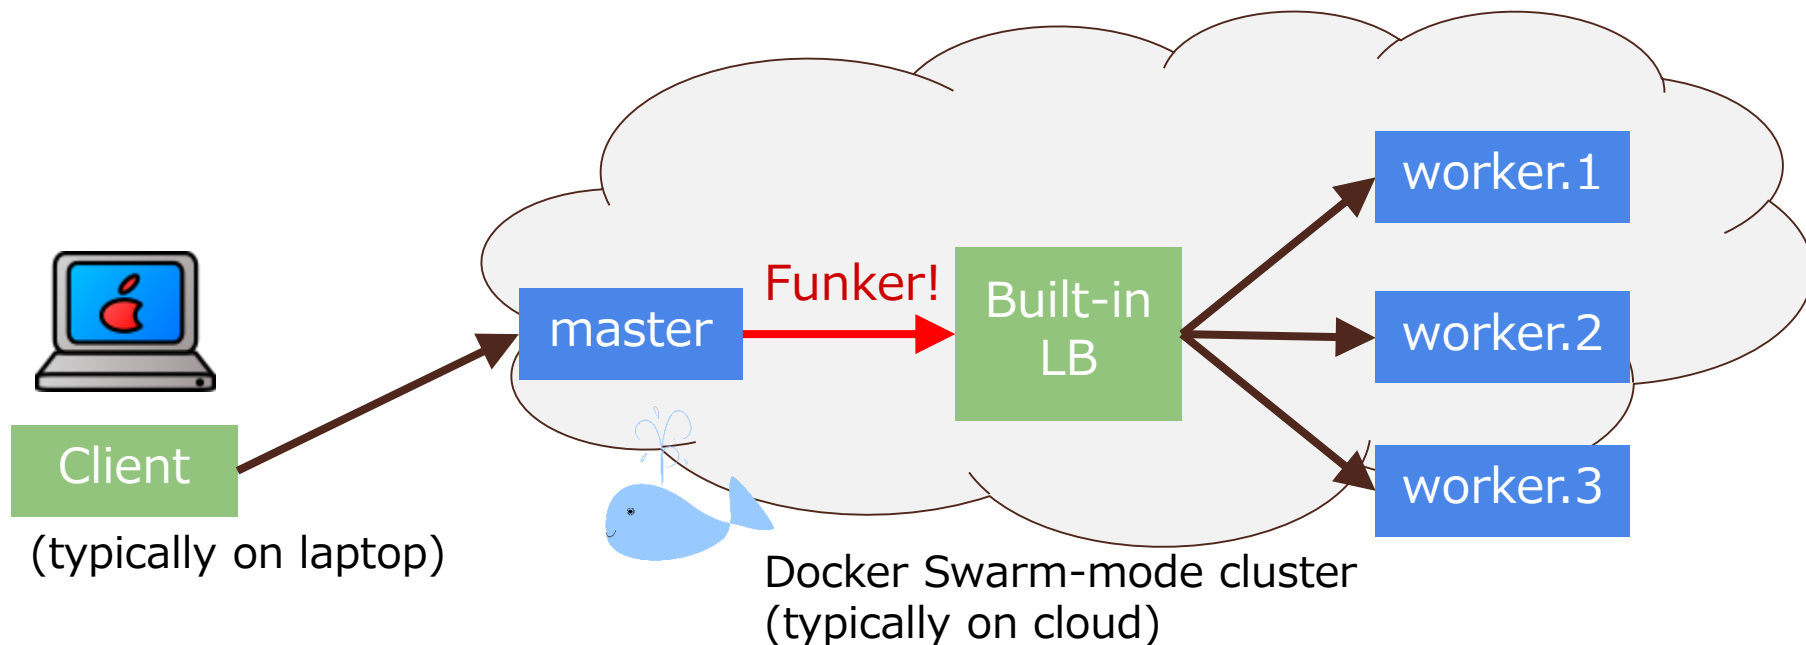
- makespans of the chunks are not uniform
- So we shuffle the chunks for "hope" of optimal scheduling
  - No guarantee for optimal schedule



# Implementation

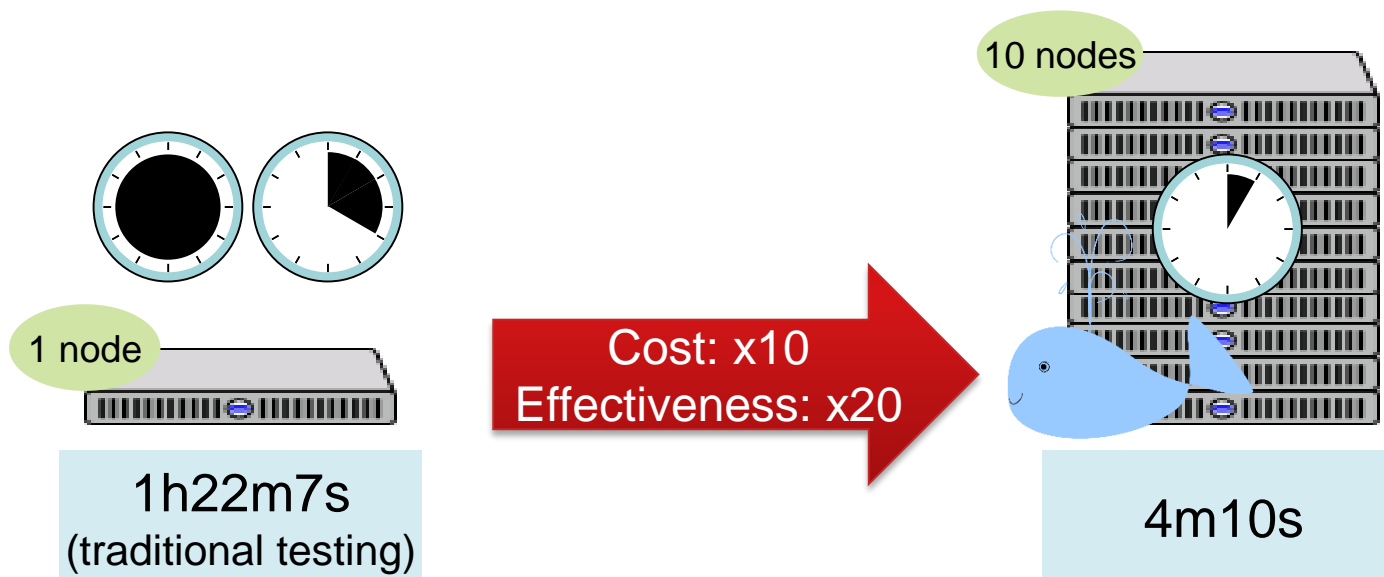
- **We use Funker ([github.com/bfirsh/funker](https://github.com/bfirsh/funker))**

- FaaS-like architecture
- Loads are automatically balanced via Docker's built-in LB
- No explicit task queue
- No `docker exec` hack
- Could be easily portable to other container orchestrators



# Experimental result

- **Evaluated my hack against the CI of Docker itself**
  - Of course, this hack can be applicable to CI of other software as well
- **Target: Docker 16.04-dev (git:7fb83eb7)**
  - Contains 1,648 test functions
- **Machine: Google Computing Engine  
n1-standard-4 instances (4 vCPUS, 15GB RAM)**



# Detailed result

Number of containers running in parallel  
(chunk size is inversely proportional to this)

Nodes	10	30	50	70
<b>1</b>	<b>15m3s</b>	N/A (more than 30m)		
<b>2</b>	12m7s	<b>10m12s</b>	11m25s	13m57s
<b>5</b>	10m16s	6m18s	<b>5m46s</b>	6m28s
<b>10</b>	8m26s	4m31s	<b>4m10s</b>	4m20s

Even with a single node,  
the result is significantly better than  
the traditional testing (1h22m7s)

Fastest  
configuration

# The code is available!



PR (merged):

<https://github.com/docker/docker/pull/29775>

```
$ cd $GOPATH/src/github.com/docker/docker
$ make build-integration-cli-on-swarm
$ ./hack/integration-cli-on-swarm/integration-cli-on-swarm \
  -replicas 50 \
  -shuffle \
  -push-worker-image your-docker-registry.example.com/worker
```

# Is it applicable to other software as well?



- **Yes, with just a few line of modifications**
- **Planning to split this hack into an independent generic test tool**

# Future work

- **"Shuffling" does not always result in optimal schedule**
- **Future work: learn the past execution history to optimize the schedule**



# Recap

- Docker Swarm-mode is effective for parallelizing CI jobs
- Some hacks for optimizing schedule

