

ORACLE®



Program

- Vocabulary
- What Is Low-Latency Java
- JVM Features
- What To Think About
- Application Analysis



Vocabulary

- Nursery == Young Space == Eden
- Nursery Collection == Young Collection == Minor GC
- Promotion
 - Move objects to Old Space
-
- Old Space == Old Generation
- Old Collection == Full Collection
- Compaction
 - Move objects together to remove fragmentation

JVM Heap Layout

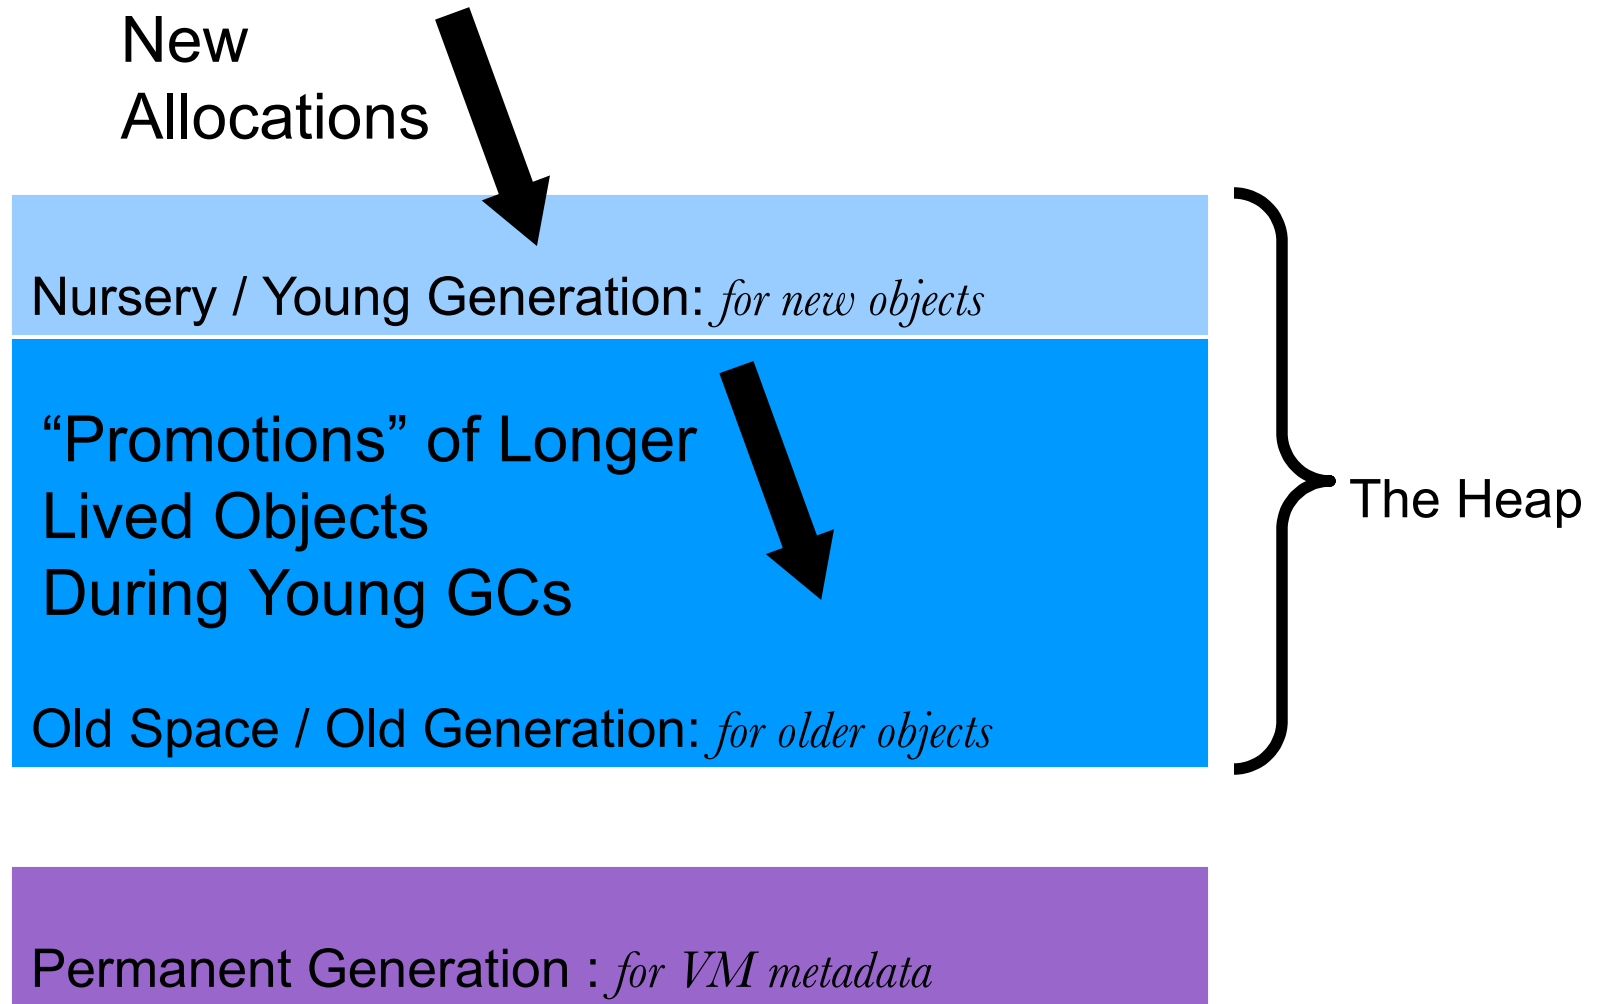
Nursery / Young Generation: *for new objects*

Old Space / Old Generation: *for older objects*

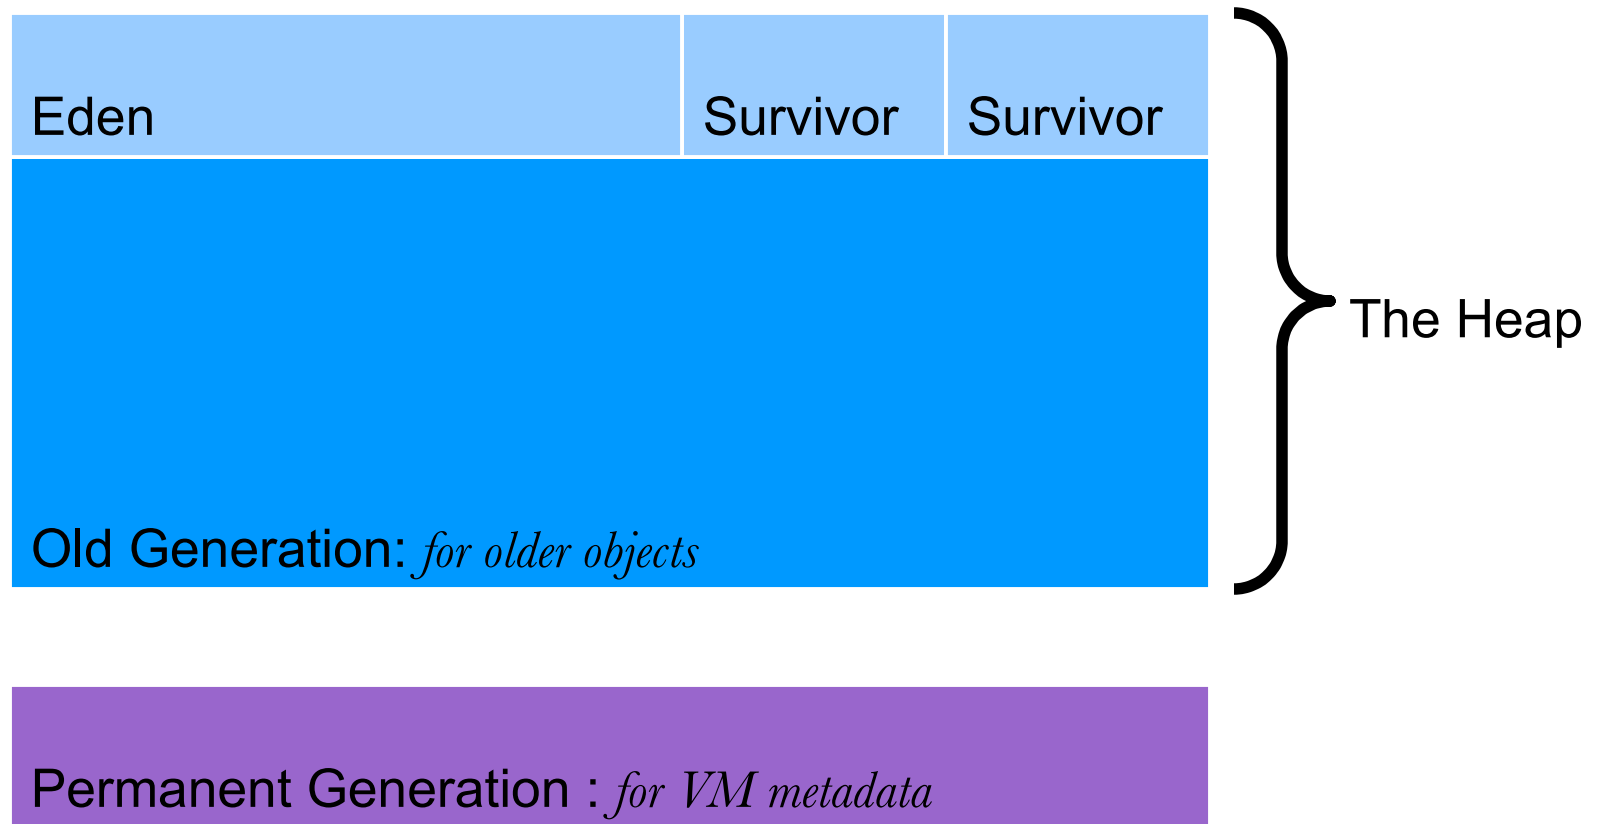
Permanent Generation : *for VM metadata*

The Heap

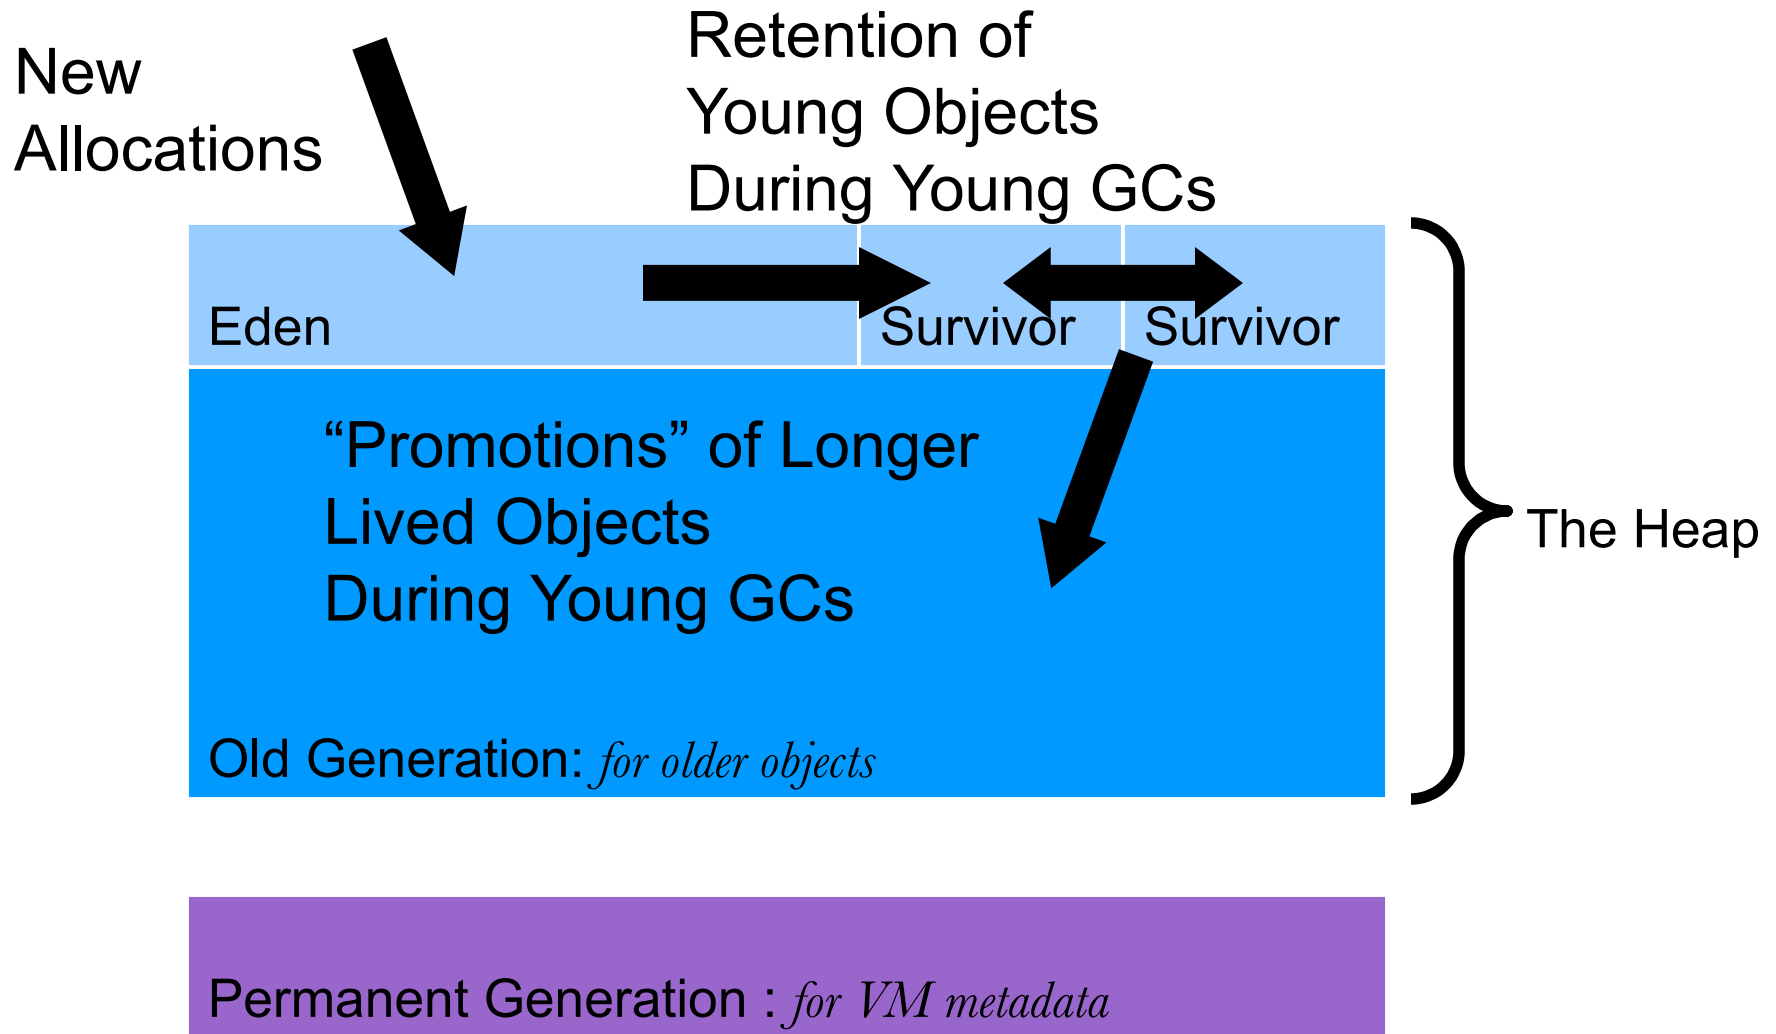
JVM Heap Layout



HotSpot JVM Heap – In More Detail...



HotSpot JVM Heap – In More Detail...




Low-Latency Java

What is Low-Latency Java

- Soft Real-Time
 - High throughput with low latency
 - No hard response time guarantees
 - No catastrophe happens at response time failures
 - But several nines response time guarantees
- Normal Java Code
 - No special APIs
 - Code tuning with focus on latency


What is Low-Latency Java

- Run standard Java code with great response times



Normal Garbage Collector

GC spikes and occasional SLA breach occurs.



Low-Latency Collector

Deterministic GC pauses, allowing guarantees of SLAs.

What is Low-Latency Java

- Response Time Requirements
 - Typical latency requirement around 5 – 50 ms
 - Max latency = transaction time + max pause time
- Typical Applications
 - Financial
 - Automatic trading
 - Matching server
 - Telecom
 - SIP applications
 - Event processing
 - RFID scanning

JVM Features

What Do the JVMs Offer

- GC Implementations
 - Generational Concurrent Collectors
 - CMS – HotSpot
 - Gencon – JRockit
 - Regional Heap Collector
 - G1 – HotSpot (In development, EA available)
 - Optimized Concurrent Collector
 - Deterministic GC – JRockit Real-Time

Generational Concurrent Collectors

- Use Case
 - Normal heap sizes
 - Reasonable amount of live data
- Nursery Collector
 - Stop the world
 - Parallel
- Old Collection
 - Mostly concurrent Mark and Sweep
 - Some short parallel pauses
 - Partial compaction

Generational Concurrent Collectors

- Tuning
 - Nursery sizing
 - Large enough to reduce frequency
 - Small enough to get adequate pauses
 - Compaction tuning
 - How large part of the heap to compact each GC
 - Large enough to avoid fragmentation
 - Small enough to get adequate pauses

Generational Concurrent Collectors

- Tuning Cont.
 - Concurrent GC threshold
 - Start concurrent GC early enough
 - Complete GC before memory become scarce
 - Tuned automatically

G1 – Regional Heap Collector

- Use Case
 - Large heaps
 - Large amount of live data
- Regional Nursery Collector
 - Stop the world
 - Parallel
- Regional Old Collector
 - Evacuation of live data
 - Remembered sets
 - No mark phase

G1 – Partial Heap Collector

- Tuning
 - Pause target
 - Lower pause target requires a more GC aware code
 - `-XX:MaxGCPauseMillis=X`
 - Nursery sizing
 - Implicit by pause target tuning
 - Evacuation tuning
 - Implicit by pause target tuning

Deterministic GC – Optimized Concurrent Collector

- Use Case
 - Normal heap sizes
 - Live data around 1/2 of the heap
- Old Collection
 - Parallel mostly concurrent Mark and Sweep
 - Heavily optimized pauses
 - Abortable compaction

Deterministic GC – Optimized Concurrent Collector

- Tuning
 - Pause target
 - Lower pause target requires a more GC aware code
 - -XXpausetarget=Xms
 - Heap sizing
 - Large enough to hold live data and free heap
 - Small enough to get adequate pauses
 - Compaction tuning
 - Implicit by pause target tuning
 - Ensure enough pause time to keep heap unfragmented

What To Think About

What To Think About – Throughput

- Potential Performance Impact
 - Bookkeeping of objects
 - Tracking new objects during a concurrent collection
 - Updating remembered sets
- Don't saturate the CPU
 - The concurrent GC threads shares the CPU

What To Think About – Allocation

- Allocation Rate
 - More allocation => More GCs => More pauses
 - Need enough time to finish concurrent phases
 - Complete GC before memory becomes scarce
 - Can still handle high allocation rates
 - Hundreds of MB per second
- Large Objects/Arrays
 - Requires free consecutive memory
 - Increases the requirement to do compaction

What To Think About – Allocation

- Avoid System.gc()
 - Let the memory system handle GCs
- Semi-Long Lived Objects
 - Increases YC times due to copying during promotion
 - Increases fragmentation in Old Space
 - Mix of short and long lived data
 - Avoid storing data from each transaction
 - Make sure tenuring thresholds is configured properly

What To Think About – Data Structures

- Understand Your Data Structures
 - Resizing Data Structures
 - Size your HashMap correctly
 - Rehashing when increasing size takes time
 - Hard to detect
 - StringBuilder/Buffer expands and copying
 - We do optimizations to try to avoid this
 - Optimal Use Case
 - ArrayList add/remove at beginning of list causes copying

What To Think About – Data Structures

- GC Friendly Data Structures
 - Mark phase iterates over live data
 - Some data structures are hard to mark in parallel

What To Think About – Reference Objects

- Use Reference Objects Sensibly
 - Soft, weak and phantom references require special processing during GC
 - Finalizers should be avoided
 - Will keep objects alive for one extra GC cycle
 - Detrimental to pause time goals

What To Think About – Profiling

- Externally Measure End-to-End Response Time
 - Internal measurements affected in the same way as the application
- GC Analysis
 - Understand your object allocation and survival patterns
 - GC pause times
 - The name of the pause is often a good hint
- Lock Profiling
 - #1 issue for scaling problems and unexplained outliers

What To Think About – Finally

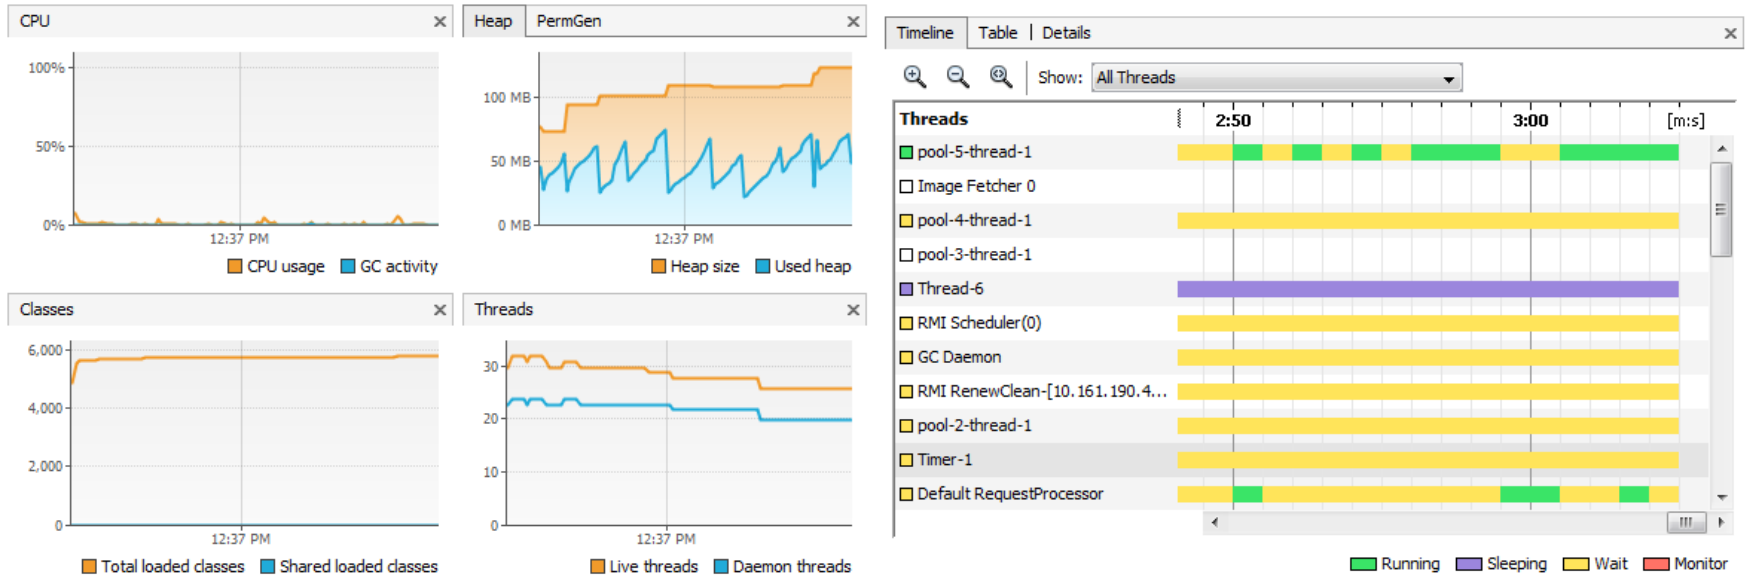
- **Don't Overdo It**
 - The JVM can handle a lot
 - Analyze to detect your current bottleneck
 - Optimize and tune

Application Analysis

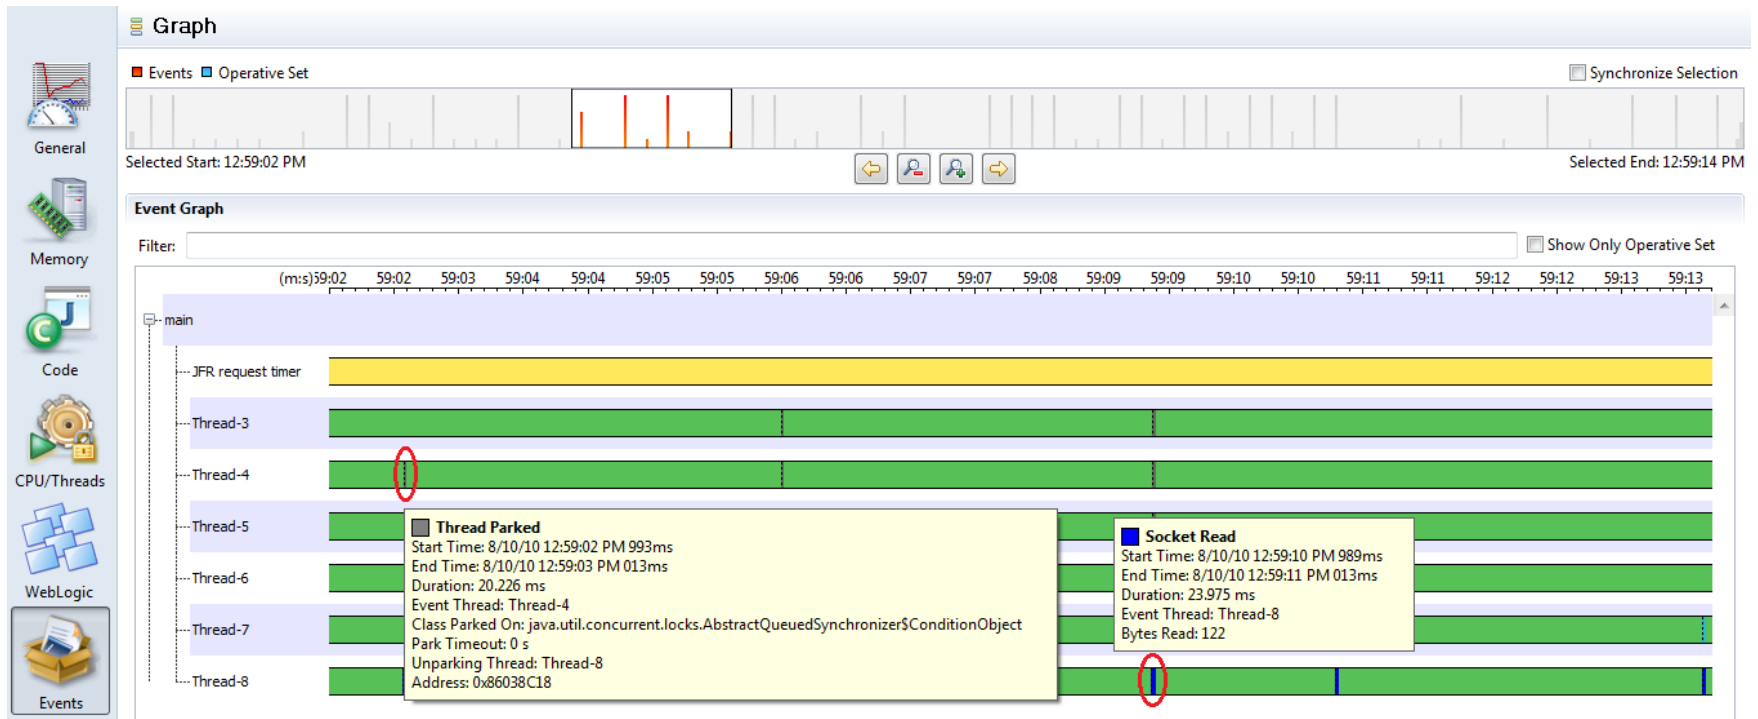
Know What to Optimize

- Analyze Your Application
 - Garbage collections
 - Pauses
 - Heap usage
 - Hot Methods
 - Lock contention
 - IO events

HotSpot VisualVM



JRockit Mission Control



Other Tools

- Java Tools
 - IBM Health Center
 - JProfiler
- Low Level Tools – Hardware Profiling
 - VTune
 - CodeAnalyst
 - Oracle Solaris Studio
 - oProfile



JavaOne and Oracle Develop

Russia 2011

April 9-10, 2011

SOFTWARE. HARDWARE. COMPLETE.