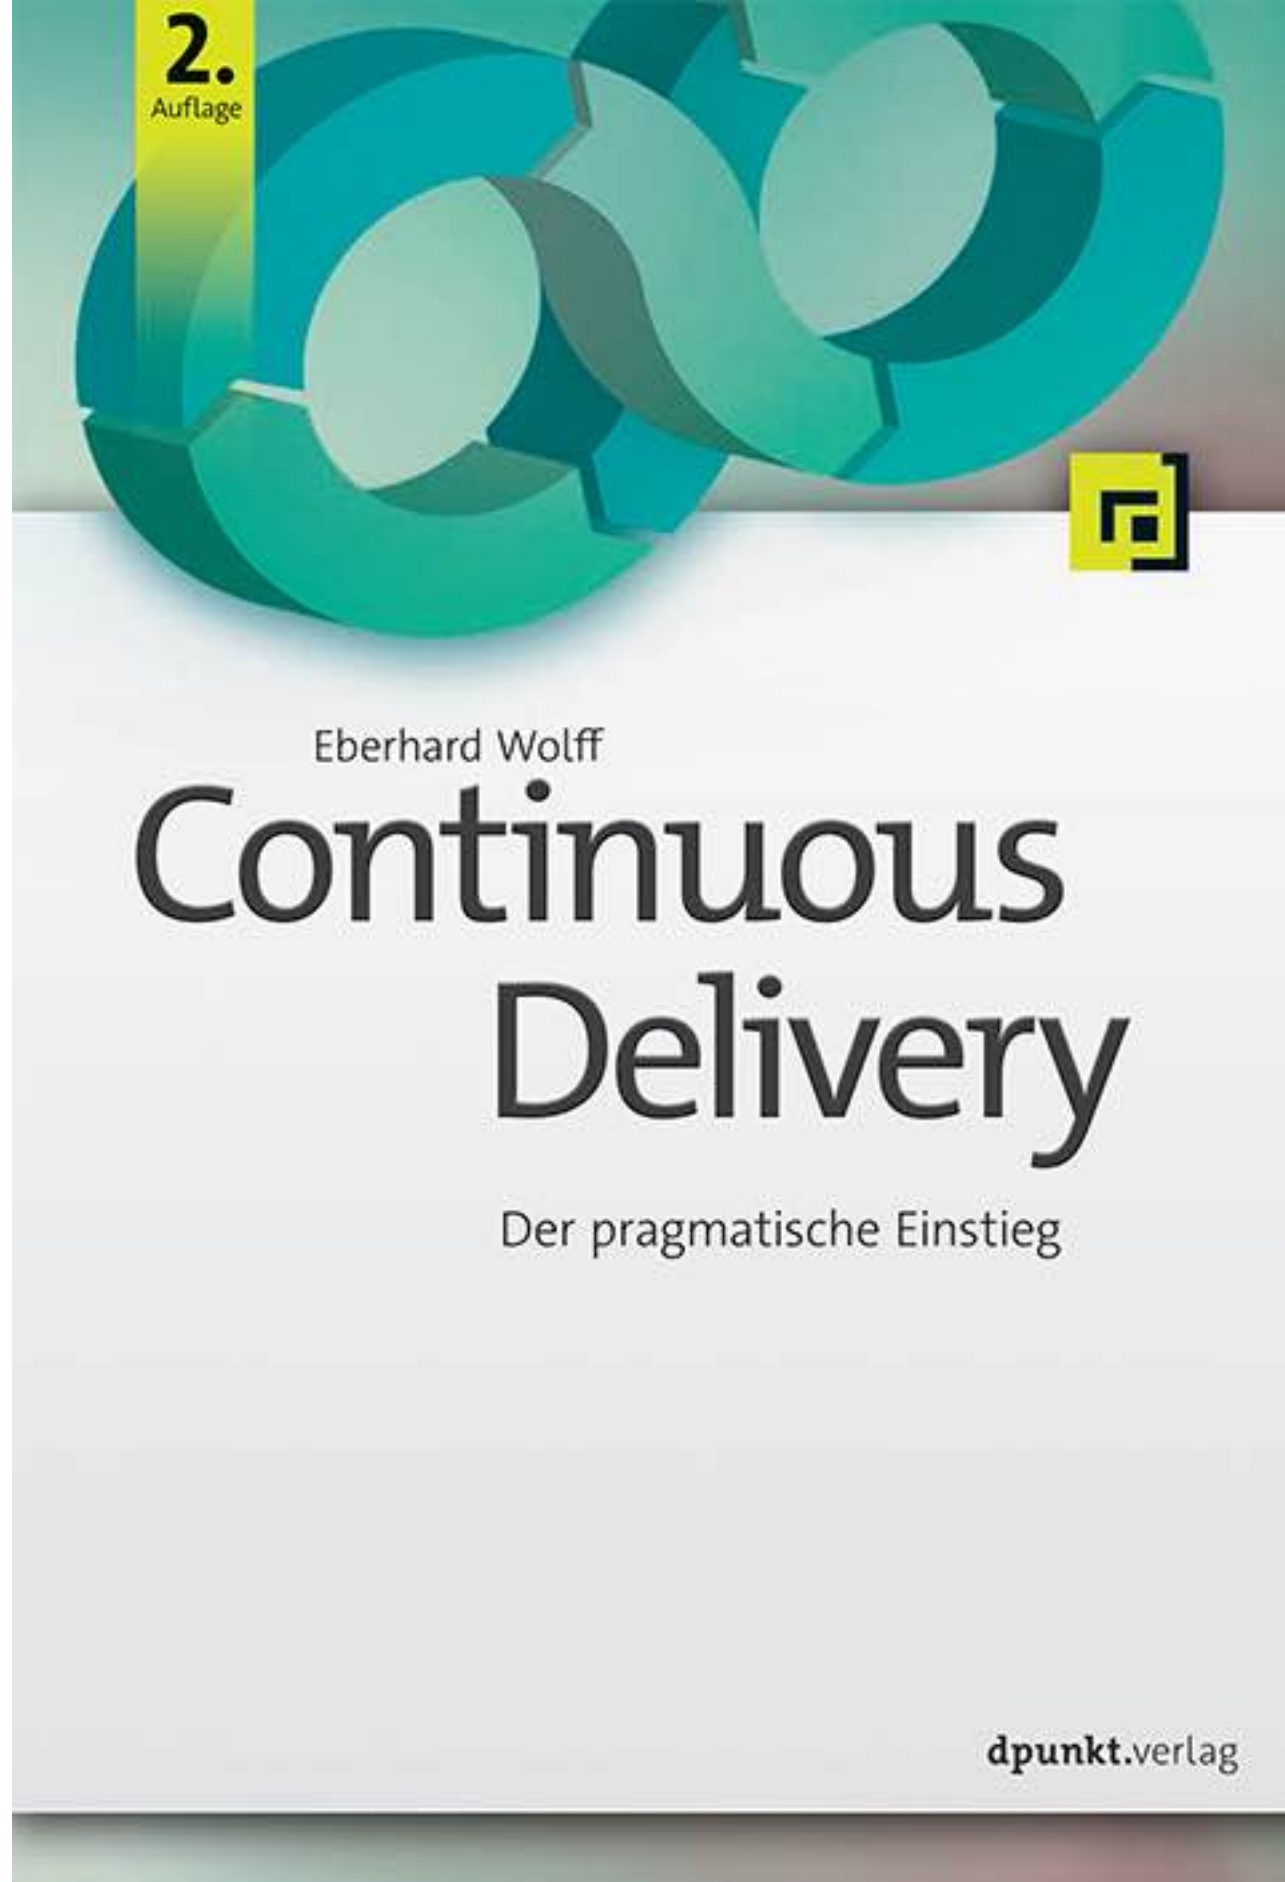


Data Architecture – Not Just for Microservices

Eberhard Wolff
@ewolff
Fellow





<http://continuous-delivery-buch.de/>



Eberhard Wolff

Microservices

Grundlagen flexibler Softwarearchitekturen

dpunkt.verlag

<http://microservices-buch.de/>



Microservices

FLEXIBLE SOFTWARE ARCHITECTURE

EBERHARD WOLFF

<http://microservices-book.com/>



Eberhard Wolff

Microservices Primer

A Short Overview

FREE!!!!

innoQ

<http://microservices-book.com/primer.html>

Volume
9

entwickler**spezial**

MICROSERVICES

DIE ANDERE ART DER MODULARISIERUNG

ARCHITEKTUR

Die Vor- und Nachteile
von Microservices

TECHNOLOGIEN

Serverlose Microservices
mit Lambda

KULTUR

Wie Microservices
Unternehmen verändern

DE € 9,80

AT € 10,80

CH sFr. 19,50

£ € 11,00



www.entwickler-magazin.de



© Stockphoto.com/Getty

+++ Architektur +++ Kultur +++ Entwicklung +++ Technologien +++ Domain-driven Design +++ Netflix +++

Classic Data Architecture

- › Centralized databases
- › ...or services that provide data
- › Ensures consistency across systems
- › ...for data model
- › ...and updates to data
- › Reuse

Classic Data Architecture

Billing

Order
Process

CRM

Order



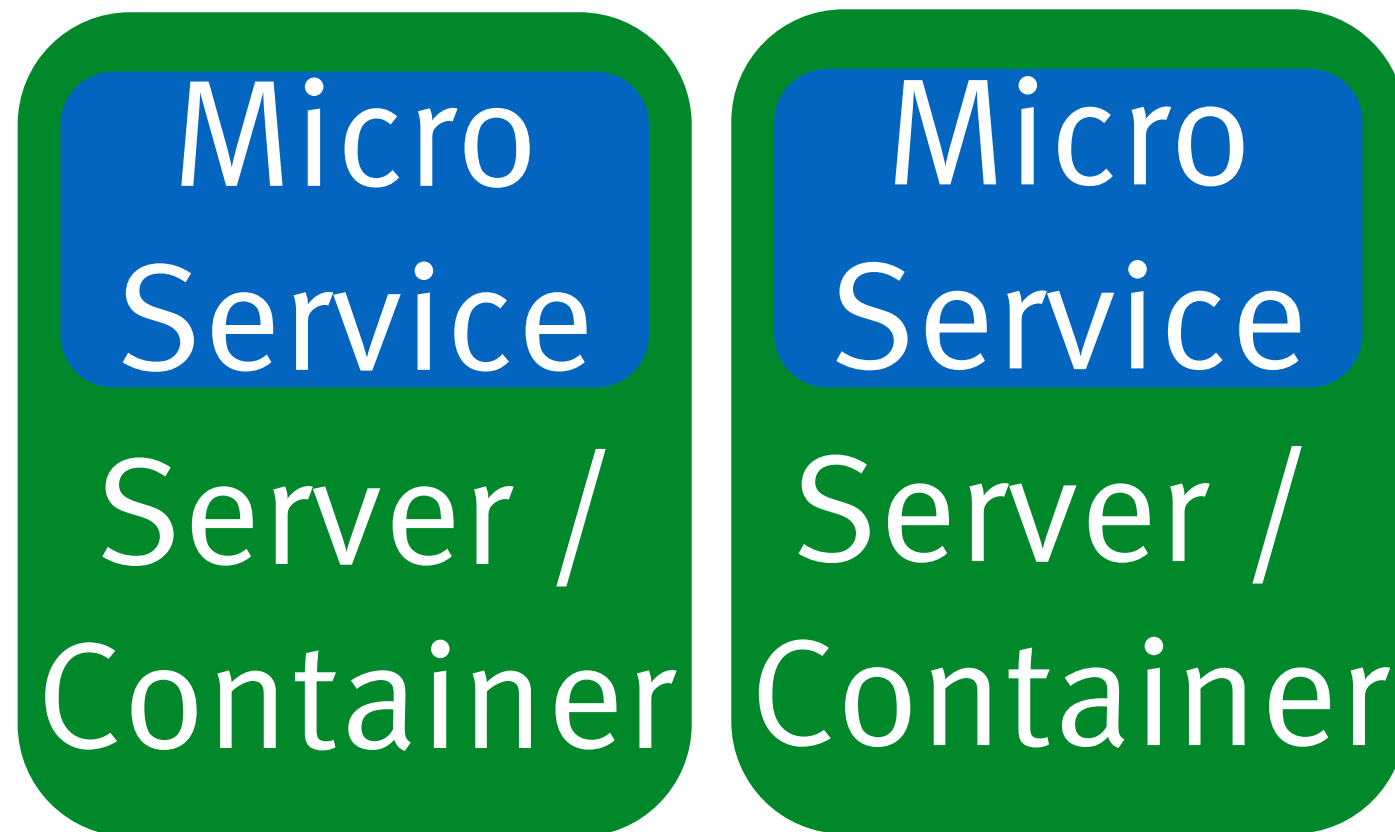
Who is using a
centralized database?

Who likes the
centralized database?


Microservices: Definition

- › No consistent definition
- › Microservices are modules
- › Independent deployment units
- › E.g. processes, Docker container
- › Microservice owned by one team

Microservices: Definition

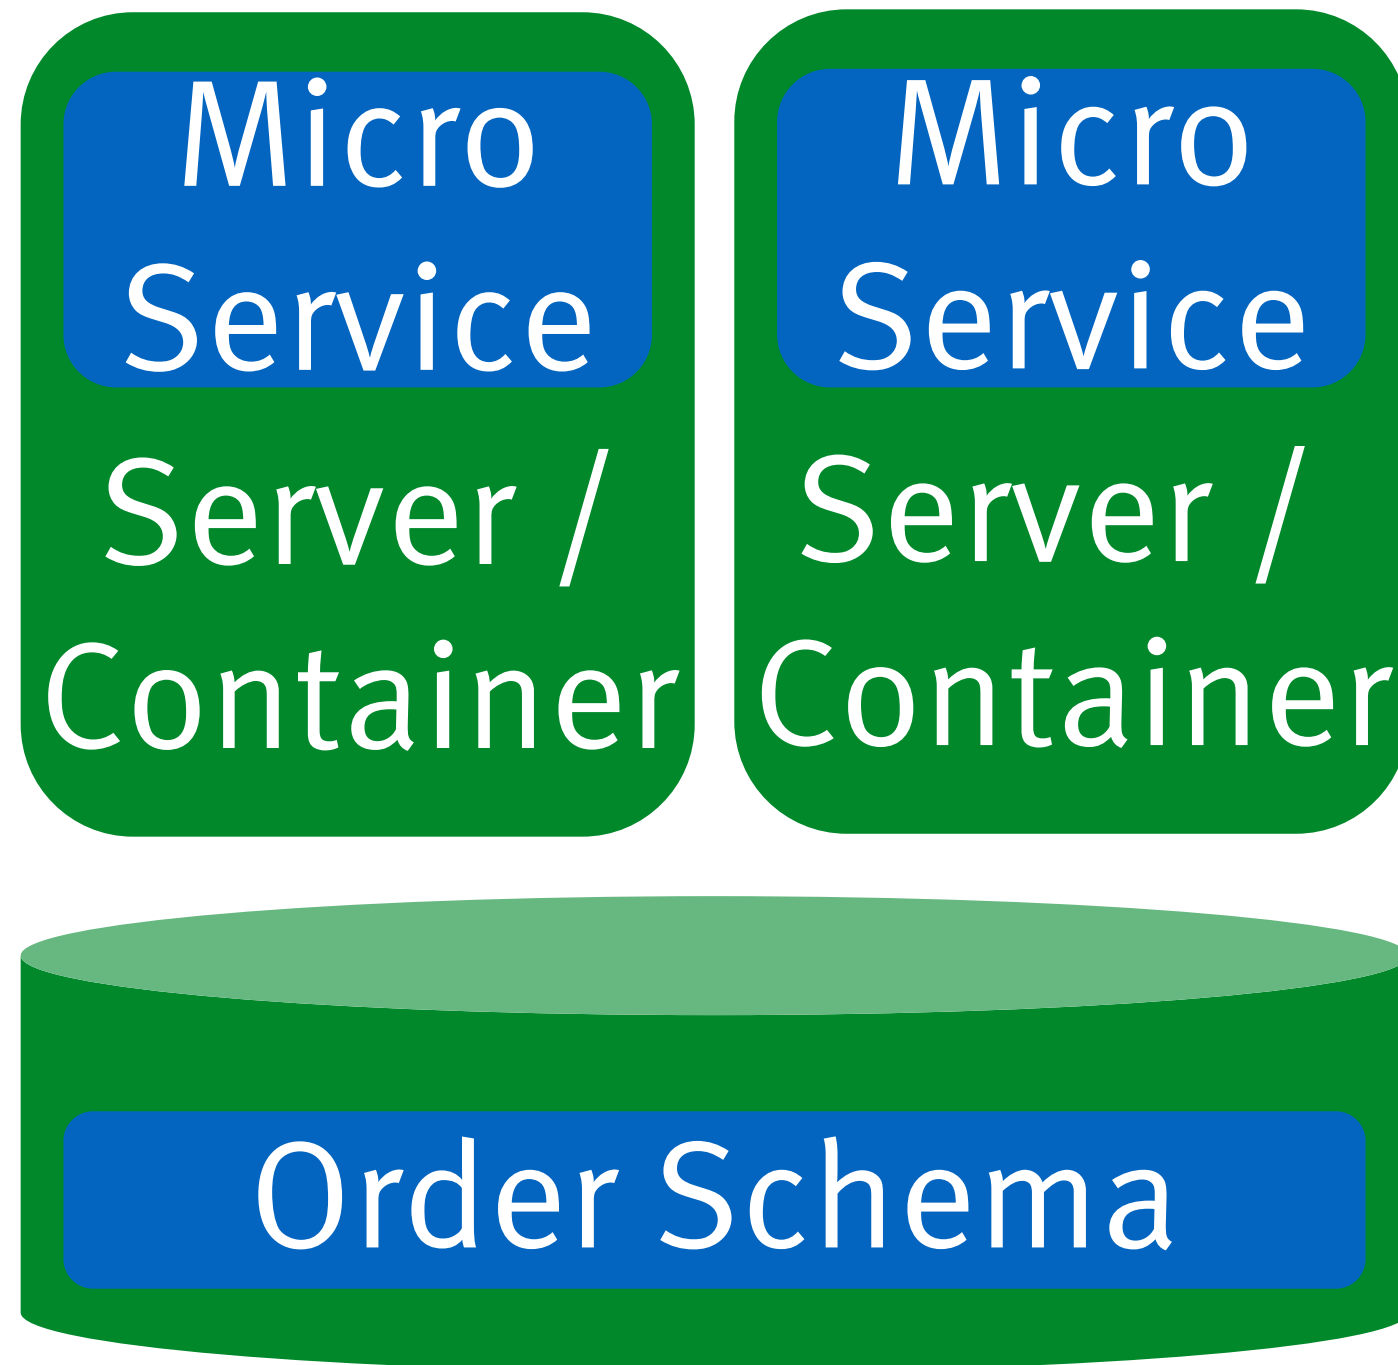


Why Microservices?

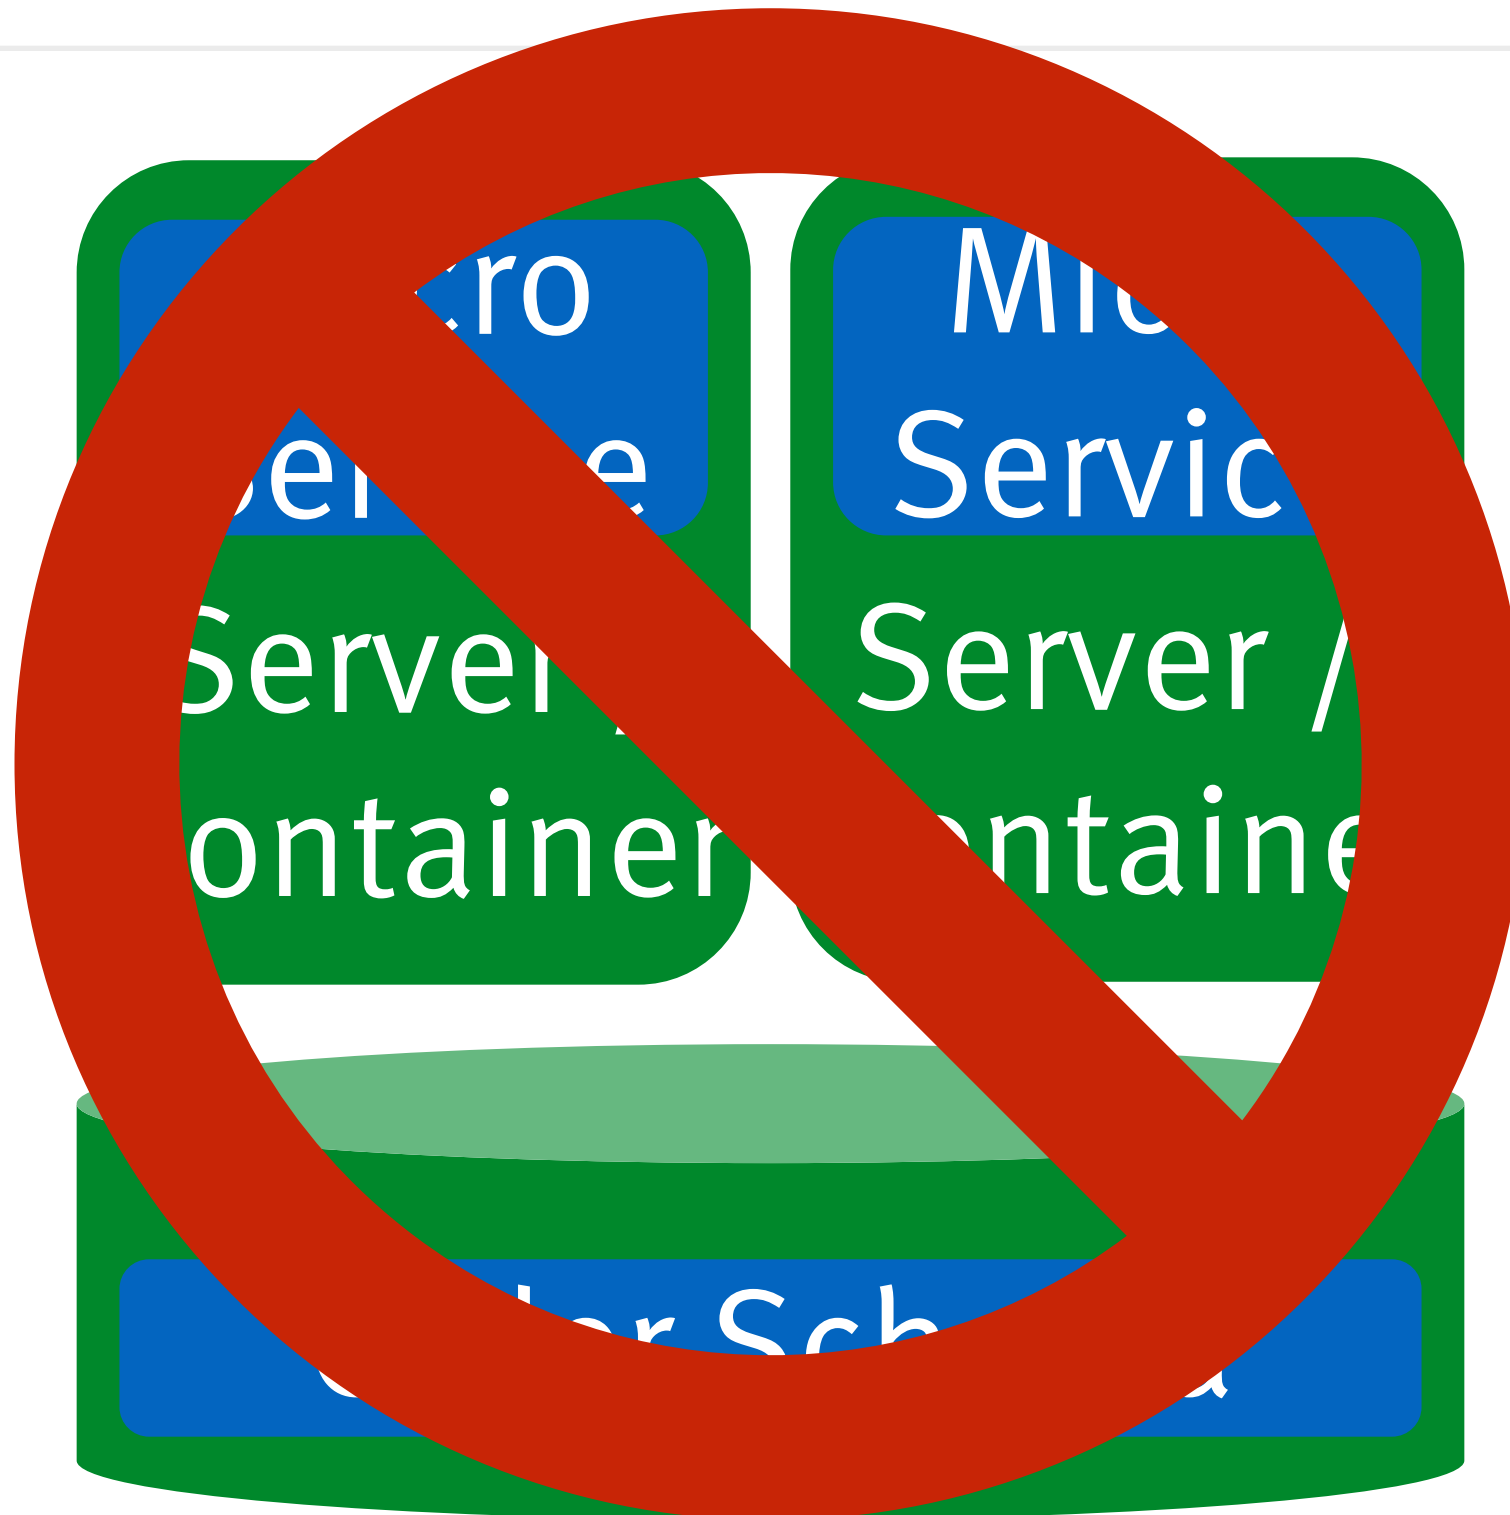
- › Develop a feature
 - › ...bring it into production
 - › ...with no coordination
 - › Independent scaling
 - › Free choice of technology
 - › Robustness
 - › Security
- 

Microservices aim for
decoupling

Microservices & Data



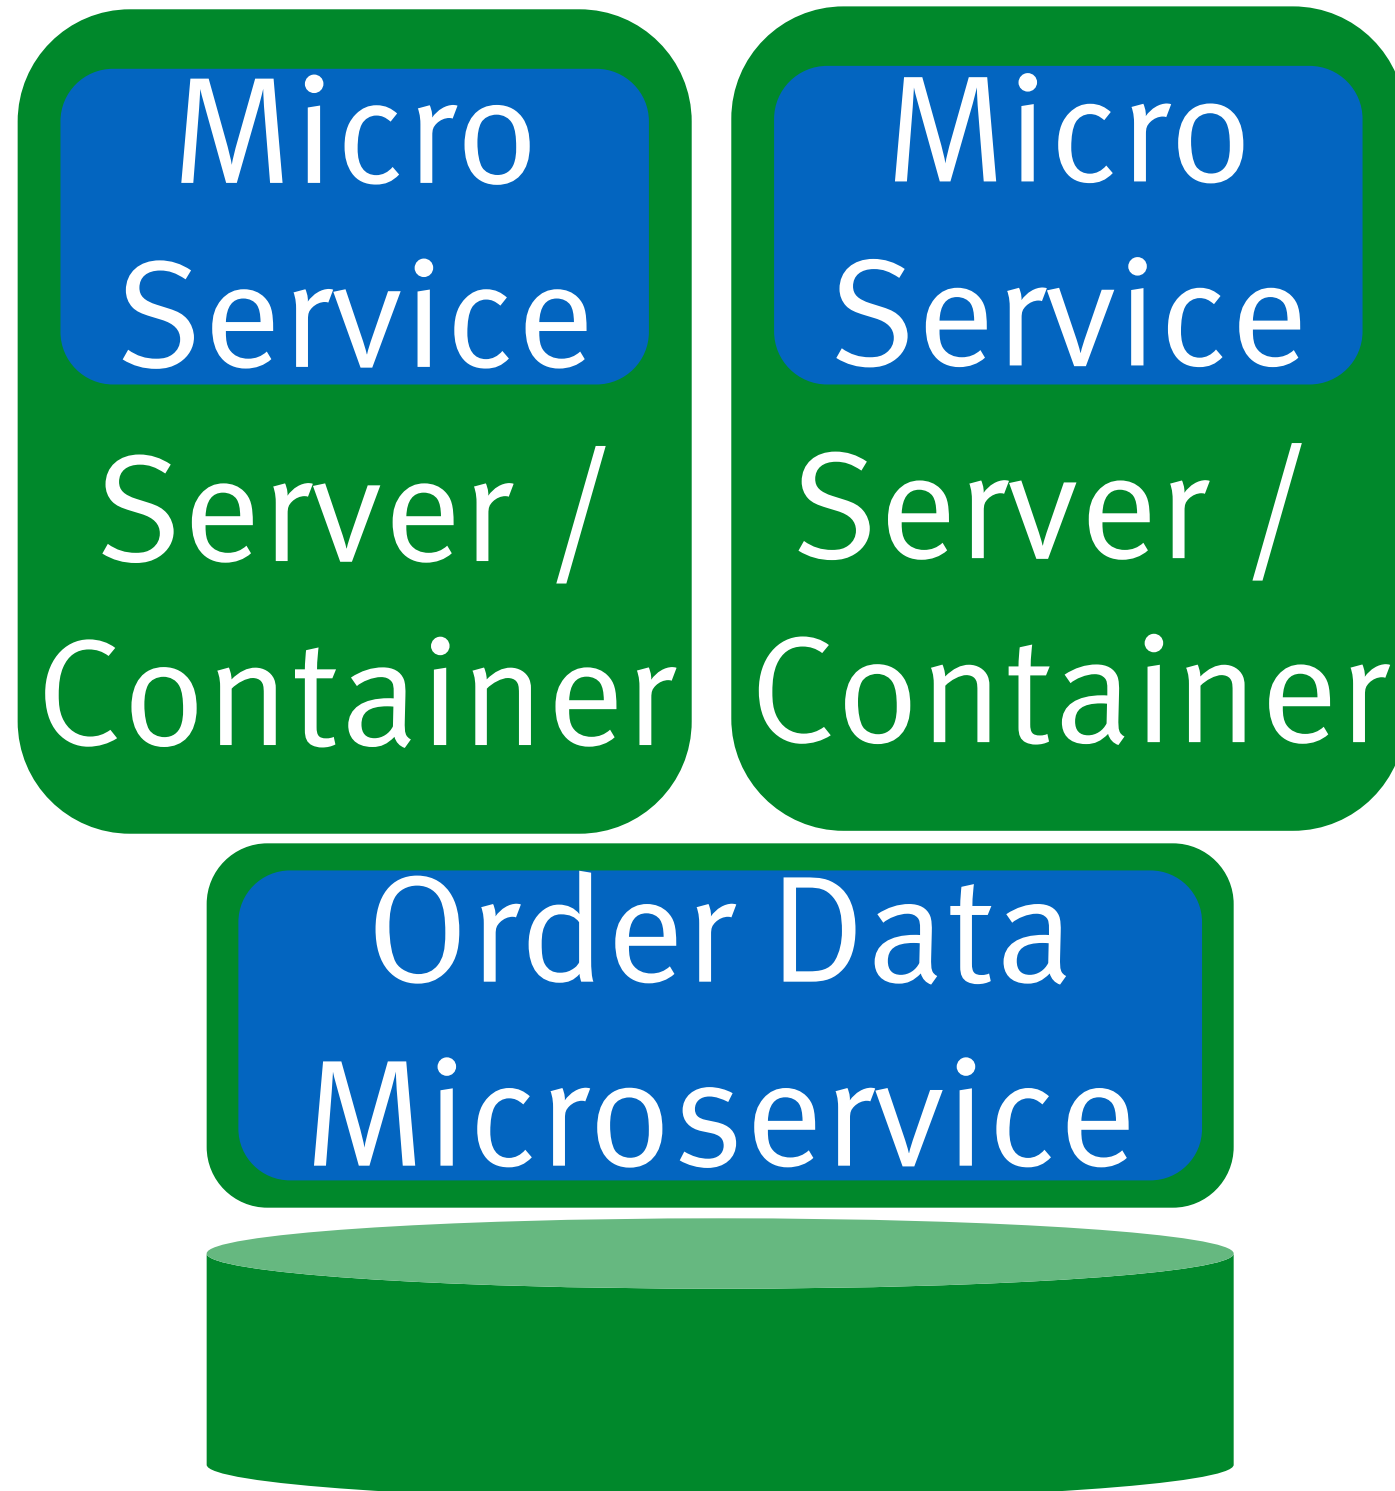
Microservices & Data



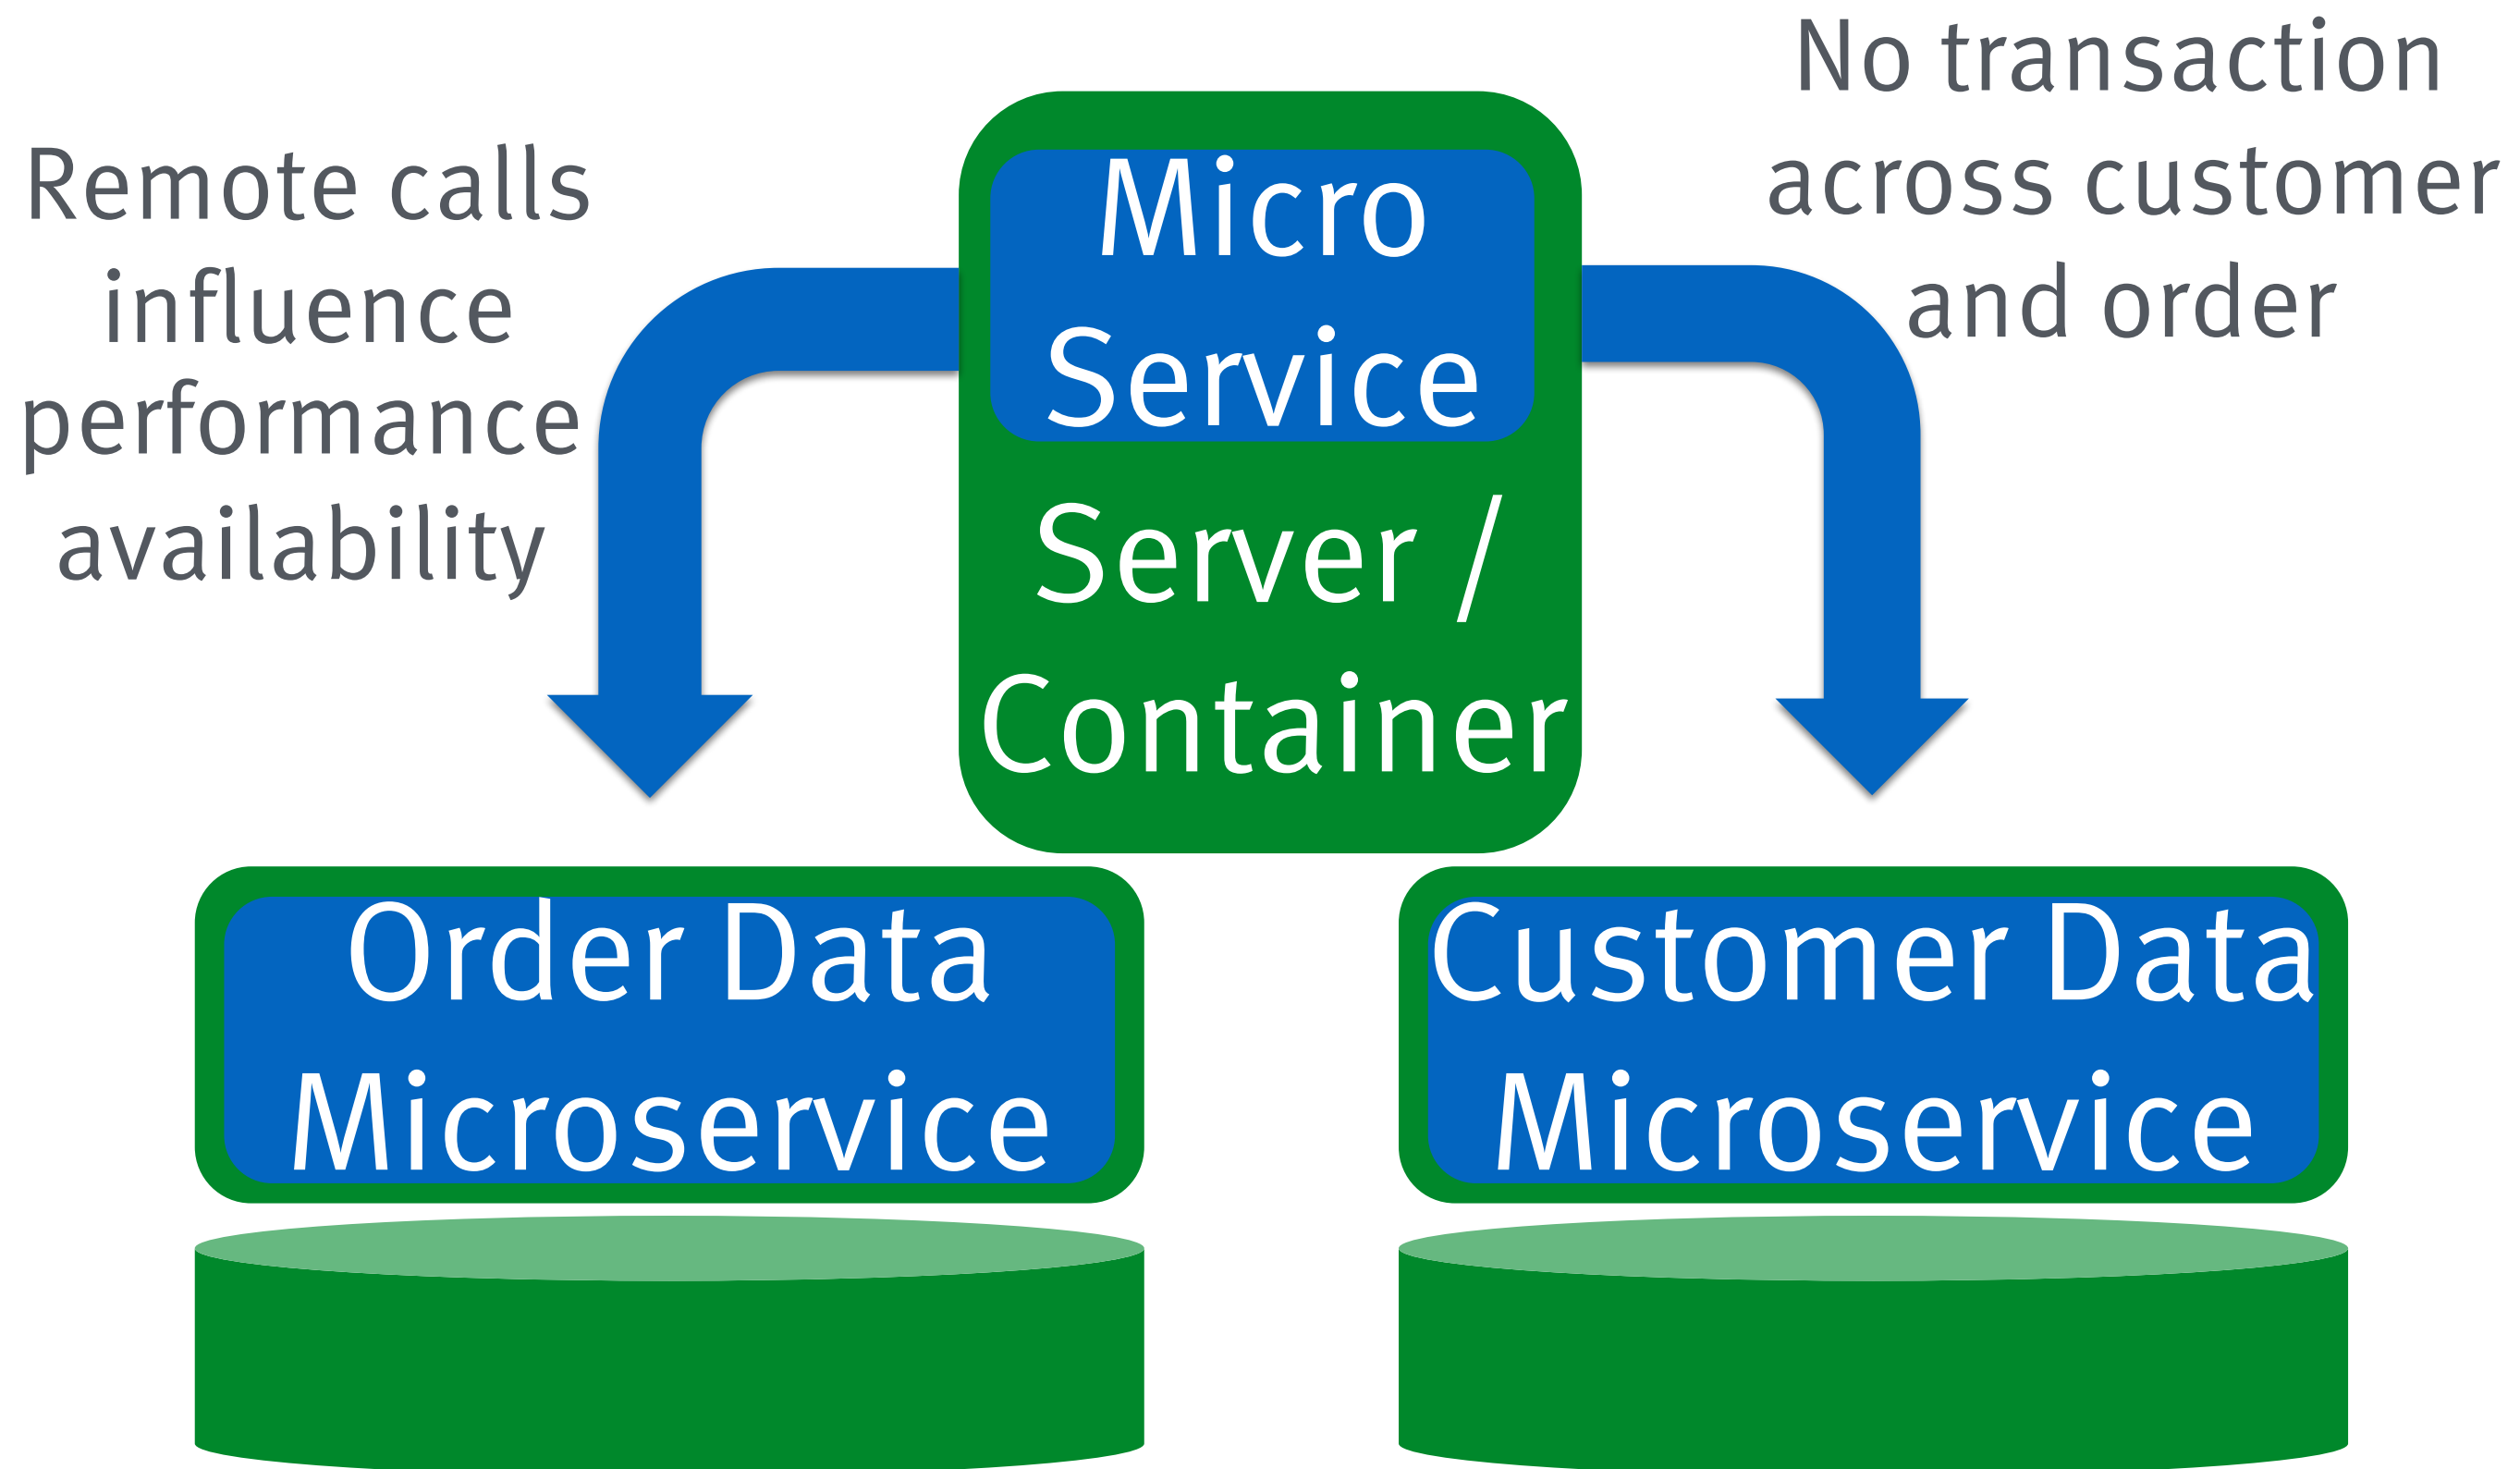
Microservices & Data

- › Decoupling for data, too
- › Separate data storage

Data Microservices



Data Microservices



Data Microservice

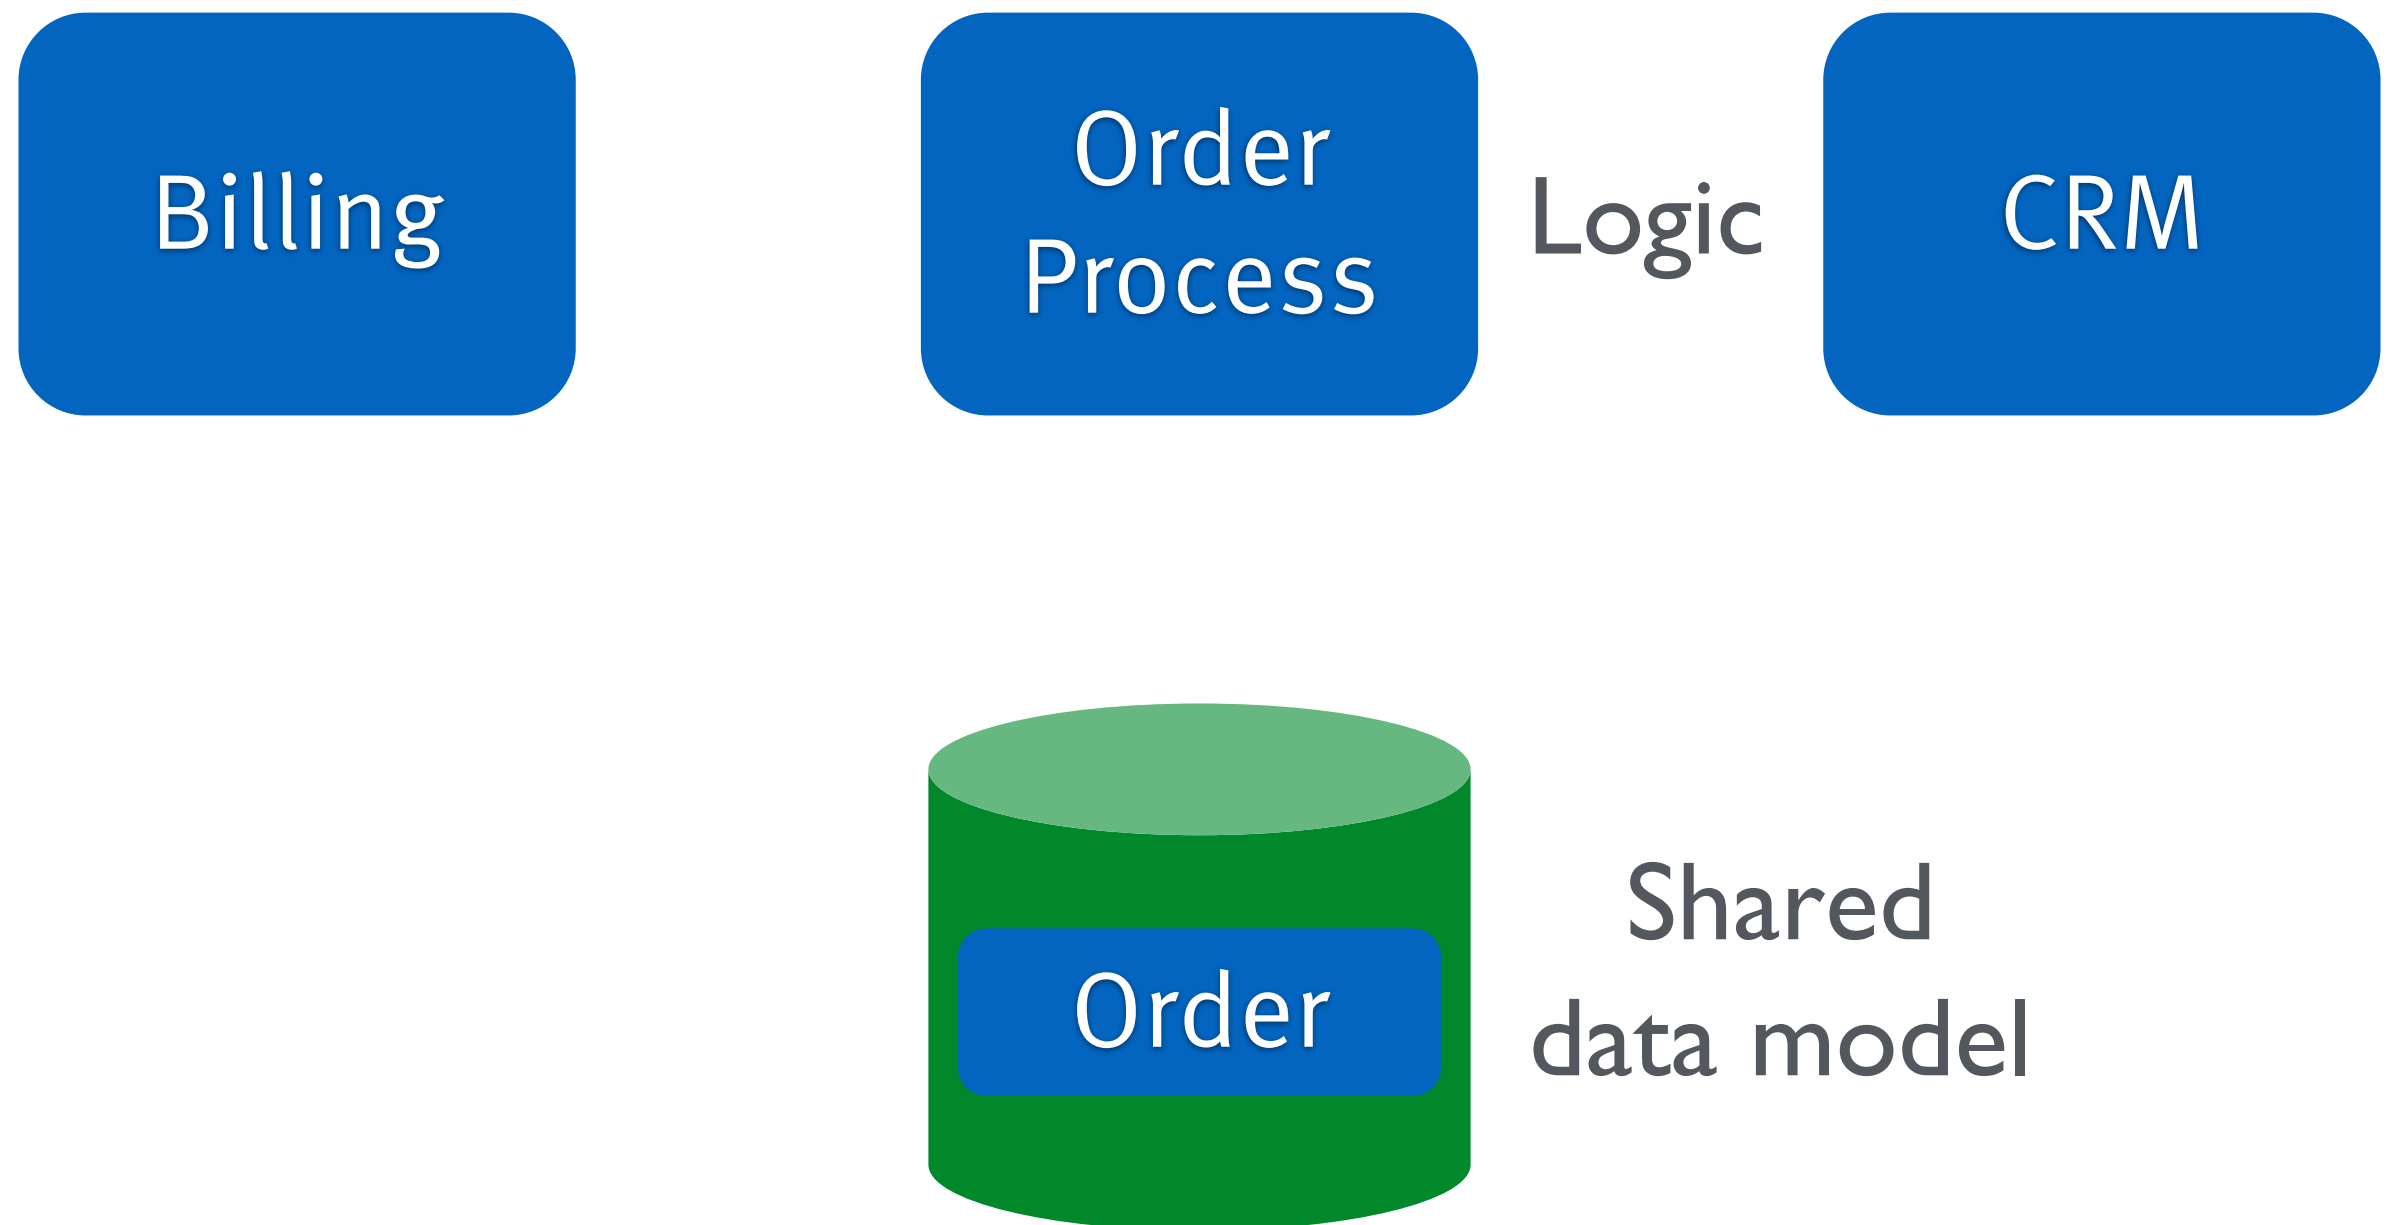
- › Change two microservices if new feature requires change to data schema
- › Transactions?
- › But: data in one place
- › No consistency issues

Data microservice
limits decoupling.

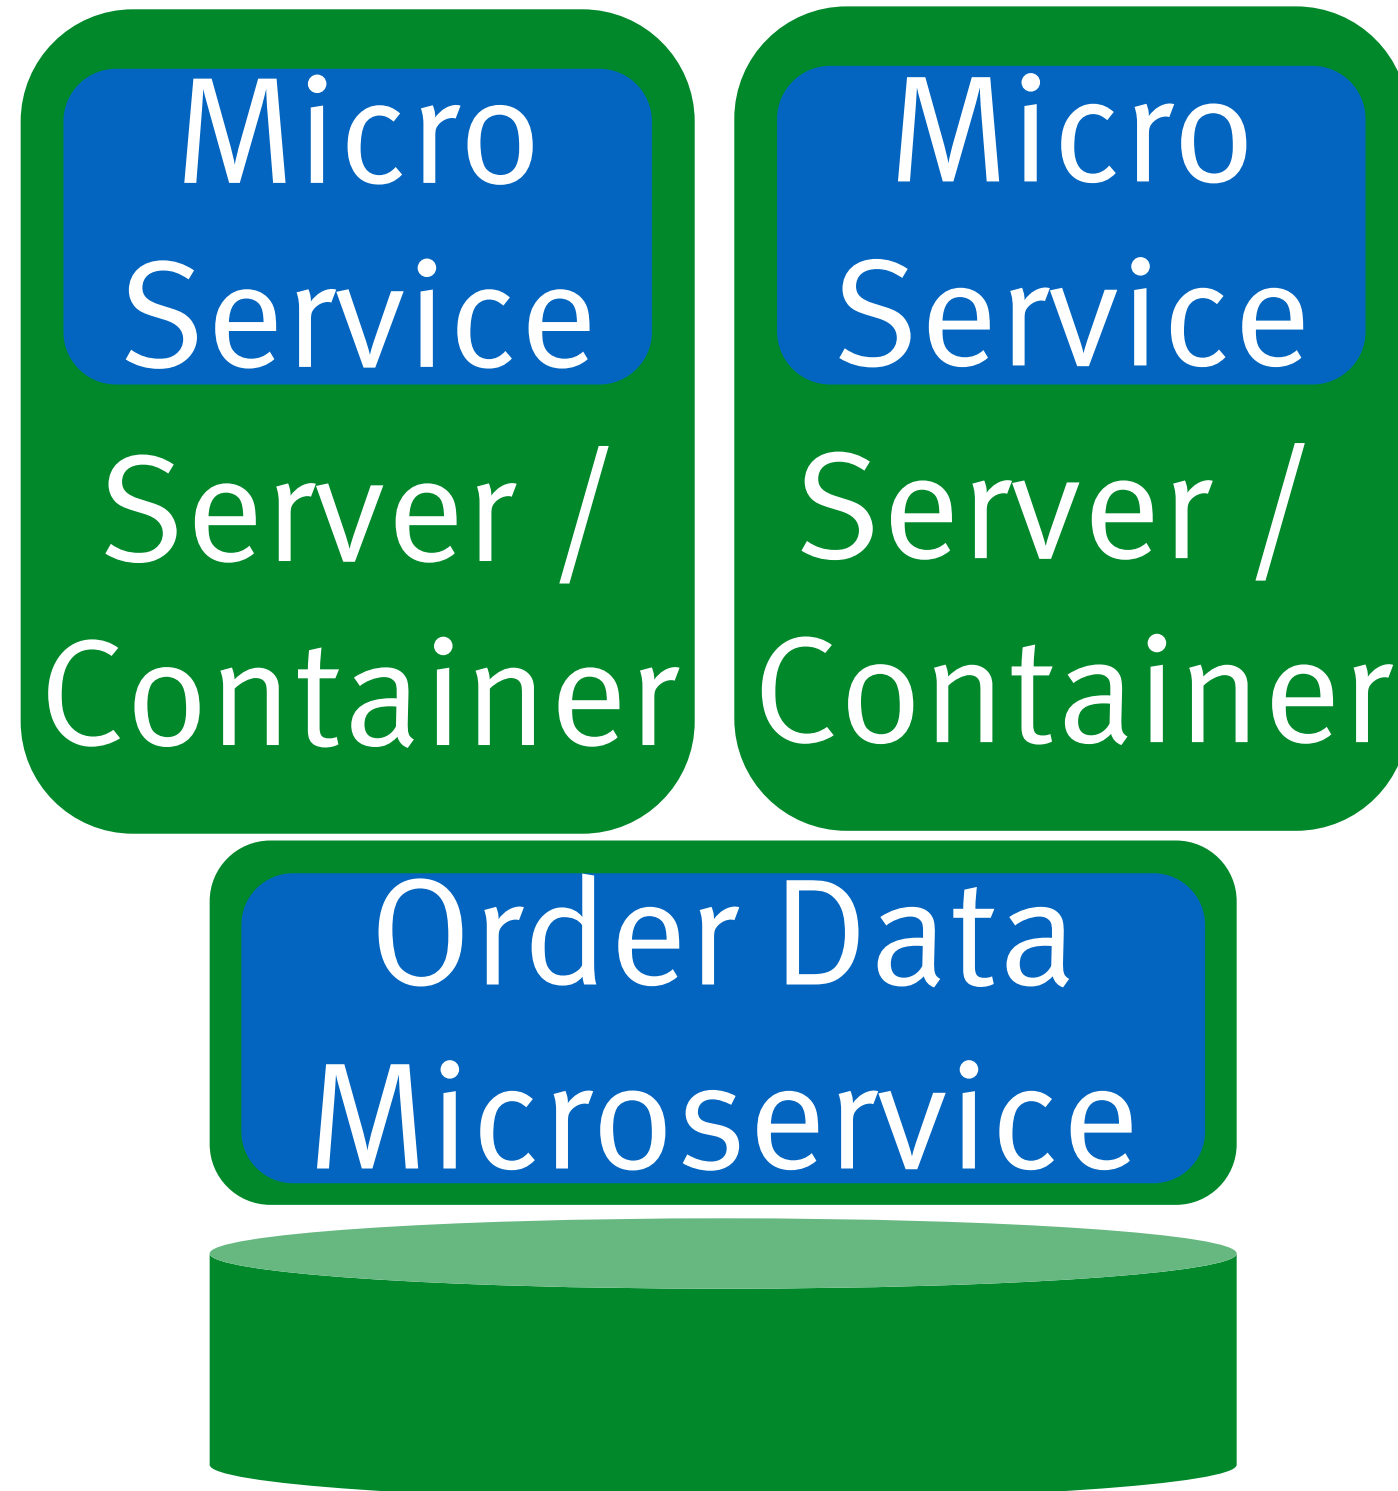
Encapsulation

- › Information hiding
- › Hide the internal data structure
- › Provide access only through a well defined interface
- › Data and databases should not be exported

Violates Encapsulation



Violates Encapsulation



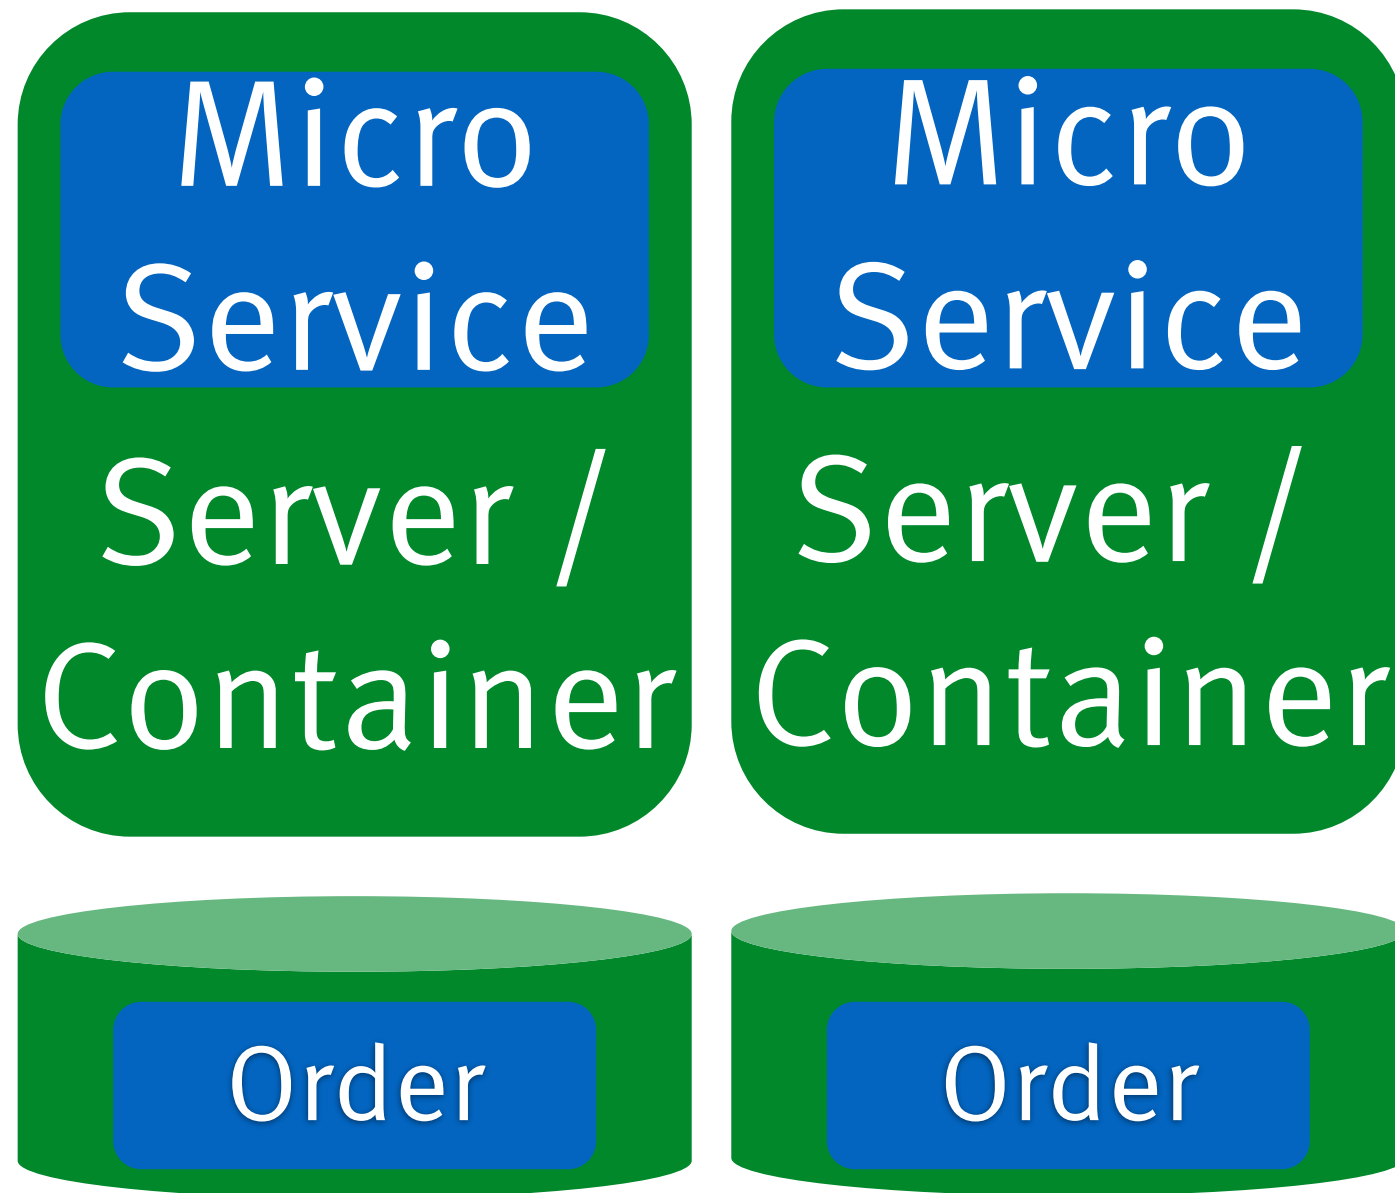
Logic

Shared
data model

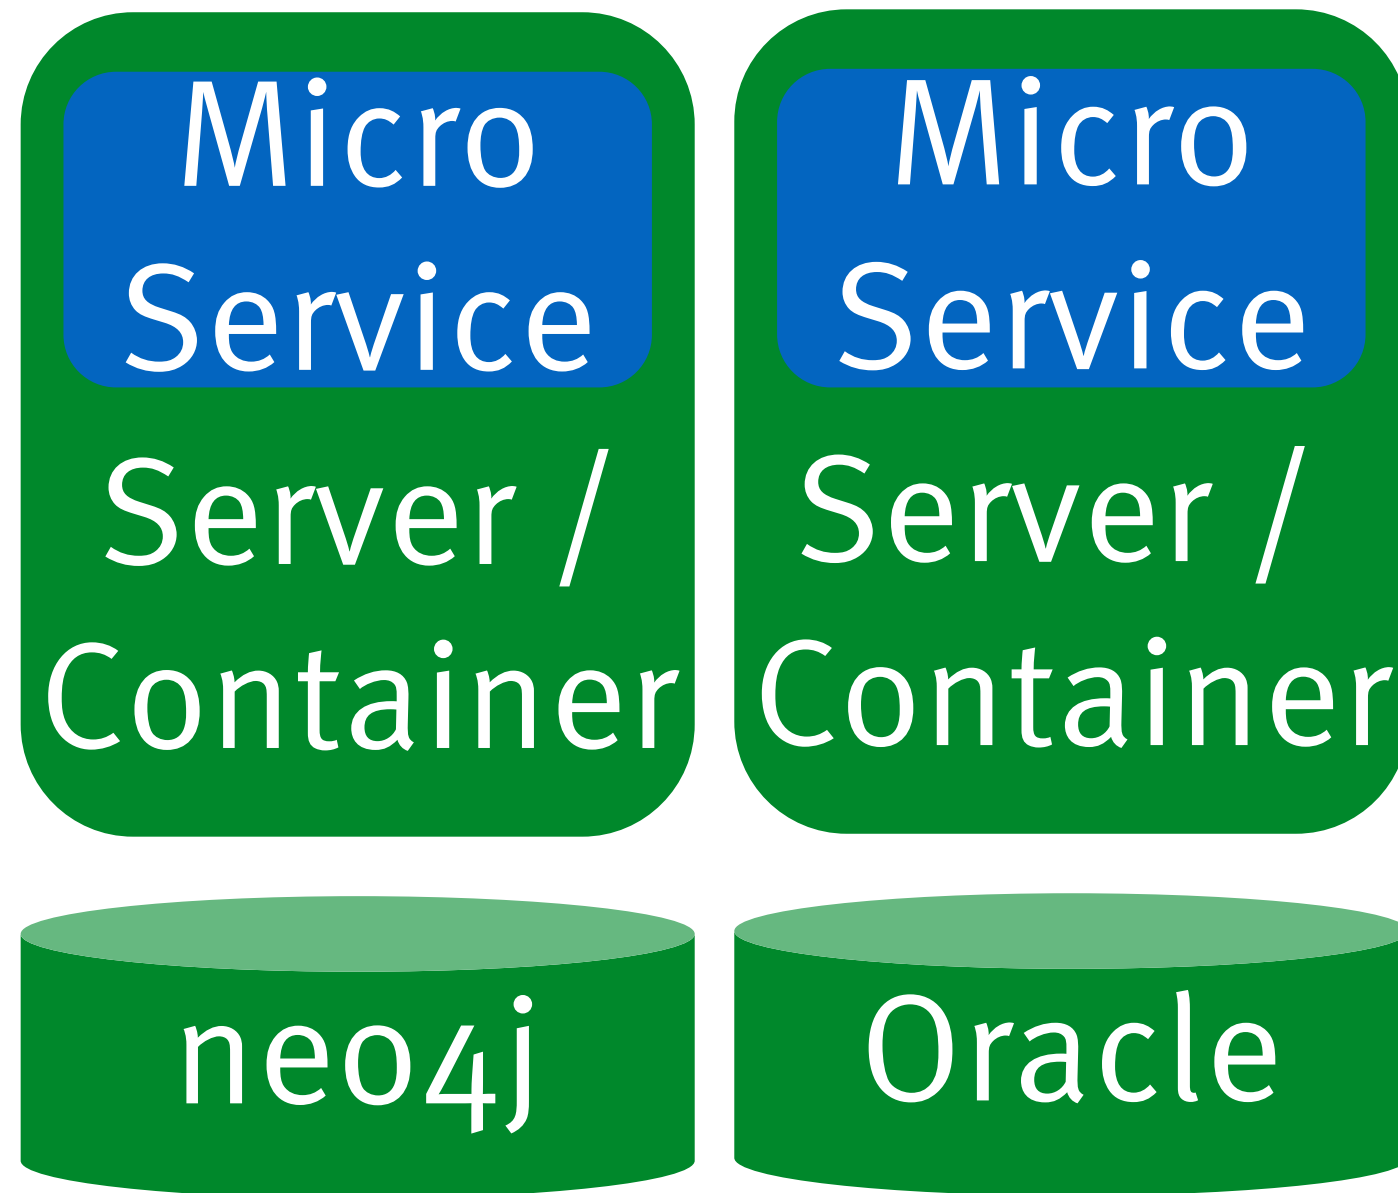
Why You Should Avoid a Canonical Data Model (Stefan Tilkov)

[https://www.innoq.com/](https://www.innoq.com/de/blog/thoughts-on-a-canonical-data-model/)
de/blog/thoughts-on-a-canonical-data-model/

Separate Databases



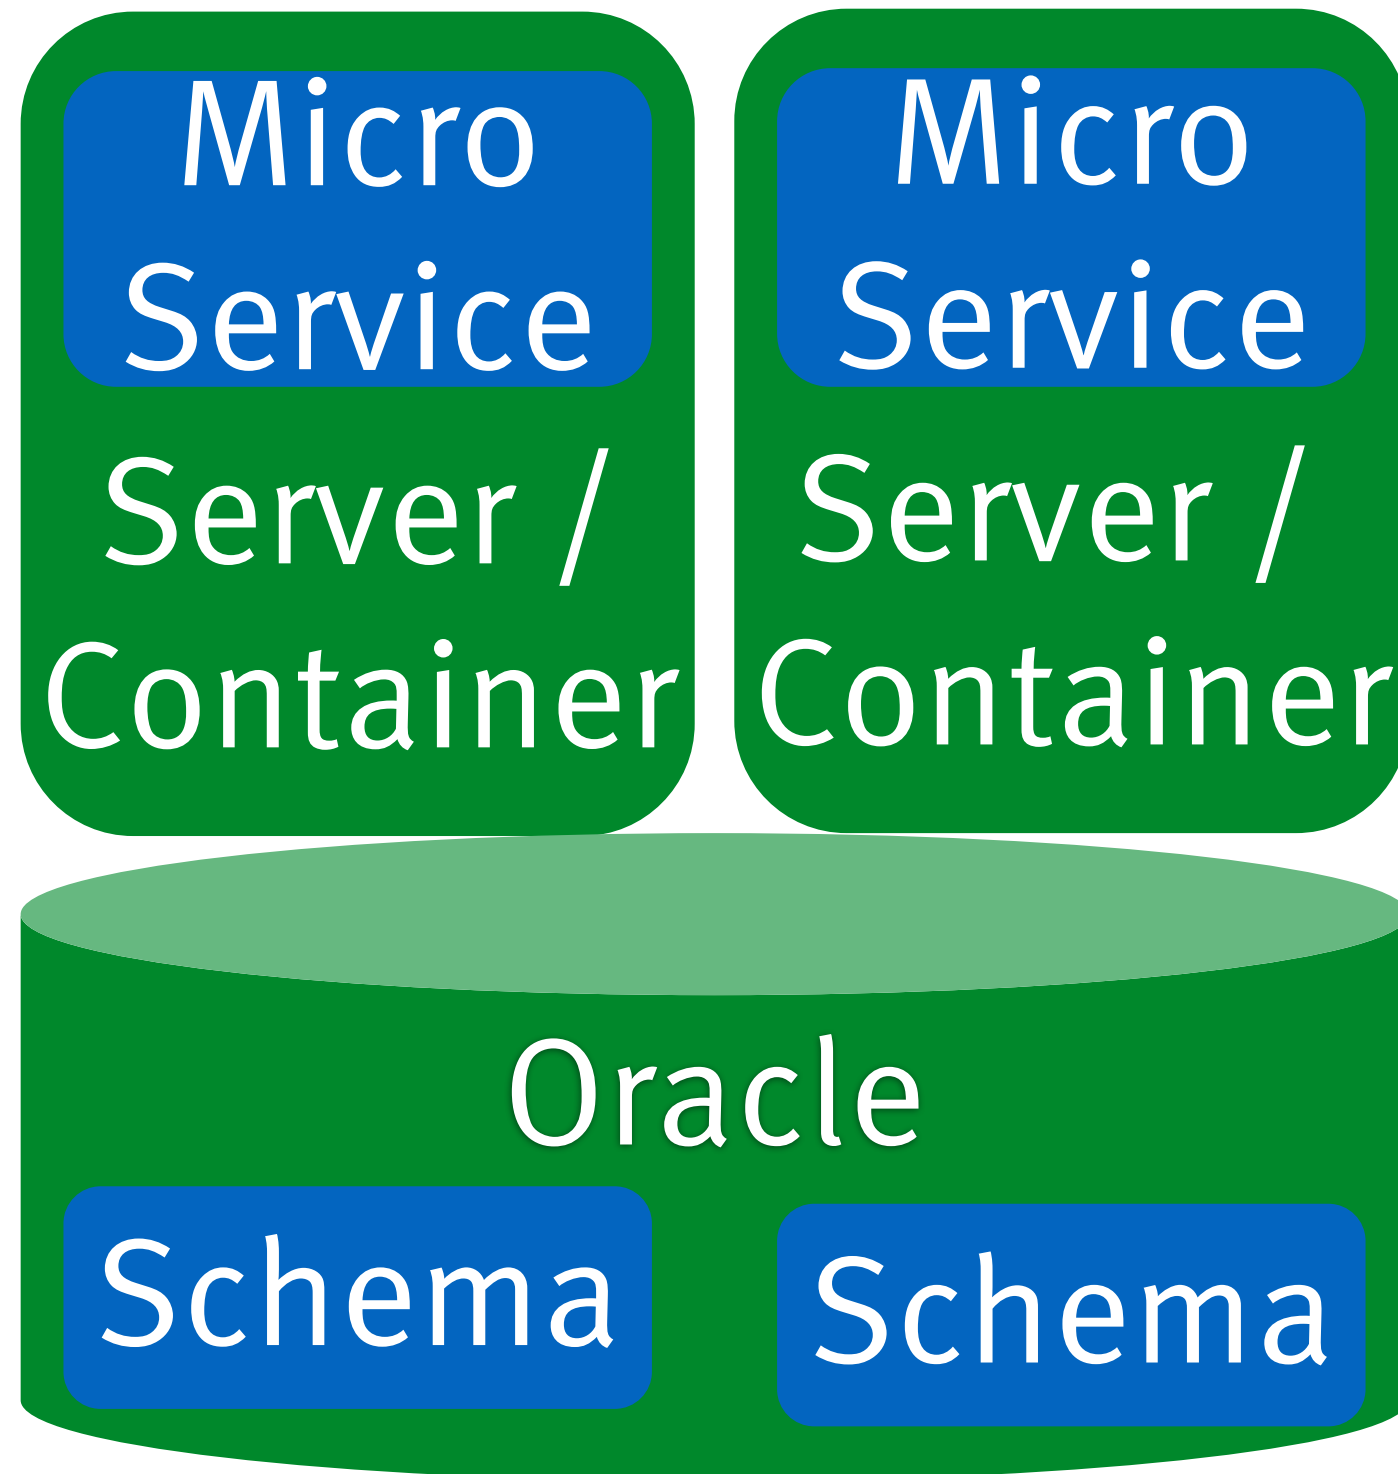
Different Databases



Different Databases

- › “Polyglot persistence”
- › Use the best tool for the job
- › Technology freedom
 - advantage of microservices
- › ...but extra effort
- › Backup, disaster recovery etc.
- › Not as easy as e.g. different frameworks

Separate Schema



Separate Schemas

- › Less effort
- › Decoupled data models
- › ...but limited independent scaling and robustness

Billing

Order
Process

CRM

Order

Order

Order

Redundancy!!!

THE END IS

NEAR





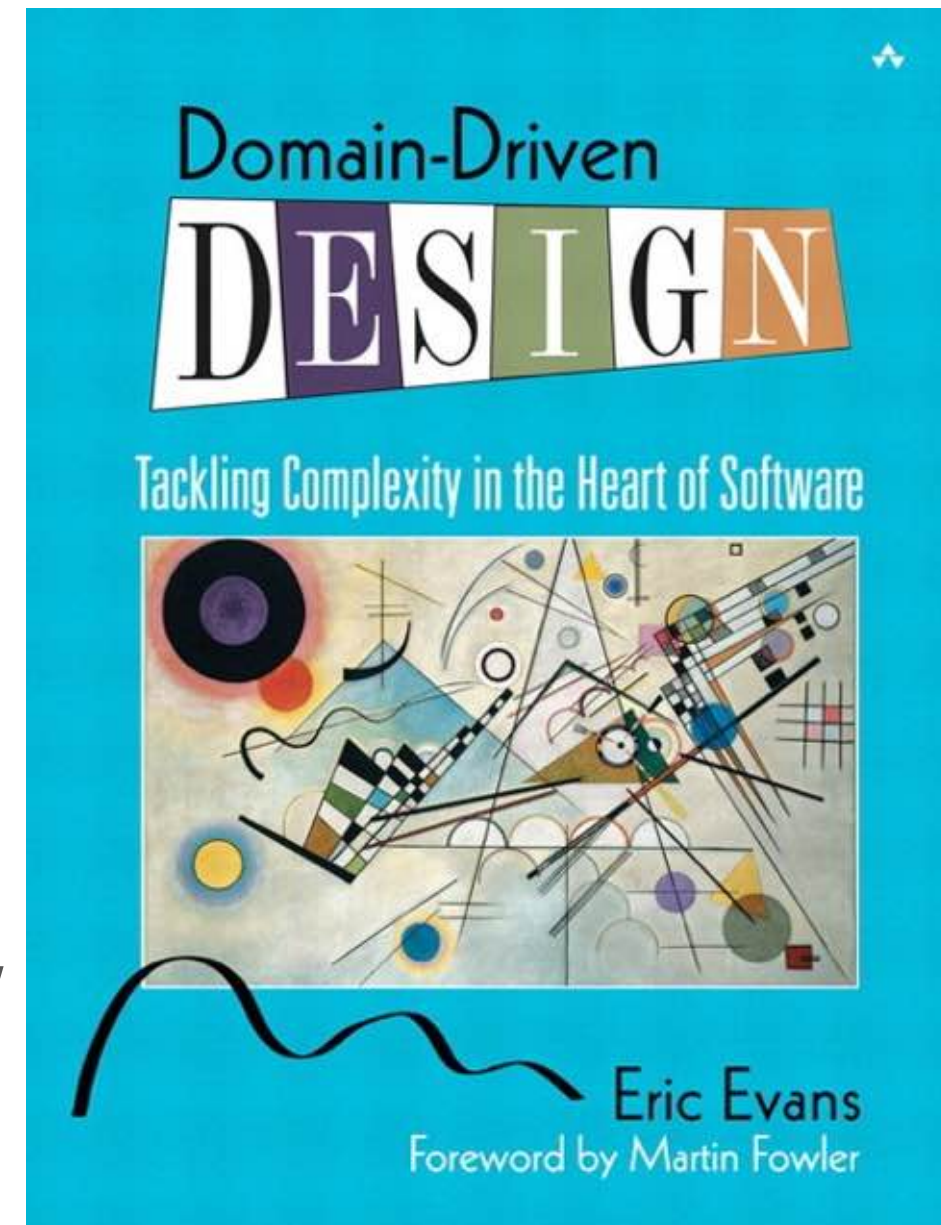
**WHAT IF I TOLD
YOU**

**THERE IS NO REDUNDANT DATA
HERE?**

Domain-driven Design

Domain-driven Design

- › 2004
- › Still very relevant
- › By Eric Evans
- › Focus on part IV
- › Free reference:
<http://domainlanguage.com/ddd/reference/>



Order

Order #

Shipping address

Tracking #

Items

Item Categories

Priority shipping

Customs #

Account #

Credit card #

...

My Domain Model
is a mess!

Bounded Context

- › Domain model is only valid for one context
- › There is no universal data model!
- › See all failed SOA attempts

Tracking

Order

Shipping
address

Tracking #

Priority
shipping

Order

Order #

Shipping address

Tracking #

Items

Item Categories

Priority shipping

Customs #

Account #

Credit card #

...

Recommendations

Order

Item

Categories

Payment

Order

Account #

Credit card #

Customs

Order

Customs #

Billing

Order
Process

CRM

Order

Order

Order

Bounded Context

- › Microservice =
BOUNDED CONTEXTS
- › Changes for new features are local
- › ...even if data models need to be changed

Billing

Order
Process

CRM

Order

Order

Order

Redundancy?

Redundancy?
Not really

Bounded Context

What about
basic data of an
order?

Strategic Design

- › How do BOUNDED CONTEXTS relate to each other?
- › Context can have relationships
- › DDD defines several relationship patterns

Shared Kernel

- › Subset of a model
- › ...that two teams share
- › Eric Evans: Including code and database
- › Microservices: Just sharing a model

Anti-corruption Layer

- › Don't let e.g. a legacy model influence a new model
- › Isolate model by additional layer
- › No need to modify the old system

Context Relationships

- › Team = Deployment Unit = BOUNDED CONTEXT
- › Context Relationships define how BOUNDED CONTEXT are used...
- › ...and how much teams need to collaborate

Coordination
Effort



Shared BOUNDED CONTEXT

SHARED KERNEL

CUSTOMER / SUPPLIER

ANTICORRUPTION LAYER

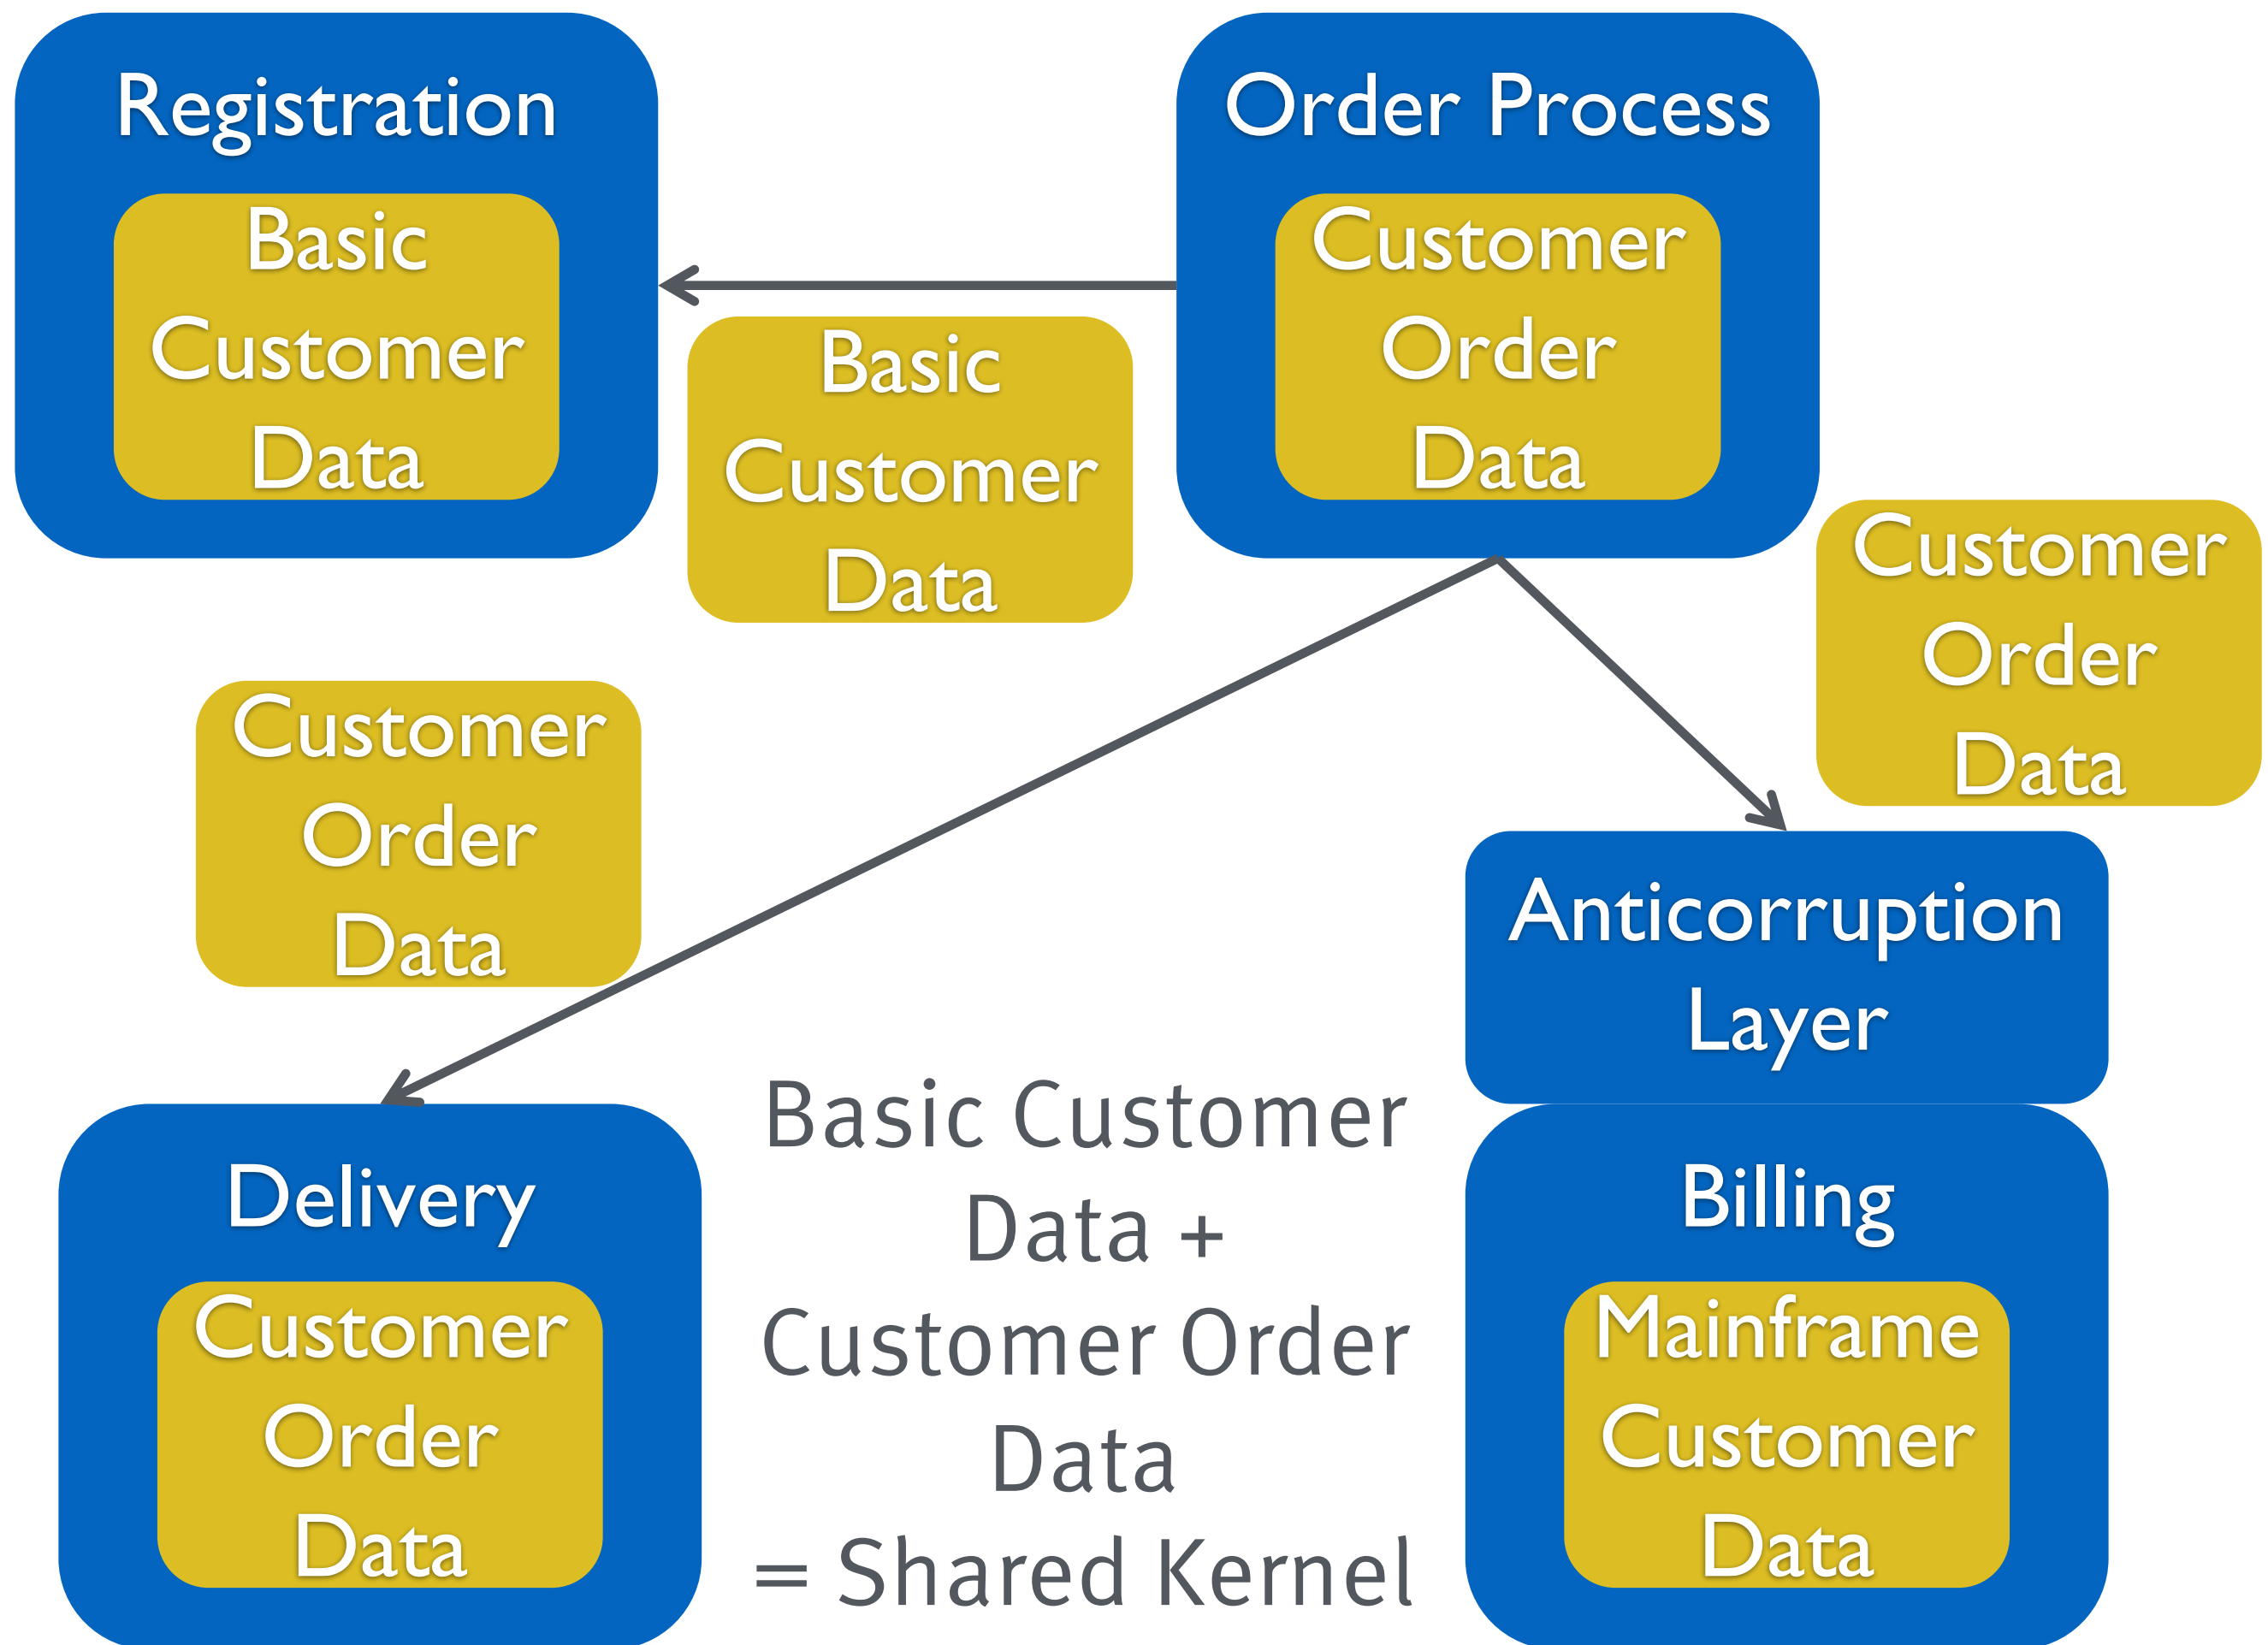
CONFORMIST

SEPARATE WAYS

Context Map

Context Map

- › Show the different BOUNDED CONTEXT
- › ...and the relation to each other
- › BOUNDED CONTEXT might be microservices
- › ...or communication links



Billing

Order
Process

CRM

Additional
data

Additional
data

Additional
data

Order
Data

Shared
Kernel
Order

Centralized Shared Kernel

- › Ensures consistency
- › ...but needs to be called for a lot of operations
- › Resilience / performance / transactions
- › Have one master as the source of truth

Billing

Order
Process

CRM

Additional
data

Shared
Kernel
Order

Additional
data

Shared
Kernel
Order

Additional
data

Shared
Kernel
Order

Decentralized Shared Kernel

- › Might be inconsistent
- › ...but all data for all requests is available in the local database
- › Better resilience...
- › ...and performance

How to Replicate Data?

Database Replication

- › Built into the database
- › Replicate schema across database instances
- › But: Microservices have separated schemas
- › Every Microservice might have different data
- › ...so database replication is not a good fit

Replication with Events

Events

- › Later addition to Domain-driven Design
- › Events with a business meaning
- › Decouple time:
Asynchronous
- › Decouple logic:
System can handle event as it pleases

New Order Event

```
graph TD; A[New Order Event] --> B[Billing]; A --> C[Order Process]; A --> D[CRM]; B --- E[(Additional data  
Shared Kernel Order)]; C --- F[(Additional data  
Shared Kernel Order)]; D --- G[(Additional data  
Shared Kernel Order)];
```

The diagram illustrates a data flow from a central 'New Order Event' to three distinct business processes: Billing, Order Process, and CRM. Each process is represented by a blue rounded rectangle. Below each process is a green cylinder representing a database or data store. Inside each cylinder are two blue rounded rectangles: 'Additional data' and 'Shared Kernel Order'. Arrows point from the 'New Order Event' bar to each of the three process boxes.

Billing

Additional
data

Shared
Kernel
Order

Order
Process

Additional
data

Shared
Kernel
Order

CRM

Additional
data

Shared
Kernel
Order

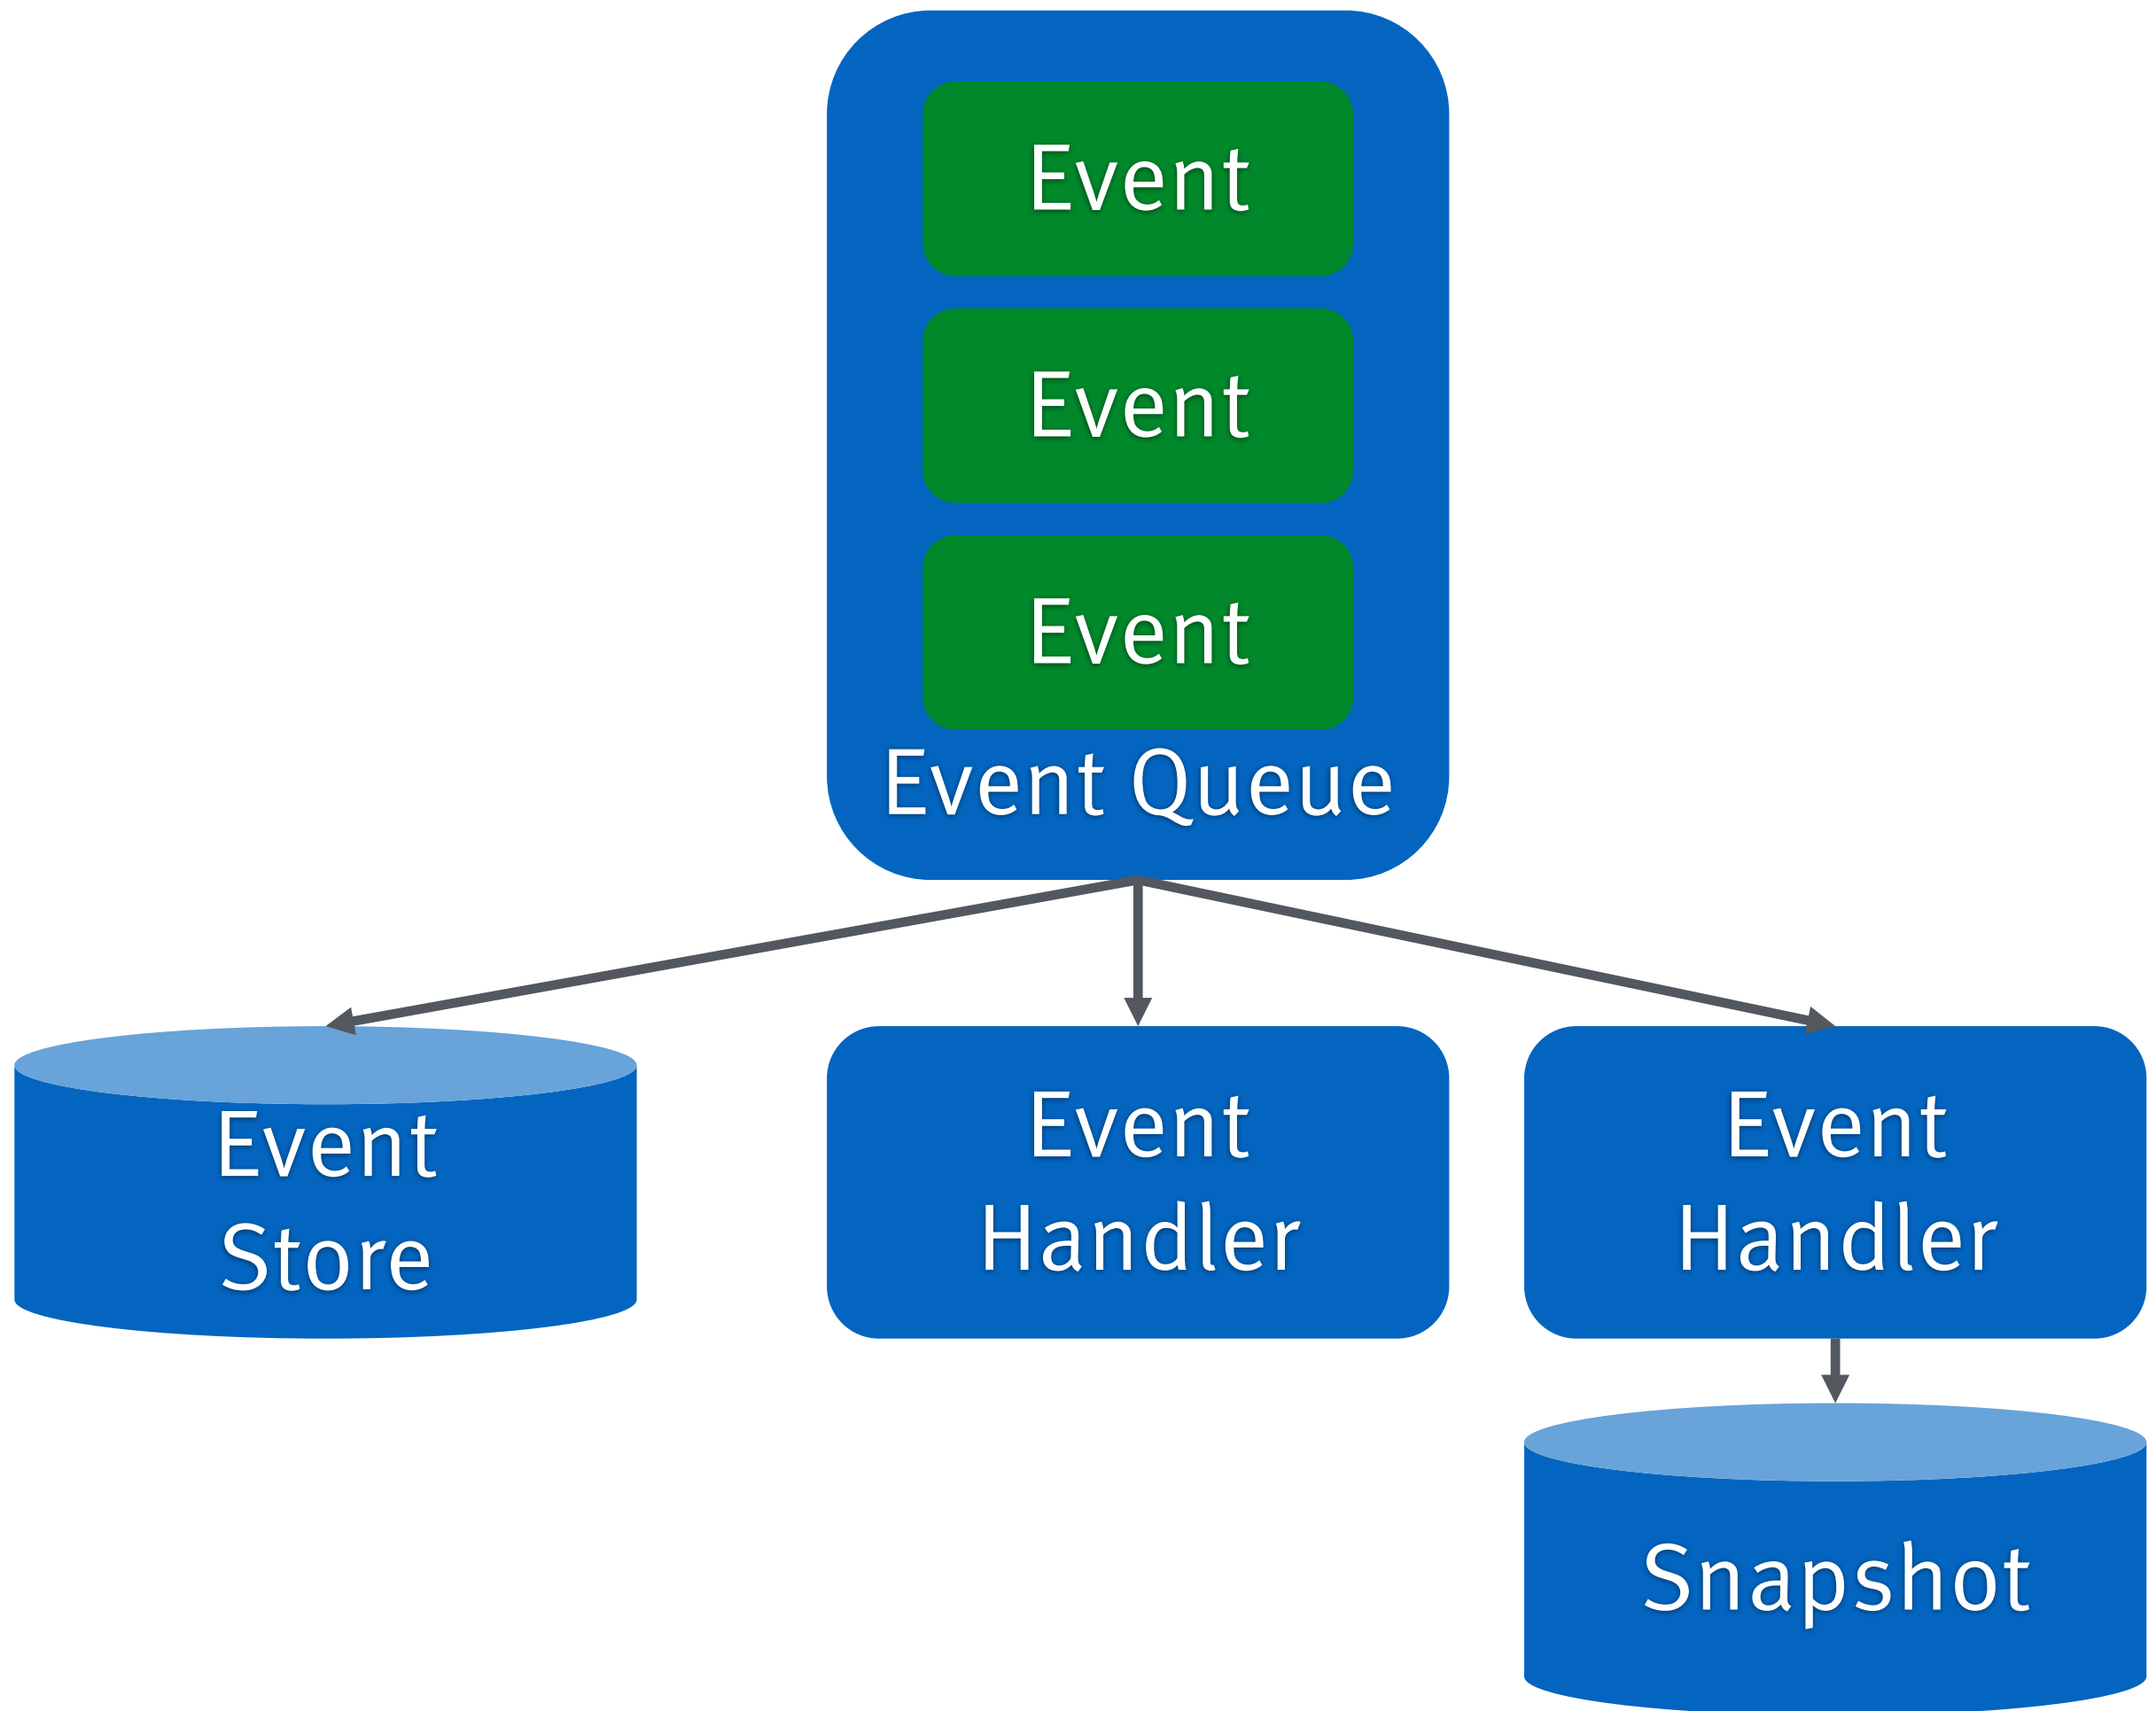
Events & Data Replication

- › Events lead to data replication
- › i.e. each system stores information it received in an event
- › Data stored in separate schema
- › Very decoupled
- › Hard to repair inconsistencies

More Fun With Events

Event Sourcing

- › Internal Structure for Microservice with events
- › Current state result of all events
- › Calculate state on the fly?



Event Sourcing

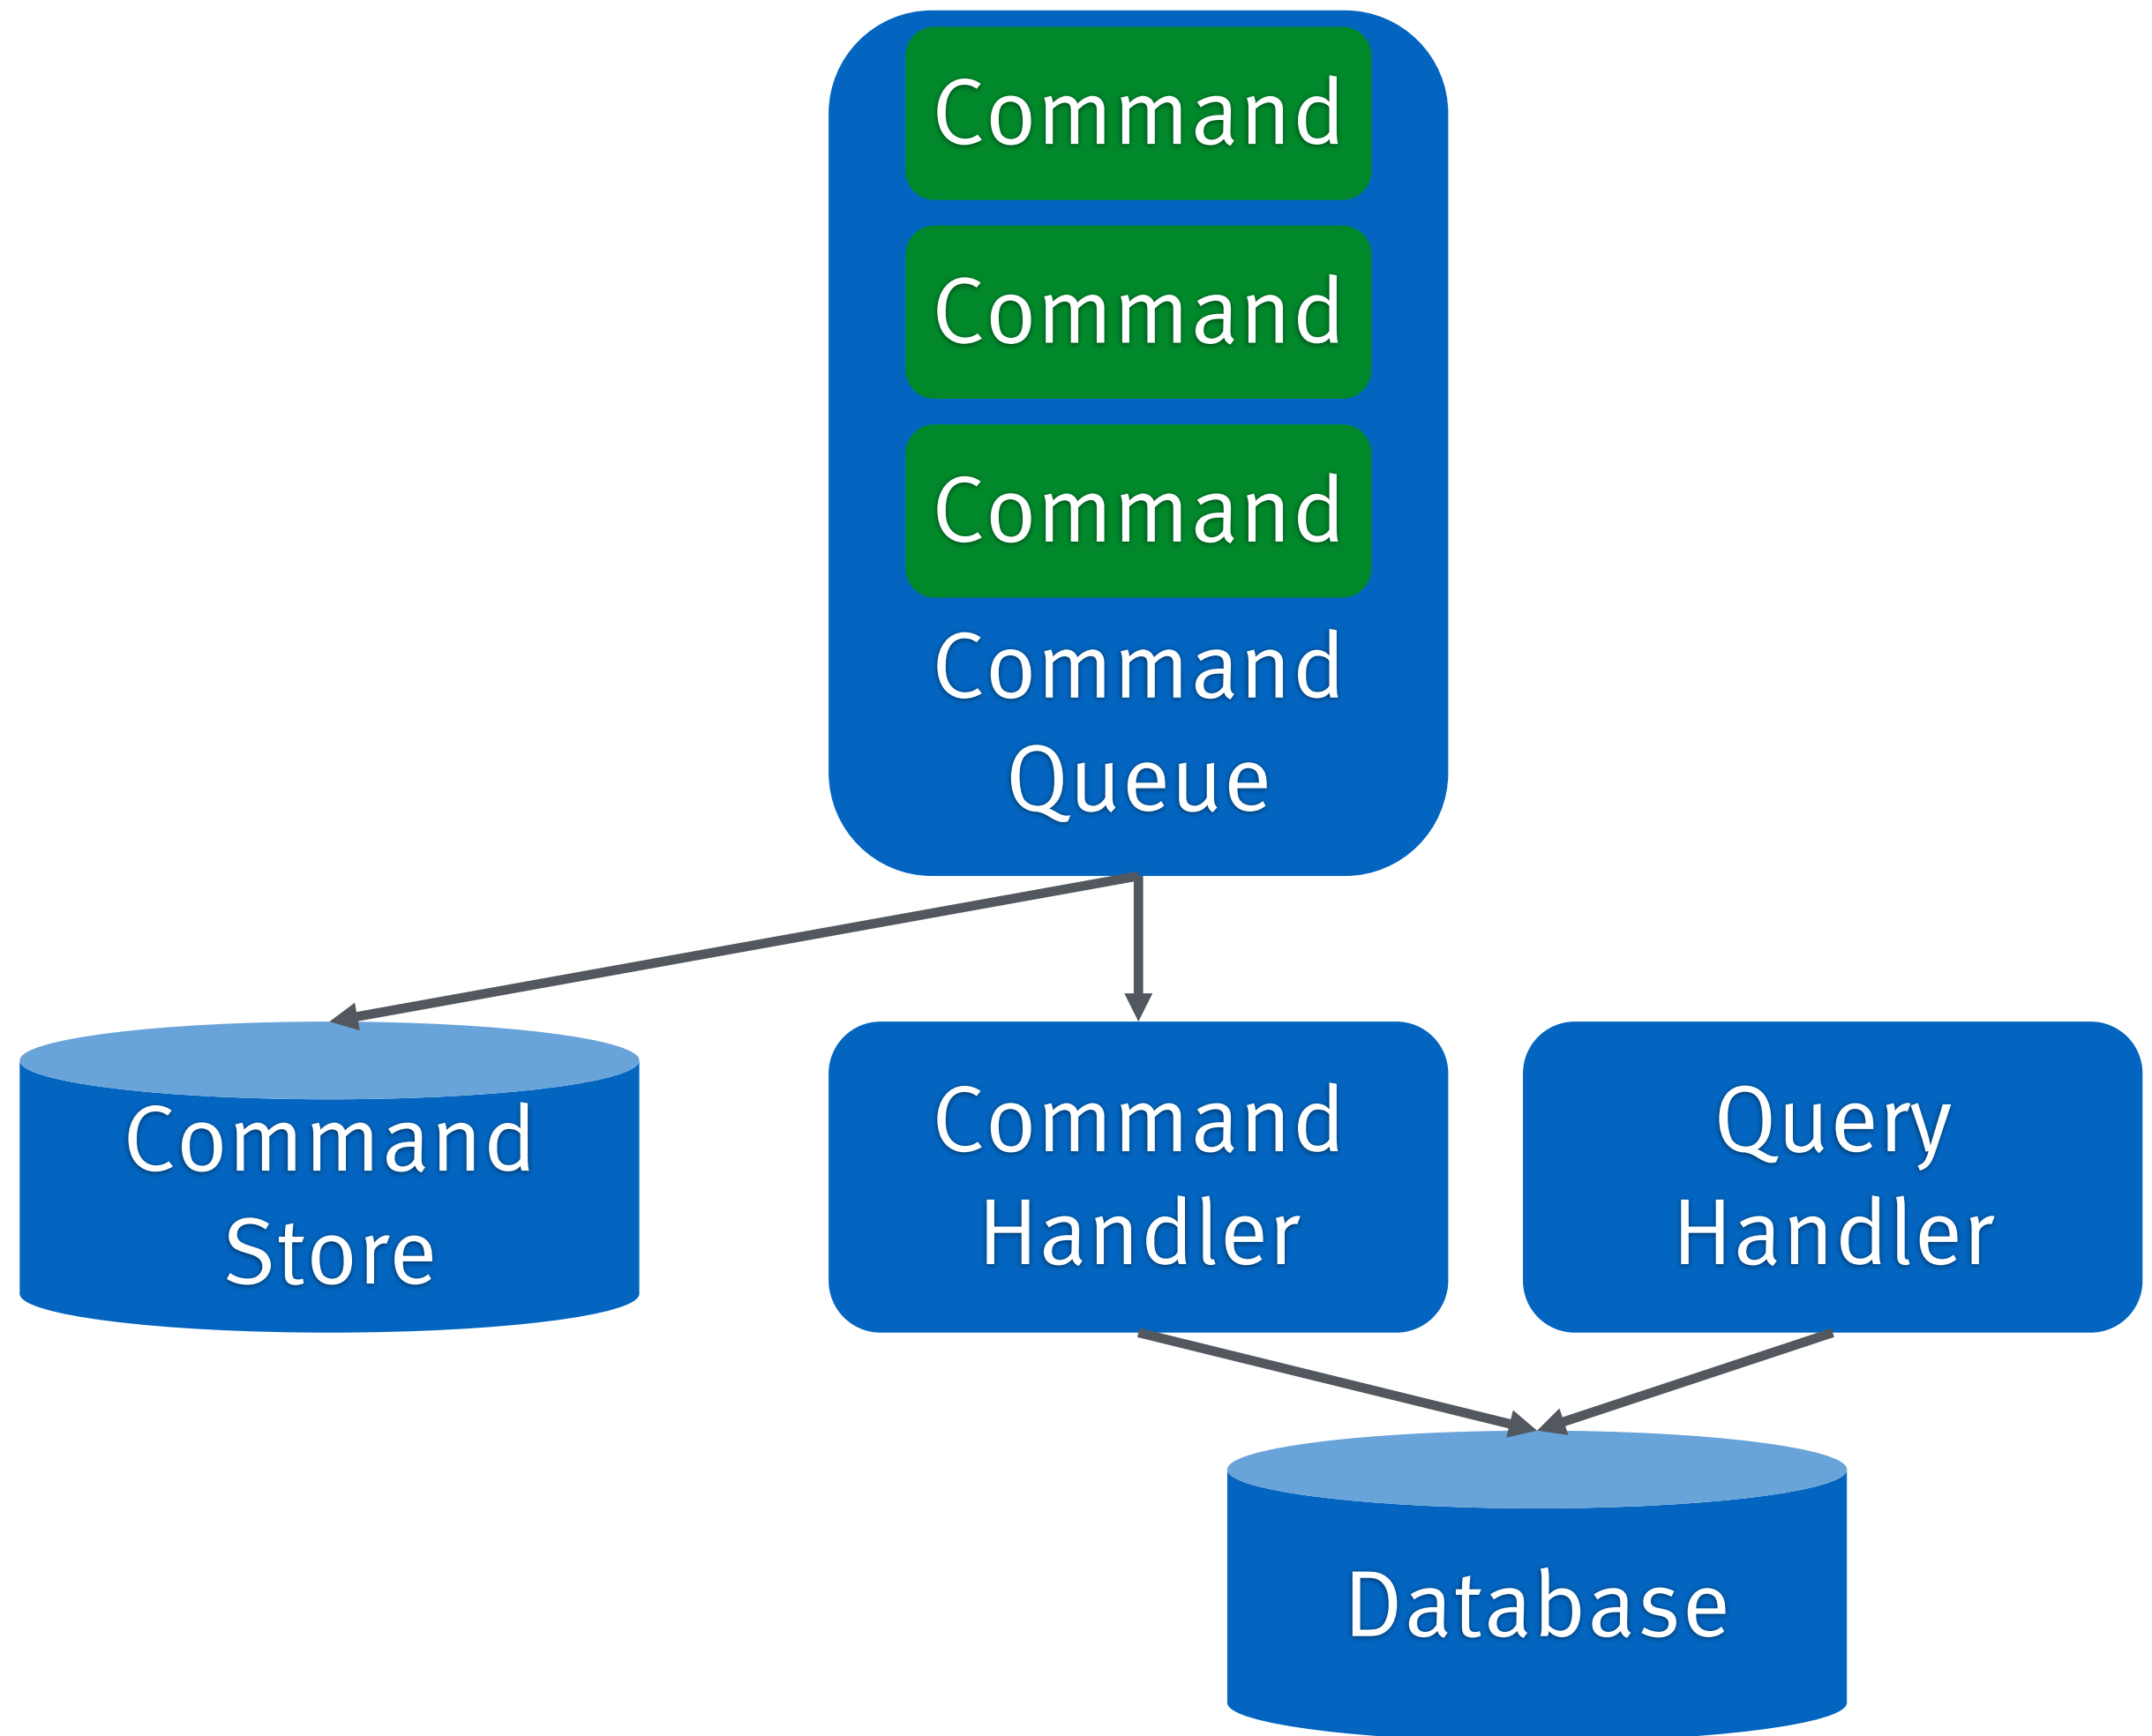
- › Event store and snapshot help to repair inconsistencies
- › Event-based architecture in microservices

CQRS

- › Command – Query Responsibility Segregation
- › Commands change data
- › Query provide data
- › Implement in separate modules
- › ...or even microservices
- › ...with potentially different BOUNDED CONTEXTS

Commands vs Events

- › Command: Change that data!
- › Event: Something has happened
- › Component decides if data should be changed



Batch Replication

Batch

- › Get all data
- › Provide API
- › ...to decouple schema
- › Copy interesting data into local database

Billing

Order
Process

CRM

Additional
data

Shared
Kernel
Order

Additional
data

Shared
Kernel
Order

Additional
data

Shared
Kernel
Order

Batch

API

API

Batch

Batch & Data Replication

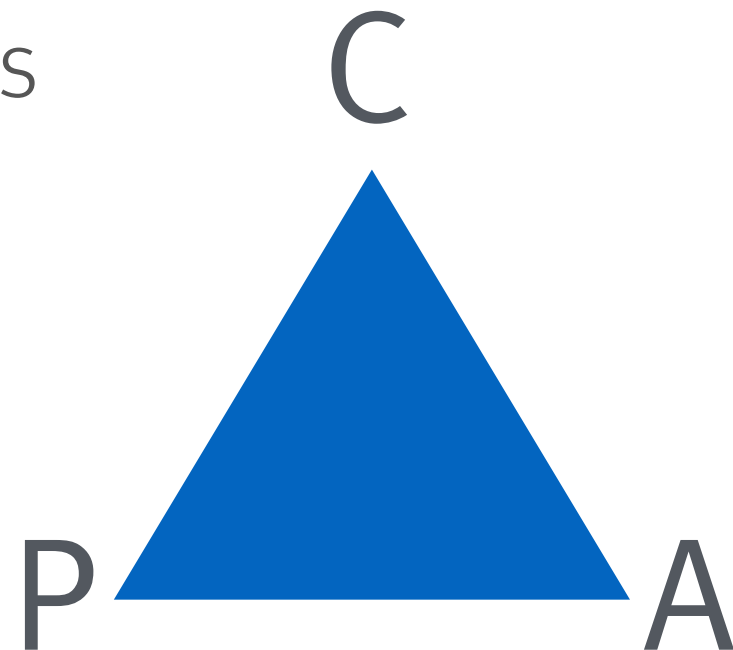
- › Easy to repair inconsistencies
- › Batch run at specific points
- › i.e. updates take time
- › Data not consistent across microservices

CAP: Challenge for Replication



CAP Theorem

- › Consistency
 - › All nodes see the same data
- › Availability
 - › Node failures do not prevent survivors from operating
- › Partition Tolerance
 - › System continues to operate despite arbitrary message loss



CAP Theorem: P

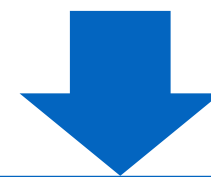
- › Network partitions do occur
- › Even with highly available network hardware
- › Also: very slow response = partition
- › Need to deal with P

CAP Theorem: C or A?

- › Node cannot access other nodes
- › Might have missed updates
- › A, not C:
Answer with a potentially wrong answer
- › C, not A:
Don't answer – the answer might be wrong



New Order Event



Billing



Order
Process



CRM



inconsistent

or unavailable

Additional
data

Shared
Kernel
Order

Additional
data

Shared
Kernel
Order

Additional
data

Shared
Kernel
Order



Conclusion

Classic:
Centralized
Database

Microservices:
private
database
decoupling

Data Microservices:
Consistent but
resilience / performance
/ transactions / decoupling?

Schema per
Microservice:
Simple infrastructure

Database per
Microservice:
Polyglot Persistence

Redundancy?

Redundant Data or
Bounded Context?

Context Map and
Context Relations

Replication

e.g. Shared
Kernel

Database
Replication

Batch

CQRS

Event
Sourcing

Events

CAP

Decentralize data!

EMail slideswjax2016@ewolff.com to get:

Slides

- + Microservices Primer
- + Sample Microservices Book
- + Sample of Continuous Delivery Book

Powered by Amazon Lambda & Microservices

