# Real-World Maintainable Software

Ten Coding Guidelines in Practice

Abraham Marín-Pérez

PEOPLECERT

```
ServletResponse resp)
xception, IOException {
ype("application/json");
                    try {
conn = DriverManager.
rty("handler.jdbcurl"));
    ResultSet results =
onn.createStatement()
       .executeQuery(
M ACCTS WHERE id="
req.getParameter(conf.
est.parametername")));
  float totalBalance = 0;
print("{\"balances\":[");
  while (results.next()) {    //
```

Better code is better business for everyone

Get it right with

Quality Software Development
Certificate in Maintainability

Save 10%*off
your exam fees

PROMOCODE: 5D6C91

* Offer expires 31th December 2016

3 easy steps to get certified!

1. Create your account at webcandidate.peoplecert.org
2. Select your exam voucher and enter PROMOCODE: 5D6C91
(offer valid until 31th December 2016)
3. Schedule and take your exam anytime, any place
(exam vouchers are valid for one year from the date of purchase)

Why get certified?

Benefits
For Organizations

- High quality code is cheaper to maintain. Initial development benefits from high maintainability from day 1
- Sustainable business needs maintainable software: Time-to-market is lower with higher maintainability
- Software quality is not left to chance: Agreement on code level specifics to achieve high quality, maintainable software

Benefits
For Individuals

- Software developers learn how to create high quality code
- Software developer teams can adopt best practices to produce maintainable software
- Software developers deliver demonstrably more value

Visit www.peoplecert.org or email communication@peoplecert.org

Real-World Maintainable Software

Ten Coding Guidelines in Practice

Abraham Marín-Pérez

Real-World Maintainable Software

by Abraham Marín-Pérez

July 2016: First Edition

Table of Contents

Foreword

It's already been 16 years that we at SIG have been helping our clients get their software systems right. For us, the journey so far has been—and remains—very rewarding, as we are privileged to work with visionary business leaders and talented software engineers.

This experience has left us with three important insights:

- Software is the DNA of our modern information society
- The moment that software gets noticed is the moment when it fails
- There are certain key but simple issues that lead again and again to the deterioration of software systems' ability to serve their intended purposes

By helping our clients address their software-related challenges, we've learned a thing or two about what it takes for a system to be future-proof and easy to maintain. We created our own model for assessing software maintainability. It's composed of a set of simple guidelines, which are based on the ISO 25010 standard and verified by TÜViT. We have consistently applied it at industry scale to more than 1,000 systems so far, across different technologies, industries, and geographies.

This report presents the theory behind our model, but it doesn't end there. Also included are client cases studies highlighting their benefits from operationalizing our model in their daily practice.

From our side, we hope you enjoy reading this report as much as we have been enjoying the SIG ride for all these years.

— Yiannis Kanellopoulos
Sr. Consultant
Software Improvement Group

Preface

Being the relatively young profession that it is, software development is still trying to figure out the best way to deliver. One of the most promising ideas of recent years comes from the software craftsmanship movement, which recommends small teams with attention to detail, risk aversion, and an appetite for continuous improvement. In teams like this, it is easy to be kept up-to-date with almost every aspect of the project, which means hidden traps and mistakes rarely go unnoticed for long. These teams consistently produce high-quality software that is easy to maintain.

Unfortunately, for better or worse, some organizations still need to manage large projects over long periods of time. In such environments, the principles of craftsmanship still apply, though one cannot hope to be kept up-to-date on every single aspect of the daily life of the project. Knowledge silos will appear, communication channels will decrease, and as a result it will be nearly impossible to assess whether staff are following a good set of best practices.

Many organizations have tried to fix this, especially from the point of view of project management. This is how, first, complex project management processes with certifications like PRINCE2 and, later, lighter processes with certifications like SCM came to be born. And, although both types of approaches achieved some level of success, they were both missing the technical side of things.

This is what initially motivated the Software Improvement Group (SIG) to create the "Ten Guidelines for Building Maintainable Software," included in the book *Building Maintainable Software* by Joost Visser (O'Reilly). The main risk of initiatives like this is that, as useful as they might seem, they could easily be archived in the depart-

ment of "Yet Another Nice Theory." This is why, in this report, I will explain how the guidelines can work in a real-life environment, considering the typical issues that every programmer faces during the course of a project, together with the hidden traps that programmers can fall into when trying to apply the Ten Guidelines.

Acknowledgments

I have always liked writing, ever since I was little. When I was in high school, I entered a regional narrative contest where I reached a modest yet satisfying third position. Then the Internet became popular and I started to write blogs and, more recently, technology news articles. This is why I was so excited when O'Reilly gave me the opportunity to write this report.

A project like this is never the product of a single person's efforts, and I'd like to thank those that have helped me along the way. First of all, I'd like to say thank you to Brian Foster, whose initial steering helped identify the best way for this report to serve its readers.

I'd also like to thank Nan Barber and Keith Conant for their reviews. Nan helped me make sure the report has a consistent style and structure, which has turned it into a more pleasant reading experience, while Keith's technical review improved the quality of the contents.

Last but not least, I'd like to thank my partner Bea for her patience and support while working on this report. Without her, this wouldn't have happened.

"How Did We Get into This Mess?"

Cape Canaveral, Florida. November 8, 1968. At precisely 9:46 a.m., the Delta E rocket ignites, propelling the Pioneer 9 spacecraft into the atmosphere. This is the fourth in a series of space missions directed at studying "space weather."

The program was highly successful: while designed to last for six months, it provided data for 35 years. The main contractor of the Pioneer 6-9 program was TRW, a company in charge not just of the construction of the spacecraft but also of the design and implementation of the software that governed it. This was during the relatively early days of the software development industry, and there weren't too many references on running software development projects. Perhaps for this reason, Winston W. Royce, one of TRW's most prominent software development project managers, published in 1970 a paper titled "Managing the Development of Large Software Systems," in which he described his views on making software projects succeed. Royce's paper was famously attributed as being the first written reference to the Waterfall Development Model, describing it as a "grandiose approach to software development."

This model took the world by storm. Companies all over the planet started to follow this methodology. Certifications were created for project managers who would be accredited as following the Waterfall Model to the letter. Teachers of computer science in universities of all countries included them in their lectures. For many decades, the Waterfall Model was adopted without question as the best and only possible way to develop software.
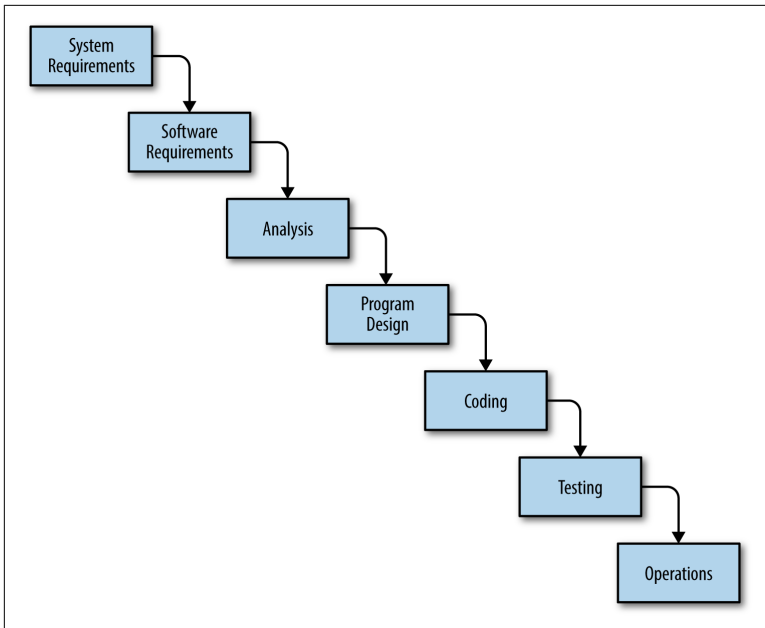
Figure 1-1. The Waterfall Development Model as described in Winston W. Royce's paper

However, in what may be a prophecy of the hunger for quick wins that would come to plague the software development industry, the early adopters of the Waterfall Model failed to properly read Royce's paper. Even though he described the Waterfall Model as the ideal way to build software, he also said that "the implementation described above is risky and invites failure." He then went on for several pages explaining the risks and downsides of his model, concluding that the only way to make it work is to run the same project at least twice so that subsequent implementations can learn from the mistakes of the previous ones.

And so it happened that software projects across the board consistently failed to meet expectations for decades. In 1994, the Standish Group published their first CHAOS report, a survey covering more than 8,000 applications, in which the overall success of software development projects throughout the industry was assessed. The results were abysmal: 31% of projects were cancelled before completion, 53% were finished with reduced functionality and/or over budget, and only 16% finished successfully according to the initial parameters.

Maybe because of these results, project managers began to worry about hitting deadlines above anything else, and potentially at the expense of long-term instability. This had the effect of increasing the costs of maintaining the code after delivery, with some sources like the IEEE estimating that maintenance causes around 60% of the total cost of the project.[1]

Step by step, failure after failure, the software development industry realized that something needed to be done differently. In 2001, the Agile Manifesto was signed, encouraging a new way of thinking about software development. Companies that began to deviate from the norm experienced the benefits. In the latest CHAOS report, results were segregated for Agile and Waterfall projects: while 29% of Waterfall-led projects are still cancelled (signifying virtually no improvement after 11 years), the failure rate of Agile-led projects is down to 9%.

It's taken time, but companies finally realize that, for a project to succeed, focus needs to be placed on the long term. Maintainability is the main issue, and for this to come at a reasonable cost it needs to be looked after from day one. It is for this reason that companies like SIG began to think about patterns and guidelines that can be applied to the everyday work of a software developer and that will assist in ensuring and assessing software maintainability. After years of experience, SIG has found a set of 10 easy-to-follow guidelines that will help keep code manageable for years to come, putting teams one step closer to success. This report will explore those guidelines, explain their applicability, and present them together with a set of real use cases that benefited from it.

You can start building maintainable code today by using these 10 guidelines.

---

1 Robert L. Glass, "Frequently Forgotten Fundamental Facts about Software Engineering" in *IEEE Software*, 2001.

The Ten Guidelines

After many years of failures, the software development industry is gradually coming to understand what makes projects succeed. Best practices have started to appear, sometimes mixed with misinformation and plain technical folklore, but through trial and error, teams around the world have started to separate the chaff from the grain.

SIG is one organization that has gone through this. And not only that, it has studied the way software development projects evolve so as to identify the difference between success and failure. After analyzing business cases for years, SIG has come up with Ten Guidelines that, if followed correctly, can make the difference between success or failure.

The Ten Guidelines are easy to understand but not necessarily easy to apply. Teams may face resistance on several fronts: sometimes from management, who may not understand the value of an investment like this, and sometimes from developers, who may take badly the fact that they are being told how to work best.

However, even if everybody is on board, the Ten Guidelines require that a level of technical expertise be applied. A lot of refactoring is needed to keep applying the guidelines overtime, and refactoring is an art that is very difficult to master. There are a number of resources that can be used to increased a developer's refactoring skills, among them the fantastic *How to Work Effectively with Legacy Code* by Michael Feathers (Prentice Hall).

In this chapter, we will briefly cover the Ten Guidelines and explain their usefulness; for more thorough coverage, the reader is advised to read *Building Maintainable Software*.

The first thing we need to note about the guidelines is that they are roughly divided into three categories. This isn't a strict division (in fact the categorization isn't present in the original source) but it's a useful way to look at the guidelines since different categories require different sets of skills:

Unit guidelines
> Write short units of code
> Write simple units of code
> Write code once
> Keep unit interfaces small

Architectural guidelines
> Separate concerns in modules
> Couple architecture components loosely
> Keep architecture components balanced
> Keep your codebase small

Enabling guidelines
> Automate tests
> Write clean code

Unit Guidelines

The first four guidelines are *unit guidelines*. In this context, a *unit* is the smallest group of code that can be executed independently; in object-oriented languages, a unit is a method within a class.

Write Short Units of Code

The first guideline indicates that our methods should be short, usually no more than 15 lines of code. This not only improves readability (fewer lines of code are easier to understand), but it also lowers the probability of hidden side effects. On top of this, a short method will have fewer variations, which means it will be easier to test.

The easiest way to apply this guideline is to move parts of the code in a method into other methods. Many IDEs will have an "extract method" function that makes this easier. Sometimes, however, the right answer is to move the code not to a different method but to a

new class—we'll see more of that when we get to the architectural guidelines.

Counting Lines of Code

Different teams may use different criteria when deciding what constitutes a line of code. In this report, we use the following:

- The signature and closing curly bracket of the method don't count. This is because these are lines that can't be removed and therefore have no bearing toward measuring the complexity of the method.

- Blank lines within the method do count. This is because, although blank lines don't have any instructions, programmers tend to add them to separate groups of lines that perform closely related tasks, which means they help indicate the complexity of the method.

- If an instruction is so long that it needs to be split into two or more lines, we count each of those lines independently. This is because we consider such instructions to represent extra complexity, and therefore it makes sense for them to contribute further to the total line count.

Choosing different criteria will obviously change the resulting number of lines, but the only effect will be that the triggering conditions for the guidelines will be met slightly sooner or later. In the end, a large method is a large method, regardless of how the lines are counted.

Let's take a look at an example. The following method contains 21 lines of code, more than the recommended limit of 15. It may not be clear what this particular method does, but that's not relevant at this point (it is part of a tool to analyze build data from a Jenkins server[1]).

```java
protected void selectBuilds(String source) {
    jenkinsClient = new JenkinsClient(source);
    List<String> allBuilds = jenkinsClient.getBuildConfigurations();
    BuildSelector buildSelector = new BuildSelector(allBuilds);
```

---

1 For further reference, see Build Hotspots on GitHub.

```
    GridPane grid = new GridPane();
    grid.setAlignment(Pos.CENTER);
    grid.setHgap(10);
    grid.setVgap(10);
    grid.setPadding(new Insets(25, 25, 25, 25));
    grid.add(buildSelector, 0, 0);

    Scene scene = new Scene(grid, 250, 400);
    m_primaryStage.setScene(scene);

    Button btn = new Button();
    btn.setText("Show me hotspots!");
    btn.setOnAction(event -> {
        List<String> selectedBuilds = buildSelector.getBuilds();
        AddDrawingToScene(selectedBuilds);
    });
    grid.add(btn, 0, 1);
}
```

We can make this method shorter by moving the creation and setup of the Grid and the Button objects into new methods, like this:

```
private GridPane createGridPane() {
    GridPane grid = new GridPane();
    grid.setAlignment(Pos.CENTER);
    grid.setHgap(10);
    grid.setVgap(10);
    grid.setPadding(new Insets(25, 25, 25, 25));

    return grid;
}

private Button createButton(BuildSelector buildSelector) {
    Button btn = new Button();
    btn.setText("Show me hotspots!");
    btn.setOnAction(event -> {
        List<String> selectedBuilds = buildSelector.getBuilds();
        AddDrawingToScene(selectedBuilds);
    });

    return btn;
}

protected void selectBuilds(String source) {
    jenkinsClient = new JenkinsClient(source);
    List<String> allBuilds = jenkinsClient.getBuildConfigurations();
    BuildSelector buildSelector = new BuildSelector(allBuilds);

    GridPane grid = createGridPane();
    grid.add(buildSelector, 0, 0);
```

```
    Scene scene = new Scene(grid, 250, 400);
    m_primaryStage.setScene(scene);

    Button btn = createButton(buildSelector);
    grid.add(btn, 0, 1);
}
```

Now we have three small methods as opposed to one big one, which makes the code easier to manipulate.

Write Simple Units of Code

The more paths of execution a method has, the more difficult it will be to reason about all of them. And when code is difficult to reason about, misunderstandings occur, and misunderstandings lead to bugs.

It's important to clarify, though, what a path of execution means. *Paths of execution* are branching points, instructions that can make the execution of the code go in one way or another. For instance, an `if` statement creates a branch of execution because, depending on the evaluation of a condition, different code will be executed. But not only that—if the condition in the `if` statement is a boolean operation involving several operators, the application of each operator will imply a new branch.

This guideline suggests that we limit branch points to a maximum of four. This will not only make the methods easier to understand but will also make them easier to test. In order to cover all different scenarios of a method, we need a number of automated tests that is at least the number of branch points plus one. Let's take a look at the following code:

```
public int getDiscount(String promoCode) {
    if(promoCode == null) {
        throw new IllegalArgumentException("promoCode");
    }

    promoCode = promoCode.trim();

    if(promoCode.length() < 5 || promoCode.length() > 8) {
        throw new IllegalArgumentException("promoCode");
    }

    if(expiredPromoCodes.containsKey(promoCode)) {
        throw new ExpiredPromoException(promoCode);
    }
```

```
    if(!availablePromoCodes.containsKey(promoCode)) {
        return 0;
    }

    return availablePromoCodes.get(promoCode);
}
```

As can probably be guessed, this method will provide the appropriate discount to apply depending on a promotional code, throwing exceptions in particular situations. This code has five branching points, which means there are six different scenarios to be considered while testing: promotional code being null, too short, or too long; promotion having expired; promotional code not matching any existing promotion (expired or not); and promotional code applied successfully.

We can reduce the number of branching points per unit by moving the validation logic to its own method, like the following:

```
public boolean isPromoCodeValid(String promoCode) {
    if(promoCode == null) {
        return false;
    }

    promoCode = promoCode.trim();

    if(promoCode.length() < 5 || promoCode.length() > 8) {
        return false;
    }

    return true;
}

public int getDiscount(String promoCode) {
    if(!isPromoCodeValid(promoCode)) {
        throw new IllegalArgumentException("promoCode");
    }

    promoCode = promoCode.trim();

    if(expiredPromoCodes.containsKey(promoCode)) {
        throw new ExpiredPromoException(promoCode);
    }

    if(!availablePromoCodes.containsKey(promoCode)) {
        return 0;
    }
```

```
    return availablePromoCodes.get(promoCode);
}
```

With the new version, we have two methods with three branching points each, which means we'll need four tests for each method to cover all cases. It may look as if we have more work to do now since we have a total of eight scenarios to cover, whereas before we only had six. However, analyzing effort this way can be deceiving. We don't *quite* have eight scenarios to cover; we have two sets of four scenarios each. This distinction is important because in software development, effort doesn't grow linearly with complexity—it grows exponentially. Therefore, it is easier to manage two sets of four scenarios each than one with six.

Write Code Once

Internet folklore has many ways to refer to this guideline, including "Stay DRY," with DRY being short for Don't Repeat Yourself, and "Don't get WET," with WET being short for Write Everything Twice. The truth is, there are so many ways to refer to this because this is one of the most powerful single sources of bugs.

It usually goes like this: A programmer, maybe due to time restrictions, decides to copy and paste a portion of code to make use of it somewhere else. Some time after that, a requirement arrives to modify that piece of code. The programmer that picks up this task, which might be the original one or a new one, doesn't remember or realize that the code that needs to be modified exists in two different places, so that programmer only applies changes to one of the copies of the code. And just like that, we have created a bug: two parts of the system that are meant to do the same thing no longer do.

But even if we manage duplication well and prevent bugs, duplicate code can still hurt a team. Whenever a task is performed, if programmers know that there is duplication in the codebase, they'll have to search for all the occurrences of the code that need to be modified and act on all of them appropriately; this is much more costly that having to change just one existing copy of the code.

The bottom line is, whenever you see duplicated code, you should refactor it into a single copy. Not only will you be saving yourself trouble and time, but also, in the process of refactoring the code, you may discover new domain concepts that fit within your overall design.

Keep Unit Interfaces Small

In the same way that *unit* means, in this context, a method in a class, *interface* here refers to the way we interact with a method; that is, the method signature. Methods with long signatures usually indicate the existence of *data clumps*: variables that always travel together and that in fact aren't particularly useful if used independently. Typical examples of data clumps are colors (expressed as their red, green, and blue components) and coordinates (expressed as their x,y components).

The way to make sure we keep interfaces small and detect these data clumps is by keeping method signatures to a maximum of four parameters. The way to apply this guideline is by bundling together two or more arguments into a new class, and then to use references to this new class. The interesting side effect is that now that we have a new class, we can start adding logic to it.

Let's consider a hotel room reservation system and, more precisely, a method to get quotes for specific rooms. Since we're only dealing with method signatures in this guideline, we won't include the body of the method:

```java
public Quote getQuote(String hotelName, RoomType roomType,
                      boolean breakfastIncluded,
                      LocalDate checkInDate,
                      LocalDate checkOutDate)
```

This signature has five parameters, one more than the guideline allows. We can fix this by bundling check-in and check-out dates into a `TimePeriod` class.

```java
public Quote getQuote(String hotelName, RoomType roomType,
                      boolean breakfastIncluded,
                      TimePeriod timePeriod) {
    // ... //
}

public class TimePeriod {
    public TimePeriod(LocalDate checkInDate,
                      LocalDate checkOutDate) {
        // ... //
    }
}
```

The interesting thing about the `TimePeriod` class is that we can easily add validation to it: ensure that check-out date is at least one day after check-in date, ensure that check-in date isn't in the past,

etc. And if we do that, we'll have validation for free whenever we need to use the pair of check-in and check-out dates. Thus, keeping unit interfaces small not only makes for simpler and more readable methods, it also helps us encapsulate concepts that better describe the domain at hand.

Architectural Guidelines

If the first four guidelines referred to characteristics that we need to measure at the unit (or method) level, the next four focus at a higher level—namely modules, components, and codebases.

In this context, a module is a collection of units; in other words, a class. Similarly, we can understand a component as an aggregation of modules (or classes) that can be grouped to offer a higher order of functionality. For many teams, a component will be something that can be developed, deployed, and/or installed independently, and will typically refer to them as a JAR file in Java, a DLL in .NET languages or, more generally, a library. However, some other teams with larger codebases will choose a bigger unit of measurement when defining what a component is, and will typically refer to them as frameworks or systems. The definition of the concept of a component will have an impact on the applicability of some of the guidelines, so teams should choose a definition carefully and potentially review it over time.

Finally, a codebase is a version-controlled collection of software; this typically means a repository in GIT-based systems, or an independent subfolder in Subversion or CVS-based systems.

As you will see, the architectural guidelines will apply to progressively broader aspects of the software, leaving behind the fine-grained details of the unit guidelines.

Separate Concerns in Modules

Modules, or classes, are meant to be representations of domain concepts; if you can't explain what a class does in a couple of simple sentences, then that class either represents more than one concept or represents a concept that is too general or abstract.

A class that holds too much responsibility will be troublesome in several ways. First, it is likely to become a *change hotspot*. Since it has so many responsibilities, it will affect a large proportion of the

business logic, and therefore the probability that it needs to be modified upon any new request will be high. Change hotspots create long (and difficult to browse) change logs and increase the probability of clashes between programmers, potentially disrupting the natural team flow.

Second, big classes have the risk of becoming a dumping ground for difficult design decisions. When new functionality needs to be added to a system and programmers are unsure about where that new functionality should go, it is not uncommon for people to choose an existing big class whose purpose is not entirely clear anyway.

Finally, big classes that concentrate a lot of logic in one place will be highly utilized by other classes in one way or another. This means we are creating a class with a high degree of change (and therefore a higher risk of accidents) and high exposure (and therefore a higher impact on accidents), and creating areas of code with high risk and impact is never a good idea.

Deciding how well a team is applying "Separate Concerns in Modules" is a little subjective. The general suggestion is to try and apply the Single Responsibility Principle, for which there is plenty of documentation— although even then some people may argue whether a particular scenario represents one single but complex principle or two independent but related principles. Some heuristics that can help are the size of the class (beyond 100 lines seems suspicious for a single principle) or the rate of public versus private methods (too many private methods may expose complexities that belong somewhere else). However, each team will have to decide what their own metrics are and how they are to be applied.

Couple Architecture Components Loosely

This guideline is similar to the previous one, but it is applied at an even higher level. With "Separate Concerns in Modules," we tried to limit the responsibilities of a class so as to limit the dependencies upon it. With "Couple Architecture Components Loosely," we try to do the same but with regards to components.

Like it happened with the previous guideline, it's a bit difficult to establish general parameters that highlight when architecture components are loosely coupled and when they aren't; the final decision

may be different from team to team. However, there are some general principles that can be applied.

First, we can draw a diagram of all the different components in our system and connect them to represent their dependencies. With this kind of diagram, we can look for components that accumulate too many incoming dependencies. For instance, in the following diagram we can see how component A is tightly coupled with the rest of the architecture, while component B isn't. Modifying component A can have repercussions in almost every other component of the system. This turns modifying component A into a risky affair, which makes it more difficult to maintain. Component B, on the other hand, has a much lower potential impact, which makes it more maintainable. In this situation, we probably should look into splitting component A into smaller, less tightly coupled components.

Figure 2-1. A component dependency diagram showing a tightly coupled component (A) and a loosely coupled one (B)

Second, we can analyze how much of each component is being exposed. For instance, if we were considering a component to be a JAR file, we would check which classes within the component are being accessed when there are calls from other components. Ideally, the exposed portion of the component will be as small as possible, since exposed classes cannot be modified safely without impact to dependent components. As an example, changing the signature of a method that is being called from a different component will

instantly cause a compilation failure unless the caller is updated at the same time; if there are many callers, this may be impractical.

Unfortunately, depending on the language there may not be an easy way to analyze the proportion of a component that is being exposed. In C#, the developer can use different access modifiers for classes that are to be available only within the component (internal) or from everywhere (public). Java, however, doesn't currently have this capability, although the new Module System in plan for Java 9 will provide it. This means that, once again, how this guideline is applied will depend on each particular team.

Keep Architecture Components Balanced

As systems grow, it may become too easy to get lost in the many facets of it, and when this happens it may become easy to miss systemic issues. This is why software needs to make sense also from a high-level point of view.

Although this sounds like a rather subjective measurement, there are objective metrics that we can apply to assess how balanced our architecture is. First of all, we can intuitively conclude that an architecture with too few components can't really describe the multiple features of a system, whereas one with too many of them will be difficult to grasp. For this reason, SIG recommends architectures designed around having nine components, with an operating margin of plus/minus three. This means that, if we have fewer than six components, we should probably look into splitting some of them, whereas if we have more than 12, we should try to consolidate them. When components are consolidated, teams may need to revisit their definition of "component" so it maps to a bigger unit of measurement.

On the other hand, the relative size of the components is also important. There probably isn't much we can infer from an architecture containing one really big component and eight tiny ones: the latter will probably be minor utility libraries, whereas the former will hold further substructures and divisions that are kept hidden. Components in an architecture should be as close in size as possible.

For most scenarios, the matter of assessing the size of a component can be as simple as counting the number of lines of code, and indeed this is what SIG recommends. However, there is an increasing number of organizations that, thanks to the dynamic capabilities pro-

vided by the Java Virtual Machine, develop different components using different JVM-compatible languages. In cases like this, counting the number of lines may be misleading, and teams may choose to count the number of files (as a proxy for the number of classes; therefore, of "concepts") or even metrics not directly related to source code, like build duration.

Keep Your Codebase Small

The vast majority of programmers will maintain that smaller codebases are easier to manage. This may sound like software development folklore, but after analyzing over 1,500 systems, SIG provided statistical evidence of it. On top of this, structuring the overall systems in several, smaller codebases as opposed to a single big one will simplify some higher-order administrative tasks: for instance, if a team needs to be split as part of an organizational restructuring, the responsibilities of the new teams can be easily decided by assigning the different codebases to each of them.

It may seem that, after applying the component guidelines above, the size of the codebase has already been taken care of. This is not necessarily true, since one could be splitting components without splitting the associated codebase. For instance, the build tool Maven, popular among Java programmers, allows the management and creation of multiple JAR files within a single project, and therefore a single codebase. In fact, given that splitting components is often easier than splitting codebases, it is advisable to try and apply this guideline by preventing unnecessary growth.

There are many things that can be done to prevent a codebase growing too big. Applying the guideline "Write Code Once" is one of them. "Couple Architecture Components Loosely" can help too, since essential components with a high number of incoming dependencies, like logging or serialization, tend to have a publicly available counterpart that could be used instead. Removing unused code is another obvious way to reduce the size of the codebase.

However, one of the most effective ways to achieve this (but frequently also the most difficult to apply) is to avoid future-proofing. *Future-proofing* is the practice of adding functionality to the codebase that hasn't been required but that people believe *might* be required at some point in the future. One common example of this is unsolicited performance optimization utilities, like caching or

connection pooling, added in case workload volumes grow to unmanageable levels. While performance is a valid concern in some situations, more often than not optimizations are added without the backup of actual data.

How to decide when a codebase is too big depends on the technology at hand, since some languages are more verbose than others. For instance, SIG sets the limit for Java-based systems at 175,000 lines of code, but this may be too much or too little when using other languages. In any case, it is important to note that the number of lines of code is just a proxy to calculate the real variable: the amount of knowledge, functionality, and effort that is contained within the codebase.

Enabling Guidelines

Up to now, we've talked about guidelines that express how the code should be structured, both at the low and high level. Also, we've talked about how the code needs to be modified whenever any of the previous guidelines isn't met. But so far we haven't talked about the inherent risk of the very application of the guidelines.

To be able to abide by the guidelines, we need to be constantly making changes to our code. However, with every change there is the risk of accidents, which means every corrective measure is an opportunity for a bug to appear. In order to apply the guidelines safely, the code needs to be easy to understand and difficult to break. The last two guidelines address these concerns, which is why I like referring to them as enabling guidelines.

Automate Tests

Automated tests require a considerable amount of investment, but the returns are almost invaluable. Automated tests can be rerun again and again almost at no cost, which means we can frequently reassess the status of the system. Thanks to this, we can safely perform as many modifications as we deem necessary, knowing that accidents that break it won't go unnoticed for long. There is plenty of literature on the topic of how to write automated tests, so we won't cover that here.

What we will cover is how to make sure that our set of automated tests is a faithful representation of the requirements of the system.

On one side, we can use a technique called Test-Driven Development (TDD), in which tests are written *before* the implementation they are testing. Moreover, when writing implementation, the programmer has to write only as little as necessary to satisfy the tests at hand. This will ensure that, if a functionality needs to be implemented, a test that verifies that functionality will be written first.

Even if we don't use TDD, or if we inherit a legacy codebase, there are ways to assess the quality of the automated tests suite. We can use code coverage analysis tools, which tag the lines of code that are being executed while running each test; this way, the tool can highlight code that isn't exercised by any test and which may constitute a gap in the testing suite. Unfortunately, code coverage analysis is not enough: one thing is saying that a particular section of code is executed while running a test, and another one is saying that the test performs the right checks after running. To go one step further, we can use tools like Jester for Java or Nester for .NET. Jester will make random modifications to the source code and then run the tests, expecting at least one of them to fail; if none of the tests fail after modifying the code, it will highlight this as a potential test gap.

Note, however, that automated tests don't entirely remove the need for manual tests. Certainly, repetitive tests that can be easily scripted don't need to be manual anymore, but qualitative assertions like usability or penetration tests will probably still need to be performed manually by an expert in the area.

Write Clean Code

Edsger Dijkstra, one of the fathers of modern programming, once said that "in programming, elegance is not a dispensable luxury but a quality that decides between success and failure". Obscure code is difficult to grasp, and therefore difficult to maintain. On the other hand, code that can be picked up by any random programmer with no problem represents the ultimate measure of maintainability.

The problem of clean code is that it's a rather ambiguous topic. It's not simply a matter of code style, indentations, or variable name lengths, it's a matter of understandability. Code should clearly show its intent and avoid any pitfalls or misgivings, so that anybody who reads it will understand it quickly and be able to modify it without making inadvertent mistakes. However, what exactly constitutes the

right set of rules to decide if the code at hand is clean or not is something that has to be agreed on by the team.

To help with this, SIG provides a basic yet functional set of rules to write clean code as part of this guideline; these rules are a great starting point for any team. If further information is needed, readers can refer to the book *Clean Code* by Robert C. Martin (Prentice Hall). Another technique that can help teams to write clean code is doing unguided but supervised peer reviews: after writing code, a programmer can share it with a peer, who will read the code and try to understand it without help; any questions that the peer needs to ask in order to understand the code represents an area where the code isn't self-evident, and therefore an area that needs to be modified.

Applying the Ten Guidelines

Apply All the Guidelines

Important as it is that these guidelines be applied to projects from day one, it is even more important that they are applied in their entirety. Applying only some of the guidelines may be even worse than not applying them at all, for this may cause a dangerous side effect: it could lead developers to think that the code is in a better state than it actually is. Let's see how this applies in practice by assuming we are happy to apply the guidelines "Write Short Units of Code" and "Write Code Once," but not "Keep Unit Interfaces Small," and see what happens when we run the code.

Applying "Write Shorts Units of Code" and "Write Code Once"

Let's assume that we are building a new system. The purpose of the system is not important right now, except to mention that this system will need a way to manage users. The first thing we create is a `User` class that simply holds the first, last, and middle name of the user.

```java
public class User {
    private final String firstName;
    private final String middleName;
    private final String lastName;

    public User(String firstName, String middleName,
            String lastName) {
```

```
        this.firstName = firstName;
        this.middleName = middleName;
        this.lastName = lastName;
    }
}
```

Adding Validation

The first validation requirement then arrives: while a middle name is optional, the first and last names are mandatory, non-blank values. We add some validation to the constructor.

```java
public class User {
    private final String firstName;
    private final String middleName;
    private final String lastName;

    public User(String firstName, String middleName,
            String lastName) {
        if (firstName == null || firstName.trim().length() == 0) {
            throw new IllegalArgumentException("firstName");
        }
        if (lastName == null || lastName.trim().length() == 0) {
            throw new IllegalArgumentException("lastName");
        }

        this.firstName = firstName;
        this.middleName = middleName;
        this.lastName = lastName;
    }
}
```

It seems clear that the validation of both first and last name is essentially the same, which means we should apply the guideline "Write Code Once" and refactor that code into its own method (showing only the constructor and the resulting new method).

```java
public User(String firstName, String middleName, String lastName) {
    ensureNotEmpty(firstName, "firstName");
    ensureNotEmpty(lastName, "lastName");

    this.firstName = firstName;
    this.middleName = middleName;
    this.lastName = lastName;
}

private void ensureNotEmpty(String field, String fieldName) {
    if (field == null || field.trim().length() == 0) {
        throw new IllegalArgumentException(fieldName);
    }
}
```

Adding Contact Details

That looks like reasonable code. At this point, the next requirement arrives: we also need the user's email address, home number, and mobile number in case we need to contact her. The resulting modified class follows.

```java
public class User {
    private final String firstName;
    private final String middleName;
    private final String lastName;
    private final String homeNumber;
    private final String mobileNumber;
    private final String emailAddress;

    public User(String firstName, String middleName,
            String lastName, String homeNumber, String mobileNumber,
            String emailAddress) {
        ensureNotEmpty(lastName, "lastName");
        ensureNotEmpty(firstName, "firstName");

        this.firstName = firstName;
        this.middleName = middleName;
        this.lastName = lastName;
        this.homeNumber = homeNumber;
        this.mobileNumber = mobileNumber;
        this.emailAddress = emailAddress;
    }

    private void ensureNotEmpty(String field, String fieldName) {
        if (field == null || field.trim().length() == 0) {
            throw new IllegalArgumentException(fieldName);
        }
    }
}
```

Validation is also required for these fields, although the validation rules differ. In this case, email address is mandatory, and we need *at least* one of home or mobile number; for simplicity, we'll assume at this point that we don't need to check the format of the strings representing the phone numbers and email address, although in a real-life scenario we would. Here is the resulting constructor:

```java
public User(String firstName, String middleName,
        String lastName, String homeNumber, String mobileNumber,
        String emailAddress) {
    ensureNotEmpty(lastName, "lastName");
    ensureNotEmpty(firstName, "firstName");
    ensureNotEmpty(emailAddress, "emailAddress");
```

```
    if ((homeNumber == null || homeNumber.trim().length() == 0)
            && (mobileNumber == null
                || mobileNumber.trim().length() == 0)) {
        String message = "homeNumber or mobileNumber is needed";
        throw new IllegalArgumentException(message);
    }

    this.firstName = firstName;
    this.middleName = middleName;
    this.lastName = lastName;
    this.homeNumber = homeNumber;
    this.mobileNumber = mobileNumber;
    this.emailAddress = emailAddress;
}
```

The check for an empty field is something that we can refactor applying "Write Code Once," which would leave the class with the following constructor and methods:

```
public User(String firstName, String middleName,
        String lastName, String homeNumber, String mobileNumber,
        String emailAddress) {
    ensureNotEmpty(lastName, "lastName");
    ensureNotEmpty(firstName, "firstName");
    ensureNotEmpty(emailAddress, "emailAddress");

    if (isBlank(homeNumber) && isBlank(mobileNumber)) {
        String message = "homeNumber or mobileNumber is needed";
        throw new IllegalArgumentException(message);
    }

    this.firstName = firstName;
    this.middleName = middleName;
    this.lastName = lastName;
    this.homeNumber = homeNumber;
    this.mobileNumber = mobileNumber;
    this.emailAddress = emailAddress;
}

private void ensureNotEmpty(String field, String fieldName) {
    if (isBlank(field)) {
        throw new IllegalArgumentException(fieldName);
    }
}

private boolean isBlank(String field) {
    return field == null || field.trim().length() == 0;
}
```

The constructor is now 15 lines long, just on the verge of the limit established by "Write Short Units of Code" but not quite enough to

force a rewrite, and there isn't any other repetition worth refactoring, which means the code is good considering the guidelines we are applying.

Adding Address

We are ready to implement our next requirement: users also need a full address, which is made up of two lines of address, a city, a postcode, and a country. The second line of address is optional, but all other fields are mandatory; again, for simplicity, we'll assume that we don't need to validate the format of the values, just that these are there. The resulting constructor follows.

```java
public User(String firstName, String middleName, String lastName,
            String homeNumber, String mobileNumber,
            String emailAddress, String addressLine1,
            String addressLine2, String city, String postcode,
            String country) {
    ensureNotEmpty(lastName, "lastName");
    ensureNotEmpty(firstName, "firstName");
    ensureNotEmpty(emailAddress, "emailAddress");
    ensureNotEmpty(addressLine1, "addressLine1");
    ensureNotEmpty(city, "city");
    ensureNotEmpty(postcode, "postcode");
    ensureNotEmpty(country, "country");

    if (isBlank(homeNumber) && isBlank(mobileNumber)) {
        String message = "homeNumber or mobileNumber is needed";
        throw new IllegalArgumentException(message);
    }

    this.firstName = firstName;
    this.middleName = middleName;
    this.lastName = lastName;
    this.homeNumber = homeNumber;
    this.mobileNumber = mobileNumber;
    this.emailAddress = emailAddress;
    this.addressLine1 = addressLine1;
    this.addressLine2 = addressLine2;
    this.city = city;
    this.postcode = postcode;
    this.country = country;
}
```

Now this constructor is 24 lines long, which is more than what the guideline "Write Short Units of Code" allows. We can make it shorter by moving all the validation logic into a new private method, resulting in the following code.

```java
public User(String firstName, String middleName,
            String lastName, String homeNumber,
            String mobileNumber, String emailAddress,
            String addressLine1, String addressLine2,
            String city, String postcode, String country) {
    validateInputs(firstName, lastName, homeNumber,
            mobileNumber, emailAddress, addressLine1, city,
            postcode, country);

    this.firstName = firstName;
    this.middleName = middleName;
    this.lastName = lastName;
    this.homeNumber = homeNumber;
    this.mobileNumber = mobileNumber;
    this.emailAddress = emailAddress;
    this.addressLine1 = addressLine1;
    this.addressLine2 = addressLine2;
    this.city = city;
    this.postcode = postcode;
    this.country = country;
}

private void validateInputs(String firstName, String lastName,
                            String homeNumber,
                            String mobileNumber,
                            String emailAddress,
                            String addressLine1, String city,
                            String postcode, String country) {
    ensureNotEmpty(lastName, "lastName");
    ensureNotEmpty(firstName, "firstName");
    ensureNotEmpty(emailAddress, "emailAddress");
    ensureNotEmpty(addressLine1, "addressLine1");
    ensureNotEmpty(city, "city");
    ensureNotEmpty(postcode, "postcode");
    ensureNotEmpty(country, "country");

    if (isBlank(homeNumber) && isBlank(mobileNumber)) {
        String message = "homeNumber or mobileNumber is needed";
        throw new IllegalArgumentException(message);
    }
}
```

Now both the constructor and the private validation method are shorter than 15 lines, so according to the guidelines we are applying we are good to go. However, by this point we can already feel that this code is not quite right. Let's take a look at the entire resulting class.

```java
public class User {
    private final String firstName;
    private final String middleName;
```

```java
    private final String lastName;
    private final String homeNumber;
    private final String mobileNumber;
    private final String emailAddress;
    private final String addressLine1;
    private final String addressLine2;
    private final String city;
    private final String postcode;
    private final String country;

    public User(String firstName, String middleName,
                String lastName, String homeNumber,
                String mobileNumber, String emailAddress,
                String addressLine1, String addressLine2,
                String city, String postcode, String country) {
        validateInputs(firstName, lastName, homeNumber, mobileNumber,
                emailAddress, addressLine1, city, postcode, country);

        this.firstName = firstName;
        this.middleName = middleName;
        this.lastName = lastName;
        this.homeNumber = homeNumber;
        this.mobileNumber = mobileNumber;
        this.emailAddress = emailAddress;
        this.addressLine1 = addressLine1;
        this.addressLine2 = addressLine2;
        this.city = city;
        this.postcode = postcode;
        this.country = country;
    }

    private void validateInputs(String firstName, String lastName,
                                String homeNumber,
                                String mobileNumber,
                                String emailAddress,
                                String addressLine1, String city,
                                String postcode, String country) {
        ensureNotEmpty(lastName, "lastName");
        ensureNotEmpty(firstName, "firstName");
        ensureNotEmpty(emailAddress, "emailAddress");
        ensureNotEmpty(addressLine1, "addressLine1");
        ensureNotEmpty(city, "city");
        ensureNotEmpty(postcode, "postcode");
        ensureNotEmpty(country, "country");

        if (isBlank(homeNumber) && isBlank(mobileNumber)) {
            String message = "homeNumber or mobileNumber is needed";
            throw new IllegalArgumentException(message);
        }
    }
```

```
    private void ensureNotEmpty(String field, String fieldName) {
        if (isBlank(field)) {
            throw new IllegalArgumentException(fieldName);
        }
    }

    private boolean isBlank(String field) {
        return field == null || field.trim().length() == 0;
    }
}
```

By this point, we have a class with no public methods (other than
the constructor) and three private methods; this gives away the feel-
ing that this class has enough complexity within it that it needs to be
sorted out in different methods, although the functionality of these
methods is not useful outside. On the other hand, the signatures of
some of those methods (the constructor and the validation method)
are so large that they span several lines. We can feel something is
missing, and that something is the application of the guideline
"Keep Unit Interfaces Small." Let's see what this class would look like
if we had applied this guideline.

Applying the Guideline "Keep Unit Interfaces Small"

"Keep Unit Interfaces Small" instructs us to keep method signatures
to a maximum of four parameters. Looking back at our example, we
didn't violate this guideline until we added the contact details, so
we'll pick it up from there.

Adding Contact Details

At this point, the constructor had six parameters and, counting vali-
dation logic, it was 15 lines long. We'll copy it next for quick refer-
ence.

```
public User(String firstName, String middleName, String lastName,
            String homeNumber, String mobileNumber,
            String emailAddress) {
    ensureNotEmpty(lastname, "lastName");
    ensureNotEmpty(firstName, "firstName");
    ensureNotEmpty(emailAddress, "emailAddress");

    if (isBlank(homeNumber) && isBlank(mobileNumber)) {
        String message = "homeNumber or mobileNumber is needed";
        throw new IllegalArgumentException(message);
    }

    this.firstName = firstName;
```

```
        this.middleName = middleName;
        this.lastName = lastName;
        this.homeNumber = homeNumber;
        this.mobileNumber = mobileNumber;
        this.emailAddress = emailAddress;
    }
```

We can bundle up all the contact details into a class of its own, ContactDetails. After this, the constructor of User will only have four parameters. Now, it makes sense that the validation of homeNumber, mobileNumber, and emailAddress is moved to the constructor of ContactDetails, which would result in the following code.

```java
public class User {
    private final String firstName;
    private final String middleName;
    private final String lastName;
    private final ContactDetails contactDetails;

    public User(String firstName, String middleName,
                String lastName, ContactDetails contactDetails) {
        ensureNotEmpty(lastName, "lastName");
        ensureNotEmpty(firstName, "firstName");

        if(contactDetails == null) {
            throw new IllegalArgumentException("contactDetails");
        }

        this.firstName = firstName;
        this.middleName = middleName;
        this.lastName = lastName;
        this.contactDetails = contactDetails;
    }

    private void ensureNotEmpty(String field, String fieldName) {
        if (isBlank(field)) {
            throw new IllegalArgumentException(fieldName);
        }
    }

    private boolean isBlank(String field) {
        return field == null || field.trim().length() == 0;
    }
}

public class ContactDetails {
    private final String homeNumber;
    private final String mobileNumber;
    private final String emailAddress;
```

```java
    public ContactDetails(String homeNumber, String mobileNumber,
                          String emailAddress) {
        ensureNotEmpty(emailAddress, "emailAddress");

        if (isBlank(homeNumber) && isBlank(mobileNumber)) {
            String message = "homeNumber or mobileNumber is needed";
            throw new IllegalArgumentException(message);
        }

        this.homeNumber = homeNumber;
        this.mobileNumber = mobileNumber;
        this.emailAddress = emailAddress;
    }

    private void ensureNotEmpty(String field, String fieldName) {
        if (isBlank(field)) {
            throw new IllegalArgumentException(fieldName);
        }
    }

    private boolean isBlank(String field) {
        return field == null || field.trim().length() == 0;
    }
}
```

This refactoring complies with "Keep Unit Interfaces Small" but not "Write Code Once" because we had to duplicate the validation helper methods in both classes. To avoid this duplication, we can move these methods into a third class, which we can call `Validator`; given that `Validator` doesn't need to handle any internal status, the new methods can be created as static. The three classes will now look pretty tidy.

```java
public class Validator {
    public static void ensureNotEmpty(String field,
                                      String fieldName) {
        if (isBlank(field)) {
            throw new IllegalArgumentException(fieldName);
        }
    }

    public static boolean isBlank(String field) {
        return field == null || field.trim().length() == 0;
    }
}

public class User {
    private final String firstName;
    private final String middleName;
    private final String lastName;
```

```java
    private final ContactDetails contactDetails;

    public User(String firstName, String middleName,
                String lastName, ContactDetails contactDetails) {
        Validator.ensureNotEmpty(lastName, "lastName");
        Validator.ensureNotEmpty(firstName, "firstName");

        if(contactDetails == null) {
            throw new IllegalArgumentException("contactDetails");
        }

        this.firstName = firstName;
        this.middleName = middleName;
        this.lastName = lastName;
        this.contactDetails = contactDetails;
    }
}
public class ContactDetails {
    private final String homeNumber;
    private final String mobileNumber;
    private final String emailAddress;

    public ContactDetails(String homeNumber, String mobileNumber,
                          String emailAddress) {
        Validator.ensureNotEmpty(emailAddress, "emailAddress");

        if (Validator.isBlank(homeNumber)
                && Validator.isBlank(mobileNumber)) {
            String message = "homeNumber or mobileNumber is needed";
            throw new IllegalArgumentException(message);
        }

        this.homeNumber = homeNumber;
        this.mobileNumber = mobileNumber;
        this.emailAddress = emailAddress;
    }
}
```

Adding Address

By this point, we can already see the effects of applying "Keep Unit Interfaces Small," but let's continue the example until the end. Let's add the address to the User class, leaving the following constructor:

```java
public User(String firstName, String middleName, String lastName,
            ContactDetails contactDetails, String addressLine1,
            String addressLine2, String city, String postcode,
            String country) {
    Validator.ensureNotEmpty(lastName, "lastName");
    Validator.ensureNotEmpty(firstName, "firstName");
```

```
    Validator.ensureNotEmpty(addressLine1, "addressLine1");
    Validator.ensureNotEmpty(city, "city");
    Validator.ensureNotEmpty(postcode, "postcode");
    Validator.ensureNotEmpty(country, "country");

    if(contactDetails == null) {
        throw new IllegalArgumentException("contactDetails");
    }

    this.firstName = firstName;
    this.middleName = middleName;
    this.lastName = lastName;
    this.contactDetails = contactDetails;
    this.addressLine1 = addressLine1;
    this.addressLine2 = addressLine2;
    this.city = city;
    this.postcode = postcode;
    this.country = country;
}
```

This constructor violates two guidelines, "Write Short Units of Code" and "Keep Unit Interfaces Small." We could bundle up all the new address fields into a new class, Address, and move the validation of relevant fields to that class. However, that would leave the constructor with five parameters, which would still violate "Keep Unit Interfaces Small." To fix this, we will also bundle the name-related attributes into a new class, FullName.

This would be the new FullName class, reusing the existing Validator:

```
public class FullName {
    private final String firstName;
    private final String middleName;
    private final String lastName;

    public FullName(String firstName, String middleName,
                    String lastName) {
        Validator.ensureNotEmpty(lastName, "lastName");
        Validator.ensureNotEmpty(firstName, "firstName");

        this.firstName = firstName;
        this.middleName = middleName;
        this.lastName = lastName;
    }
}
```

This would be the new `Address` class:

```java
public class Address {
    private final String addressLine1;
    private final String addressLine2;
    private final String city;
    private final String postcode;
    private final String country;

    public Address(String addressLine1, String addressLine2,
                   String city, String postcode, String country) {
        Validator.ensureNotEmpty(addressLine1, "addressLine1");
        Validator.ensureNotEmpty(city, "city");
        Validator.ensureNotEmpty(postcode, "postcode");
        Validator.ensureNotEmpty(country, "country");

        this.addressLine1 = addressLine1;
        this.addressLine2 = addressLine2;
        this.city = city;
        this.postcode = postcode;
        this.country = country;
    }
}
```

And, finally, this would be the end state of the `User` class:

```java
public class User {
    private final FullName fullName;
    private final ContactDetails contactDetails;
    private final Address address;

    public User(FullName fullName, ContactDetails contactDetails,
                Address address) {
        if (fullName == null) {
            throw new IllegalArgumentException("fullName");
        }

        if (contactDetails == null) {
            throw new IllegalArgumentException("contactDetails");
        }

        if (address == null) {
            throw new IllegalArgumentException("address");
        }

        this.fullName = fullName;
        this.contactDetails = contactDetails;
        this.address = address;
    }
}
```

The `User` class can be further improved by applying "Write Code Once" and removing duplication, which would leave it as follows:

```java
public class User {
    private final FullName fullName;
    private final ContactDetails contactDetails;
    private final Address address;

    public User(FullName fullName, ContactDetails contactDetails,
                Address address) {
        ensureNotNull(fullName, "fullName");
        ensureNotNull(contactDetails, "contactDetails");
        ensureNotNull(address, "address");

        this.fullName = fullName;
        this.contactDetails = contactDetails;
        this.address = address;
    }

    private void ensureNotNull(Object fieldValue,
                               String fieldName) {
        if (fieldValue == null) {
            throw new IllegalArgumentException(fieldName);
        }
    }
}
```

If we now compare the end result with the one we obtained without applying the guideline "Keep Unit Interfaces Small," we can see a big difference. In the former case, at the end of the coding exercise we only had one class, `User`, which was clearly bigger than desirable and had too many variables and moving parts; it's easy to see that, if the first version of `User` were to be modified, mistakes could be inadvertently made. However, just by applying "Keep Unit Interfaces Small," the end result was dramatically different: instead of one class, we had five, all of them small and obvious in their own way. If a programmer was to make a mistake in one of these five classes, it would stand out much more clearly and therefore the probability of somebody catching it would be higher.

This example illustrates just how important it is that the Ten Guidelines are not cherry-picked. They were designed to complement each other, and their benefits can only be achieved if applied together; as we just saw, leaving out even only one of them could have a dramatic effect on the quality of the end product.

Getting Value from the Ten Guidelines

There is a common joke in software development that says "the first 90% of a software development project costs as much as the other 90%." What this joke shows is that while the vast majority of requirements in a project tend to be relatively straightforward, there is always a small fraction of requirements that represent unusual edge cases and abnormalities that are incredibly hard to code, up to the point that this small fraction requires as much energy and dedication as the rest of the project, frequently running the project over budget.

The same can happen when applying the Ten Guidelines to our software. If we attempt to apply all the guidelines strictly to the letter, we'll find ourselves spending a lot of time and effort on a few outliers. However, while we rarely have the option to drop requirements that are too difficult to implement, we do have the option to give ourselves some room for flexibility when applying the guidelines. We must remember that the Ten Guidelines are here to help us make projects more maintainable in the future, and therefore to help us reduce the cost of writing software. If we let the guidelines govern all of our decision-making, we'll find ourselves rewriting the code so much that we'll drive the cost up instead of down.

There needs to be, therefore, a threshold of tolerance when applying the guidelines to make sure we get value out of them. What's more, the particular threshold may be different for different teams, or even for the same team over time. For instance, a team of junior programmers may have a hard time refactoring code to make it fully compliant, which means a higher tolerance may be needed to keep a

good balance of short- and long-term benefits; however, an experienced team with a higher knowledge of design patterns and other refactoring techniques can manipulate the code much more effectively, which means we can aim at a higher level of quality. Similarly, a team that inherits some legacy codebase may need to start with a high tolerance when applying the guidelines so as to avoid paralysis by refactoring, although as the code is improved, this tolerance can be progressively reduced.

SIG, the creators of the Ten Guidelines, have embedded this wisdom into the SIG/TÜViT Evaluation Criteria (or EC for short), its quality-scoring system for code. EC is based on the Ten Guidelines but includes a number of tolerance thresholds that are benchmarked and calibrated every year against a set of software systems across industries. This allows teams to monitor the quality of their codebase in an effective manner, correcting those violations that could cause a greater impact to the code but leaving the harmless ones.

Flexibility is a double-edged sword, though. Every project experiences moments of pressure when deadlines loom nearer than one would desire, and in those moments stakeholders may be tempted to argue that the guidelines are just slowing the team down and that higher flexibility is needed. This usually causes unnecessary confrontation between stakeholders and team members. The important thing to note in these situations is that those arguing in favor of relaxing the application of the guidelines do so not because they don't appreciate the long-term benefits these provide but because they haven't been made aware of their value. For organizations to fully benefit from the Ten Guidelines, everybody needs to be fully aware of them: the Ten Guidelines have to be instilled into the core values of the organization.

Not Too Much, Not Too Little: Just Right

The previous section showed how being too strict about the Ten Guidelines could grind the team to a halt, whereas being too lax could be the equivalent of cutting corners. However, reaching the right balance can only be done by taking into account many other variables beyond the guidelines. In other words, while the guidelines need to be applied in their entirety to be effective, they also need to be applied within a context.

SIG, with its EC scoring system, provides a way to measure how much a codebase is adhering to the guidelines. But the score is only the starting point. Whenever changes in a codebase make the score fall below the agreed-upon threshold, it is not enough to send the code back to the programmers and expect them to rewrite it until the score is back up. A conversation needs to happen so as to find the reason for which the score went down; maybe a key member of the team has been out sick, maybe work has turned onto a particularly complex domain, or maybe programmers have been plain sloppy. Only after finding the root cause we can understand what actions need to be taken to bring the score back up: lower expectations until out sick members are back, further training on the domain at hand, or stronger leadership.

The most important thing to take into account is that the Ten Guidelines, and their associated EC score, are not a measure of success by themselves, but rather a way to achieve success. Only then we will be able to ascertain the right level of flexibility for a given project, team, and time.

Ten Real-World Use Cases

Perhaps the most important takeaway of this report is the fact that the Ten Guidelines aren't just a theoretical exercise. In fact, the guidelines were created after examining what made software projects successful and maintainable, and once created, SIG applied them successfully to real-life projects.

In this chapter we will see 10 business cases where the Ten Guidelines were applied successfully.

Interamerican Greece

Interamerican Greece, or IAG, is the top insurance company in Greece. In 2010, it had 1,400 employees and worked with 1,800 intermediaries to provide insurance to more than one million Greeks. But by that time, the market was beginning to change. On one side, customers wanted to buy insurance directly from the insurer via the Internet; on the other side, intermediaries wanted to have more e-commerce options. IAG responded to the challenge to become the first company in Greece to sell insurance over the Internet.

However, the increased reliance on IT to drive business exposed a significant risk: much of IAG's IT infrastructure had suffered from years of underinvestment—some of the systems were 40 years old. Failures couldn't be completely ruled out, and these would be directly visible to the client, which would damage their reputation. IAG needed a new strategy.

The in-house programmers had already migrated some older systems to a newer Java platform, and the main question at IAG was whether this was the right approach for the other systems. SIG analyzed the migrated systems according to the Ten Guidelines using their EC score system, giving a score of three out of five. This wasn't as high as desirable, but it was enough to prove that the strategy was sound. Therefore, IAG set out to migrate further systems using their in-house teams but with the assistance of SIG to aim at a score of four out of five. Thanks to monthly reviews and constant communication with staff, the values of the Ten Guidelines were transmitted and applied, helping IAG's software nearly achieve the maximum score, five out of five.

Alphabet International

Alphabet International, part of the BMW Group, provides leasing and business mobility solutions in Europe. After acquiring ING Car Lease in 2011, their market share boosted dramatically: by 2012, the combined company had more than 490,000 cars under contract in 19 countries.

One of the main challenges in the acquisition of ING Car Lease, and in any other merger or acquisition, is consolidating the IT systems of both companies. Both Alphabet International and ING Car Lease had their own applications to manage lease contracts, some of them in-house and some of them provided by partners. SIG was brought in to analyze those applications, highlighting the ones that were too tied to operations in a particular region (and therefore not suitable for a global operation), those that didn't align with long-term objectives, and those that would become part of Alphabet International's main IT assets. This knowledge was instrumental for Alphabet International to lay out their long-term technological strategy, and provided a benchmarking framework to assess the quality of applications as they were being modified, created, or phased out.

Port of Rotterdam Authority

The Port of Rotterdam is the largest port in Europe and, until 2004, was also the busiest in the world. It is managed by the Port of Rotterdam Authority, a public company owned by the Municipality of Rotterdam and the Dutch State. Managing port activities is incredi-

bly complex, and until recently this was done by a collection of multiple, independent systems connected with varying degrees of luck.

In order to remain competitive, the Port of Rotterdam Authority realized that they needed a new, fully integrated solution, one that could be managed more effectively and efficiently. They named their new system HaMIS, the Harbour Master Management and Information System. Although they wanted to create this solution in-house, they realized that managing complex IT systems is not part of their core business; this added the long-term goal of externalizing management of HaMIS to a third-party provider at some point.

SIG was hired to assist in the creation and externalization of HaMIS. The Ten Guidelines, implemented in the EC scoring system, were used with a double objective. On one hand, the objective was to ensure that HaMIS provided the same level of functionality as the previous, scattered solutions. On the other hand, they also needed to provide empirical data that demonstrated the high quality of HaMIS, which could be used by the Port of Rotterdam Authority to negotiate a better deal when externalizing the management of HaMIS. What's more, once management had been transferred to a third party, the EC scoring system could be used as part of the Service Level Agreement to ensure that the third party kept quality high.

Care Schadeservice

Care Schadeservice is the largest car repair organization in the Benelux region. Repairing cars is a labor-intensive activity, which means repair costs can easily skyrocket for even the simplest breakdown. For this reason, Care Schadeservice aimed at finding as many efficiencies as possible, creating a streamlined workflow that minimized waste. Care Schadeservice initially used off-the-shelf IT products to manage their workflow, but being the biggest company in the care repair sector, they were the first to notice the limitations of standardised products. So in 2004 they set out to create their own in-house solution, CareFlow, with a second generation of the platform coming out six years later.

SIG was involved in the creation of both generations, although in different ways. During the creation of CareFlow v1, Care Schadeservice was part of a holding company listed in the stock market, which required them to include an objective assessment of any software

developed in-house; SIG, with their EC scoring system based on the Ten Guidelines, provided such assessment.

When CareFlow v2 was needed, Care Schadeservice decided that they wanted to outsource the creation and maintenance of the system but with the requirement of obtaining a product of certifiable quality. The EC scoring system was perfect for this. Care Schadeservice contracted NetRom to perform the work, and with the help of SIG established a minimum score of three out of five for the software.

Vhi Ireland

Vhi Ireland is a leading private health insurance company. Like many other companies, Vhi Ireland had a combination of systems that had evolved over decades. On top of this, in many occasions the systems had to be modified by external developers and subcontractors, which contributed to a lack of cohesiveness in the codebase.

The situation was such that Vhi Ireland had doubts regarding the suitability of their existing IT systems for their long-term objectives; their main concern was that there could be technical issues that slowed down or even blocked development of further projects. In order to objectively measure the state of affairs, Vhi Ireland contacted SIG to perform a full assessment of their software. The end result was the creation of a new online platform, which was implemented while instilling the quality values of the Ten Guidelines to developers and senior management.

Rabobank International

Rabobank International needed a new application for foreign currency transactions by bank employees and large clients. They analyzed existing packages but didn't find any suitable ones, so they turned to a custom-built application. Nexaweb was contracted to create and maintain the new foreign currency application: RITA.

Unfortunately, RITA didn't meet expectations. There were constant complaints from the trading floor regarding the stability of the application, and stakeholders were unhappy about the slow turnaround of new features. Rabobank International and Nexaweb had agreements specific to the architecture and code quality, but these

were too abstract to be discussed effectively. A more specific way to measure quality was needed.

With the help of SIG, the underlying problems of RITA were slowly unearthed. Fixing the architectural problems provided stability to the platform, while improving the quality of the code made it easier to work with, which lead to shorter release cycles. This, in turn, lifted pressure on each of the individual releases, since a feature missing a particular release would soon be included in the next one. This ultimately increased confidence in the platform, allowing it to be deployed globally.

Ministry of Infrastructure and Environment in the Netherlands

The Ministry of Infrastructure and Environment in the Netherlands is responsible for, among other things, the electronic annual environmental report (e-MJV). This report includes, for each company in the Netherlands, the total amount of harmful substances they have released into the water and air. These reports are then sent to the National Institute for Public Health and the Environment, and from there to the European emission register. This way, the Netherlands can make sure that they comply with international environmental agreements such as the Kyoto Protocol and the European environmental agreements.

Up until 2000, all the reports were obtained through paper forms. Since this was too inefficient, a desktop application was created so that companies could report these details electronically. The desktop application worked well for a few years, but then its architectural flaws became apparent: making sure that the right version of the application was being used on all computers was administratively costly. A web-based second generation was therefore commissioned.

However, this second version was rushed out due to legislative deadlines, which affected the end result. The application was slow, the interface wasn't user-friendly, and the system was buggy. Many changes were needed to bring the application to acceptable levels, and there were doubts about whether the state of the code would allow for such changes in a timely manner.

In order to understand what the best way to improve the application was, SIG was contacted to analyze e-MJV and provide an assessment

of weaknesses. Thanks to the Ten Guidelines, SIG was able to indicate that the code was of average quality (three out of five), and summarize the main issues within it. A decision was made to address the most important ones until the code quality achieved a score of at least four out of five, the level at which e-MJV could be considered maintainable. After this point, working with e-MJV became much easier, changes were performed faster, and the future of e-MJV seemed safer.

ProRail

ProRail is responsible for the maintenance and operation of the rail network in the Netherlands. The Netherlands might be a small country, but it boasts an incredibly busy rail network: 6,500 kilometres of track, 3,000 crossings, 4,500 kilometres of overhead lines, 8,600 railroad switches, and 390 railways stations. On top of this, the Netherlands is home to Rotterdam, where the largest port in Europe is located. This indicates the level of passenger and freight traffic that the Dutch rail network has to support.

Being responsible for the maintenance and operation of the rail network means that ProRail needs to control every single aspect of the infrastructure. In order to achieve this, ProRail makes use of technical drawings that describe each aspect of the infrastructure: a signalling drawing shows the different signals, an overhead drawing shows the structure of overhead power lines, and so forth. The administration of these drawings was a source of problems for ProRail for two reasons: one, most of the technical drawings were managed manually on paper; two, the different types of designs risked imposing restrictions on each other and often these restrictions weren't apparent until the installation was being performed, leading to delays and extra costs.

To fix this, ProRail made use of the services of LOXIA, who developed a new software to manage the drawings electronically. However, the responsibilities of this new software were far-reaching: security, safety, operational costs—there was a lot at stake. ProRail needed to ensure the software was of the highest quality, and SIG was employed to help with the quality assessment. SIG's score system based on the Ten Guidelines acted not only as a way to measure the quality of the software, but the score itself served as a motivating catalyst among staff. Now ProRail can easily share electronic designs

with operators to assist them in creating timetables for their trains while employees keep a vigilant eye over code quality for reliable results.

ING Bank

ING Bank is a global bank with a presence in more than 40 countries and a workforce comprising more than 75,000 people. For any company this big, constantly adapting to new practices and techniques is a matter of survival. ING Bank had already performed a transformation towards Agile development, but then they needed to take it one step further to institutionalize code quality beyond the scope of the team.

Work started with their Mijn ING website, a product with which more than 40 teams are involved. Each team curated its own code to make sure the quality was good from a local perspective, but there were fears of cross-team inefficiencies. To solve this, SIG was brought in to perform a holistic analysis. As we have seen, the Ten Guidelines also cover overarching architectural concerns, and this enabled SIG to pinpoint weaknesses that were invisible to the individual teams. Thanks to this, ING Bank was able to provide higher transparency regarding the quality of their software and, therefore, a better vision for senior management on the overall status.

Nykredit

Nykredit is the largest credit company in Denmark. Mortgages and commercial banking activity make up the core business of Nykredit, although they also offer personal pension plans and insurance. They also lead the way in Internet banking, having received the award for "the most digital company in Denmark" in 2010 and 2011.

By 2010, it became apparent that they had outgrown their IT infrastructure. The monolithic IT solution that Nykredit had hitherto used was aimed at small- and medium-sized banks, and it was time to move to a new, more tailored solution. Nykredit decided to use Finacle by Infosys, but with an important number of customizations. They were planning to employ 150 people over five years for this project, and they needed a way to make sure their investment would go in the right direction.

SIG helped Nykredit create their own quality assessment tool: Nykredit Quality Tooling, or NQT. NQT analyzes all the Java code written by or for Nykredit, highlighting portions that violate the Ten Guidelines or other quality metrics established as part of their agreement; this flagged code is reviewed by software architects before a final report is produced for senior management. Thanks to this, Nykredit has been able to gradually deprecate their old platforms and substitute them with the new software.

About the Author

Abraham Marín-Pérez is an independent Java programmer, author, public speaker, and Agile consultant. He helps organizations achieve their objectives through a number of varying challenges, both technical and non-technical. He also helps run the London Java Community, and contributes as a Java Editor at InfoQ.